


Priority queue ADT

part 1

CS 146 - Spring 2017



Today

- Dijkstra's algorithm
- The minimum spanning tree problem
- The cut property for MSTs
- Prim-Dijkstra-Jarnik algorithm
- Kruskal's algorithm

```

map bfs(graph G, vertex s) {
    dist = new map()
    queue = new FIFOqueue()

    for every vertex v in G
        dist.put(v, +inf)
    dist.put(s, 0)
    queue.enqueue(s)

    while queue not empty {
        v = queue.dequeue()
        for each neighbor w of v {
            if dist.get(w) == +inf {
                dist.put(w,
                    dist.get(v) + 1)
                queue.enqueue(w)
            }
        }
    }
    return dist
}

```

```

map dijkstra(weighted-graph G, vertex s)
    dist = new map()
queue = new priorityQueue()

    for every vertex v in G
        dist.put(v, +inf)
    dist.put(s, 0)
queue = new priorityQueue(dist)

    while queue not empty {
        v = queue.extractMin()
        for each neighbor w of v {
            if w should be updated {
                dist.put(w,
                    dist.get(v) + weight(v,w))
queue.decreaseKey(w)
            }
        }
    }
    return dist
}

```

Dijkstra's
algorithm:
first steps
(incomplete)

```

map dijkstra(weighted-graph G, vertex s) {
    dist = new map()


    for every vertex v in G
        dist.put(v, +inf)
    dist.put(s, 0)
    queue = new priorityQueue(dist)

    while queue not empty {
        v = queue.extract-min()
        for each neighbor w of v {
            if dist.get(w) > dist.get(v) + weight(v,w) {
                dist.put(w, dist.get(v)+ weight(v,w))
                queue.decreaseKey(w)
            }
        }
    }
    return dist
}

```

Dijkstra's algorithm

Example: Dijkstra's algorithm



```

map dijkstra(weighted-graph G, vertex s) {
    dist = new map()

    for every vertex v in G
        dist.put(v, +inf)
    dist.put(s, 0)
    queue = new priorityQueue(dist)

    while queue not empty {
        v = queue.extractMin()
        for each neighbor w of v {
            relax(v->w)
            if dist.get(w) > dist.get(v) + weight(v,w) {
                dist.put(w, dist.get(v)+ weight(v,w))
                queue.decreaseKey(w)
            }
        }
    }
    return dist
}

```

updates dist to w
via path through edge v->w


```
map dijkstra(weighted-graph G, vertex s) {  
    dist = new map()
```

```
    for every vertex v in G  
        dist.put(v, +inf)
```

```
    dist.put(s, 0)
```

```
    queue = new priorityQueue(dist)
```

```
    while queue not empty {
```

```
        v = queue.extractMin()
```

```
        for each neighbor w of v {
```

```
            if dist.get(w) > dist.get(v) + weight(v,w) {
```

```
                dist.put(w, dist.get(v)+ weight(v,w))
```

```
                queue.decreaseKey(w)
```

```
            }
```

```
        }
```

```
    }
```

```
    return dist
```

```
}
```

← once per vertex

← once per edge
**over the entire
algorithm**

(not just inside loop)

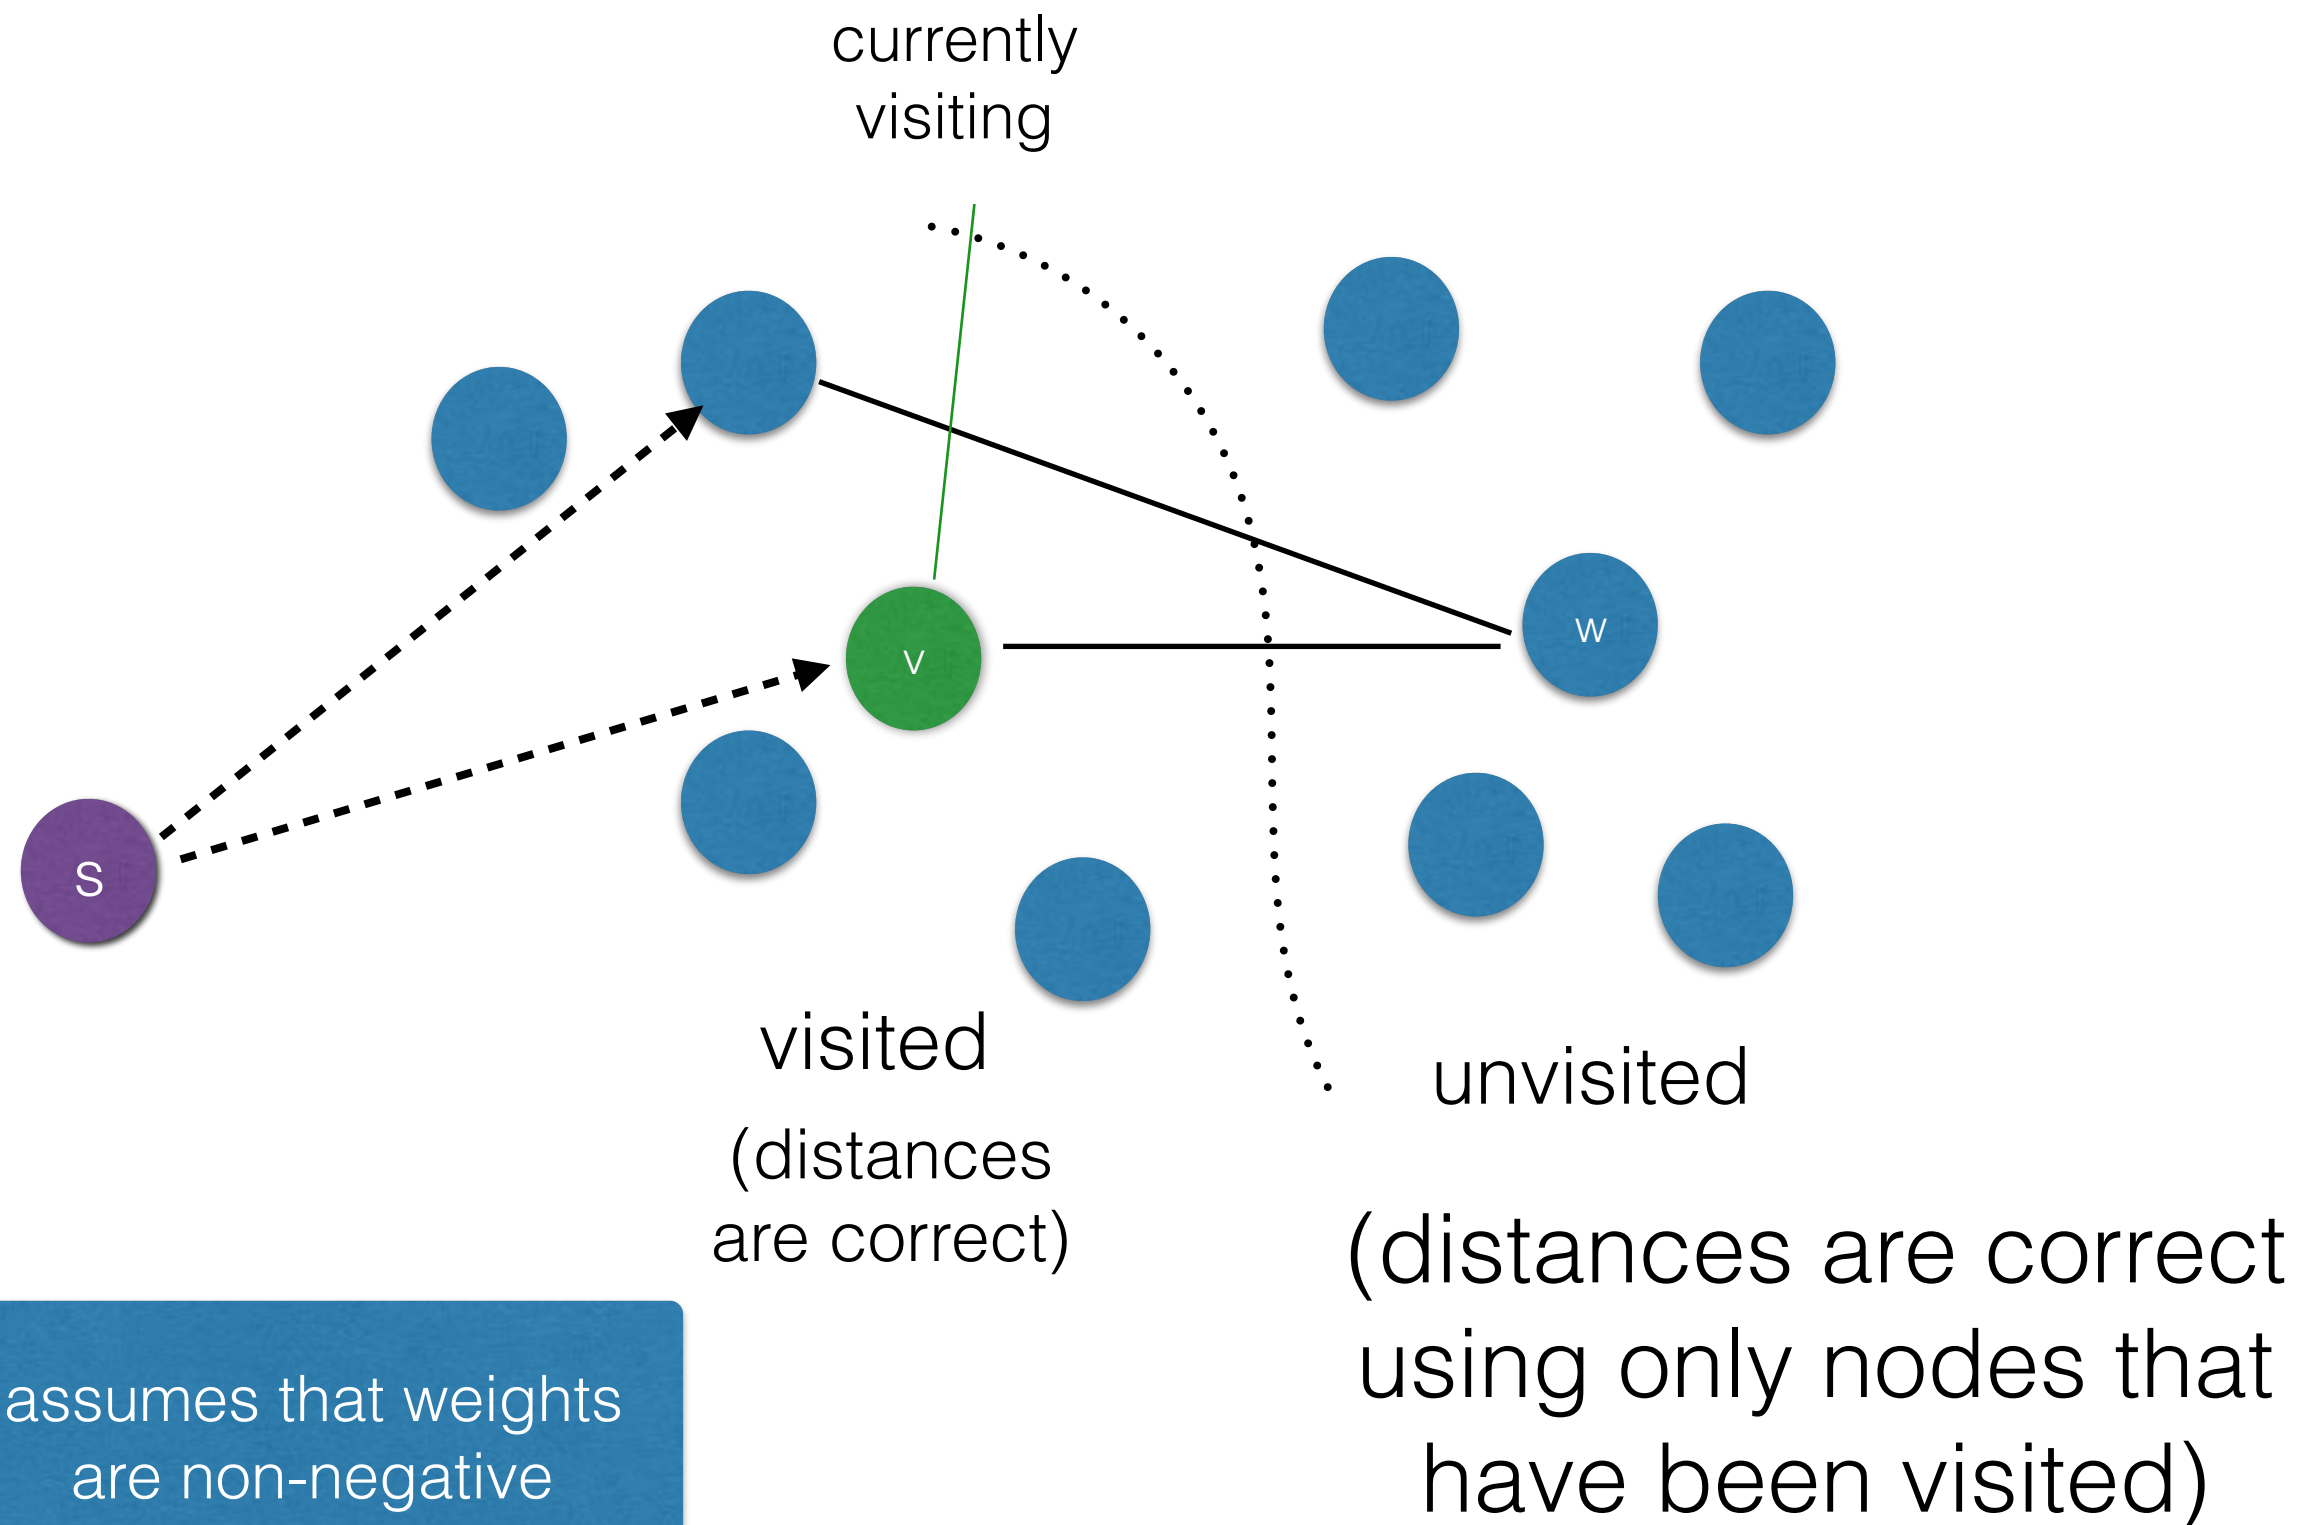
Running time of Dijkstra's algorithm

$$\begin{aligned} &\leq V + 4E \times \text{time}(\text{dictionary op}) \\ &\quad 1 \times \text{time}(\text{queue.makeQueue}) \\ &\quad \mathbf{V} \times \text{time}(\text{queue.extractMin}) \\ &\leq \mathbf{E} \times \text{time}(\text{queue.decreaseKey}) \end{aligned}$$


up to a constant factor, #queue ops = #dict ops,
but dictionary ops are constant time with hash table
running time is dominated by queue operations

$$O(T(\text{makeQueue}(V)) + V \times T(\text{extractMin}(V)) + E \times T(\text{decKey}(V)))$$

Why is Dijkstra's algorithm correct?




Shortest paths from B



The shortest paths from a vertex to all other nodes **form a tree**.

Why?

B's shortest path tree



```


map augmented-dijkstra(weighted-graph G, vertex s) {
    dist = new map()
    prev = new map()    ← maps nodes to previous node on
                        source's shortest path tree
    for every vertex v in G {
        dist.put(v, +inf)
        prev.put(v, null)
    }
    dist.put(s, 0)
    queue = new priorityQueue(dist)

    while queue not empty {
        v = queue.extract-min()
        for each neighbor w of v {
            if dist.get(w) > dist.get(v) + weight(v,w) {
                dist.put(w, dist.get(v)+ weight(v,w))
                queue.decreaseKey(w)
                prev.put(w, v)
            }
        }
    }
}


return prev

```


Retrieving the actual path
step 1: record shortest path tree



Retrieving the actual path
step 2: reconstruct path to some target node



Minimum spanning trees



What is the cheapest possible power grid that will connect all the cities?


Observation

- The solution cannot contain cycles
- b/c removing an edge from this cycle would reduce the cost w/o compromising connectivity
- The solution must be a tree which we shall call...

Minimum spanning tree

- Input: undirected connected weighted graph $G = (V, E)$
- Output: a tree $T = (V, E')$ with $E' \subseteq E$ that minimizes

$$\text{weight}(T) = \sum_{e \in E'} \text{weight}(e)$$




how many
can you find?

Observation

- minimum spanning trees are not unique!
- a graph can have more than one
- but all will be the same weight, obviously...

Problem:
find an algorithm for finding the
minimum spanning tree of a graph



Ideas


1. from a single vertex, grow a tree by repeatedly adding a minimum-edge weight connecting a vertex not already in the tree.
2. construct tree by repeatedly adding the next lightest edge that doesn't produce a cycle

Beware!

The greedy approach doesn't always produce the best solution.



What's a graph cut?




Cut = a **partitioning** of the vertices into 2 sets

aka splitting or division

formally, a pair $(X, V \setminus X)$
where $X \subseteq V$, both non-empty


What's a graph cut?



Edges crossing the cut: what we are really after

Cut property

“For a given cut, the lightest edge across the cut, if unique, is part of all MSTs.”












if not, the edge can be swapped out


Prim Dijkstra Jarnik's algorithm

idea: from a single vertex, grow a tree by repeatedly adding a minimum-edge weight connecting a vertex not already in the tree.

Jarnik (1930), Prim (1957), Dijkstra (1959)



Why is PDJ's algorithm correct?



At each iteration,
PDJ picks the lightest edge crossing the cut
between the vertices already in MST and the rest.

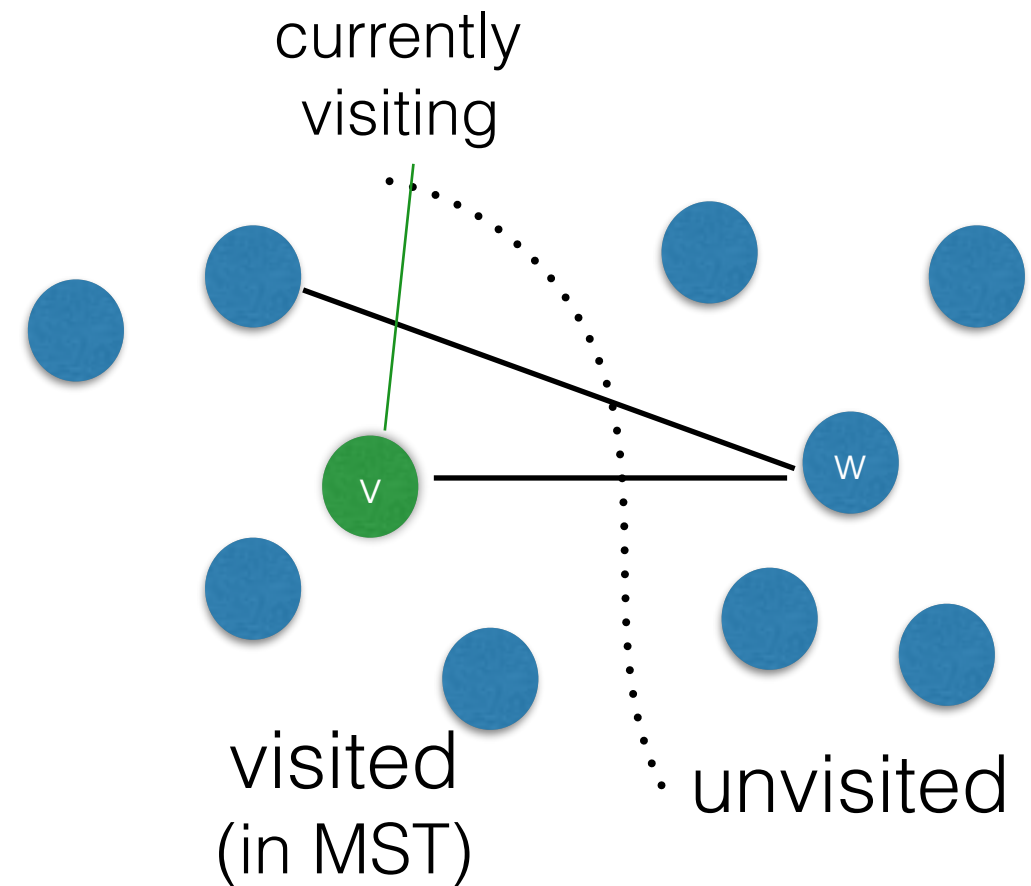
By cut property, this edge must be in MST.

```

map primDijkstraJarnik(weighted-graph G, vertex s) {
  cost = new map()
  prev = new map()
  for every vertex v in G {
    cost.put(v, +inf)
    prev.put(v, null)
  }
  dist.put(s, 0)
  queue = new priorityQueue(cost)

  while queue not empty {
    v = queue.extract-min()
    for each neighbor w of v {
      if cost.get(w) > weight(v,w) {
        cost.put(w, weight(v,w))
        queue.decreaseKey(w)
        prev.put(w, v)
      }
    }
  }
  return prev
}

```



update to cheapest edge
between explored and w

```

map primDijkstraJarnik(weighted-graph G, vertex s) {
    cost = new map()
    prev = new map()
    for every vertex v in G {
        cost.put(v, +inf)
        prev.put(v, null)
    }
    dist.put(s, 0)
    queue = new priorityQueue(cost)

    while queue not empty {
        v = queue.extract-min()
        for each neighbor w of v {
            if cost.get(w) > weight(v,w) {
                cost.put(w, weight(v,w))
                queue.decreaseKey(w)
                prev.put(w, v)
            }
        }
    }
    return prev
}

```



running time?

Kruskal's algorithm

idea: construct tree by repeatedly add the next lightest edge
that doesn't produce a cycle

More precisely...

forest



no cycle

disconnected

3 connected
components

tree



no cycle

connected

1 connected
component

general graph




cycles

connected

1 connected
component

tree




no cycle
connected
1 connected
component

general graph



cycles
connected
1 connected
component



cycles
disconnected
3 connected
components









Kruskal's algorithm

More precisely...


idea: construct tree by

starting with a forest of single-vertex components

repeatedly adding to the forest the next lightest edge
that doesn't produce a cycle




Why is Kruskal's algorithm correct?




pretend the green edge has weight 4

At each iteration, Kruskal picks the lightest edge that does not create a cycle.

Why is Kruskal's algorithm correct?



Put all other components to one side or the other, we have a cut.



The edge chosen by Kruskal is the lightest across this cut.

By cut property, this edge must be in MST.

```
map Kruskal(weighted graph G) {  
    mst = new Graph() with same vertices as G and no edges  
    sortedEdges = sort edges of G in order of increasing weight  
    for each edge (u,v) in sortedEdges {  
        if u and v are in different connected components of mst  
            add (u,v) to mst  
    }  
    return mst  
}
```


$E \log E$

```
map Kruskal(weighted graph G) {  
    mst = new Graph() with same vertices as G and no edges  
    sortedEdges = sort edges of G in order of increasing weight  
    for each edge (u,v) in sortedEdges {  
        if u and v are in different connected components of mst  
            add (u,v) to mst  
    }  
    return mst  
}
```

loops E times

$E \log E + E \times \text{Time}(\text{determine if in same component of MST})$

How to determine if u and v are in different components?

Approach 1

do DFS (tree traversal) on MST with u as starting vertex.

- $O(V + E)$ each time
- $O(E \log E + E (V+E)) = O(E^2)$ total time

How to determine if u and v are in different components?

Approach 2

use auxiliary special-purpose data structure called **union-find** to keep track of disjoint sets

- $\text{find}(x)$: find the designated representative element of the set to which x belongs
- $\text{union}(x,y)$: combine the sets to which x and y belong
- $\text{makeset}(x)$: create a singleton set containing x only

some operations may run slow, some fast, but k of these operations run in $O(k \alpha(n))$ where n is the number of elements.

```

map Kruskal(weighted graph G) {
    mst = new Graph() with same vertices as G and no edges
    ds = new UnionFind()
    for v in G
        ds.makeset(v)
    sortedEdges = sort edges of G in order of increasing weight
    for each edge (u,v) in sortedEdges {
        if u and v are in different connected components of mst
        add (u,v) to mst
        if ds.find(u) is not ds.find(v) {
            add (u,v) to mst
            ds.union(u,v)
        }
    }
    return mst
}

```

```

map Kruskal(weighted graph G) {
    mst = new Graph() with same vertices as G and no edges
    ds = new UnionFind()
    for v in G
        ds.makeset(v)
    sortedEdges = sort edges of G in order of increasing weight
    for each edge (u,v) in sortedEdges {
        if u and v are in different connected components of mst
        add (u,v) to mst
        if ds.find(u) is not ds.find(v) {
            add (u,v) to mst
            ds.union(u,v)
        }
    }
    return mst
}

```

$V \alpha(V)$

$E \log E$

$\leq 3E \alpha(V)$

Time is $O(E \log E + (V+3E) \alpha(V)) = O(E \log V)$

Priority queues: recap

- used in greedy algorithms where the best possible move is processed first
 - Dijkstra's alg to solve single source shortest paths
 - PDJ alg to solve the MST problem
- alternative approach: sort before processing
 - Kruskal's alg to solve the MST problem