

Operating systems

Memory: speeding up address translation
free memory management


Review

What is paging good for?


Today

- Speeding up address translation
- Managing free memory

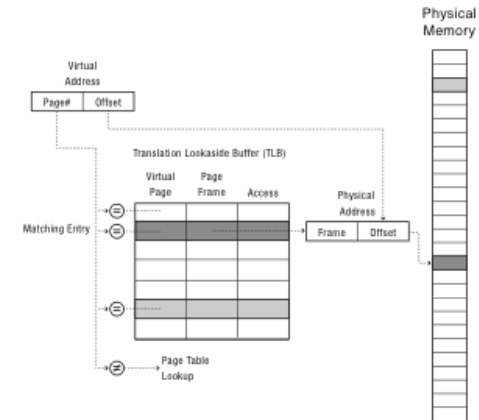
add reg1, reg2 ← fetch instruction
mult reg1, 2 ← fetch instruction



Translation lookaside buffer (TLB)

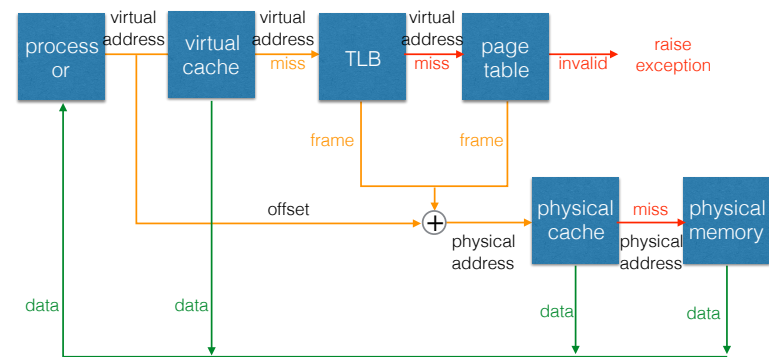


Inside the TLB



$$\text{cost of address translation} = \text{cost of TLB lookup} + \text{cost of full translation} \times \text{probability of a miss}$$

Caches galore




- Caching is a standard technique for speeding up data lookup
- Challenge #1: how to maintain consistency?
- Challenge #2: what to keep in cache, what to kick out?


- How do you minimize memory fragmentation?



Hole coalescing on a release



Holes as linked lists



Choosing a hole to allocate

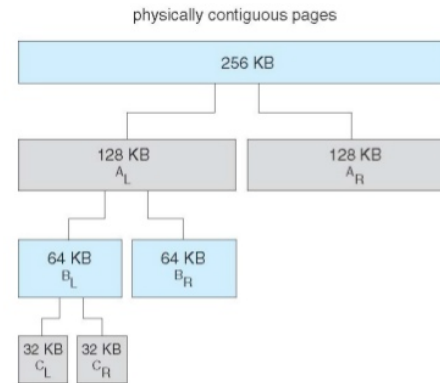
Issues to consider

- utilization
- external fragmentation
- search time

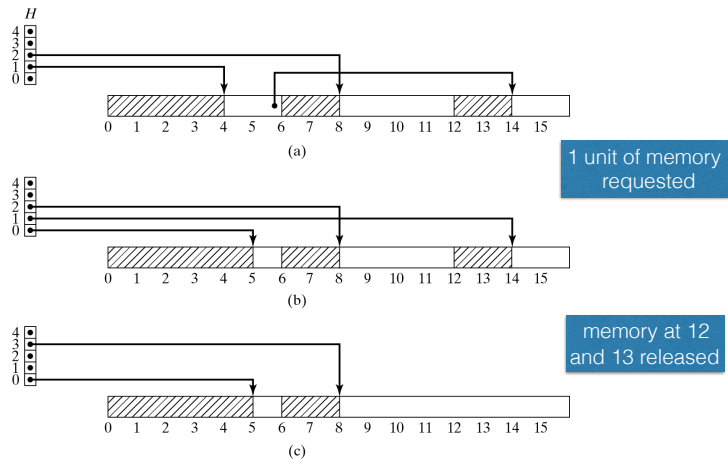
Allocation strategies

- first-fit
- next-fit
- best-fit
- worst-fit


Buddy system



Buddy system: example



Slab allocation

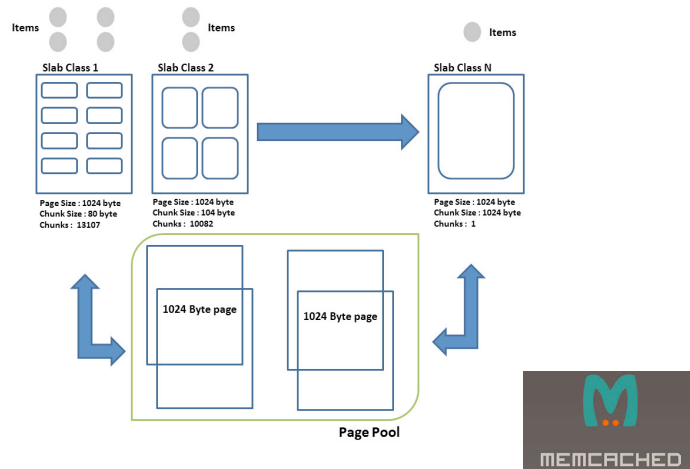


slab: one or more physically contiguous pages

cache: one or more slabs, one cache per kernel data structure type



Slab allocation




Managing insufficient memory


- Memory compaction
- Swapping
- Overlays

Memory compaction

consolidating smaller holes
dispersed throughout memory
into a single larger hole




Swapping



selecting a process and
temporarily evicting it
to secondary storage

Overlays

different portions of the program
replace each other in memory
as execution proceeds.



When does free memory need to be managed?





- OS level
 - (early on) in segmented memory systems
 - to manage kernel memory
- program level
 - in languages such as C and C++: malloc/free, new/delete to manage the heap
 - Java has **new** but no **delete**, why?
 - when you know the usage pattern of your application and handling it yourself is faster: eg memcached

GARBAGE COLLECTION




- reference counting
- mark-sweep
- mark-compact
- copying-collector


GARBAGE COLLECTION



- **reference counting**
- mark-sweep
- mark-compact
- copying-collector



GARBAGE COLLECTION



- reference counting
- **mark-sweep**
- mark-compact
- copying-collector

Recap: memory management

- **efficient and flexible memory use**
 - paging and segmenting system
 - multi-level page tables -> sparse addressing
 - shared code
- **security and isolation**
 - branch and bound
 - read/write access
- **speeding up data retrieval**
 - caching: TLB, virtual and physical caches
- **managing free memory (depends on the actual use)**
 - not a problem if chunks are uniform size
 - techniques: coalescing, hole selection, buddy system, slab allocation
 - language support for managing heap: malloc/free, garbage collection