

Operating Systems

Memory: caching

Last time...

Last time...

- How do you minimize memory fragmentation?



When does free memory need to be managed?

- OS level
 - (early on) in segmented memory systems
 - to manage kernel memory
- program level
 - in languages such as C and C++: malloc/free, new/delete to manage the heap
 - Java has **new** but no **delete**, why?
 - when you know the usage pattern of your application and handling it yourself is faster: eg memcached


Recap: memory management

- **efficient and flexible memory use**
 - paging and segmenting system
 - multi-level page tables -> sparse addressing
 - shared code
- **security and isolation**
 - branch and bound
 - read/write access
- **speeding up data retrieval**
 - caching: TLB, virtual and physical caches
- **managing free memory (depends on the actual use)**
 - not a problem if chunks are uniform size
 - techniques: coalescing, hole selection, buddy system, slab allocation
 - language support for managing heap: malloc/free, garbage collection

Caching

- why? (main idea, examples)
- how? (challenges)


Why? caching speeds up data lookups




get (data):
 if data in cache: **cache hit!**
 return data
 else: **cache miss**
 get it from data source
 save data to cache

average fetch time =
 $\text{Pr}(\text{hit}) \times (\text{latency from cache}) +$
 $\text{Pr}(\text{miss}) \times (\text{latency from source})$

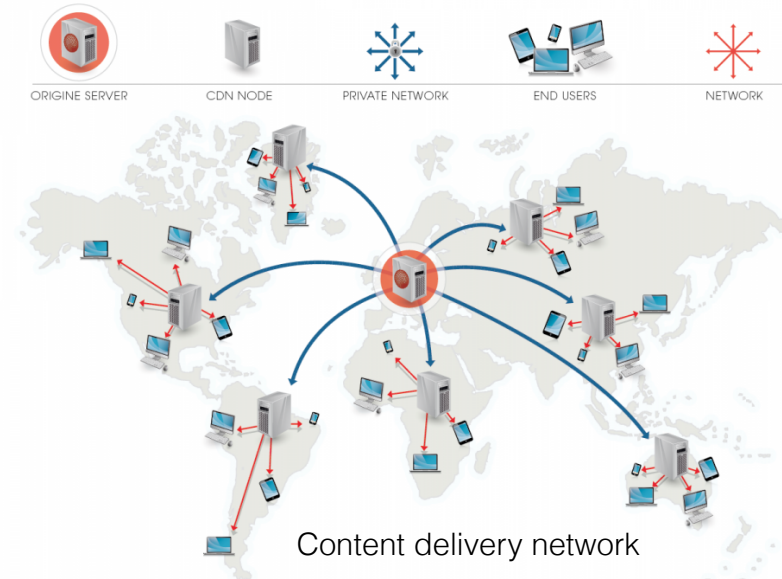
Extended Memory Hierarchy



Source: http://www.ts.avnet.com/uk/products_and_solutions/storage/hierarchy.html



caching in a network

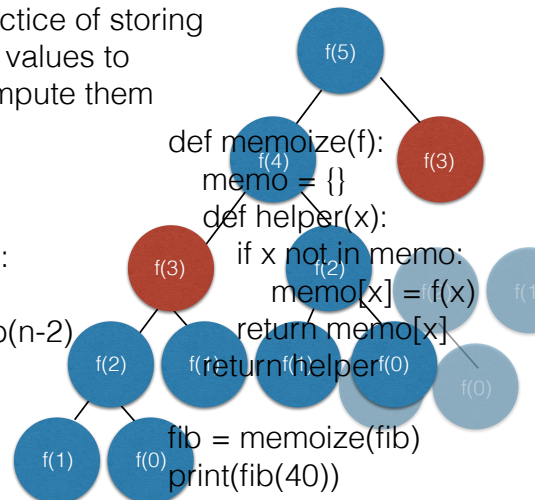


Source: <https://4cornermedia.com/wp-content/uploads/2016/03/CDN-1.png>

Memoization

in algorithms, the practice of storing
previously computed values to
avoid having to recompute them

```
def fib(n):
    if n == 0 or n == 1:
        return 1
    return fib(n-1) + fib(n-2)
```




Examples of caching


- hardware level:
 - TLB: cache for addresses
 - L1-L3: data caches
- OS level: demand paging
- software level: memoization
- network level
 - browser, webcache, server cache, database cache
 - content delivery networks

Caching: challenges


- efficient lookups (today)
- cache coherence (a little today, mostly later)
- replacement policy (next lecture)



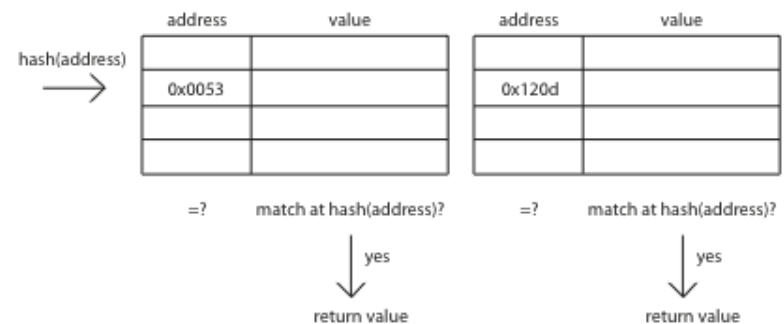
Lookups at the hardware level:
fully associative cache



Lookups at the hardware level:
direct mapped



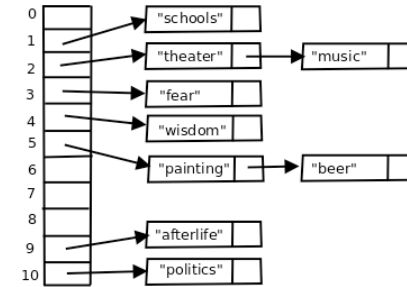
Lookups at the hardware level:
set associative



Hash tables


- $O(1)$ lookups on average

How do you resolve collisions?



hash chaining


How do you resolve collisions?



Open addressing: store values directly in the array


Cuckoo hashing

Easy Insertion




Cuckoo hashing

Inserting With 1 Conflict




Cuckoo hashing

Inserting With 2 Conflicts



Cuckoo hashing

Infinite Loop



Cuckoo hashing

- Pagh and Rodler, 2001
- expected amortized $O(1)$
- $O(1)$ lookups in worst case, rather than expected case




Efficient lookups

- CPU cache
 - fully associative cache
 - direct mapped cache
 - set associative cache
- software: hash table
 - hash chaining
 - open addressing: linear probing, **cuckoo hashing**


Cache coherence

- how to detect and repair stale data?

2 writing policies

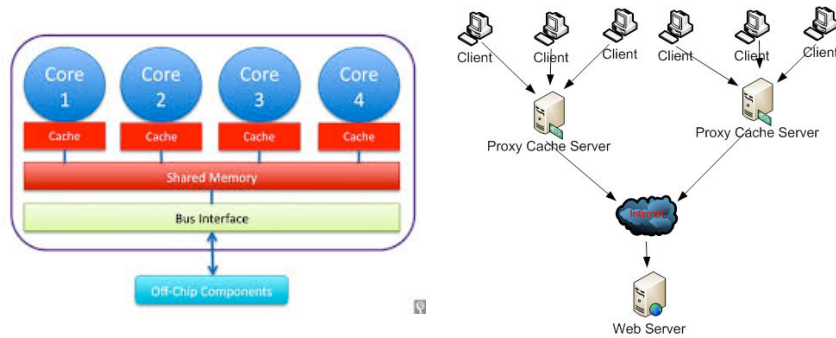


The problem with multiple caches



What time should we go to the movies?

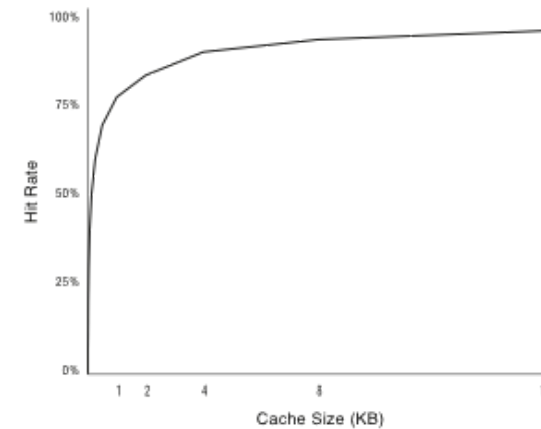
Caching of shared memory




Sources: http://www.codeproject.com/KB/web-cache/ExploringCaching/cache_array.jpg
<http://embedded-computing.com/articles/conquering-concurrency-bugs-multicore-systems/>

Working set

the amount of memory a process requires at any given time




Zipf distribution



Recap

- caching speeds up data lookups if the data is likely to be re-requested again
- data structures for $O(1)$ -lookup
 - set-associative (hardware)
 - hash tables (software)
- cache coherence
 - with 1 cache: buffer writes back to memory
 - with multiple caches: things get complicated
- next time: replacement policy



Acknowledgements

- <http://ospp.cs.washington.edu/slides.html>
- cuckoo hashing explanation: Bridget Howell, Joe Whitney: <https://www.youtube.com/watch?v=OBuGqu2d4v4>