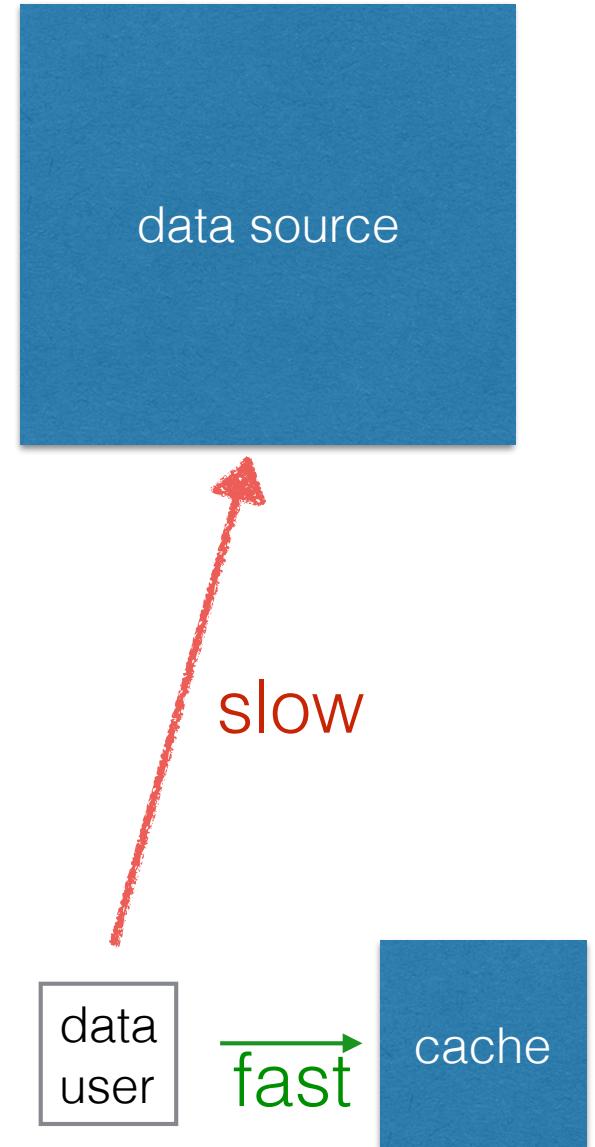


Operating Systems

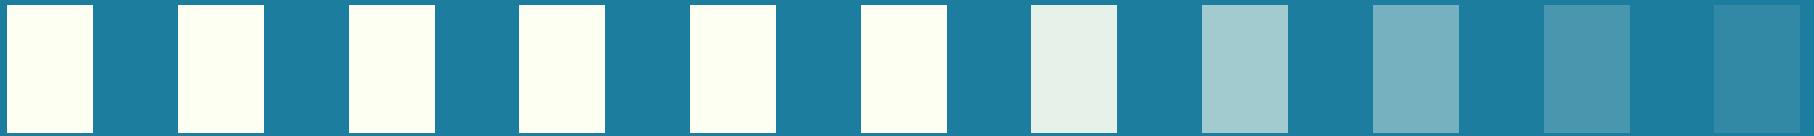
Memory: replacement policies

Last time

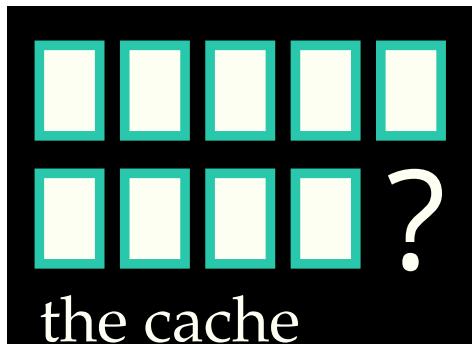
- **caching speeds up data lookups if the data is likely to be re-requested again**
- **data structures for O(1)-lookup**
 - set-associative (hardware)
 - hash tables (software)
- **cache coherence**
 - with 1 cache: buffer writes back to memory
 - with multiple caches: things get complicated
- today: **replacement policy**



THE PROBLEM



input: a sequence of requests for objects



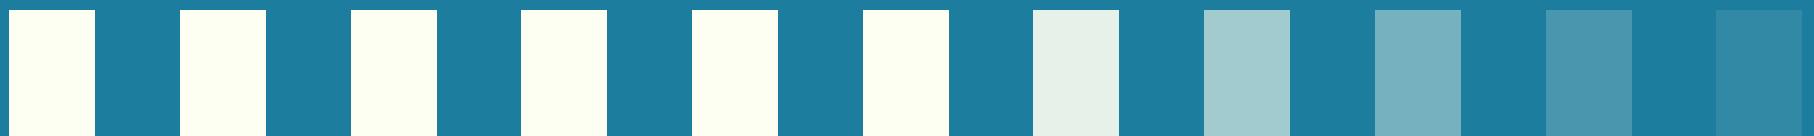
GOAL
minimize number of cache misses

clairvoyant

LRU

FIFO
LFU

BASIC ALGORITHMS



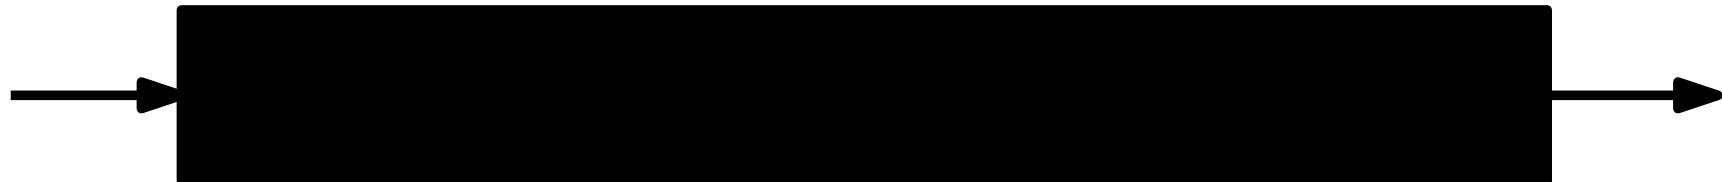
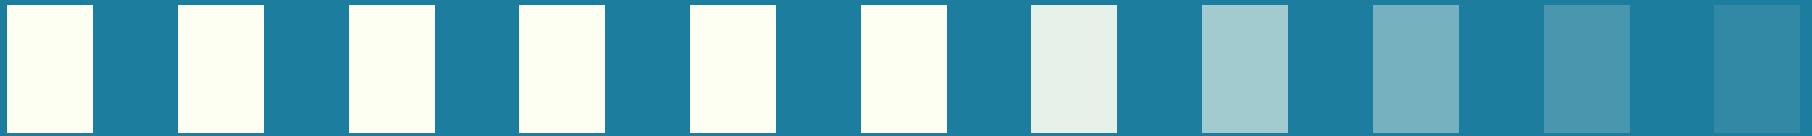
BASIC ALGORITHMS

LRU

clairvoyant

LFU

FIFO



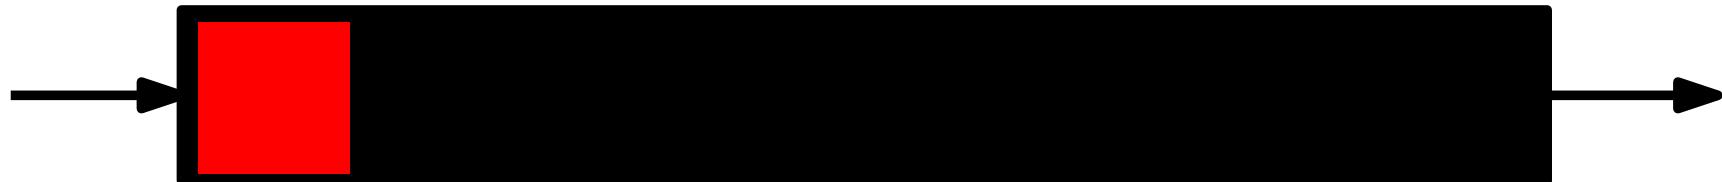
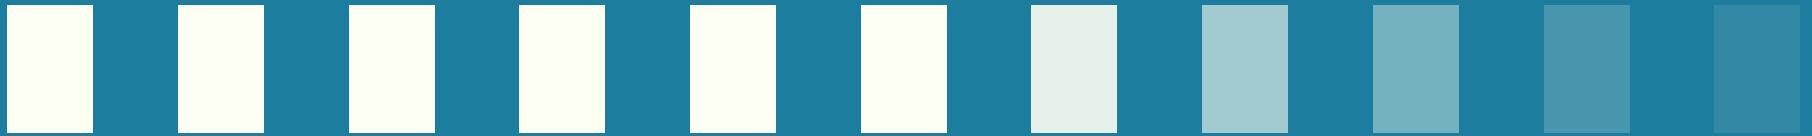
BASIC ALGORITHMS

LRU

clairvoyant

LFU

FIFO



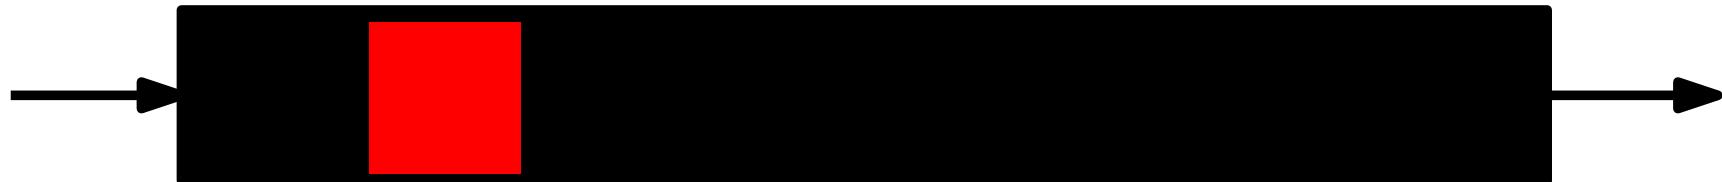
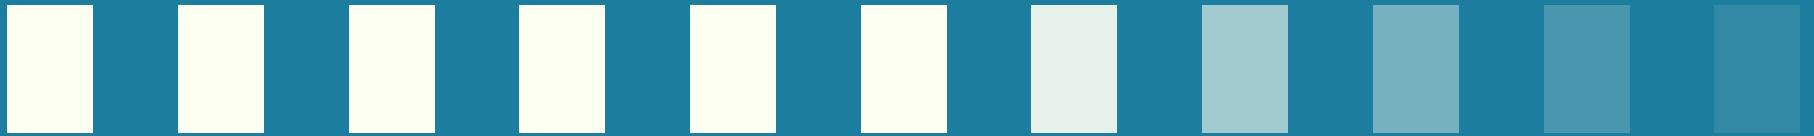
BASIC ALGORITHMS

LRU

clairvoyant

LFU

FIFO



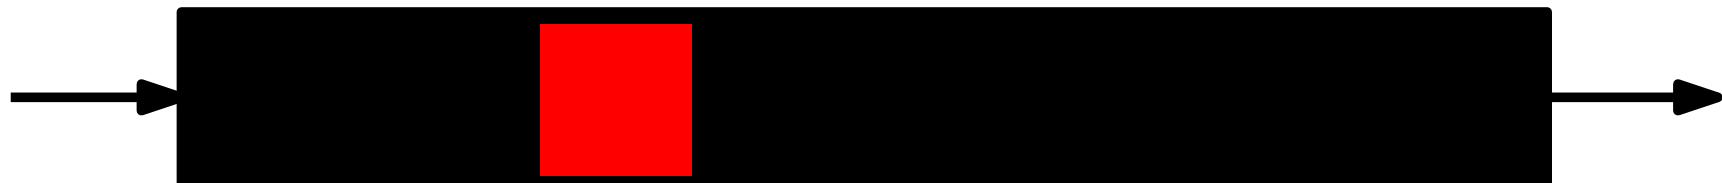
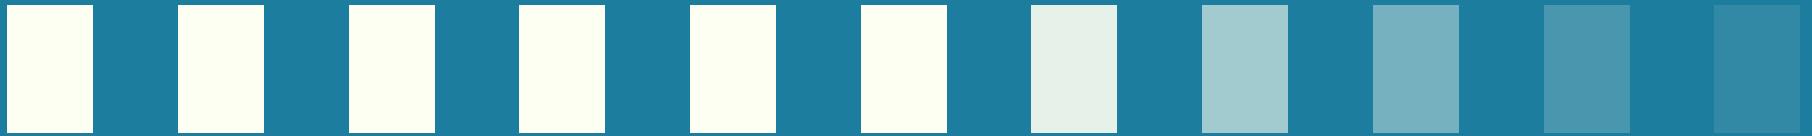
BASIC ALGORITHMS

LRU

clairvoyant

LFU

FIFO



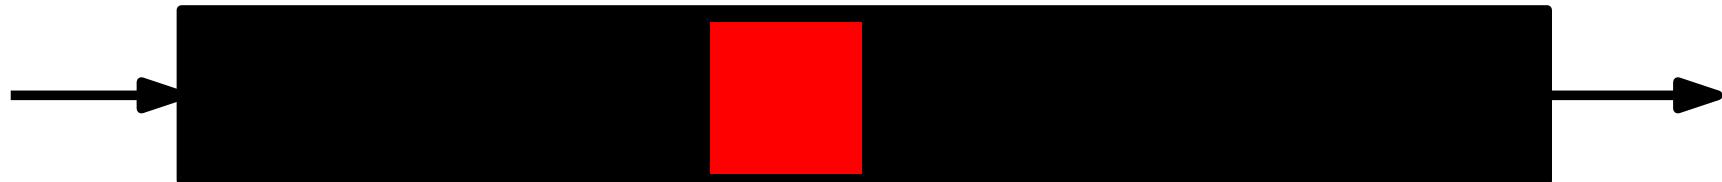
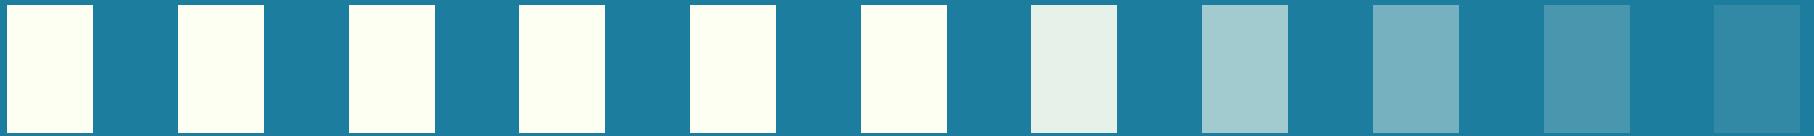
BASIC ALGORITHMS

LRU

clairvoyant

LFU

FIFO



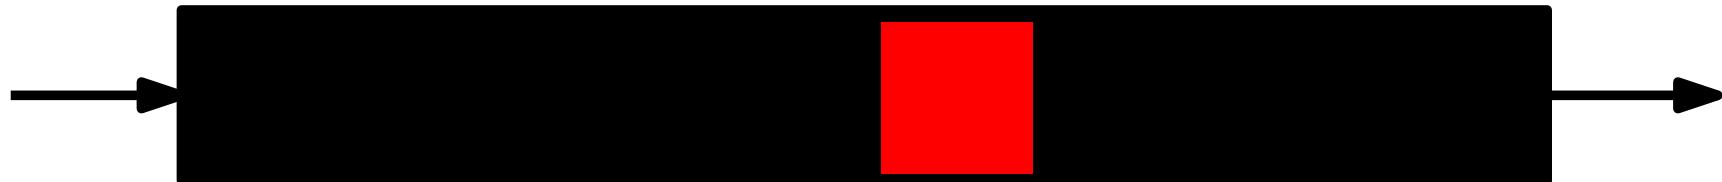
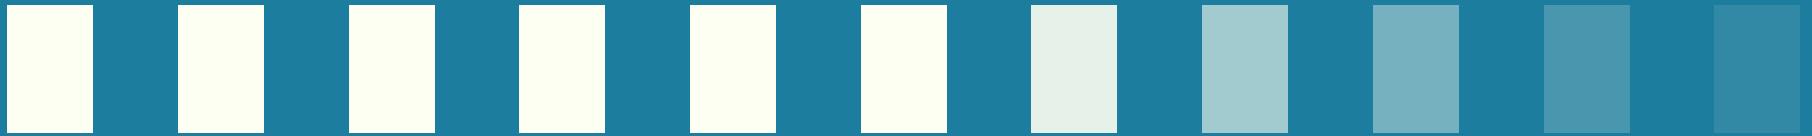
BASIC ALGORITHMS

LRU

clairvoyant

LFU

FIFO



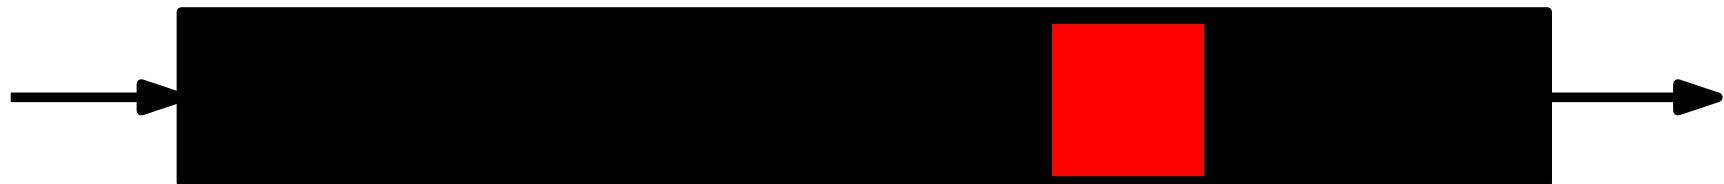
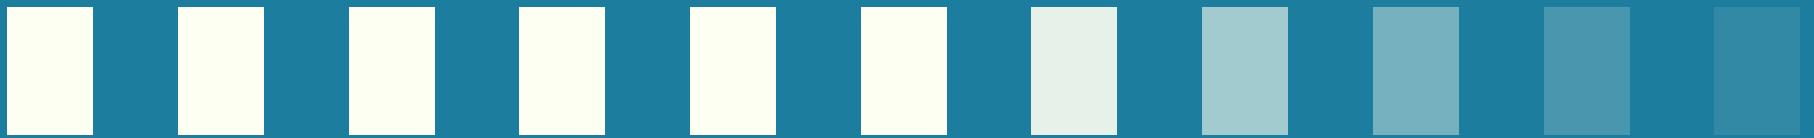
BASIC ALGORITHMS

clairvoyant

LRU

FIFO

LFU



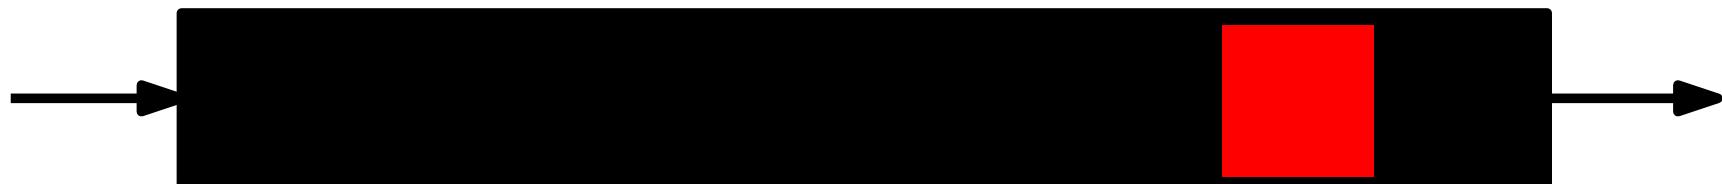
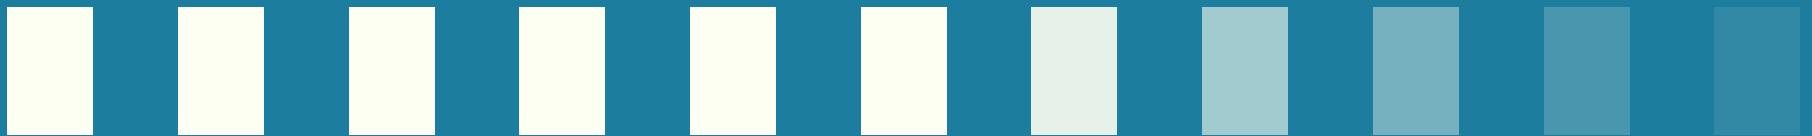
BASIC ALGORITHMS

clairvoyant

LRU

FIFO

LFU



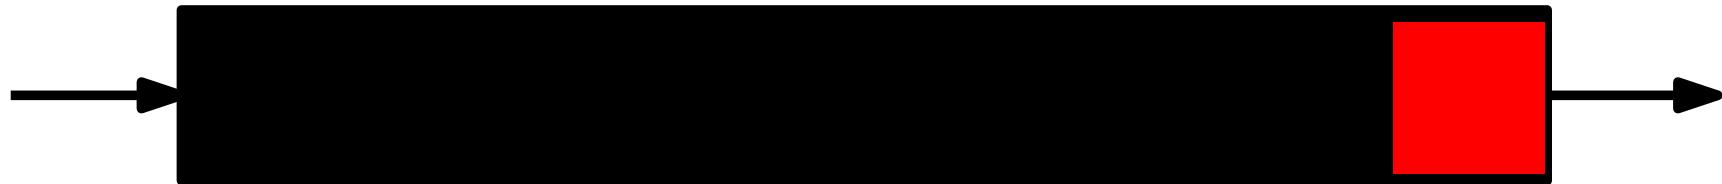
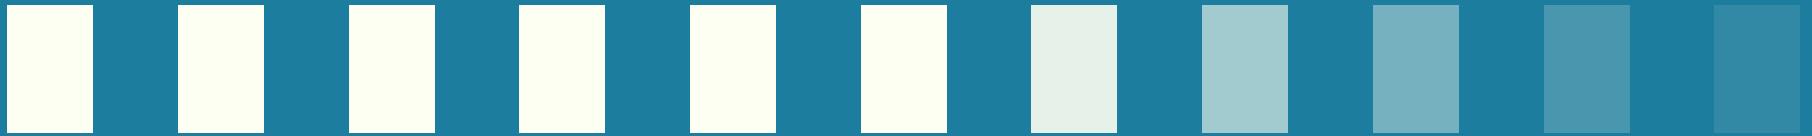
BASIC ALGORITHMS

LRU

clairvoyant

LFU

FIFO



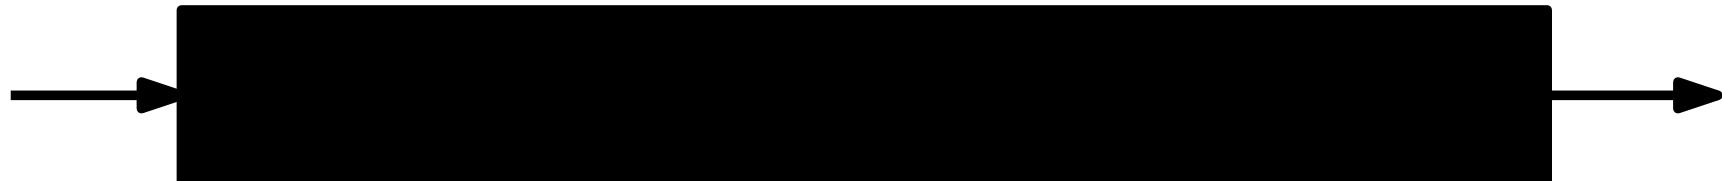
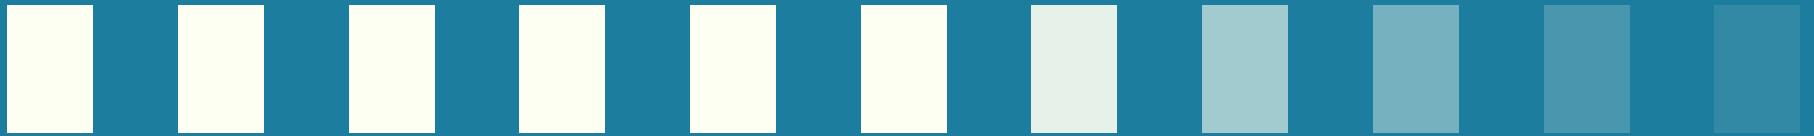
BASIC ALGORITHMS

clairvoyant

LRU

LFU

FIFO



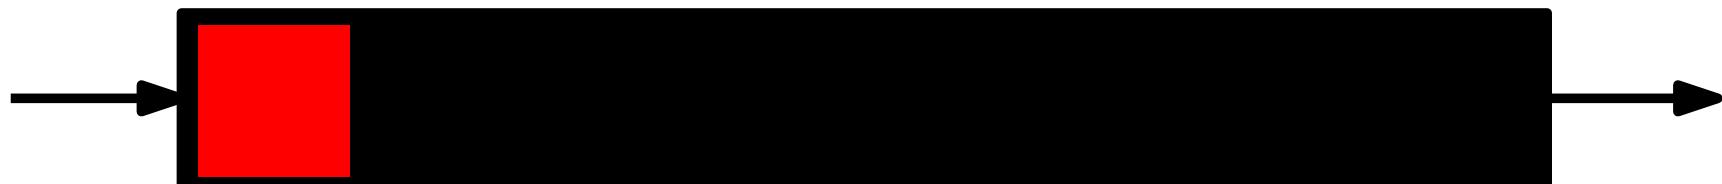
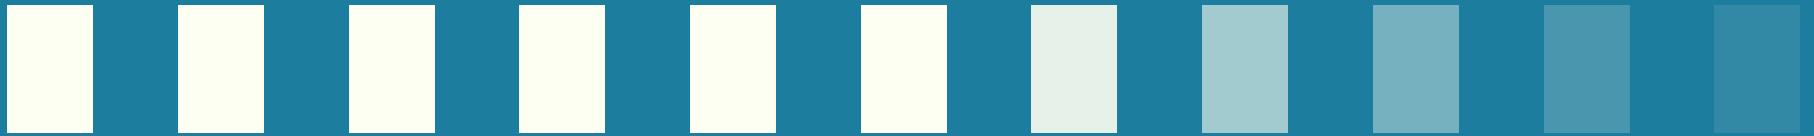
BASIC ALGORITHMS

clairvoyant

LRU

LFU

FIFO



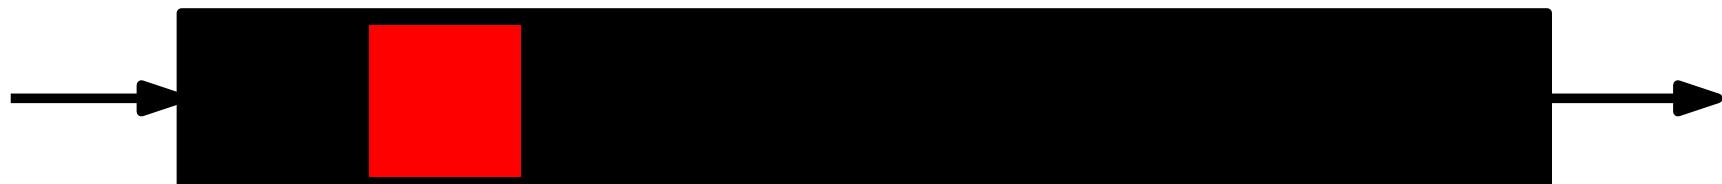
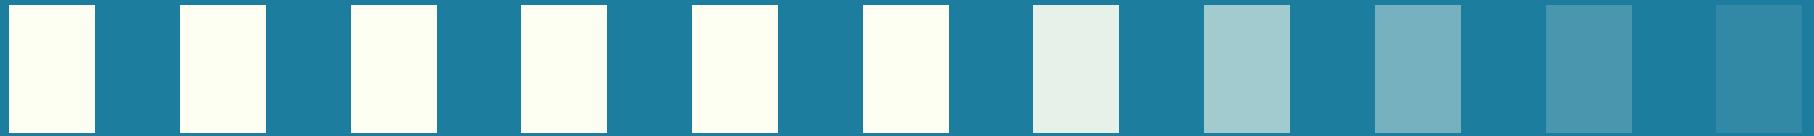
BASIC ALGORITHMS

clairvoyant

LRU

LFU

FIFO



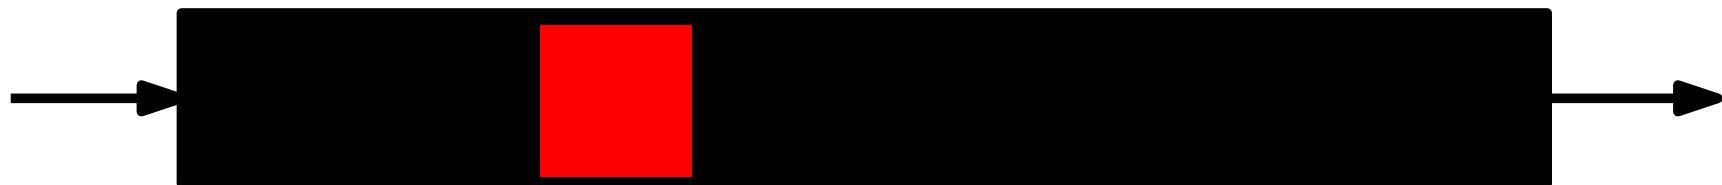
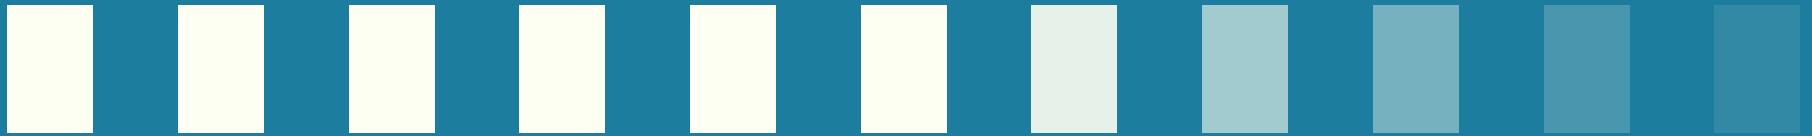
BASIC ALGORITHMS

clairvoyant

LRU

LFU

FIFO



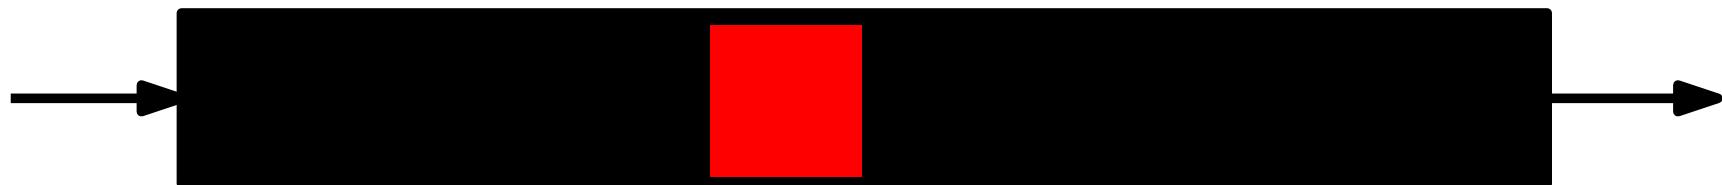
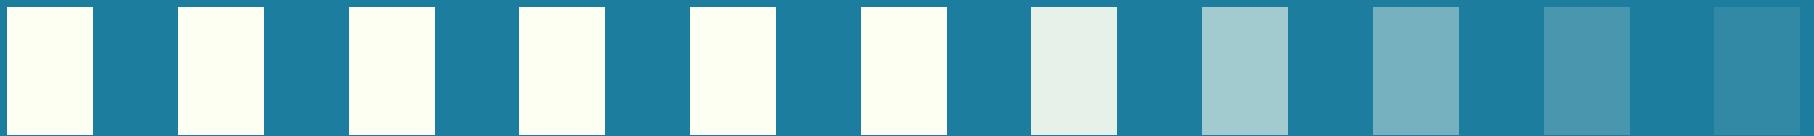
BASIC ALGORITHMS

clairvoyant

LRU

LFU

FIFO



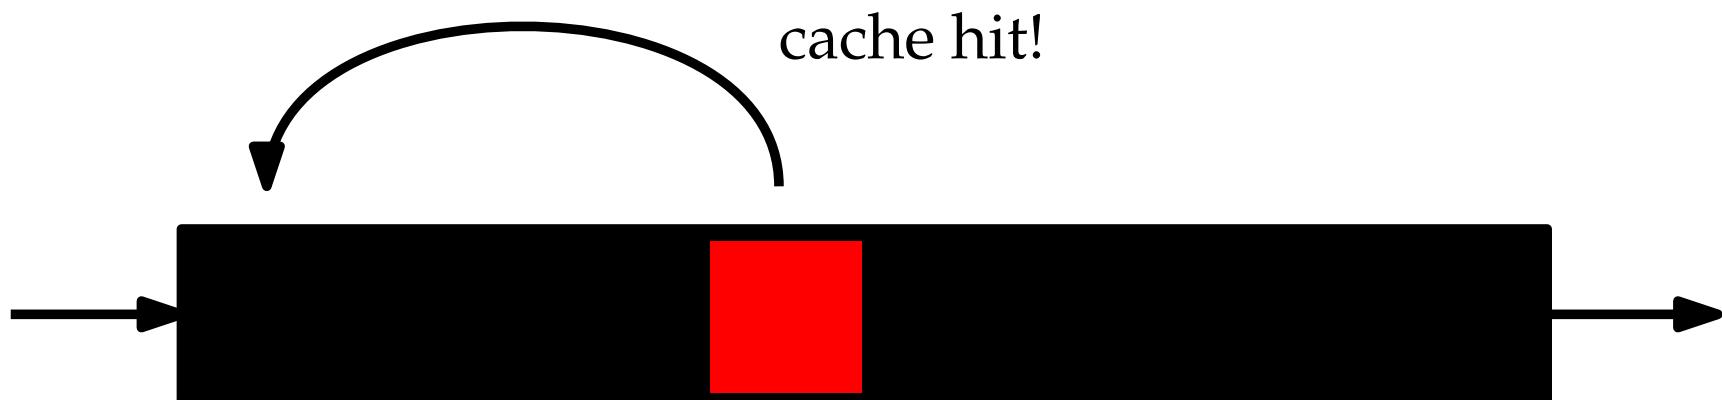
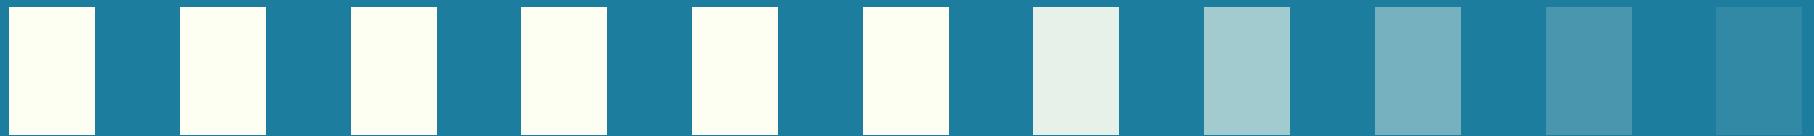
BASIC ALGORITHMS

clairvoyant

LRU

LFU

FIFO



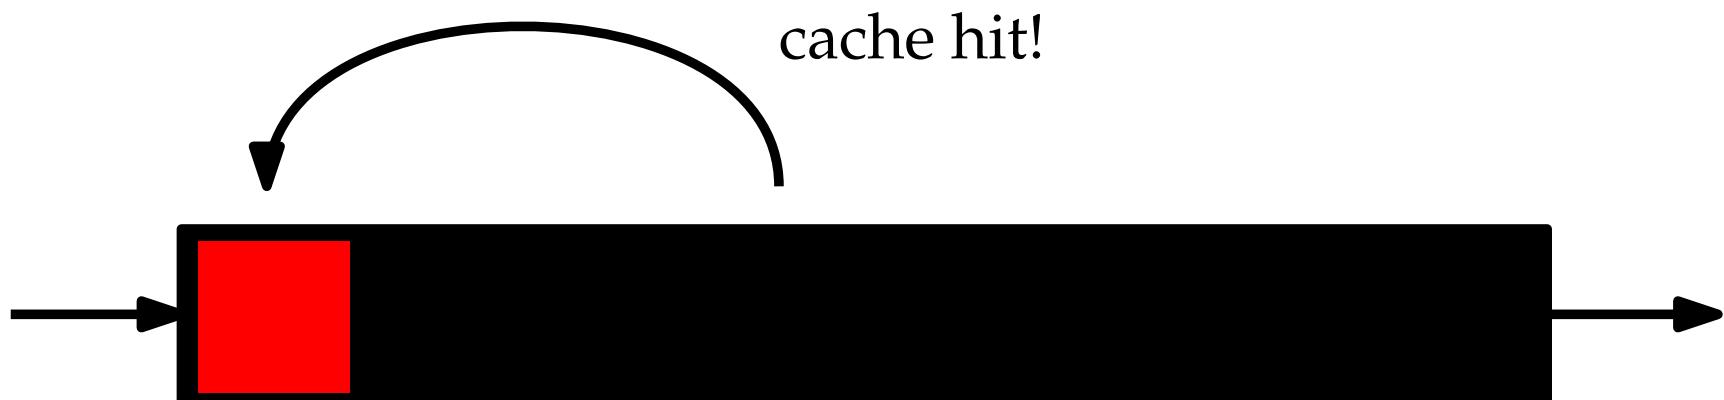
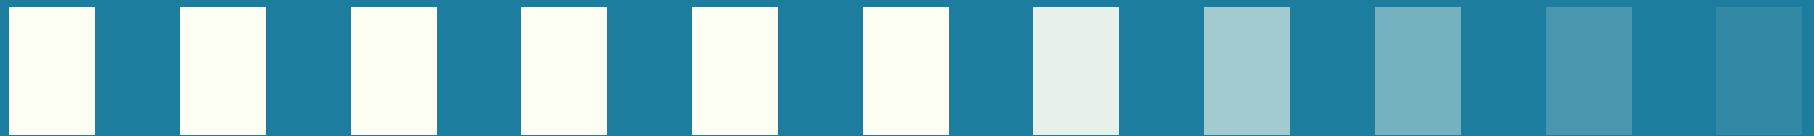
BASIC ALGORITHMS

clairvoyant

LRU

LFU

FIFO



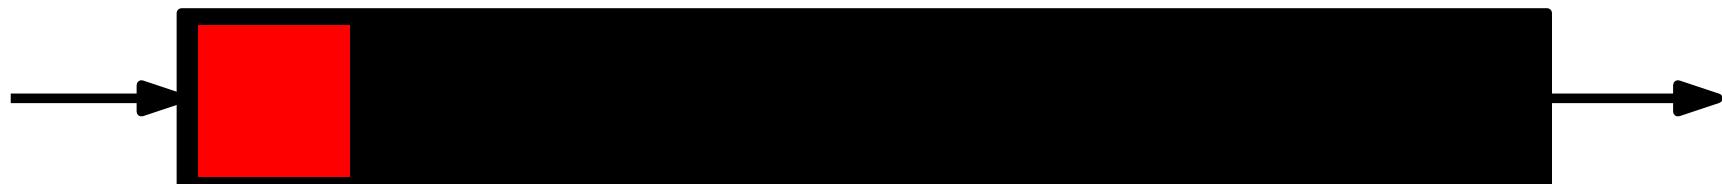
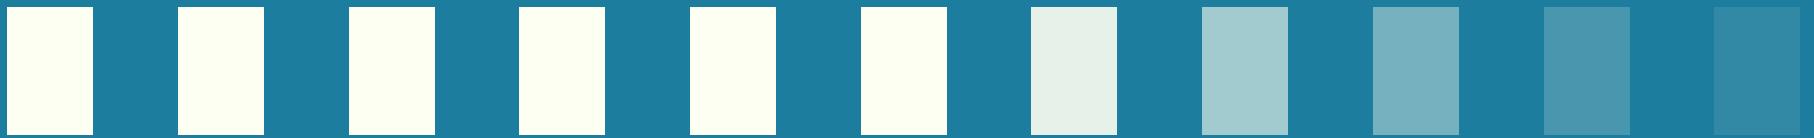
BASIC ALGORITHMS

clairvoyant

LRU

LFU

FIFO



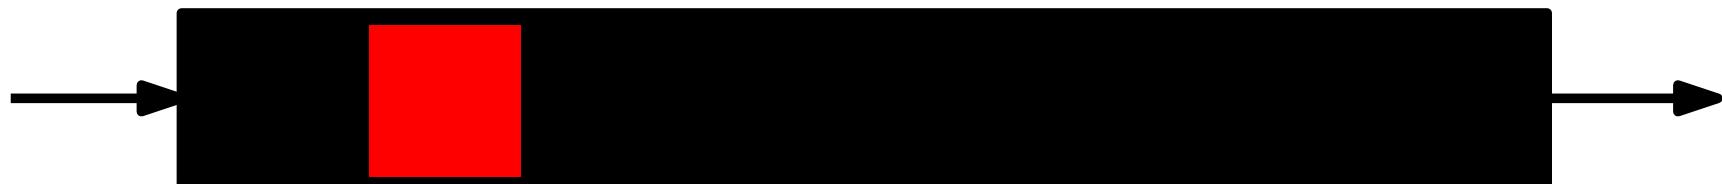
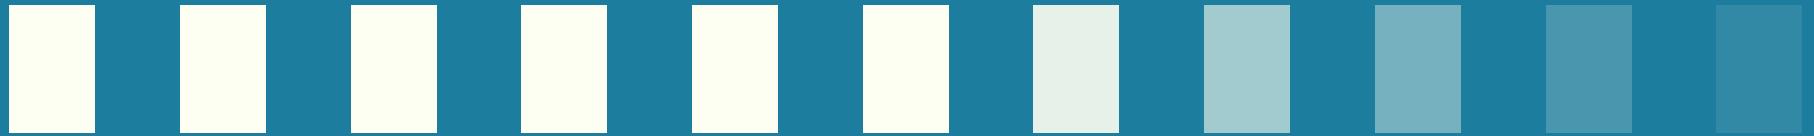
BASIC ALGORITHMS

clairvoyant

LRU

LFU

FIFO



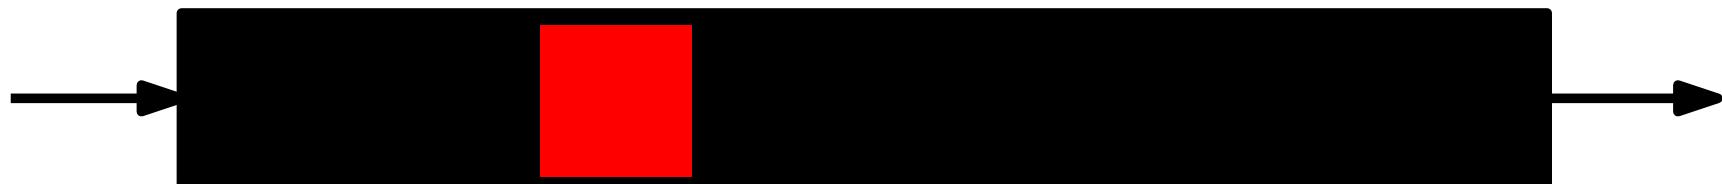
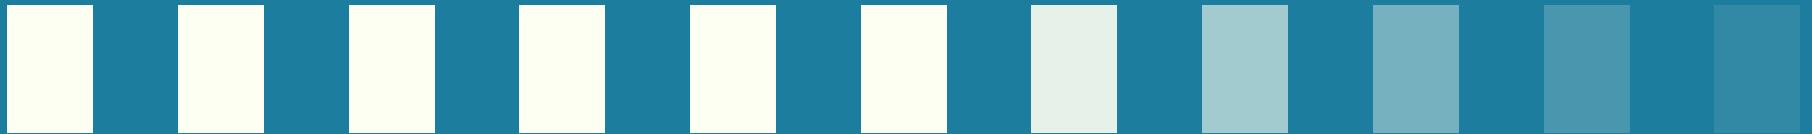
BASIC ALGORITHMS

clairvoyant

LRU

LFU

FIFO



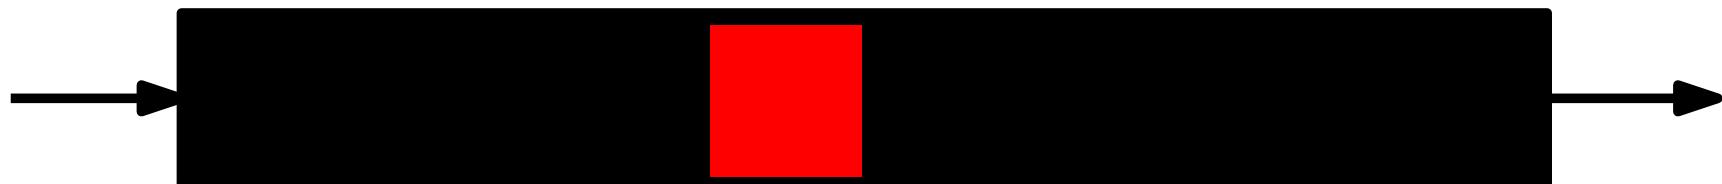
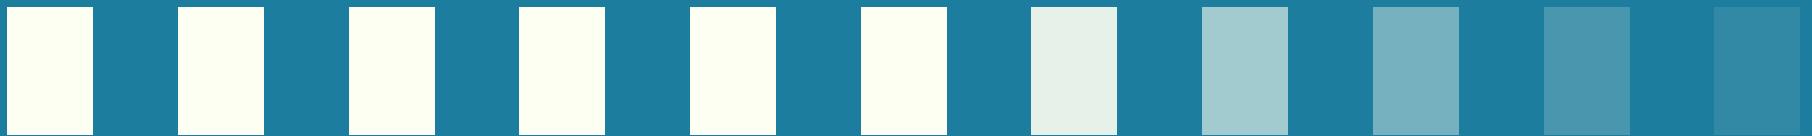
BASIC ALGORITHMS

clairvoyant

LRU

LFU

FIFO



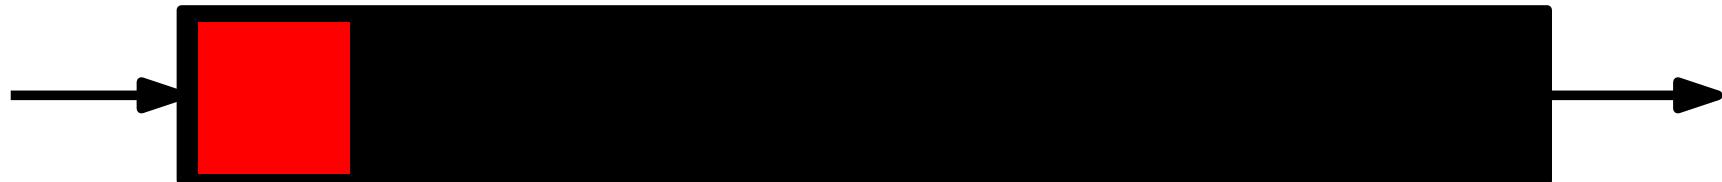
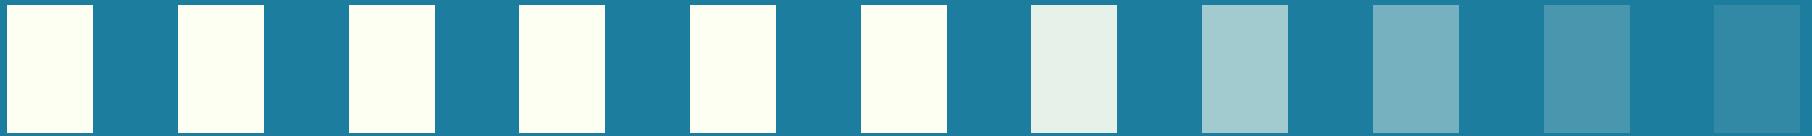
BASIC ALGORITHMS

clairvoyant

LRU

LFU

FIFO



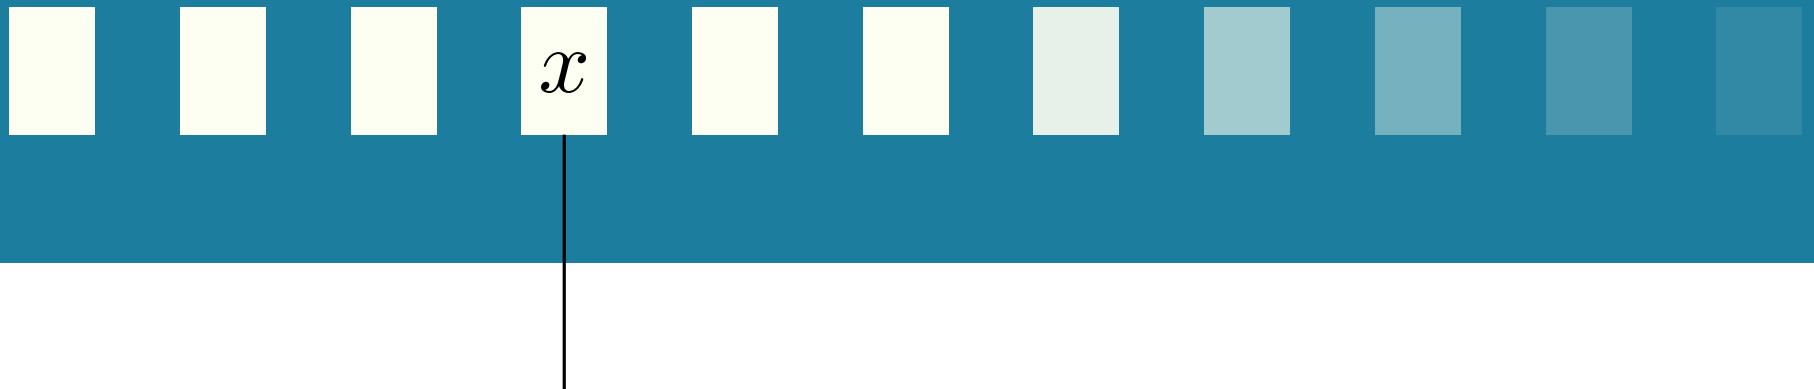
BASIC ALGORITHMS

clairvoyant

LRU

LFU

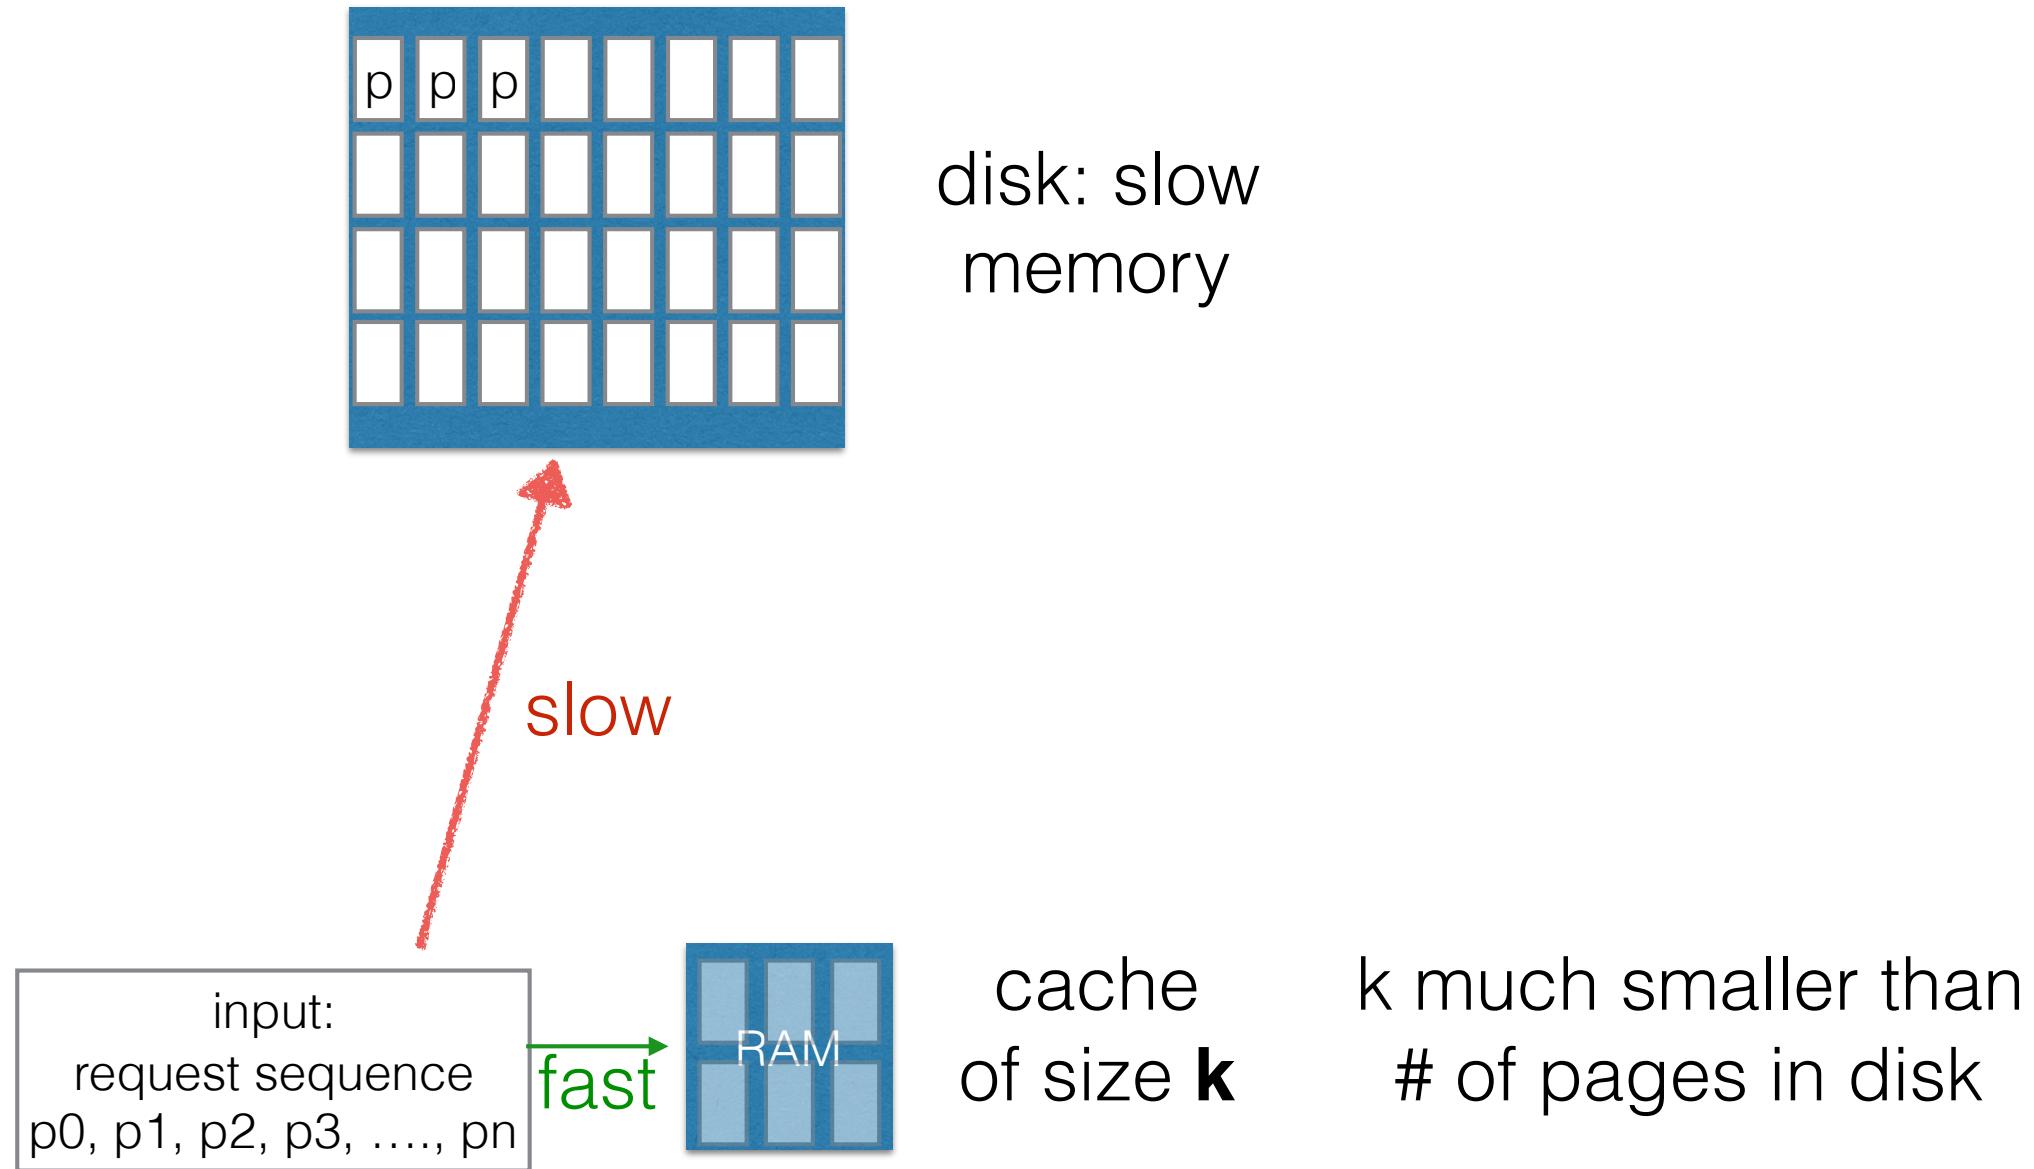
FIFO



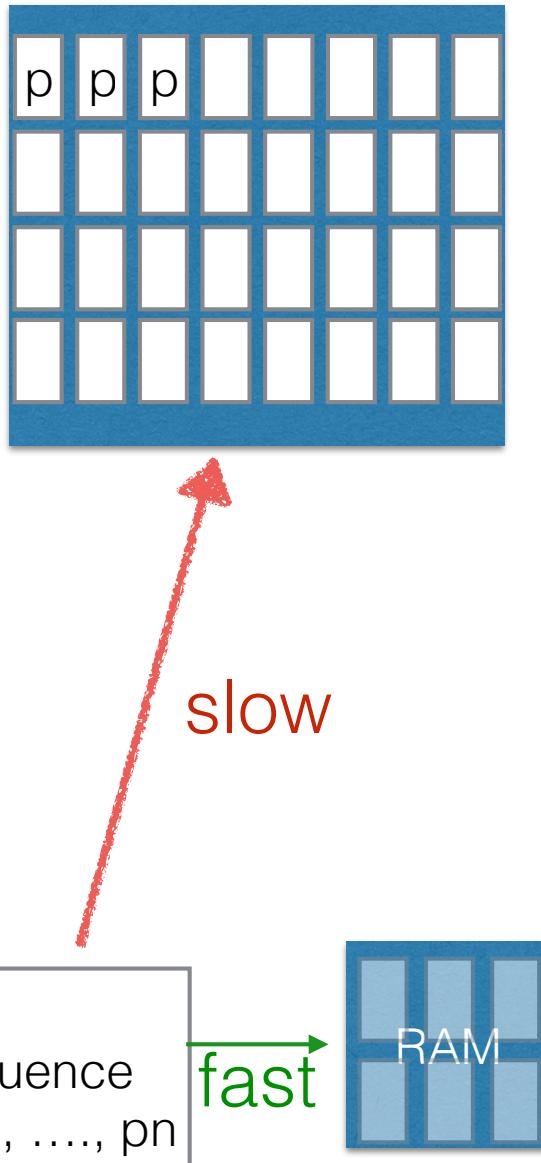
$\text{frequency}(x) = \text{number of requests}$

evict object with smallest frequency

Model of virtual memory



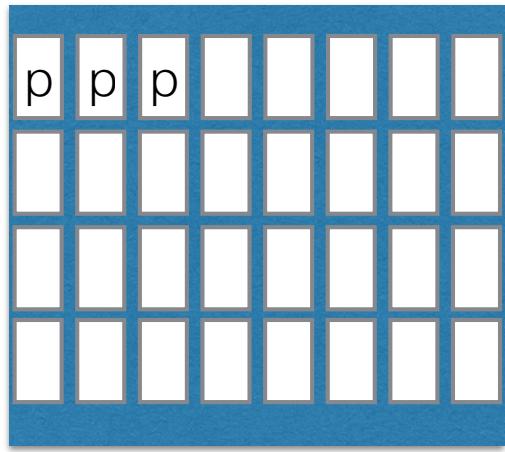
Model of virtual memory



goal

minimize total cost of
bringing pages from disk

Paging problem



input

$s = p_0, p_1, p_2, p_3, \dots, p_n$

output

for each p_i , which p to evict from cache

goal

minimize # misses



LFD is optimal

- proposed by Belady
- LFD: longest forward distance
- synonyms: OPT, MIN, clairvoyant, Belady's optimal alg
- proof given in class

LFD proof of optimality

- let ALG be any paging algorithm, s input seq. Will construct a seq of algorithms: ALG = ALG0, ALG1, ALG2, ... ALGn = LFD st
- $\text{cost(ALG0, s)} \geq \text{cost(ALG1, s)} \geq \text{cost(ALG2, s)} \geq \dots \geq \text{cost(ALGn)} = \text{LFD}$, $n = \text{len}(s)$
- how? by induction: suppose ALGi is determined
- construct ALG(i+1) as follows: for the first $i-1$ requests, do the same eviction decisions as ALGi. On request i, ALGi evicts the page currently in its cache that will be requested furthest into the future, call this page v. ALG(i) may or may not evict this same page. If it does, let ALG(i+1) = ALG(i) for the future too, so then done.
- Otherwise ALG(i) evicts page $u \neq v$. So after processing i, both caches are the same, but ALG(i) has v and ALG(i+1) has u, where u is requested before v.
- On future requests, ALG(i+1) mirrors the same decisions as ALG(i) except if ALG(i) evicts v, then ALG(i+1) evicts u. Since u will be requested first, ALG(i) will incur a miss when u is requested, at which point both may or may not be identical (off by 1). If v is requested in the future, then the costs will be identical.

Belady's anomaly

- increasing the cache size may result in more misses
- hw: FIFO suffers from this phenomenon
- hw: LRU doesn't

Competitive analysis

Which replacement policy is better?

$$\text{competitive ratio of ALG} = \max_{\text{sequences } s} \frac{\text{misses(ALG, } s\text{)}}{\text{misses(OPT, } s\text{)}}$$

ALG is said to be **k-competitive** if its competitive ratio is $\leq k$

LRU is k-competitive

for any input sequence s :

$$\text{misses}(\text{LRU}, s) \leq k \text{ misses}(\text{OPT}, s)$$

Proof sketch: divide s into k -phases, where each phase consists of references to exactly k distinct pages
then $\text{misses}(\text{LRU}, s) \leq k$
and $\text{misses}(\text{OPT}, s) \geq 1$

FIFO is k-competitive

for any input sequence s :

$$\text{misses}(\text{FIFO}, s) \leq k \text{ misses}(\text{OPT}, s)$$

Proof: homework

LFU is not competitive

there exists a sequence of input sequences s_1, s_2, s_3, \dots

$$\lim_{i \rightarrow \infty} \frac{\text{misses(LFU, } s_i)}{\text{misses(OPT, } s_i)} = \infty$$

Proof: let $s_i = p_0 \times i+1, p_1 \times i+1, \dots p_{(k-1)} \times i+1, [p_k, p_{(k+1)}] \times i$

then $\text{misses(LFU, } s_i) = k-1 + 2i$

and $\text{misses(OPT, } s_i) = k$

LIFO is not competitive

there exists a sequence of input sequences s_1, s_2, s_3, \dots

$$\lim_{i \rightarrow \infty} \frac{\text{misses(LIFO, } s_i\text{)}}{\text{misses(OPT, } s_i\text{)}} = \infty$$

Proof: homework

A timeline of paging analysis

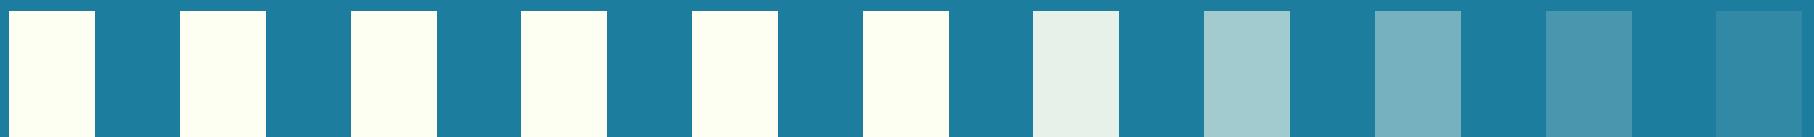
- Belady's paper shows that **MIN is optimal**
 - **competitive analysis**, Sleator and Tarjan
 - 1995: 1st analysis of paging with **locality of reference**, Borodin et al. using access graphs
 - 1996: strongly competitiveness, Irani et al.
 - 1999: LRU is better than FIFO
 - 2000: markov paging, Karlin et al.
 - 2007: relative worst order, Boyar et al
 - 2009: relative dominance, Dorrigiv et al
 - 2010: parameterized analysis, Dorrigiv et al
 - 2012: access graph model under worst order analysis, Boyar et al
 - 2013: access graph model under relative interval analysis, Boyar et al.
 - 2015: competitiveness in terms of locality parameters, Albers et al.
- over the years:
 - alternative performance measures proposed
 - various frameworks for locality of reference proposed

LRU

LFU

FIFO

STRENGTHS AND WEAKNESSES

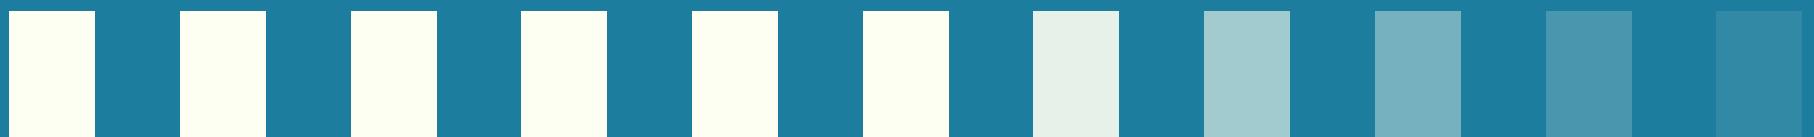


good if popularity changes quickly
unpopular objects stay too long

LRU

FIFO
LFU

STRENGTHS AND WEAKNESSES



good if popularity changes quickly

unpopular objects stay too long

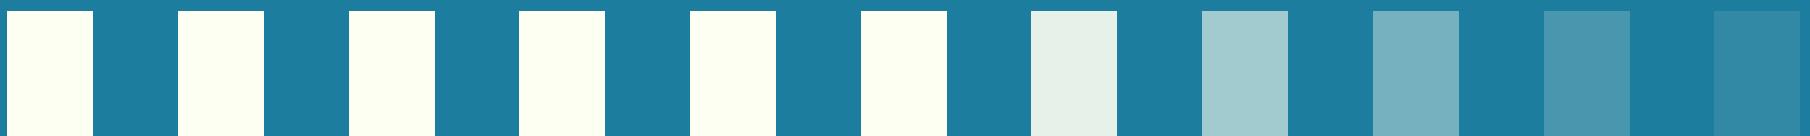
LRU

FIFO

LFU

good if popularity is static
suffers from cache pollution over time

STRENGTHS AND WEAKNESSES



good if popularity changes quickly
unpopular objects stay too long

LRU

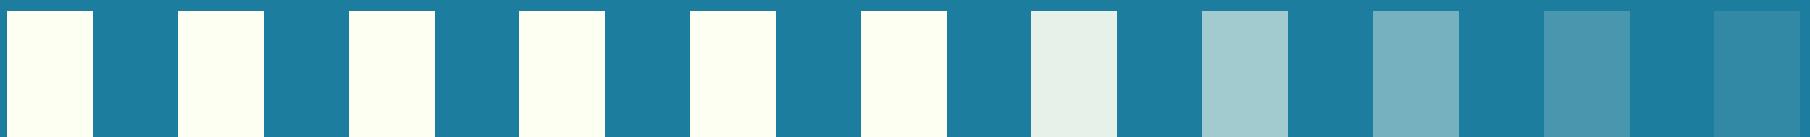
good when
fragmentation
is an issue

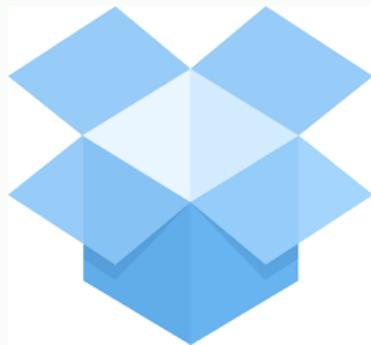
FIFO

LFU

good if popularity is static
suffers from cache pollution over time

STRENGTHS AND WEAKNESSES



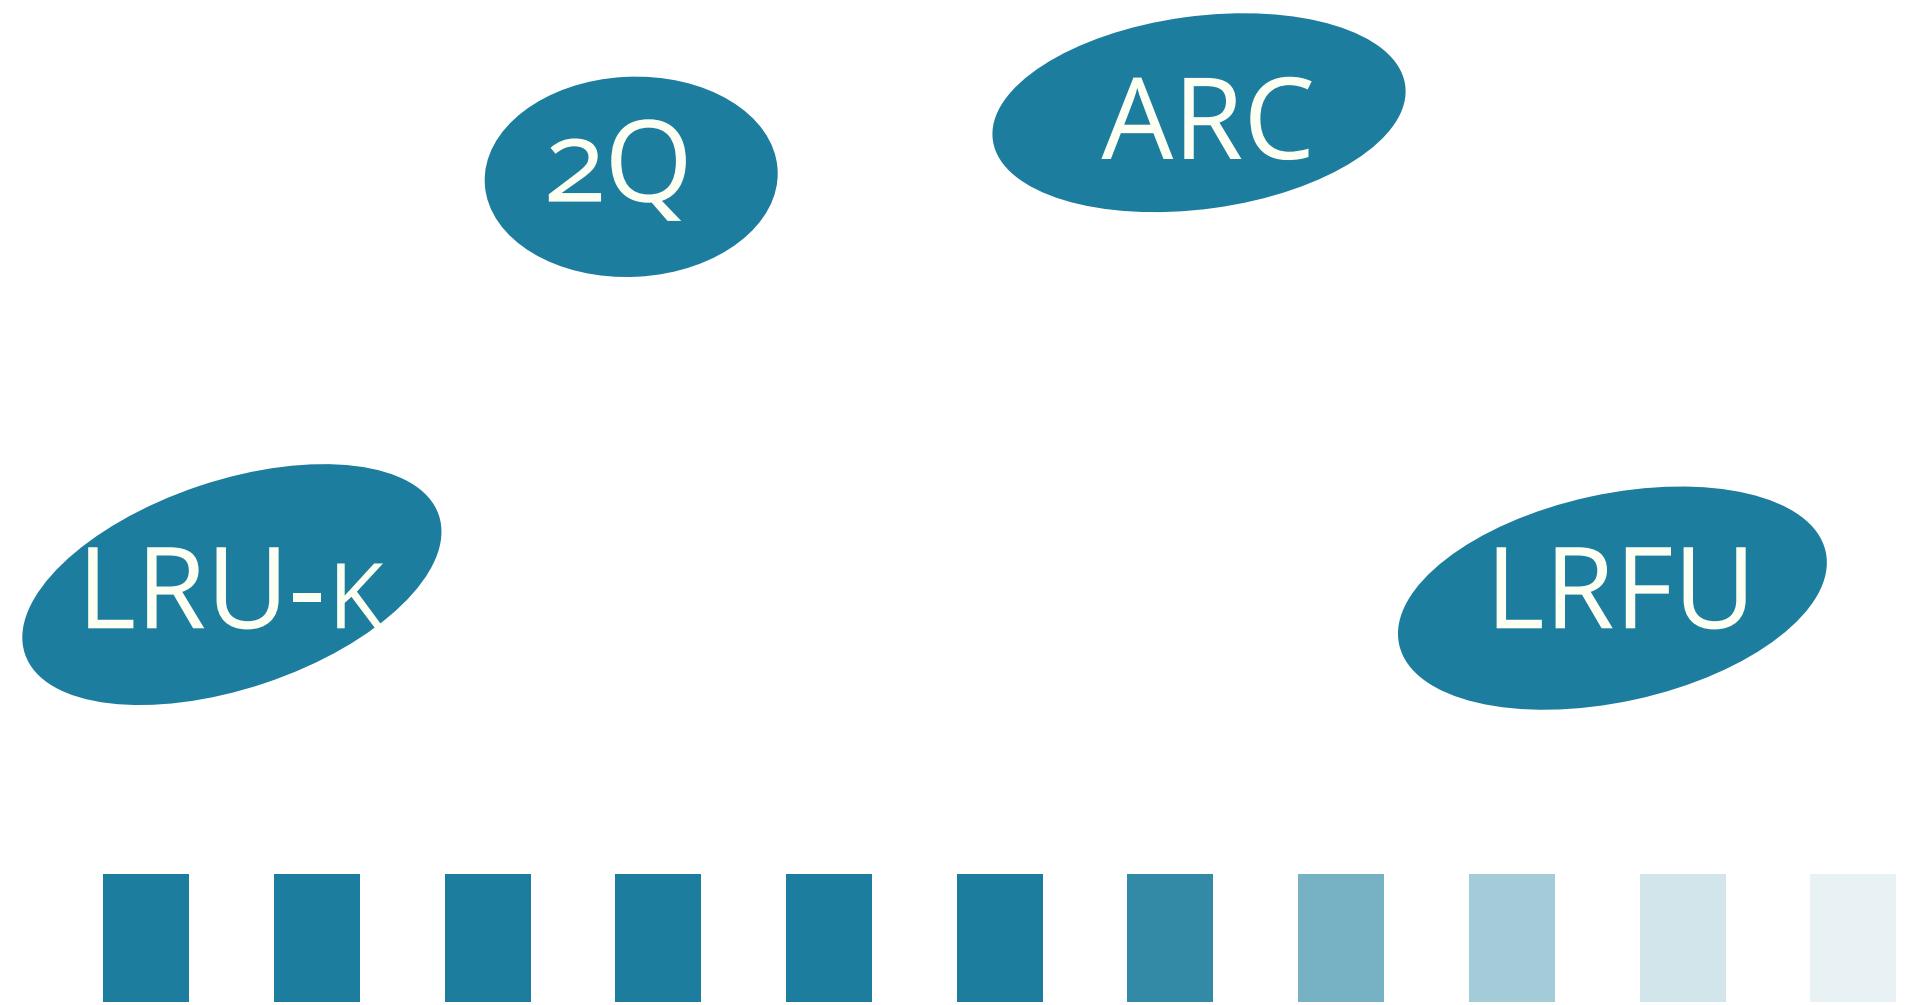


LRU

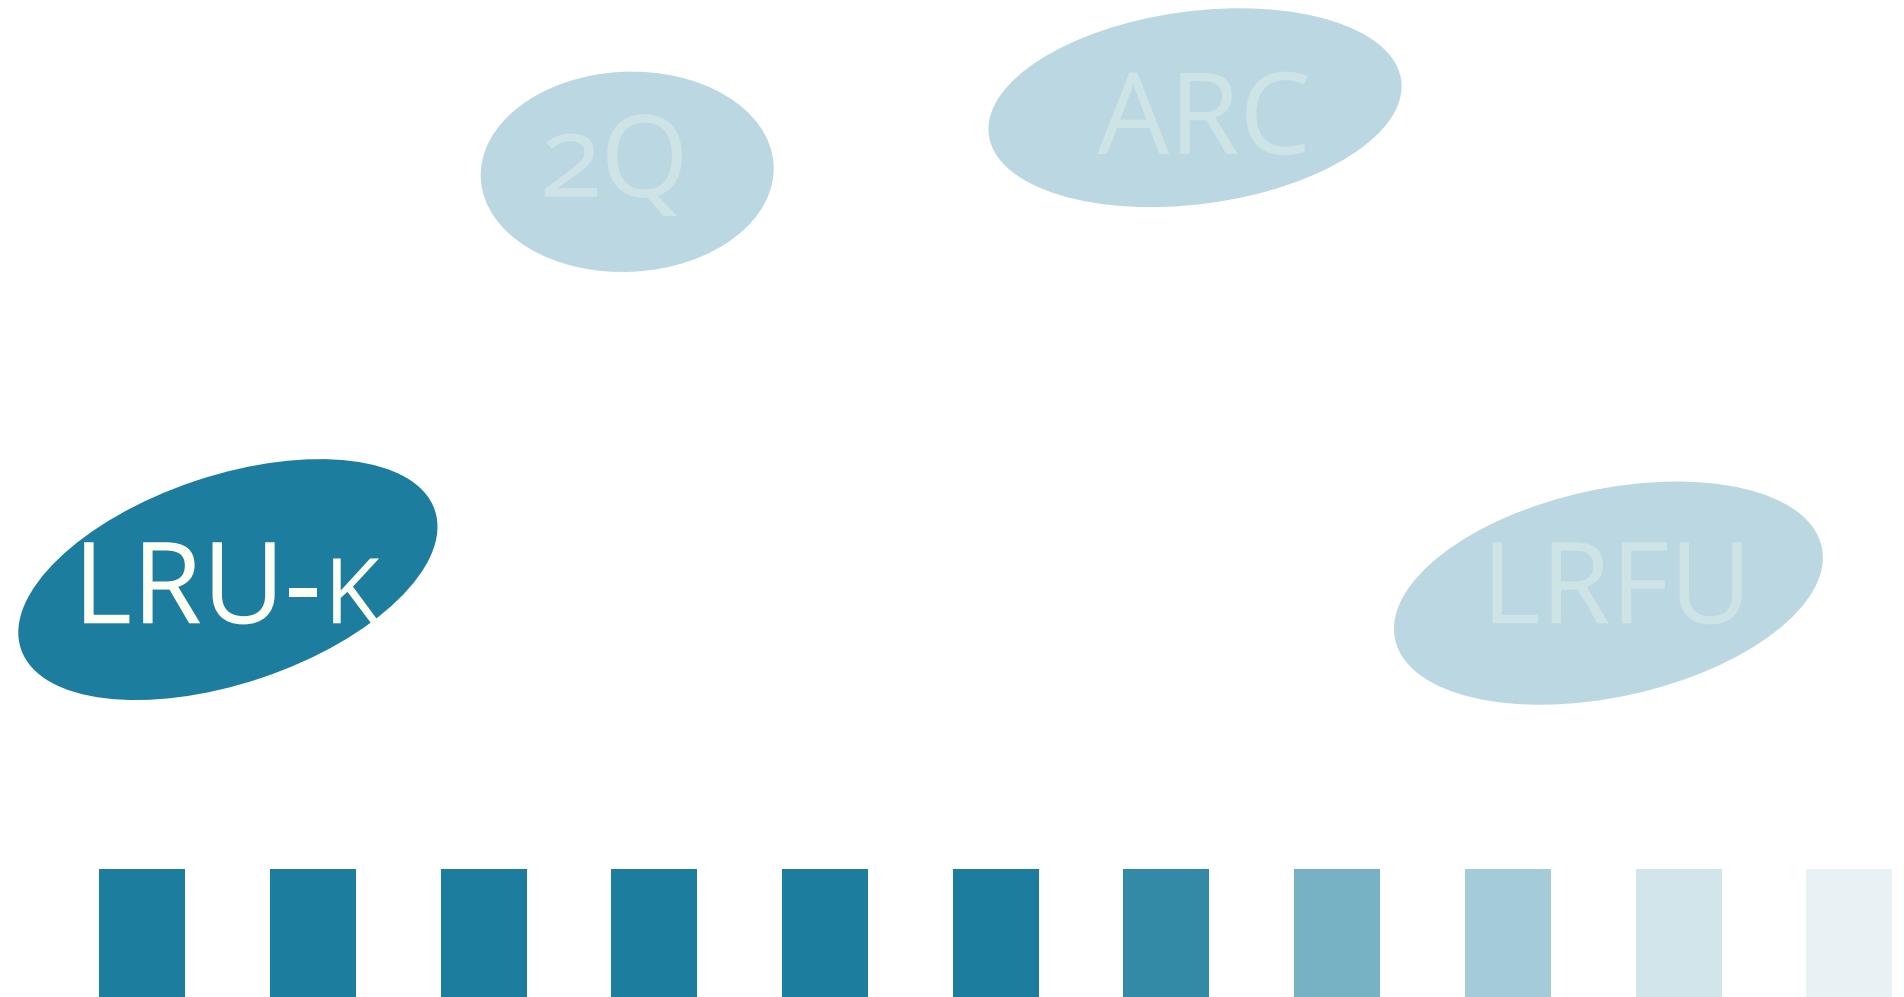
FIXING IT

admission control

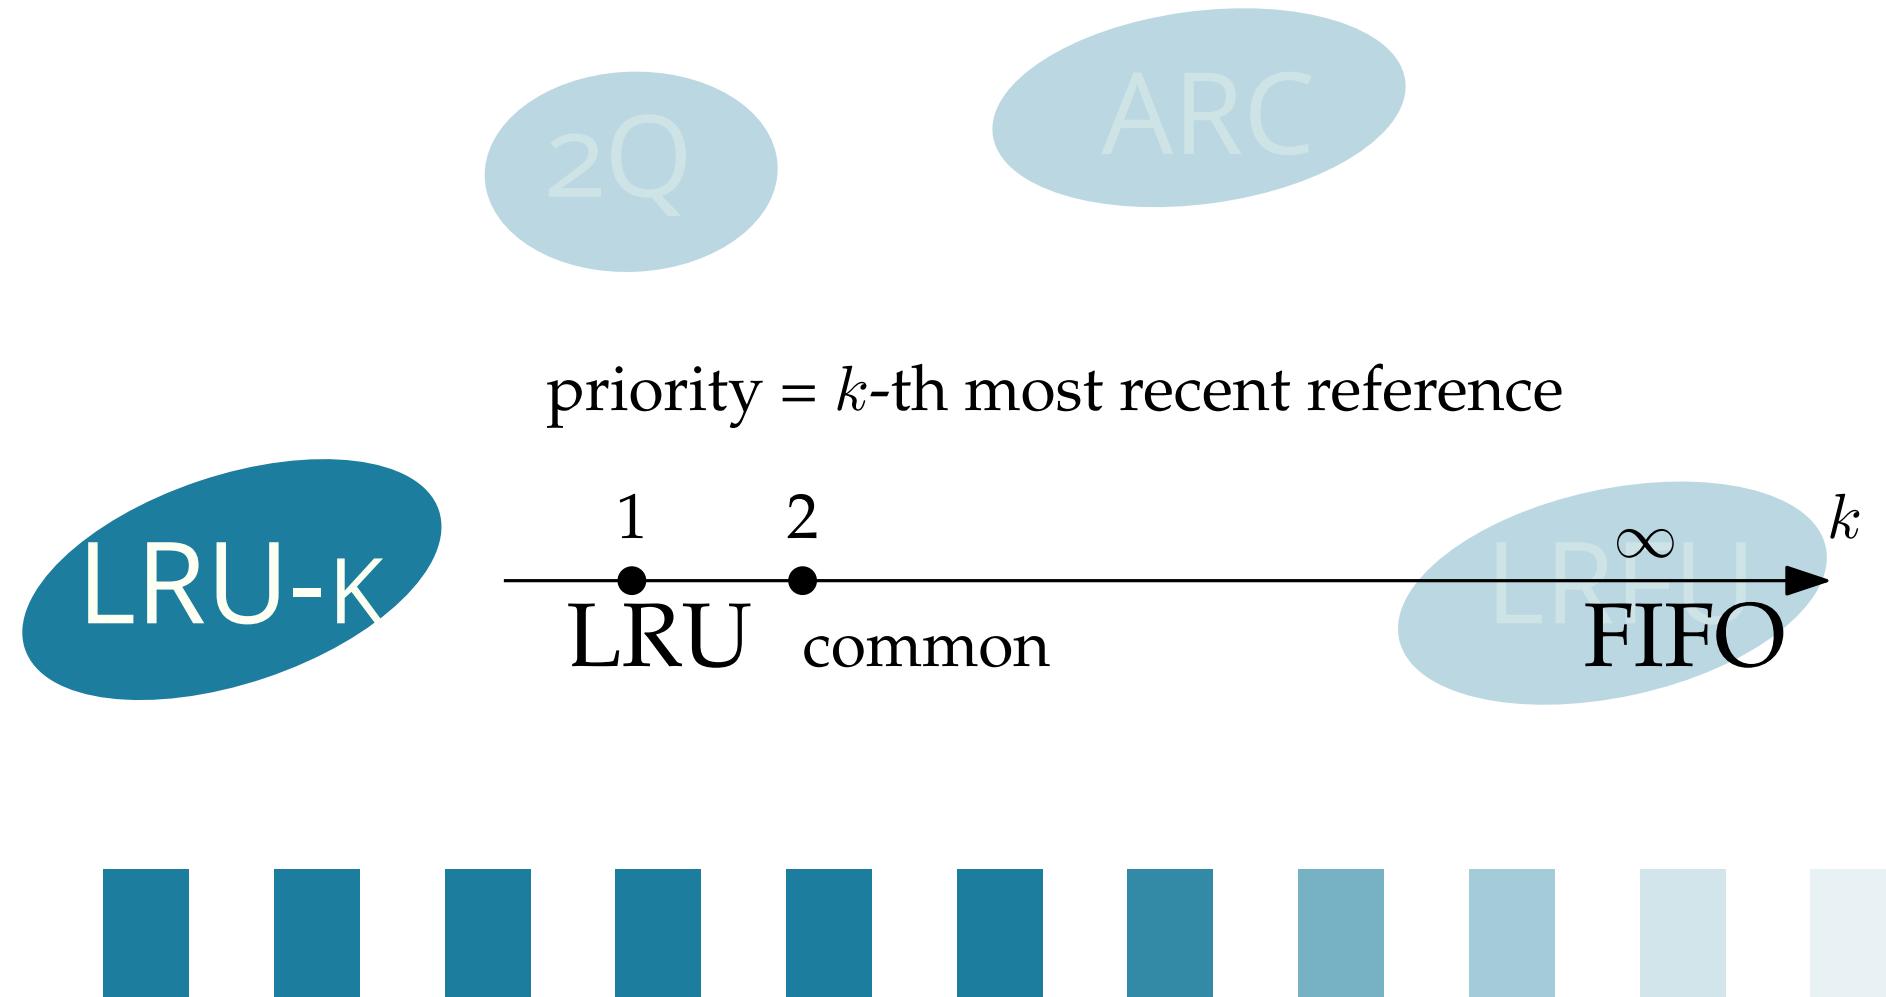
find a middle ground between LRU and LFU



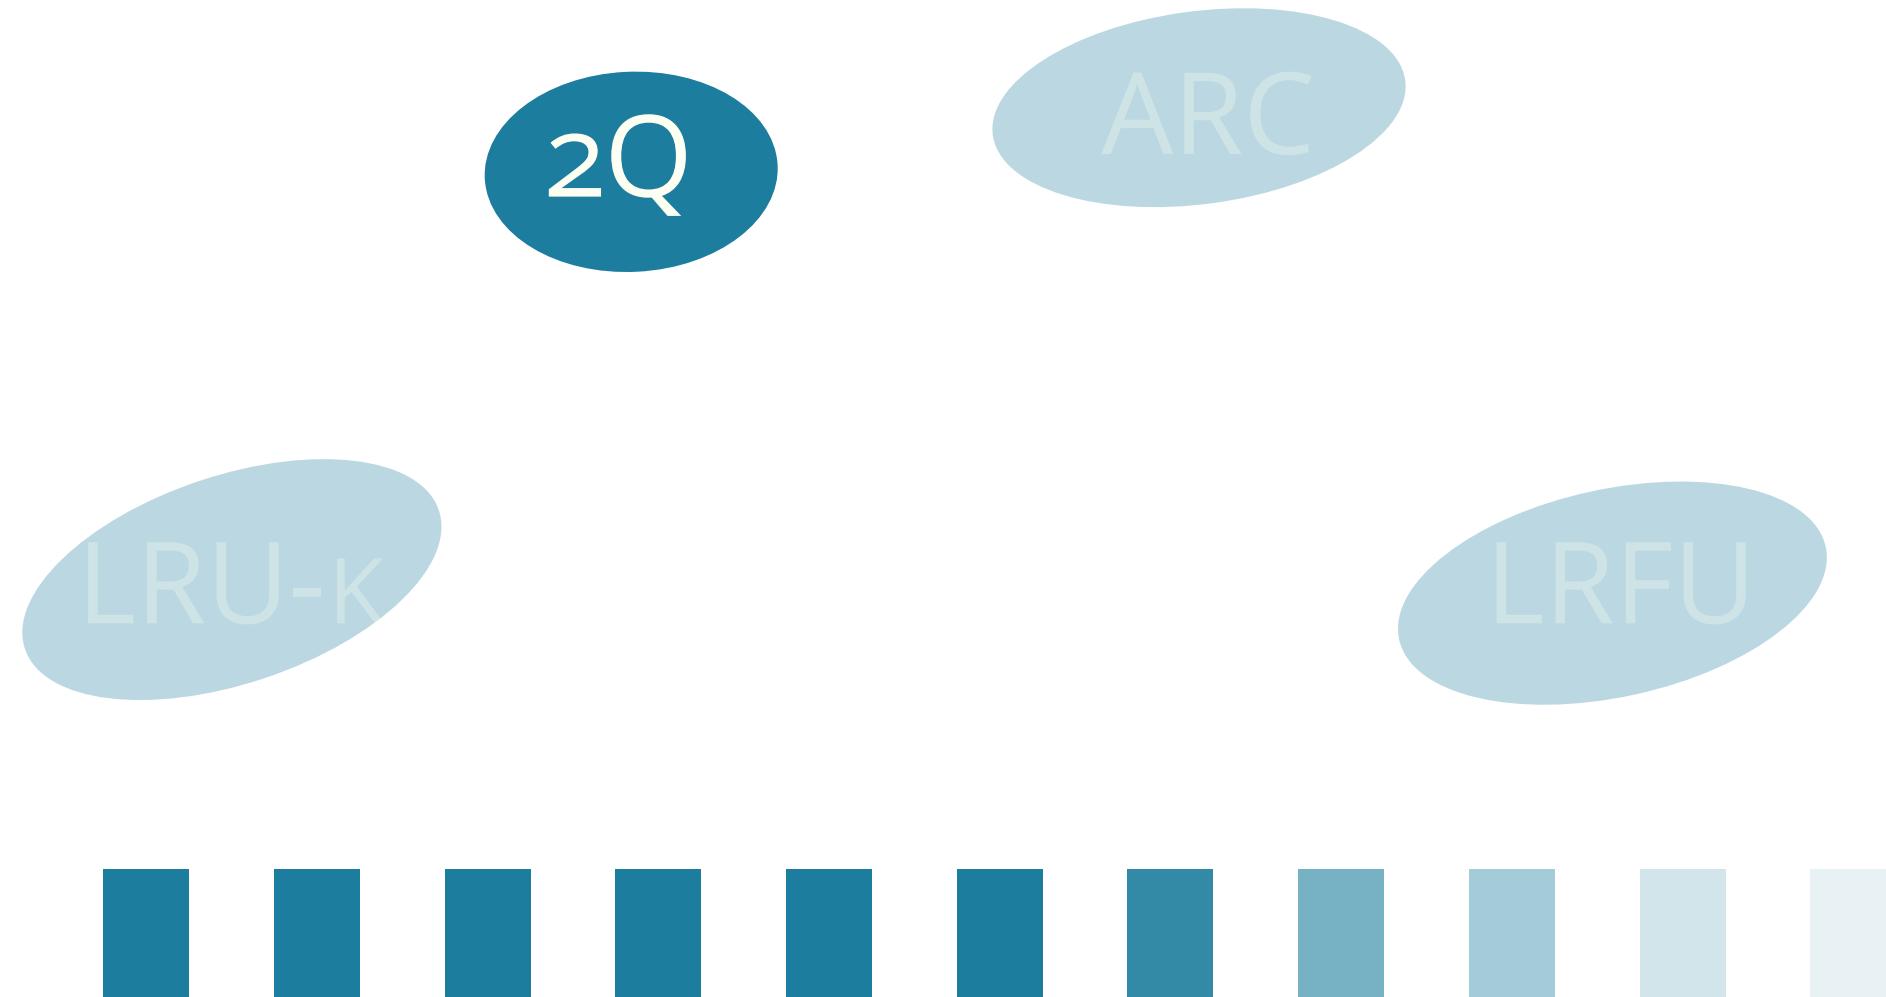
LRU-LFU HYBRID EVICTION ALGORITHMS



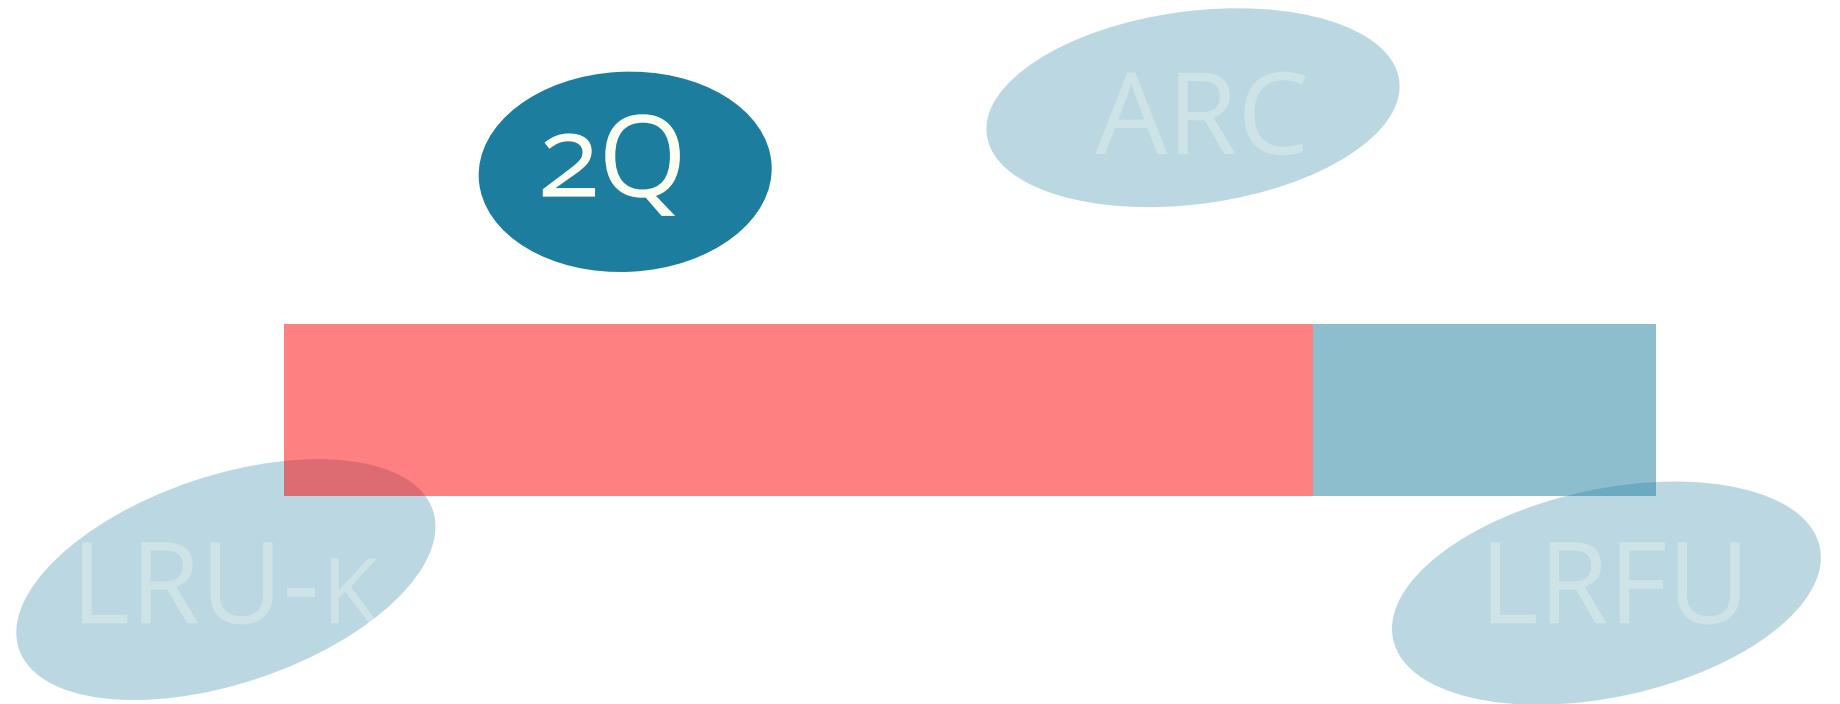
LRU-LFU HYBRID EVICTION ALGORITHMS



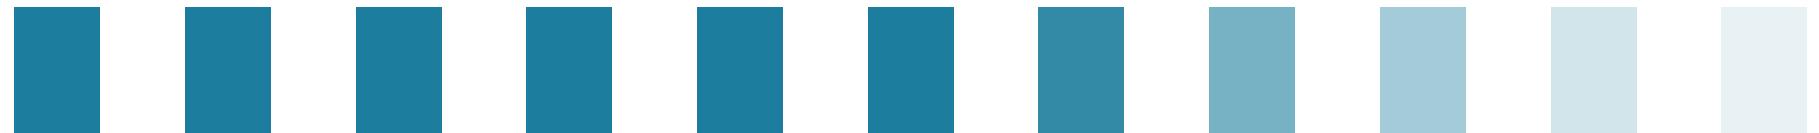
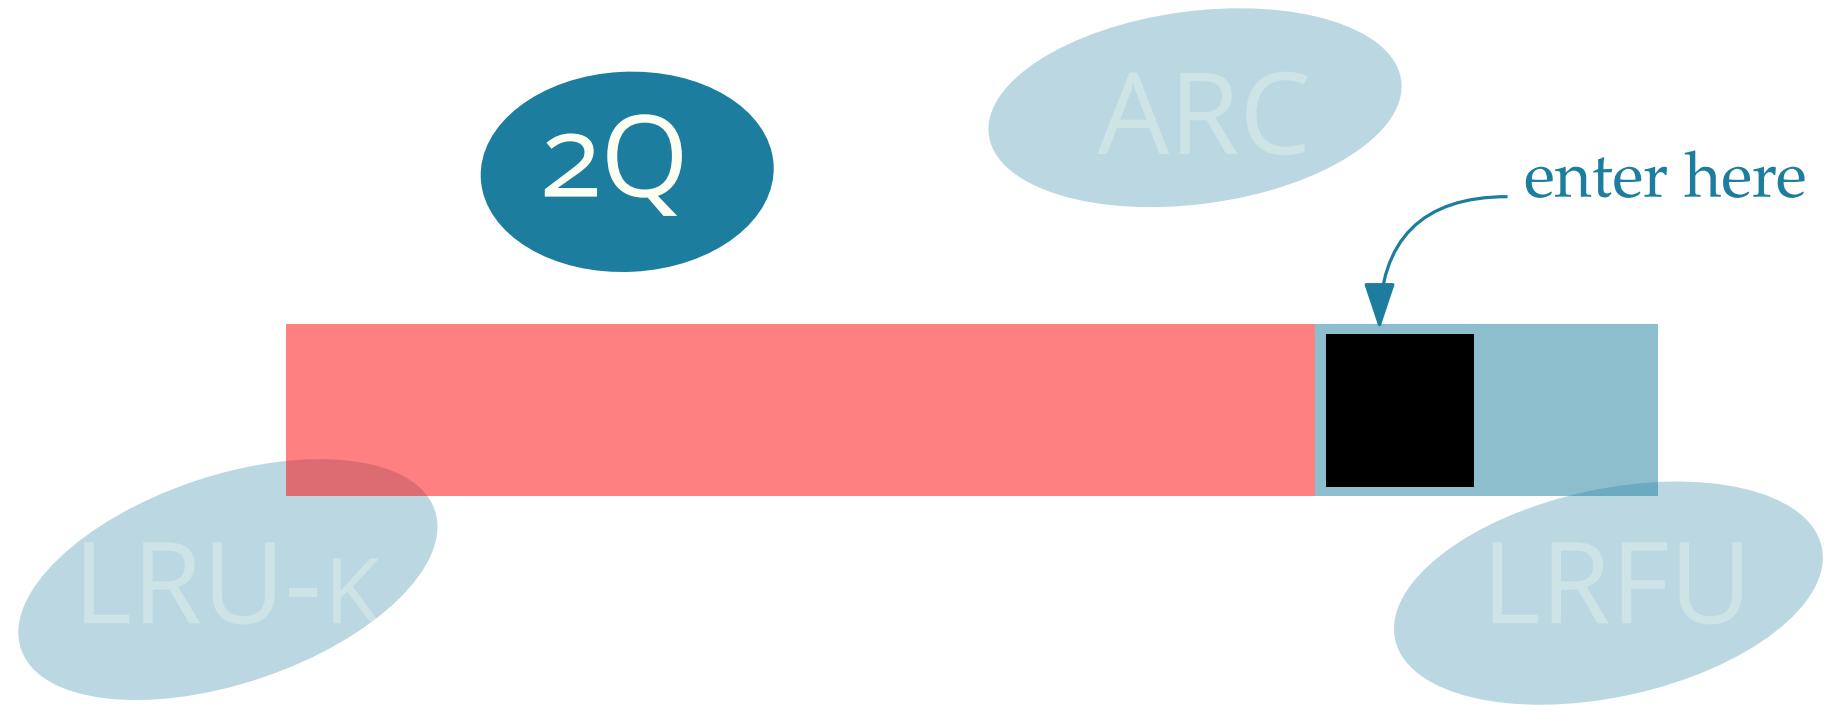
LRU-LFU HYBRID EVICTION ALGORITHMS



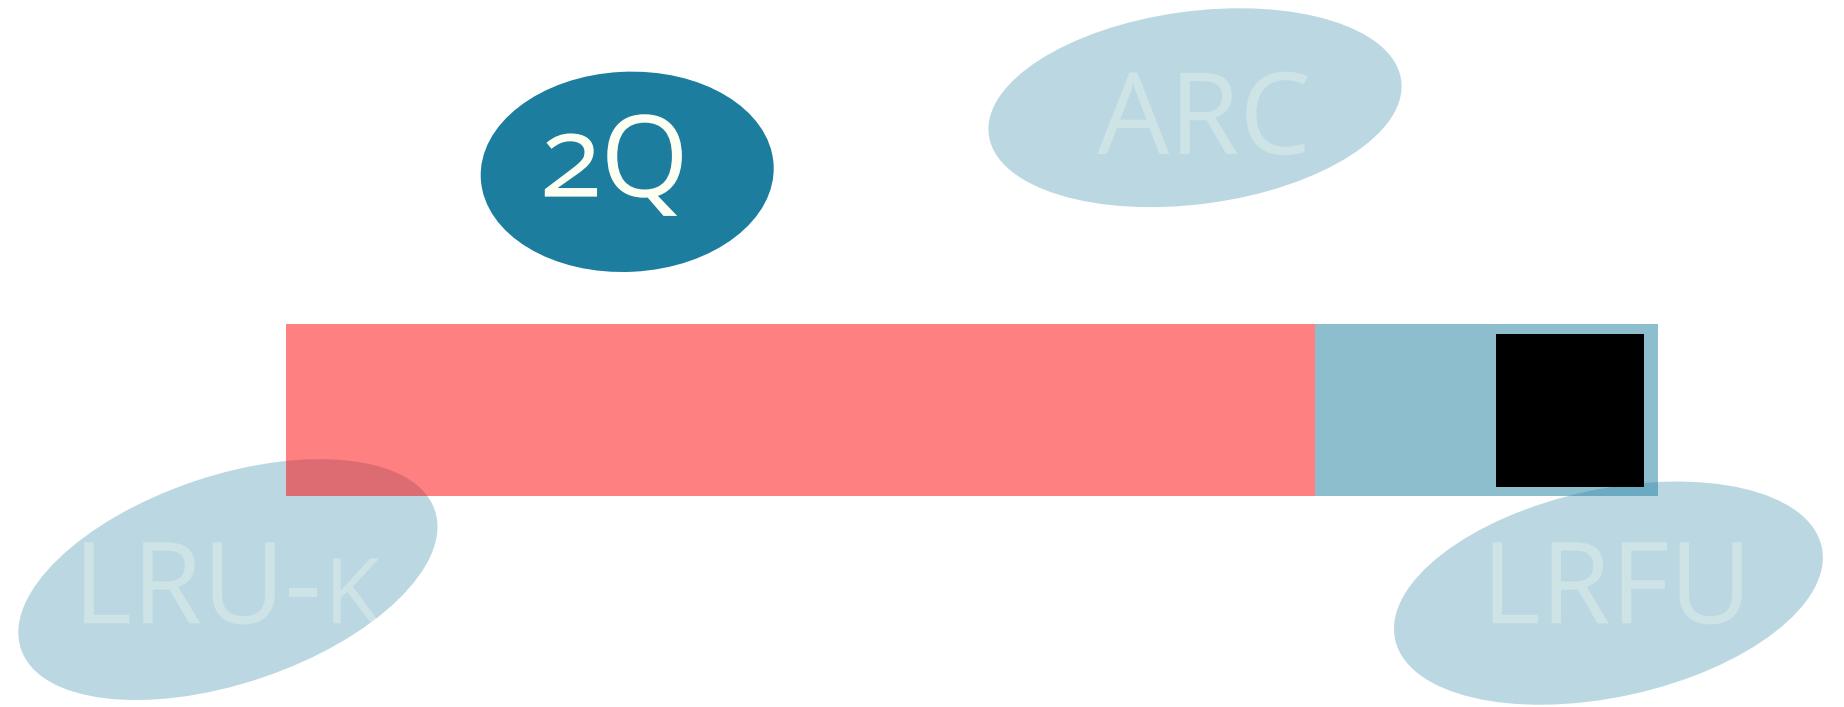
LRU-LFU HYBRID EVICTION ALGORITHMS



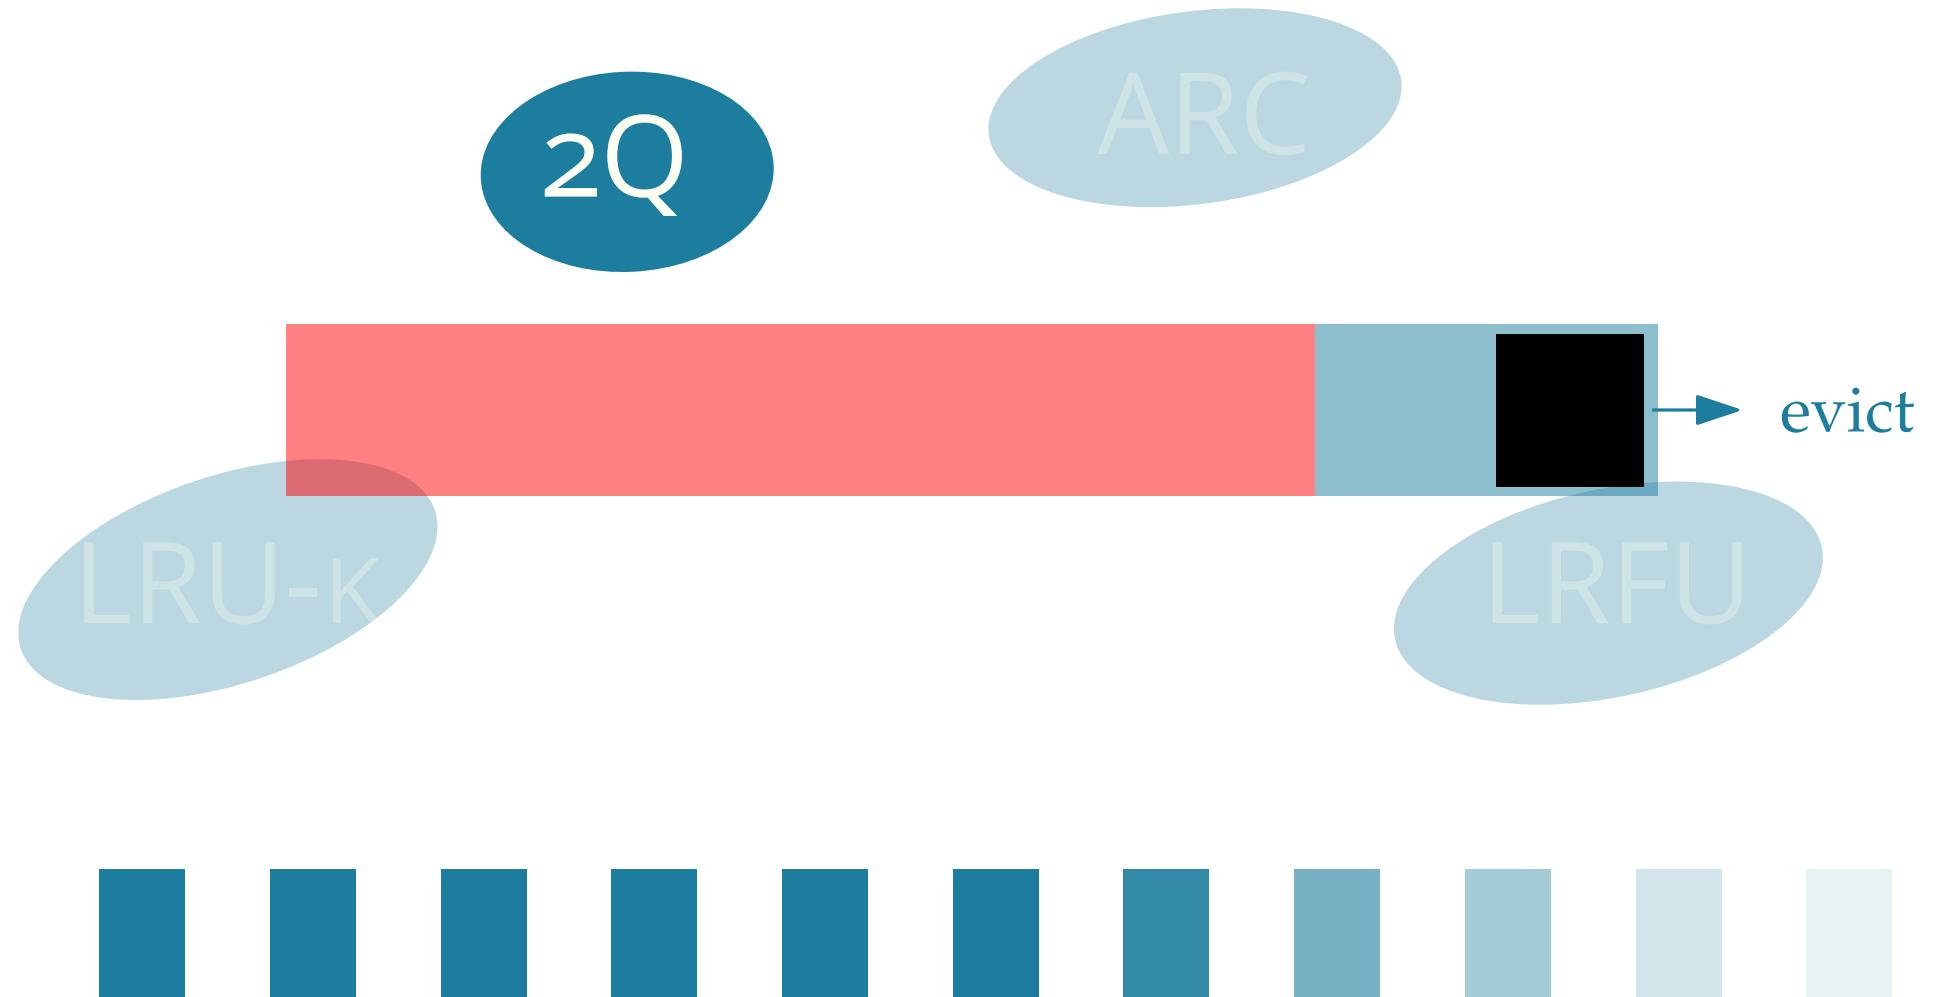
LRU-LFU HYBRID EVICTION ALGORITHMS



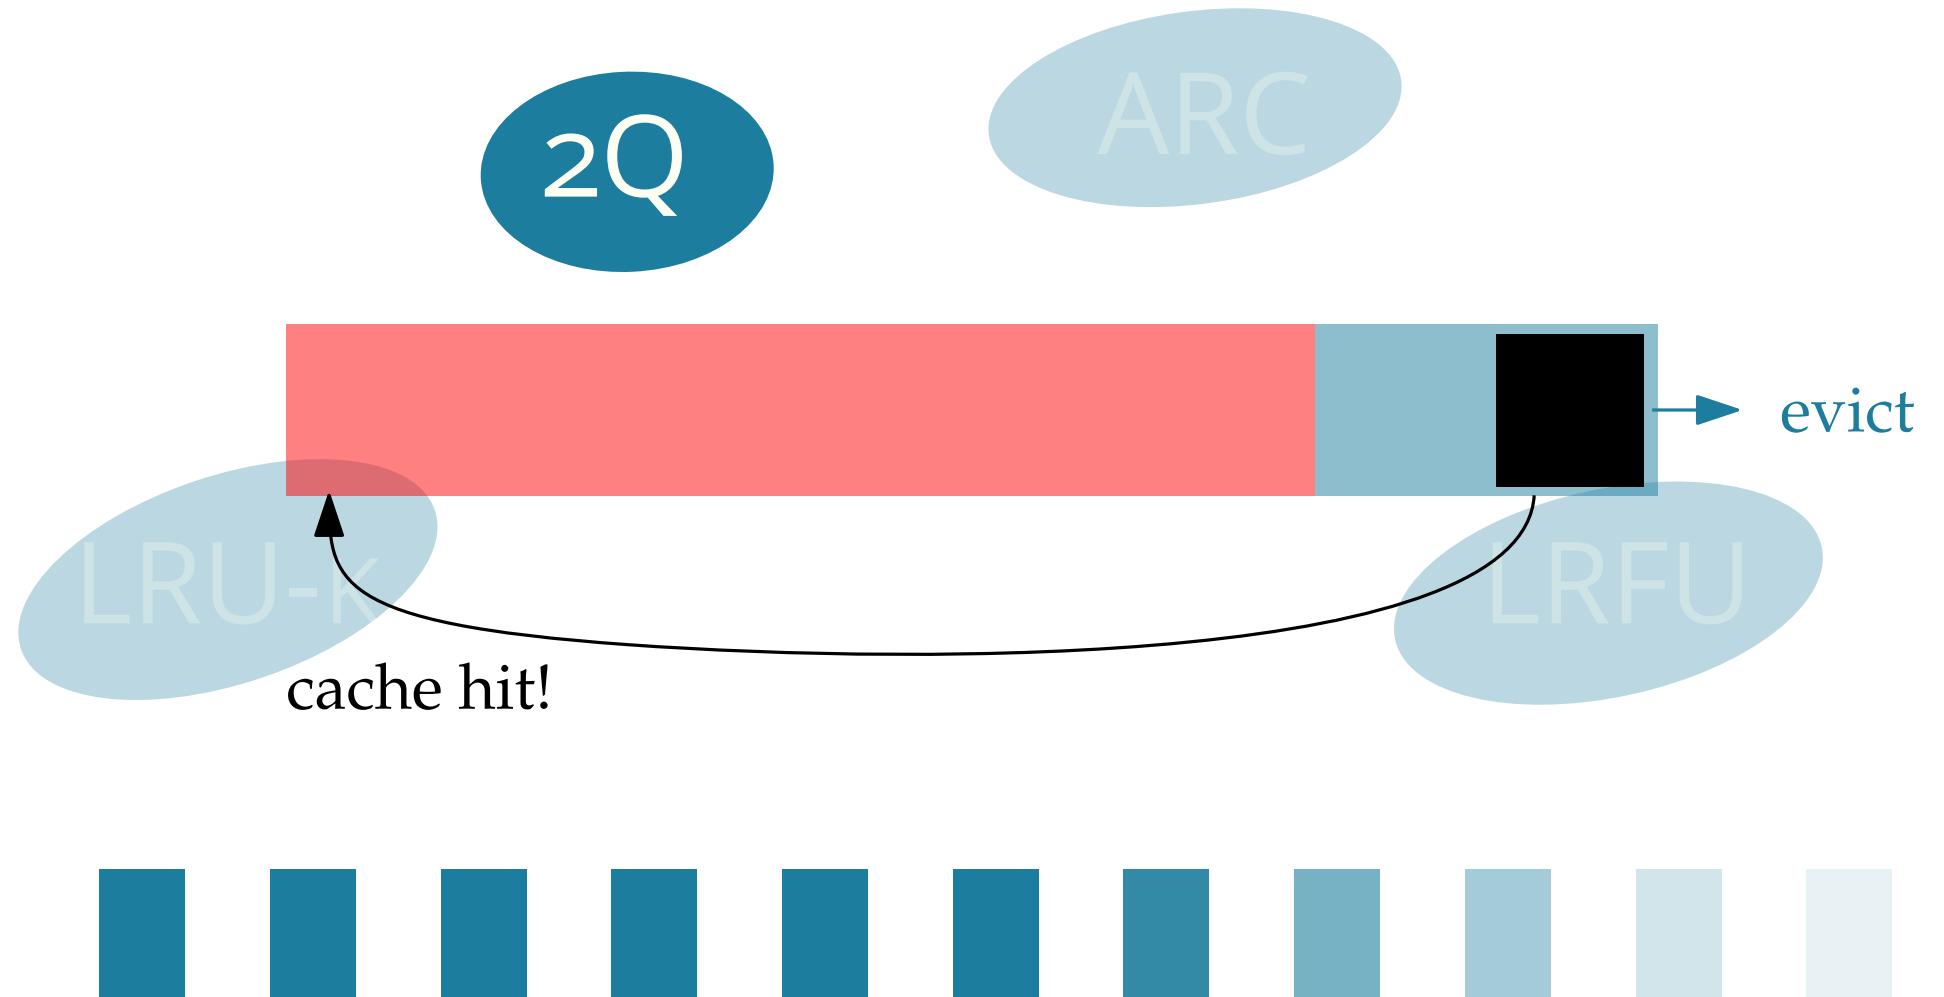
LRU-LFU HYBRID EVICTION ALGORITHMS



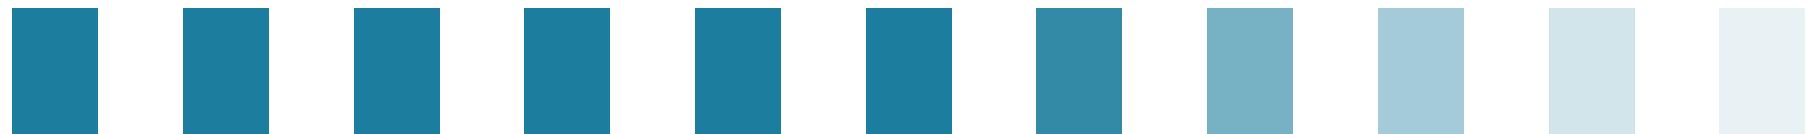
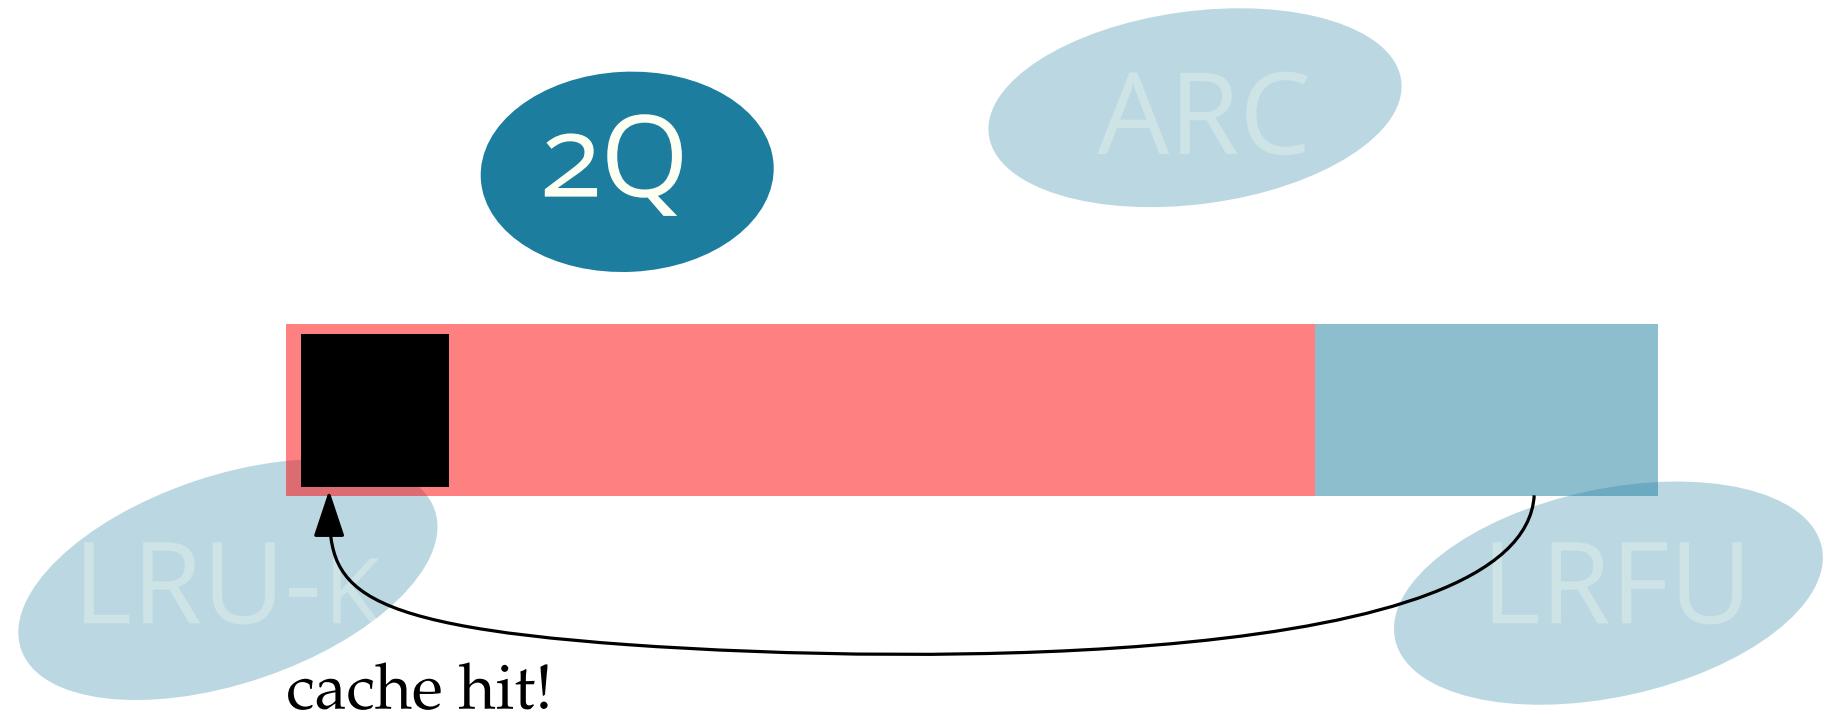
LRU-LFU HYBRID EVICTION ALGORITHMS



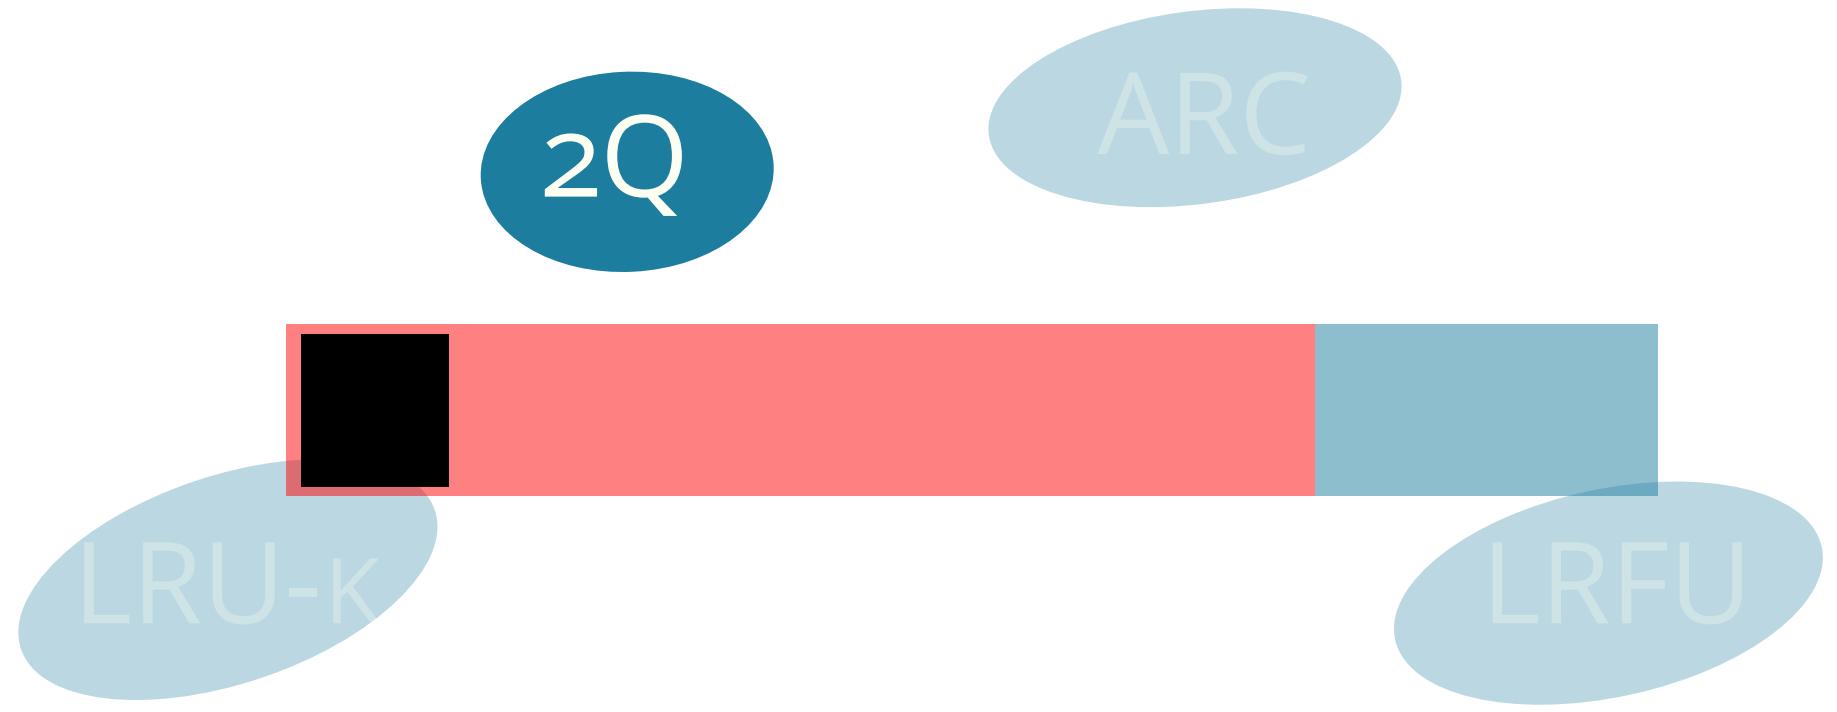
LRU-LFU HYBRID EVICTION ALGORITHMS



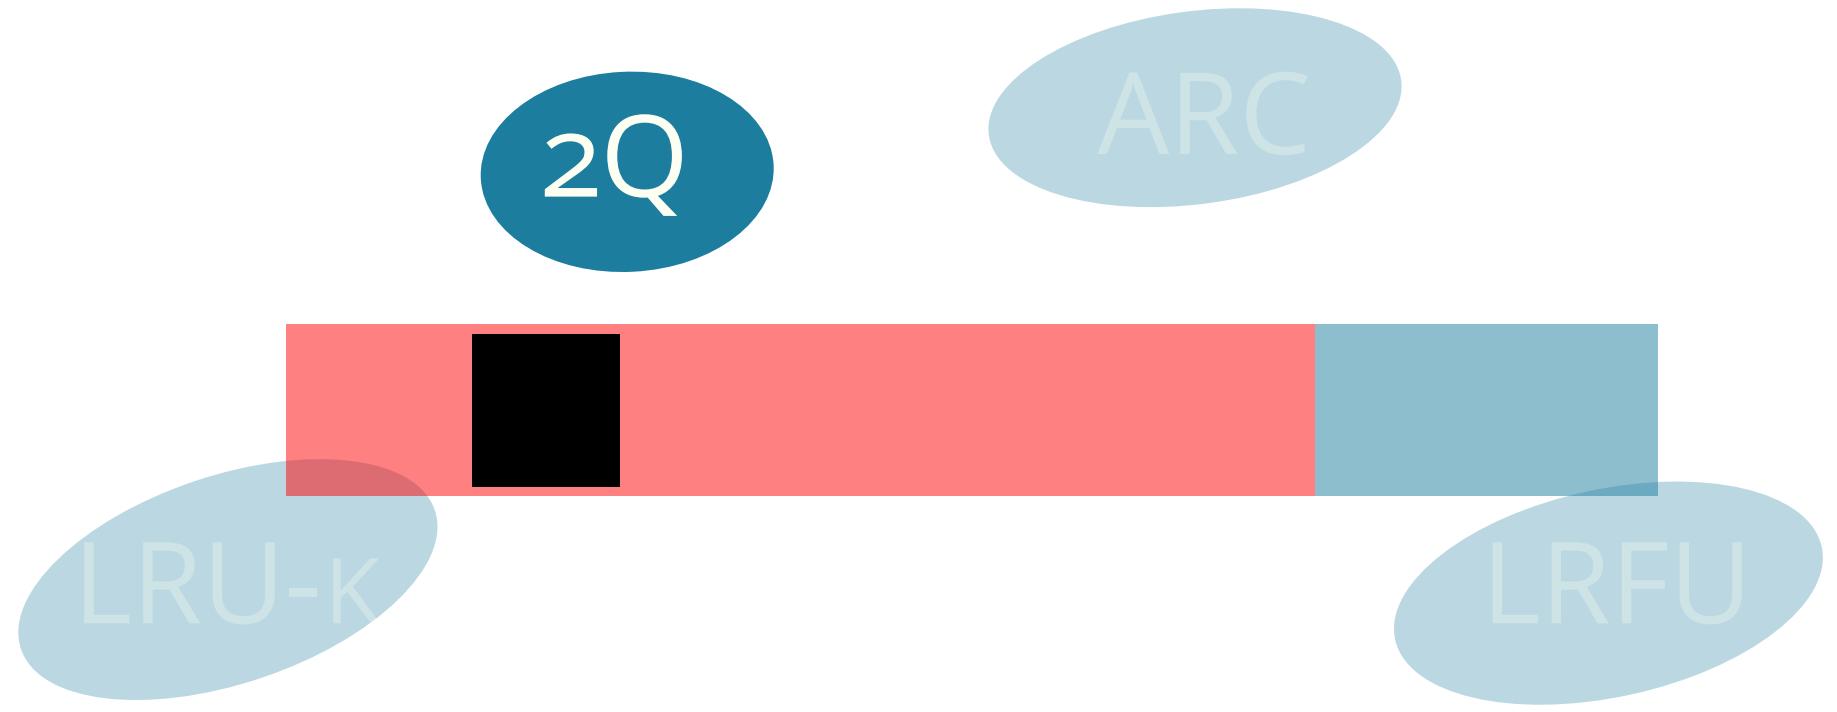
LRU-LFU HYBRID EVICTION ALGORITHMS



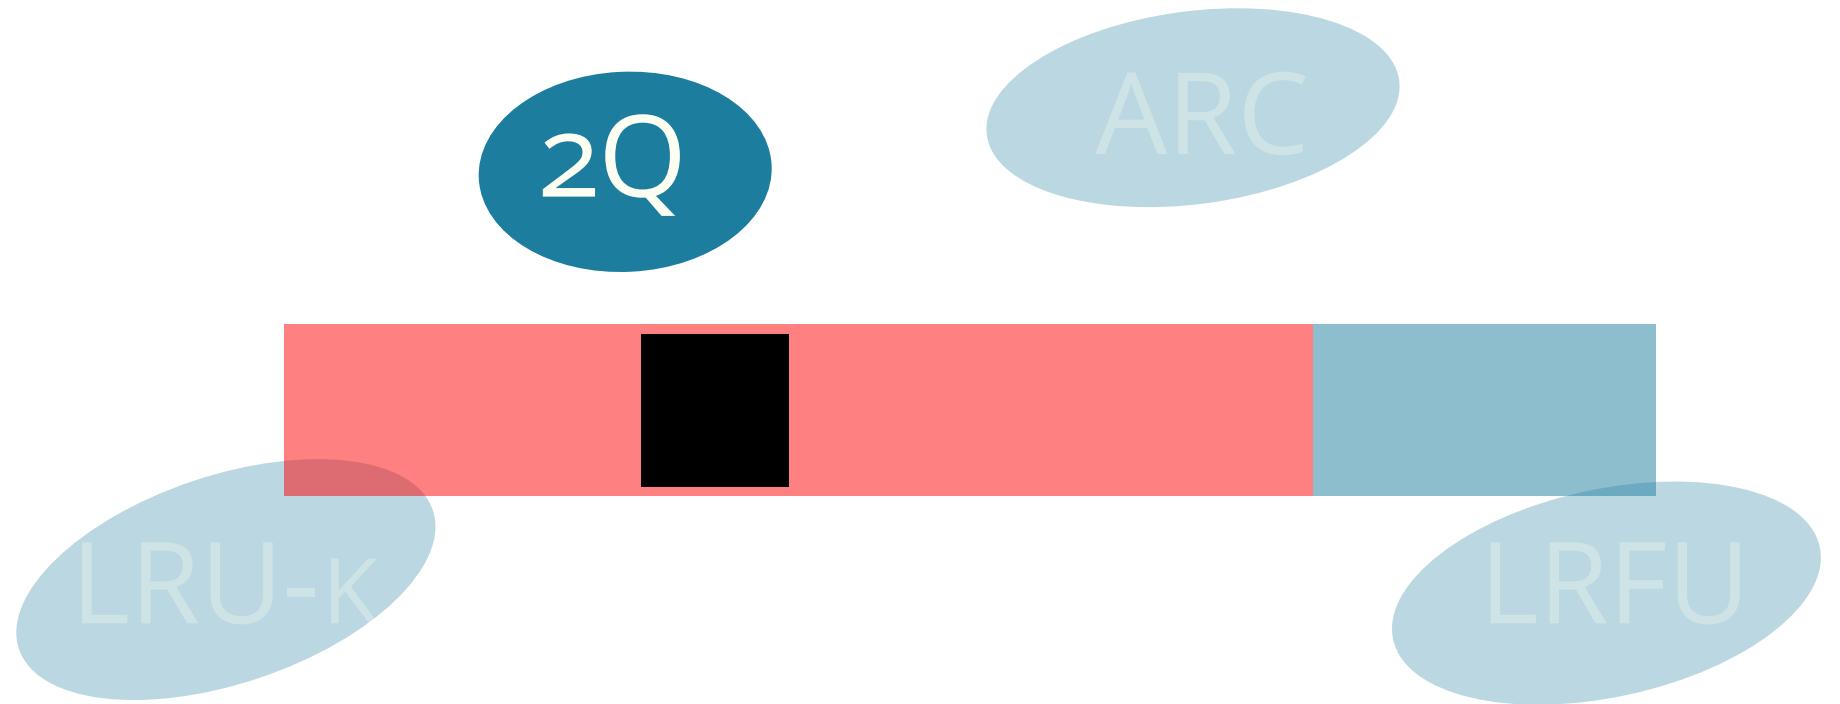
LRU-LFU HYBRID EVICTION ALGORITHMS



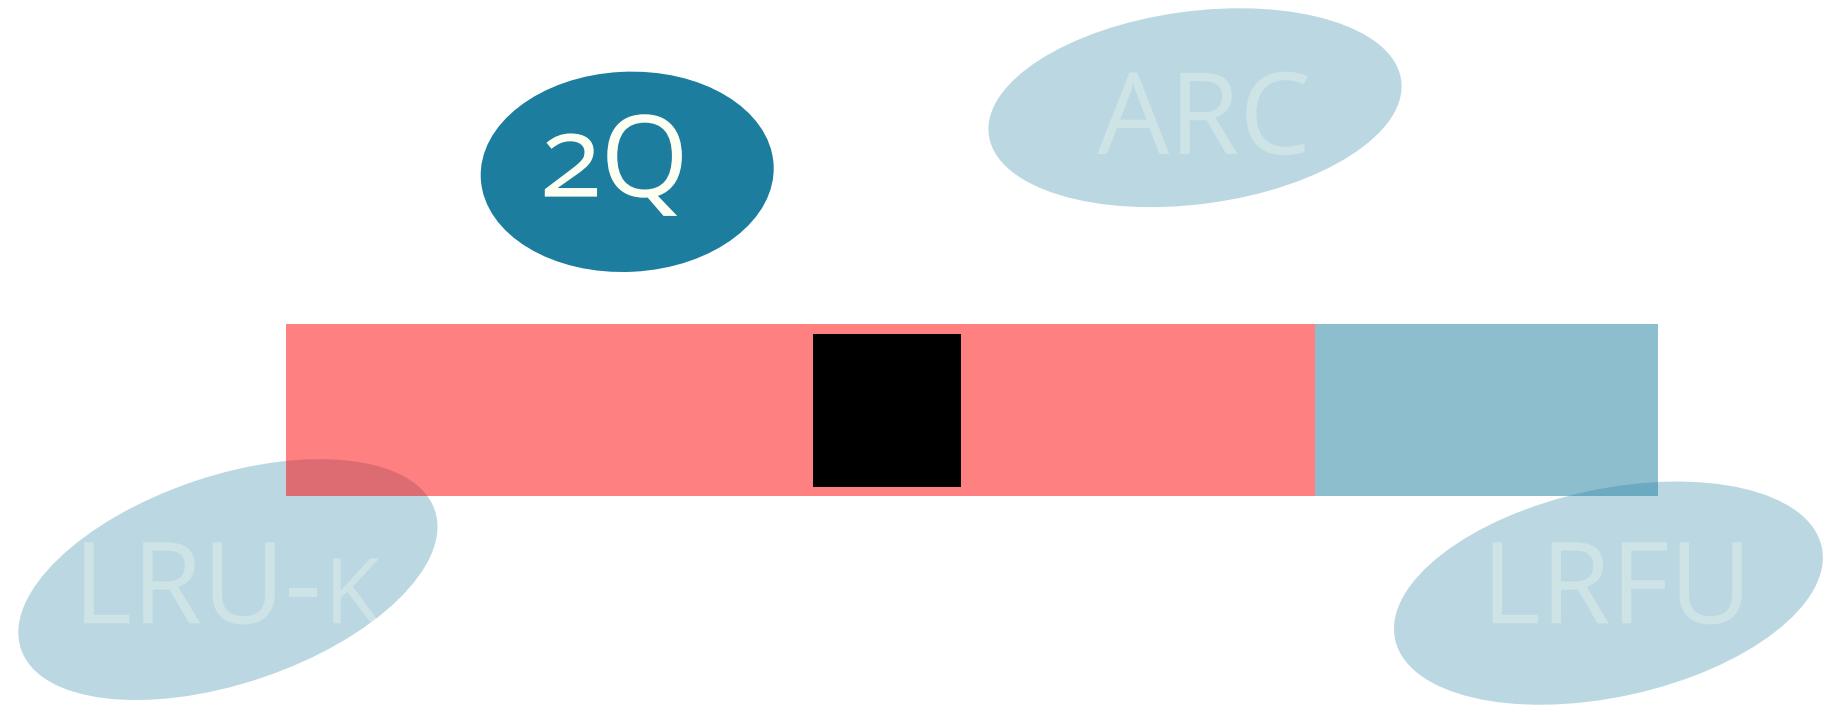
LRU-LFU HYBRID EVICTION ALGORITHMS



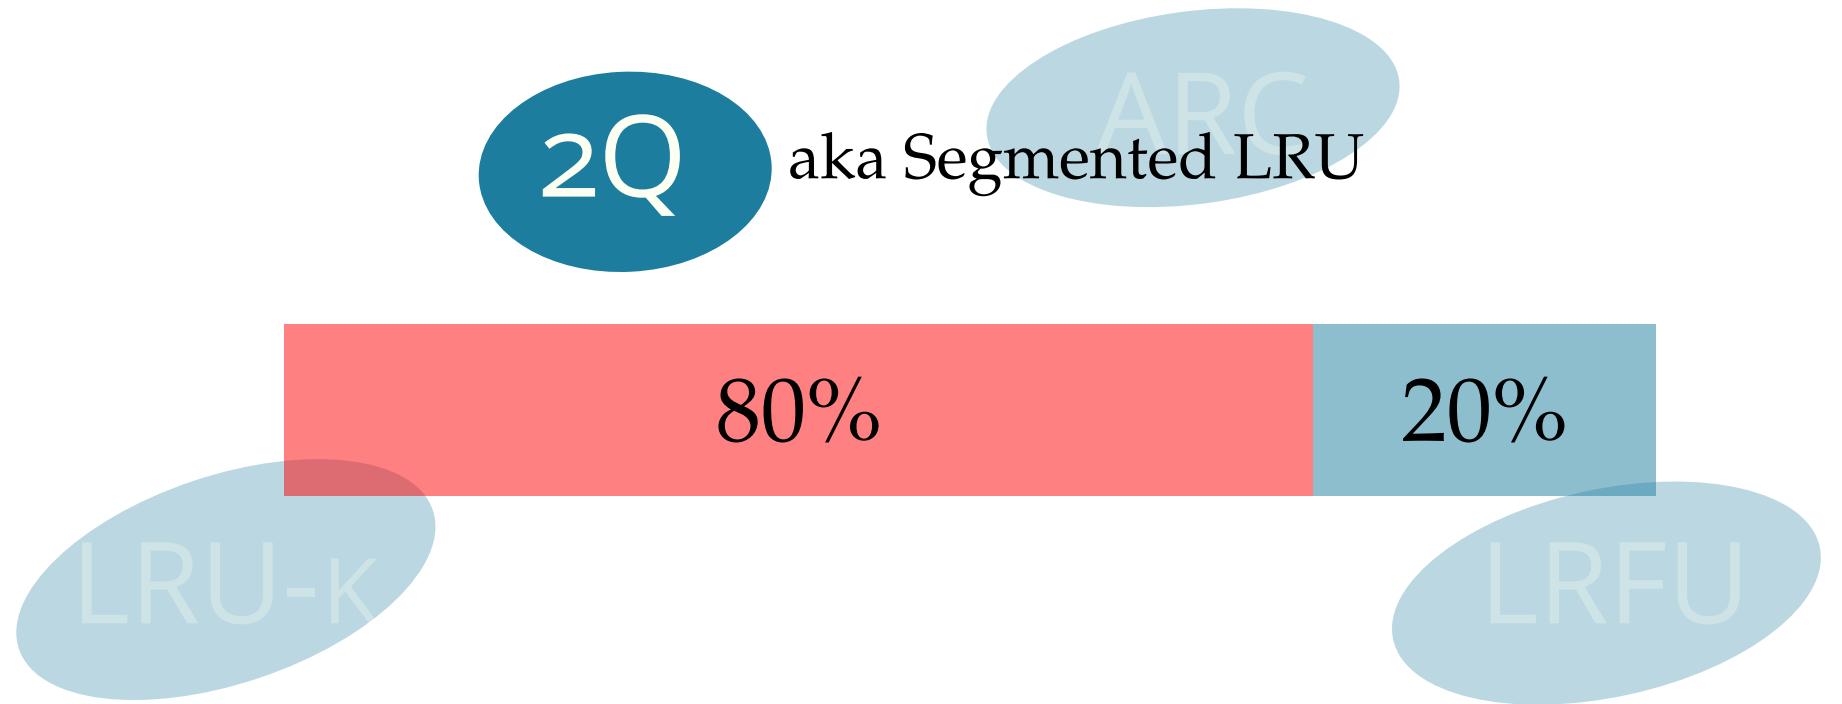
LRU-LFU HYBRID EVICTION ALGORITHMS



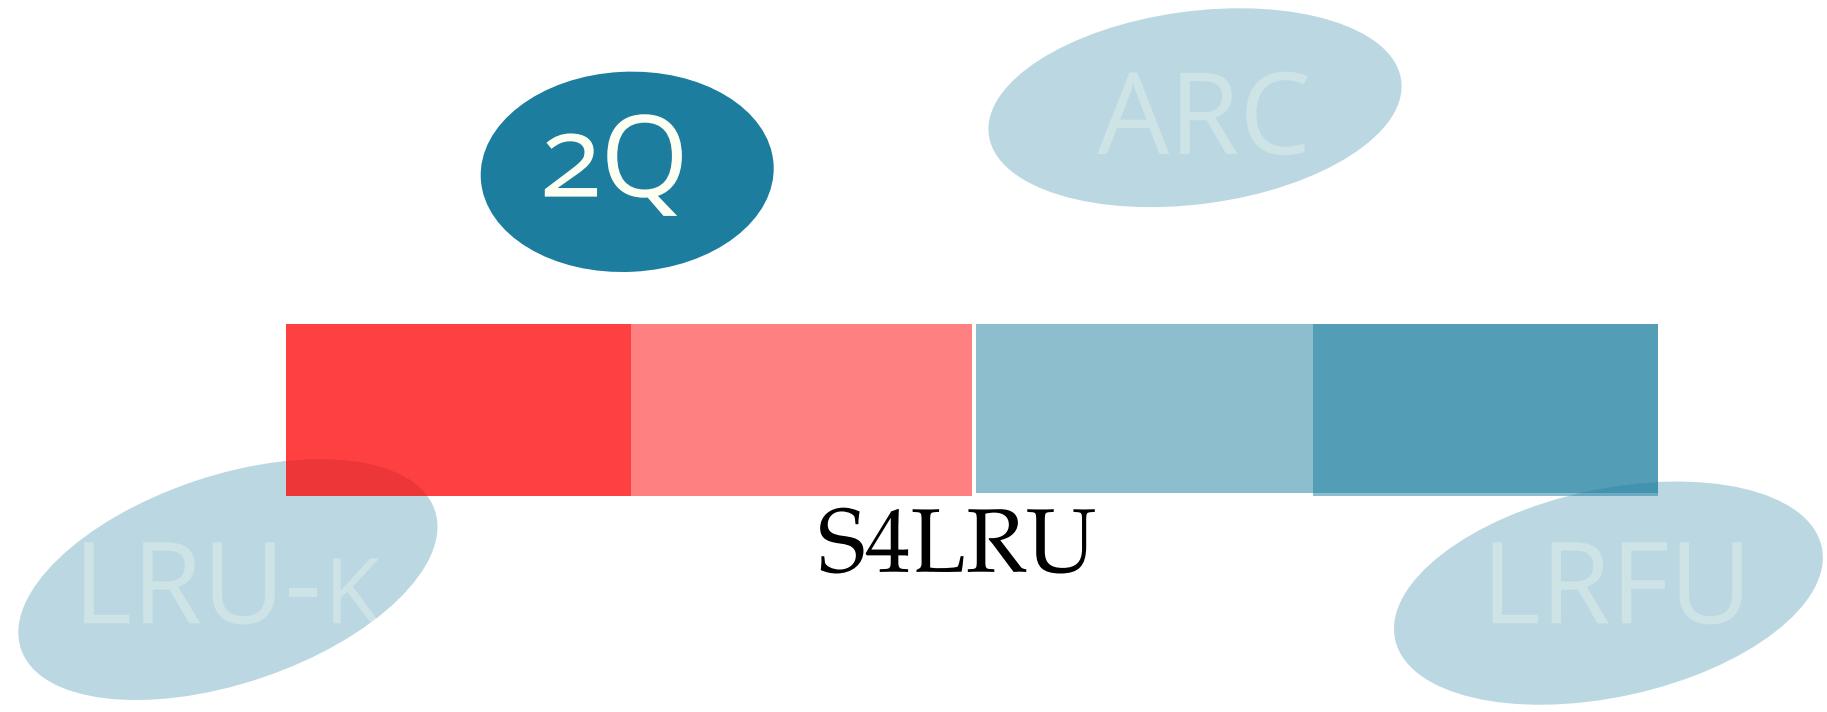
LRU-LFU HYBRID EVICTION ALGORITHMS



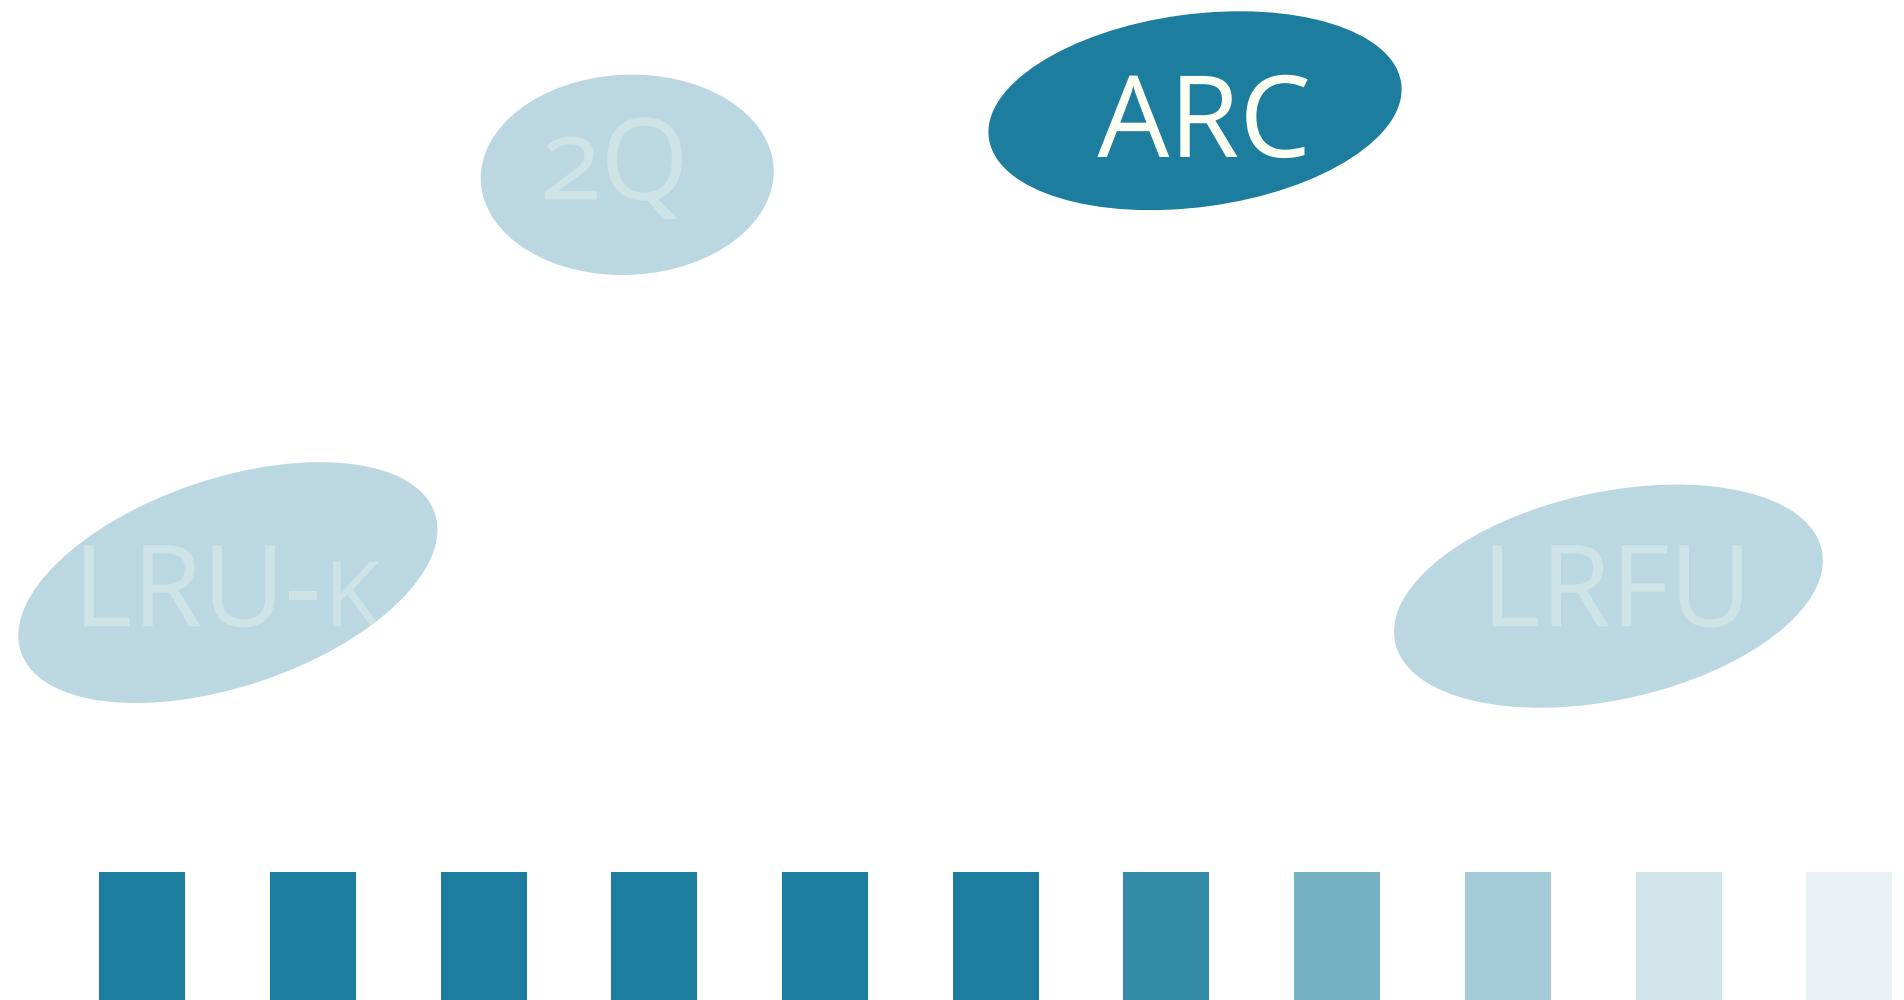
LRU-LFU HYBRID EVICTION ALGORITHMS



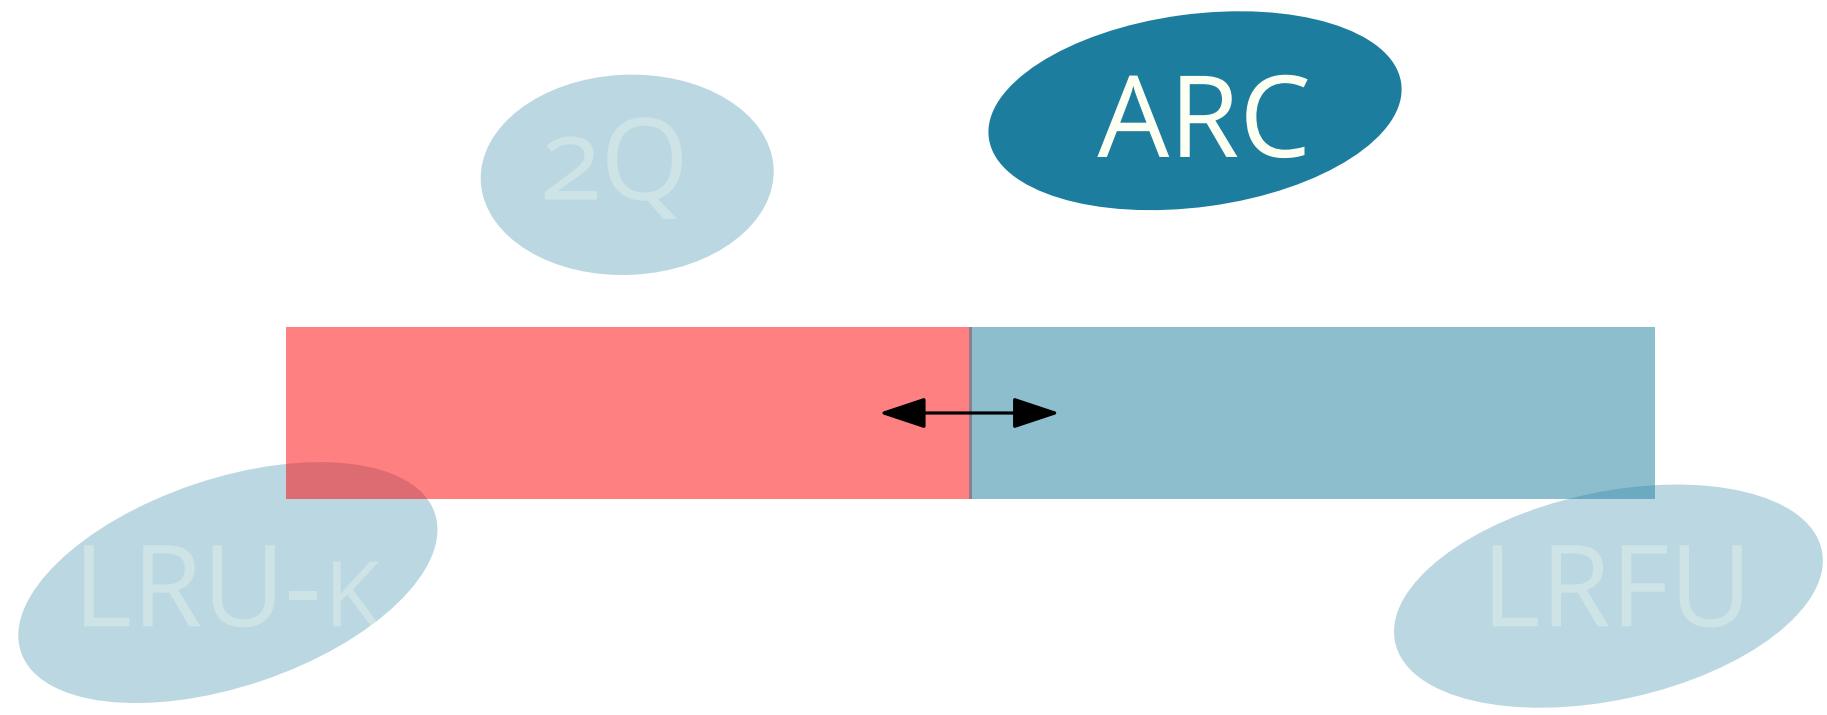
LRU-LFU HYBRID EVICTION ALGORITHMS



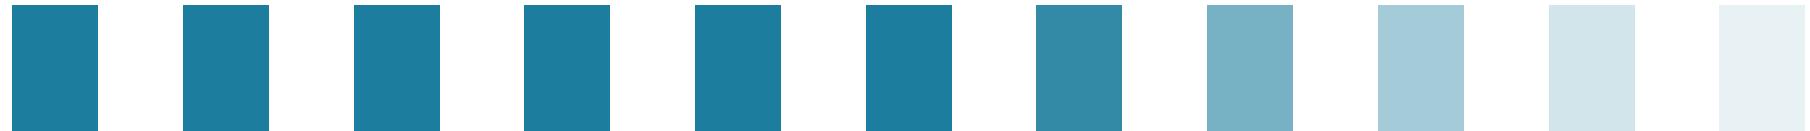
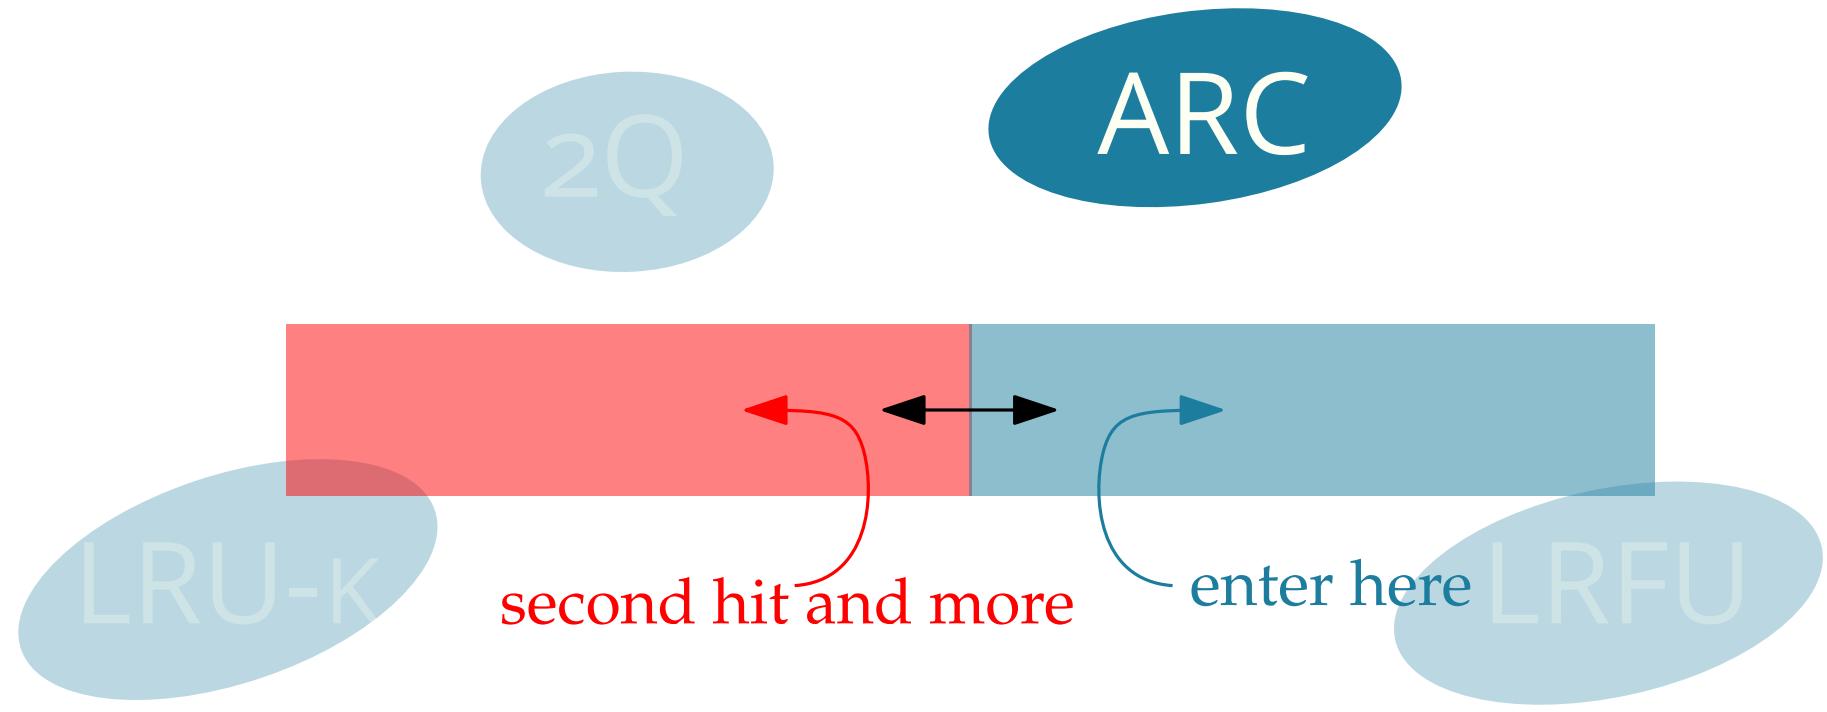
LRU-LFU HYBRID EVICTION ALGORITHMS



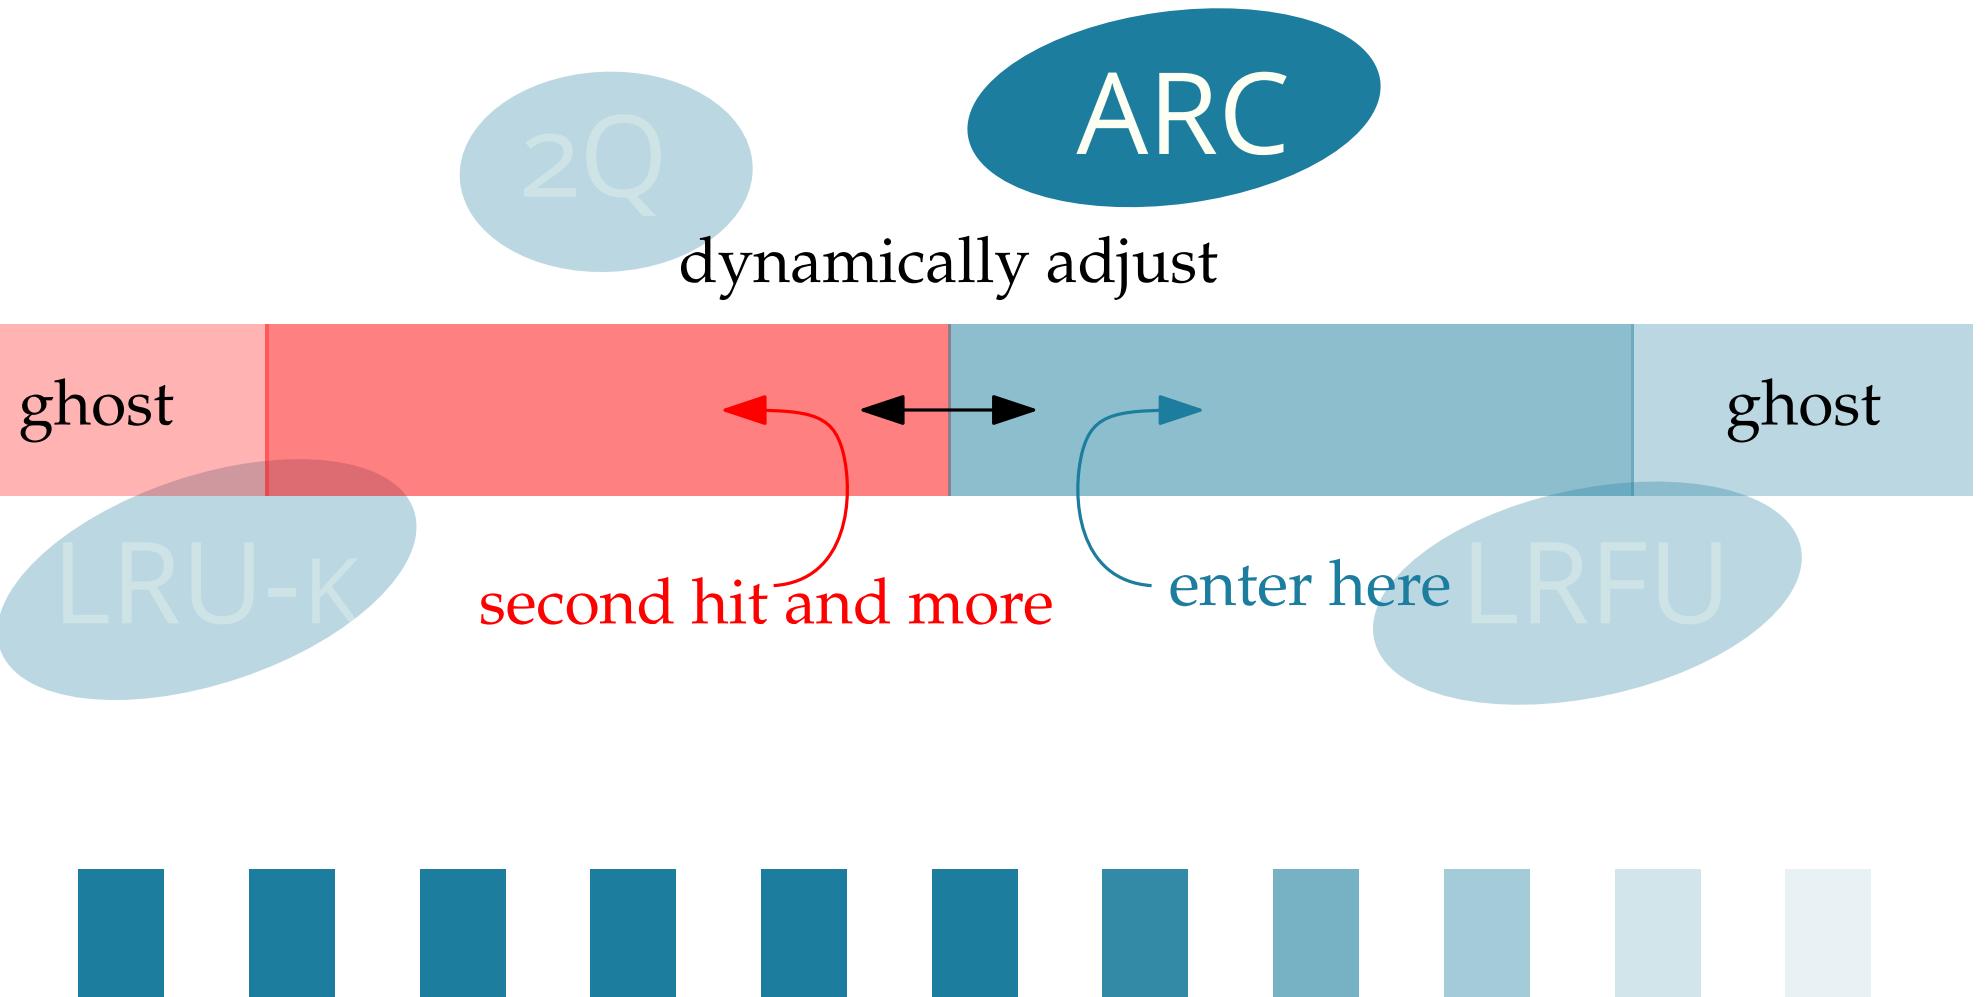
LRU-LFU HYBRID EVICTION ALGORITHMS



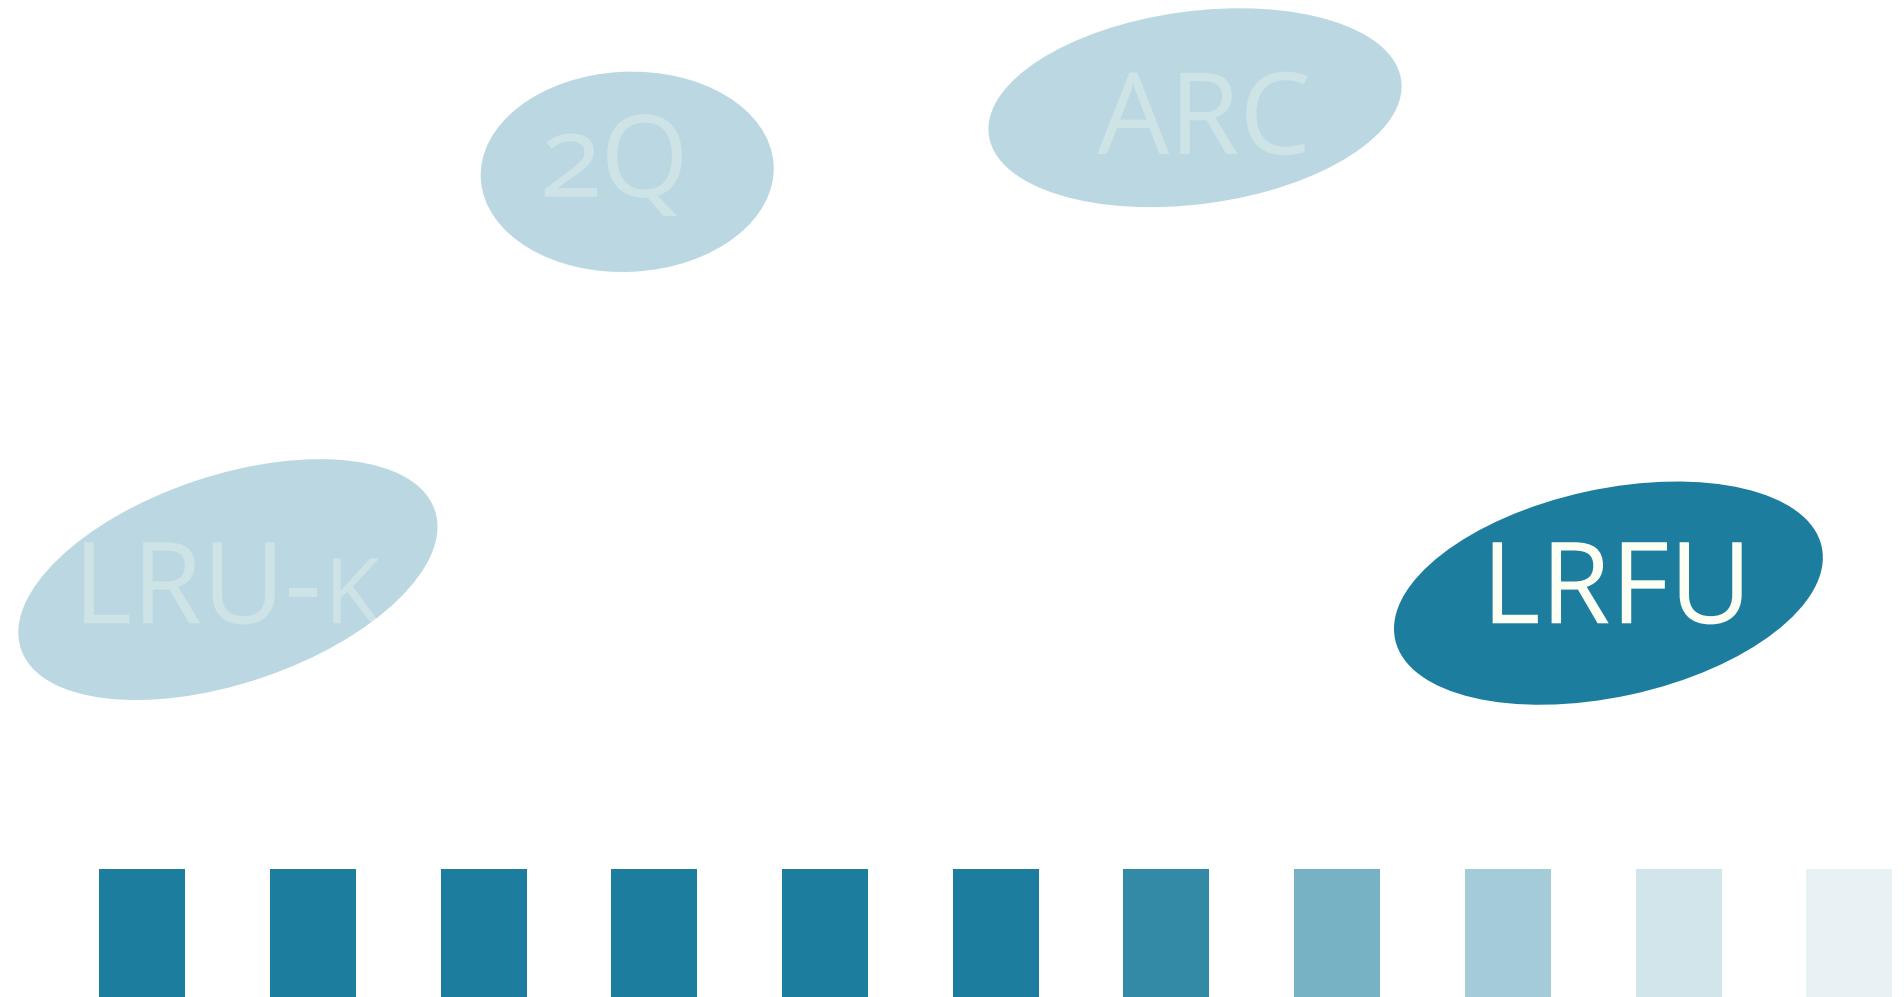
LRU-LFU HYBRID EVICTION ALGORITHMS



LRU-LFU HYBRID EVICTION ALGORITHMS



LRU-LFU HYBRID EVICTION ALGORITHMS



LRU-LFU HYBRID EVICTION ALGORITHMS

combined recency and frequency

$$\text{CRF}(x) = \sum_i \alpha^{-t(x,i)}$$

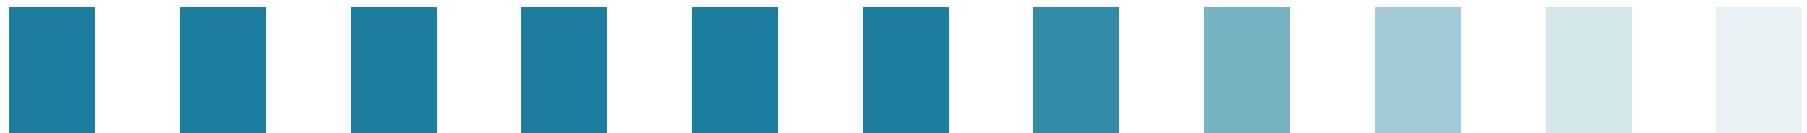
2Q

ARC

time since last i -th request to x

LRU-K

LRFU



LRU-LFU HYBRID EVICTION ALGORITHMS

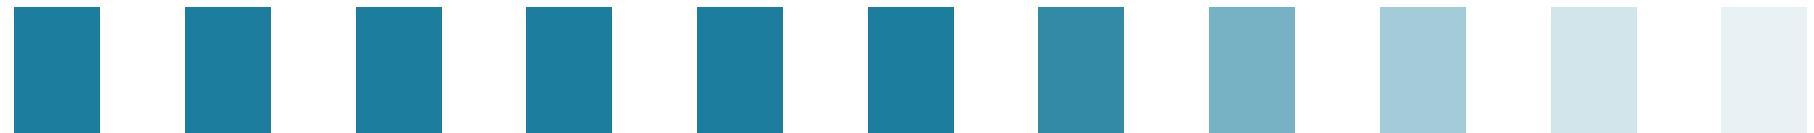
combined recency and frequency

$$\text{CRF}(x) = \sum_i \alpha^{-t(x,i)}$$

2Q

ARC

time since last i -th request to x



LRU-LFU HYBRID EVICTION ALGORITHMS

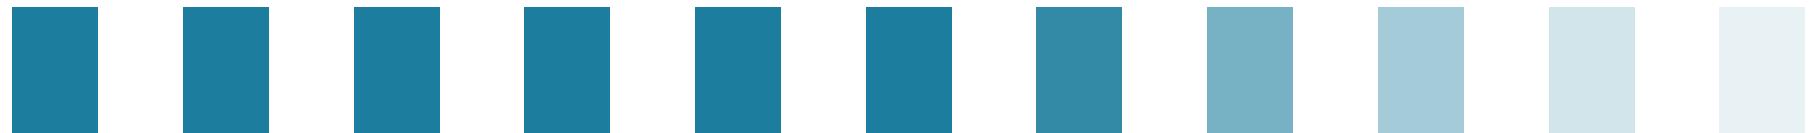
combined recency and frequency

$$\text{CRF}(x) = \sum_i \alpha^{-t(x,i)}$$

2Q

ARC

time since last i -th request to x



LRU-LFU HYBRID EVICTION ALGORITHMS



HYBRID SOLUTIONS TEND TO DO WELL

Evaluation of Caching Strategies Based on Access Statistics of Past Requests, Hasslinger, Ntougias, MMB & DFT 2014

LRU ← CACHING POLICIES → LFU

exact ————— TOWARDS PRACTICALITY ————— approximate

UNDERSTANDING THE WORKLOAD

LRU ← → LFU

handles bursts in popularity well
the colder can push out the warmer

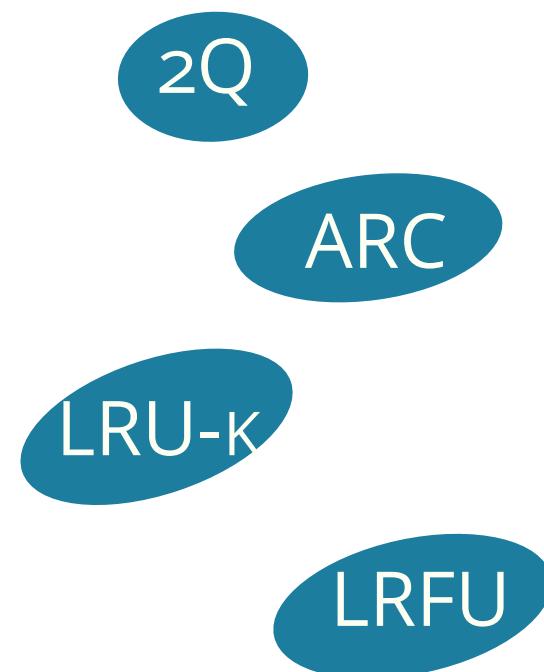
suffers from cache pollution
optimal if popularity is static

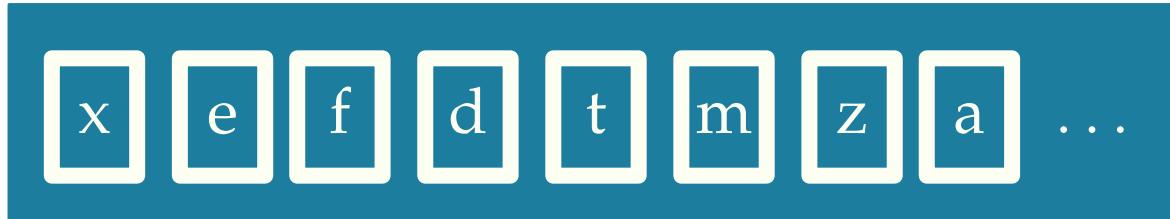
2nd hit caching
score-based LRU
SIZE

LRU ← → LFU

handles bursts in popularity well
the colder can push out the warmer

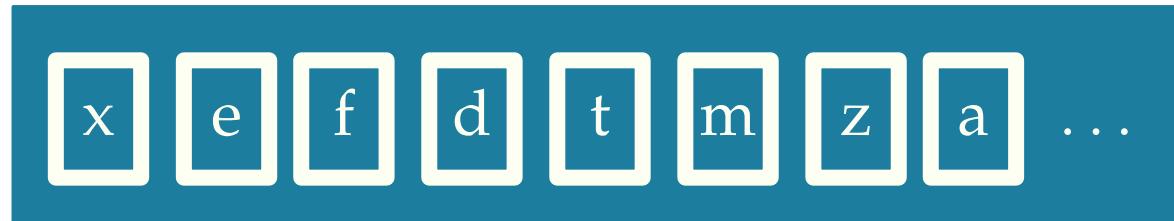
suffers from cache pollution
optimal if popularity is static



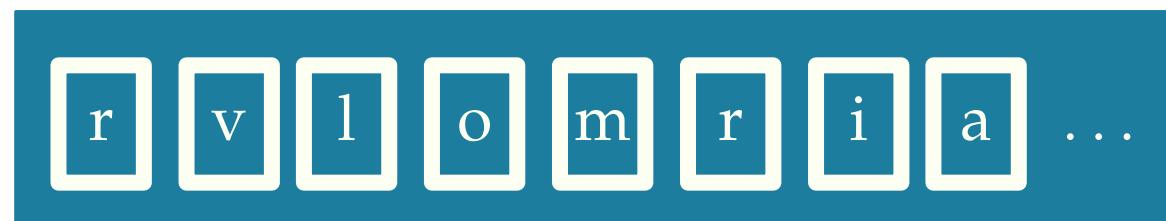


evict this guy first

what a caching algorithm boils down to

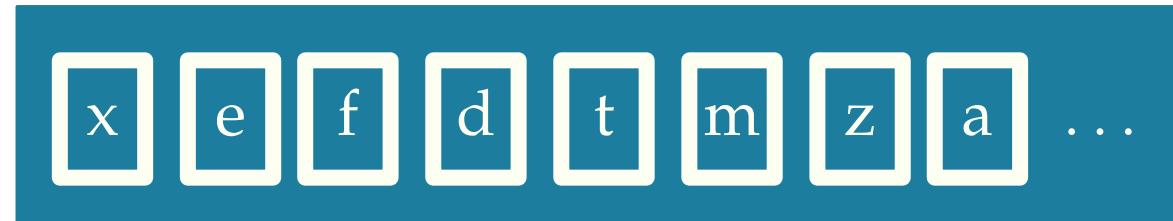


↑
evict this guy first

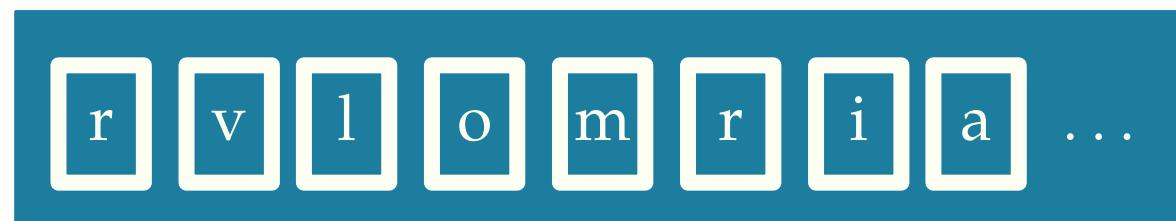


priority 5.1 14 99.99 3 π 2.8 e 2000

what a caching algorithm boils down to



↑
evict this guy first

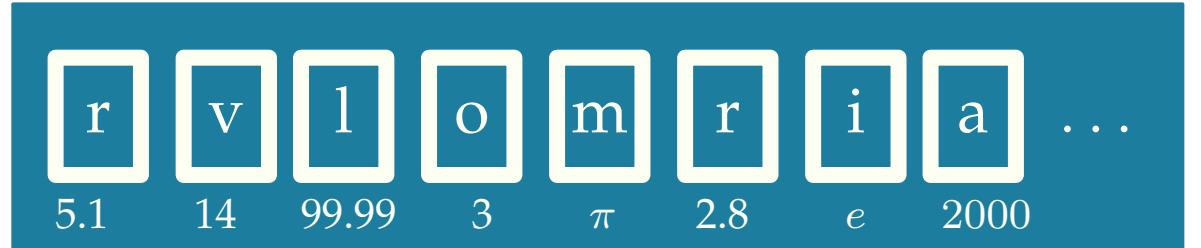


priority 5.1 14 99.99 3 π 2.8 e 2000

↑
evict this guy

what a caching algorithm boils down to

caching algorithm template



handle-request(object o):

if o is in cache:

 update priority(o)

 (for all p in cache: update priority(p))

else:

 while there is not enough room in cache for o:

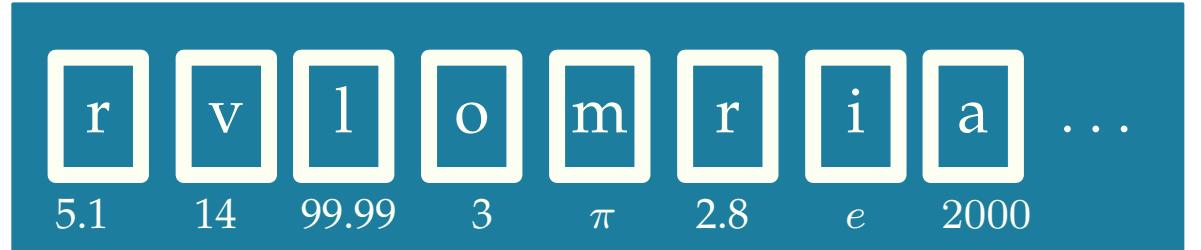
 find and evict object in cache with smallest priority

 bring o in cache

 compute priority(o)

 (for all p in cache: update priority(p))

caching algorithm template



handle-request(object o):

if o is in cache:

 update priority(o)

 (for all p in cache: update priority(p))

else:

 while there is not enough room in cache for o:

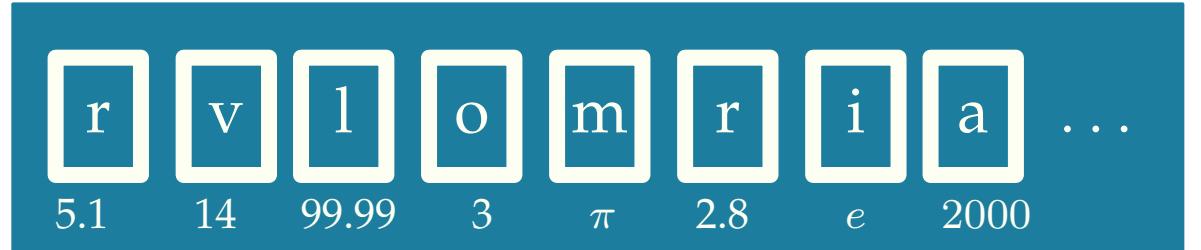
 find and evict object in cache with smallest priority

 bring o in cache

 compute priority(o)

 (for all p in cache: update priority(p))

caching algorithm template



handle-request(object o):

if o is in cache:

update priority(o)

(for all p in cache: update priority(p))

else:

while there is not enough room in cache for o:

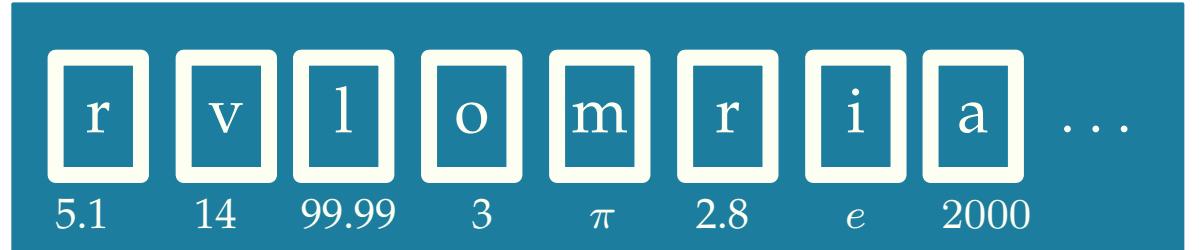
find and evict object in cache with smallest priority

bring o in cache

compute priority(o)

(for all p in cache: update priority(p))

caching algorithm template



handle-request(object o):

if o is in cache:

update priority(o) O(1)

(for all p in cache: update priority(p))

else:

while there is not enough room in cache for o:

find and evict object in cache with smallest priority

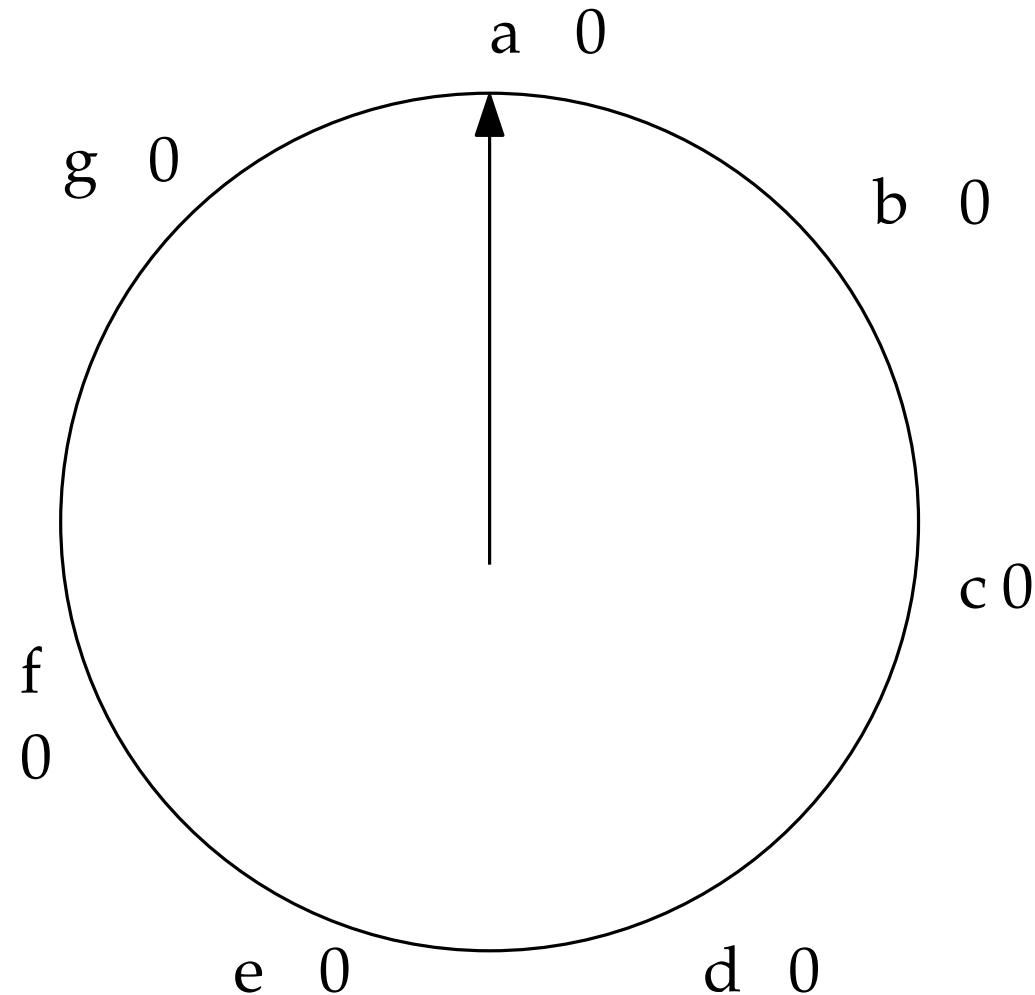
O(1)

bring o in cache

compute priority(o) O(1)

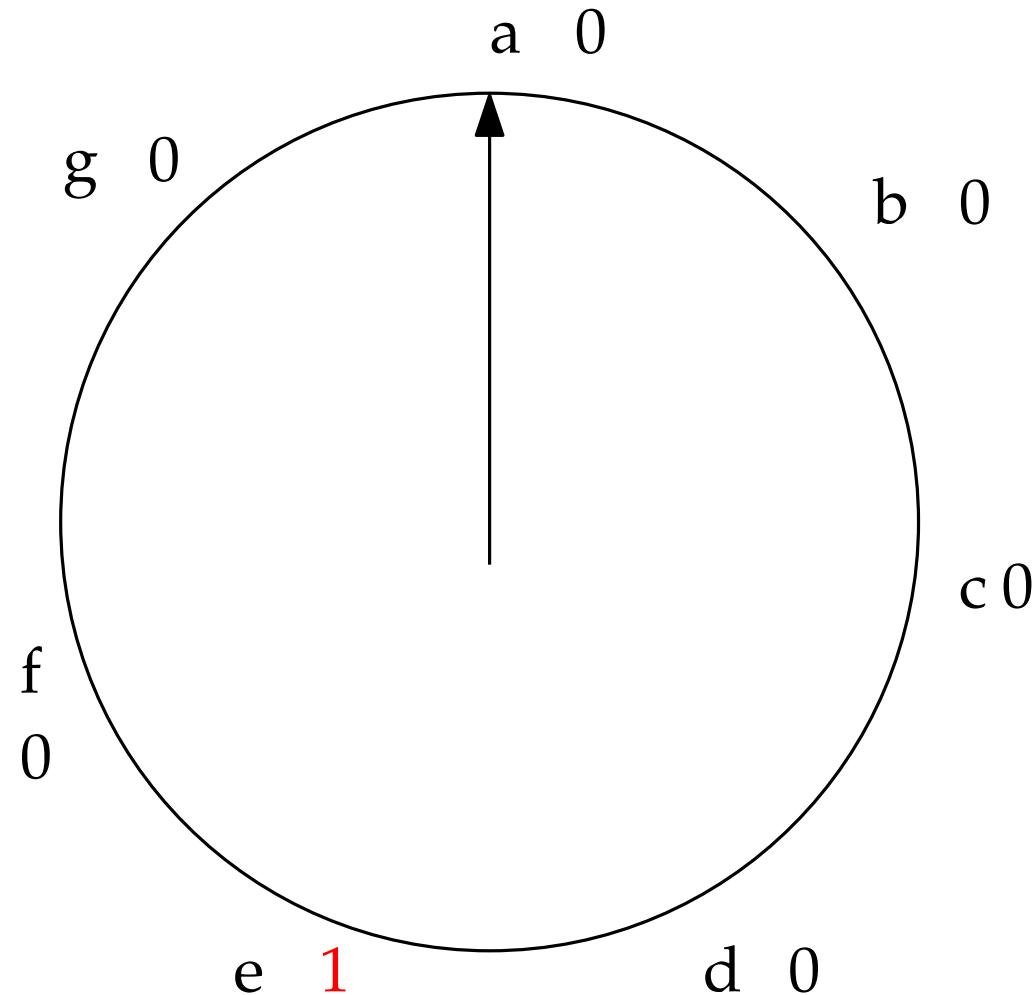
(for all p in cache: update priority(p))

clock: an approximation of LRU



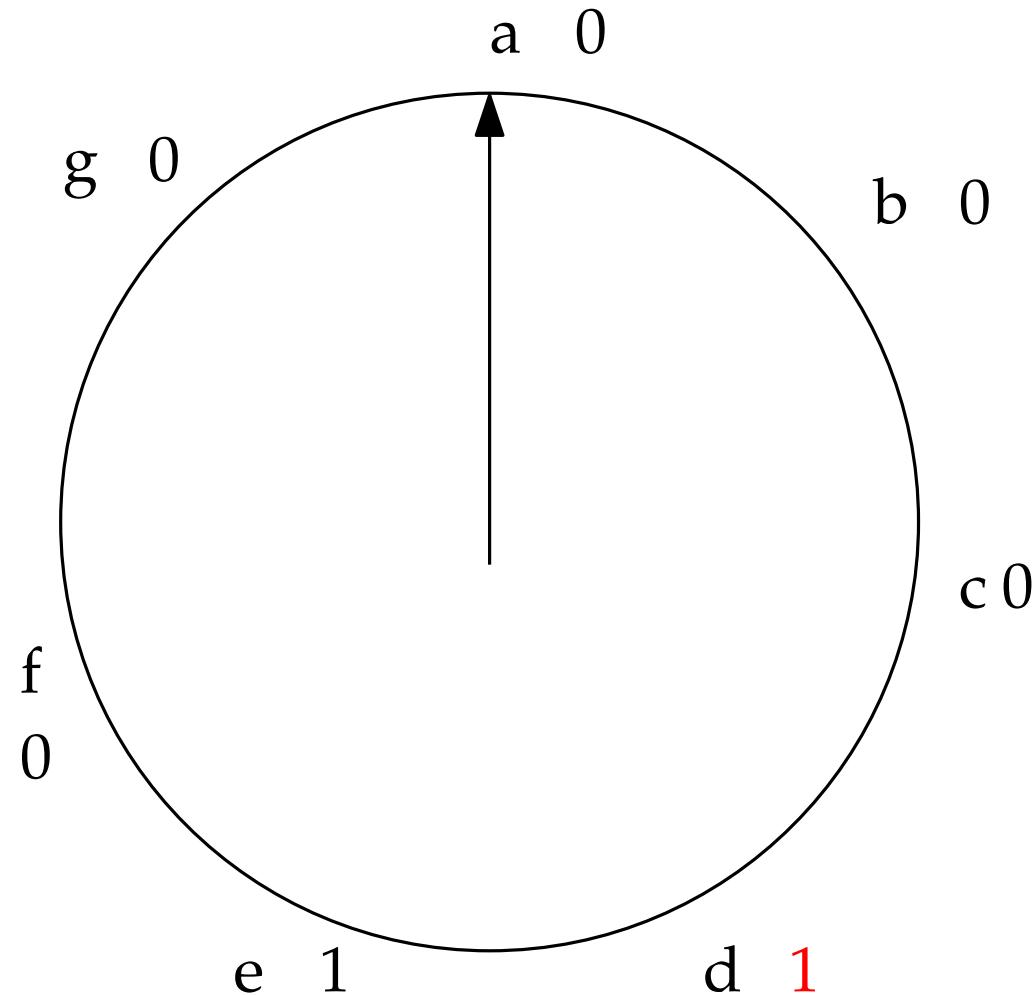
e d a e b h

clock: an approximation of LRU



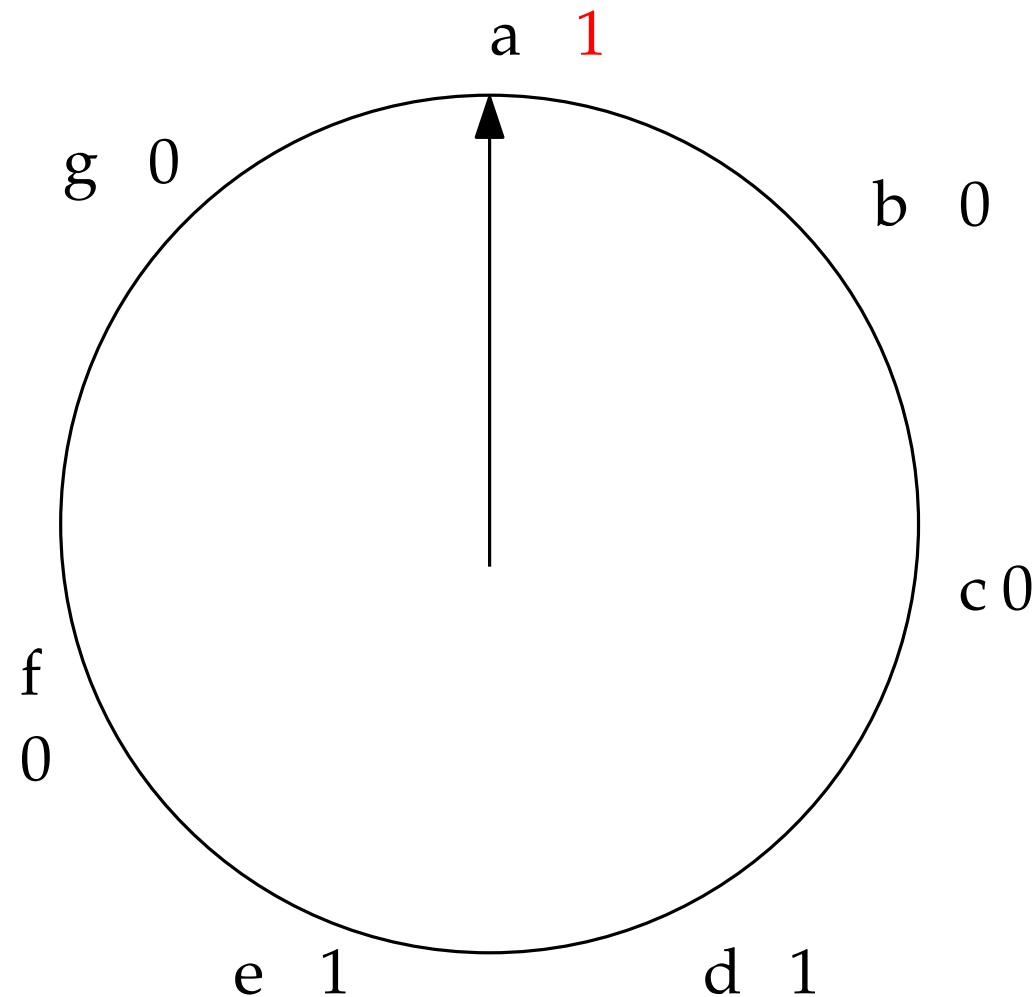
e d a e b h

clock: an approximation of LRU



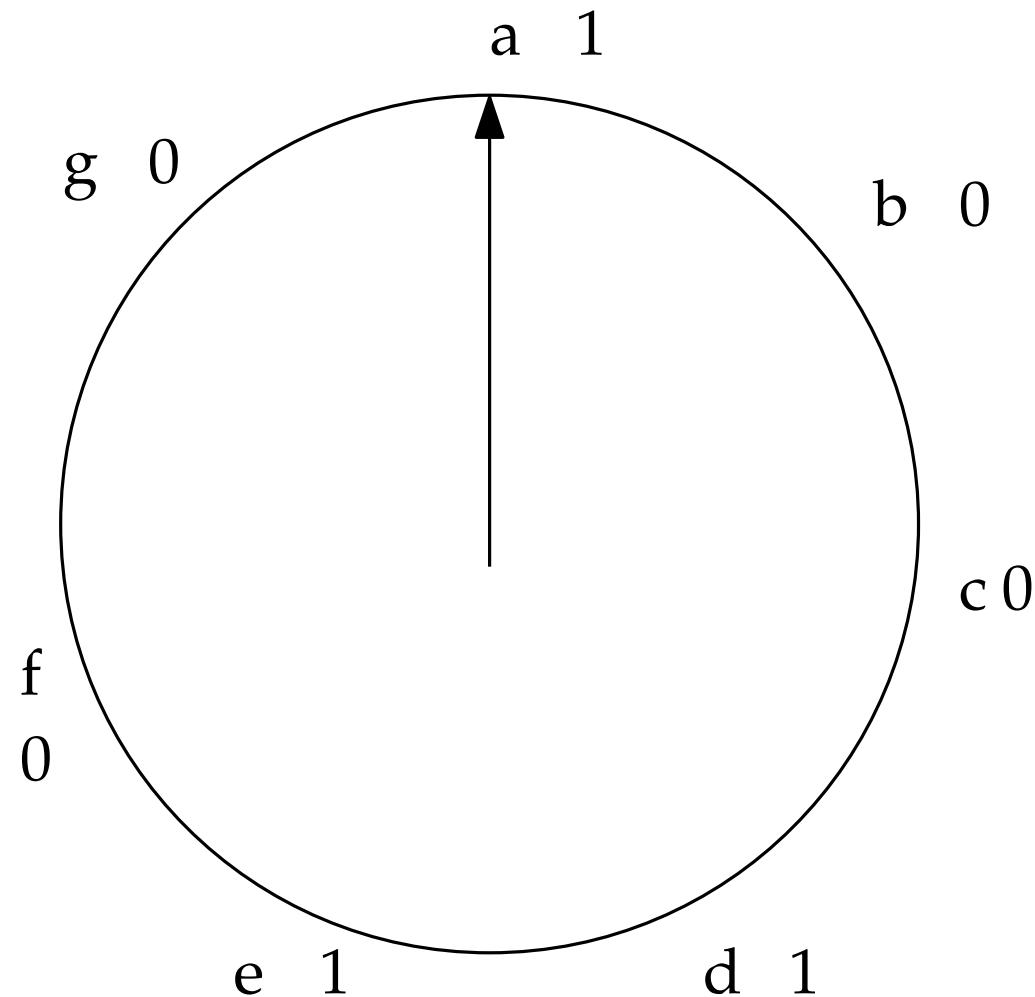
e d a e b h

clock: an approximation of LRU



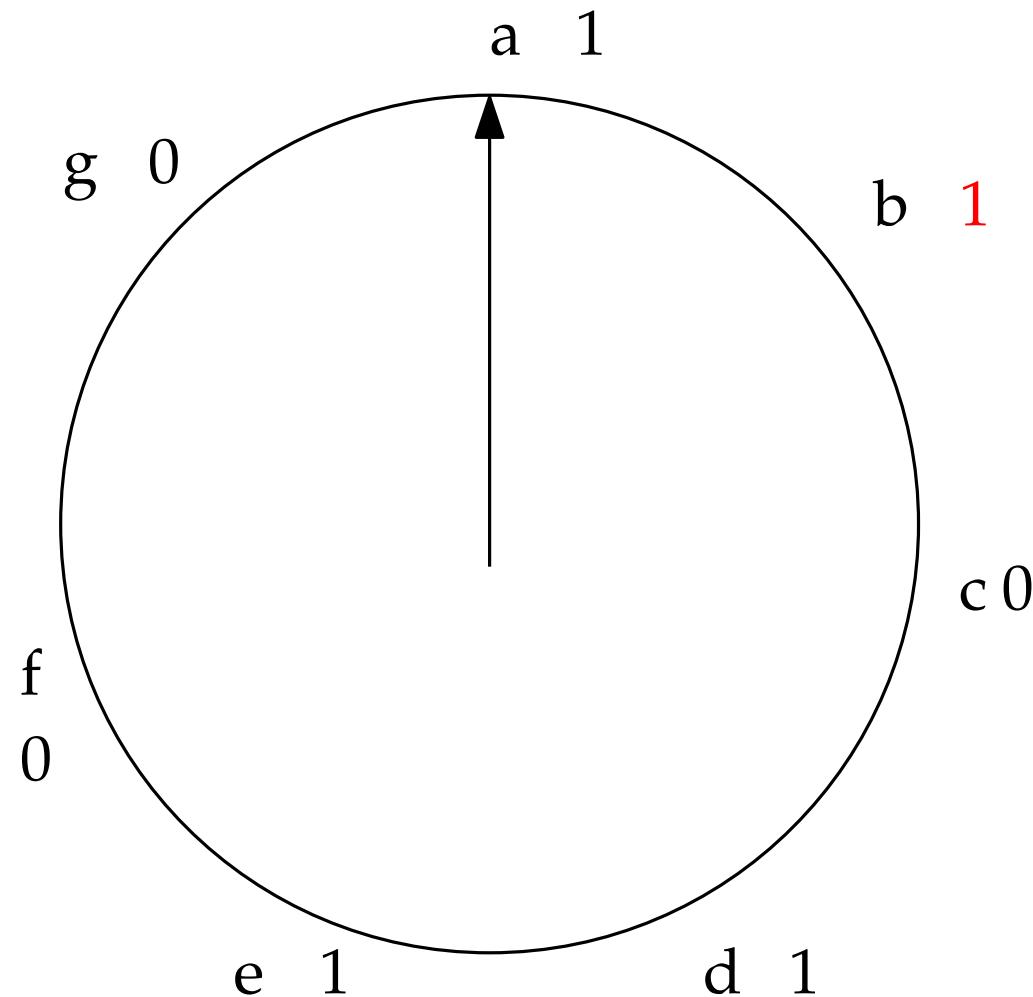
e d a e b h

clock: an approximation of LRU



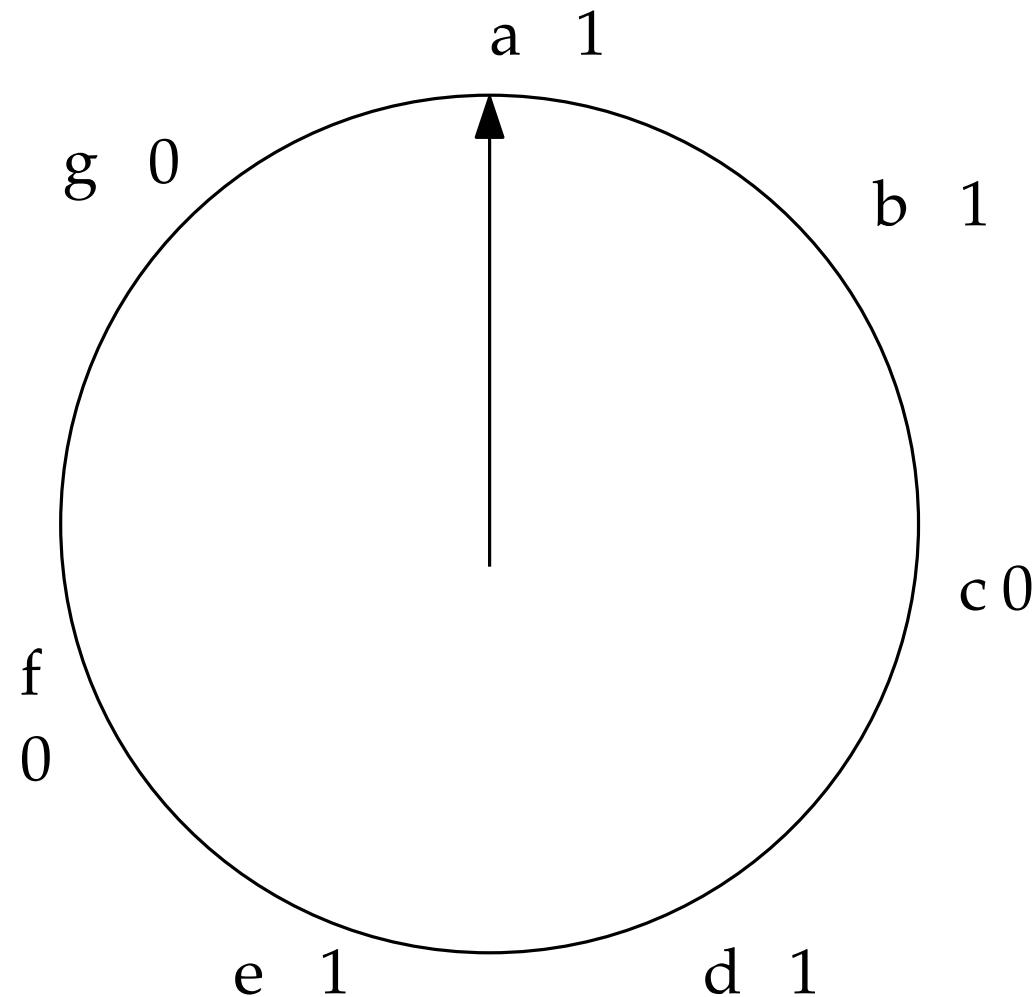
e d a e b h

clock: an approximation of LRU



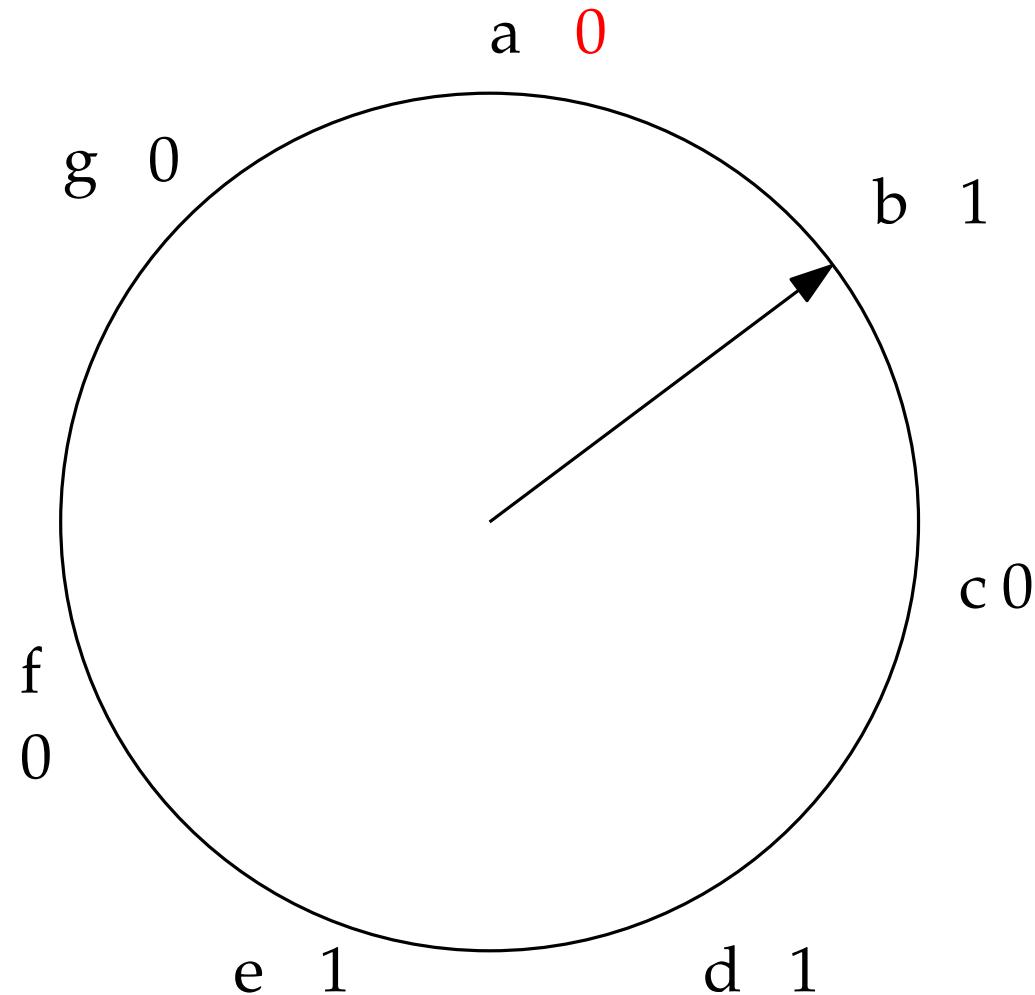
e d a e b h

clock: an approximation of LRU



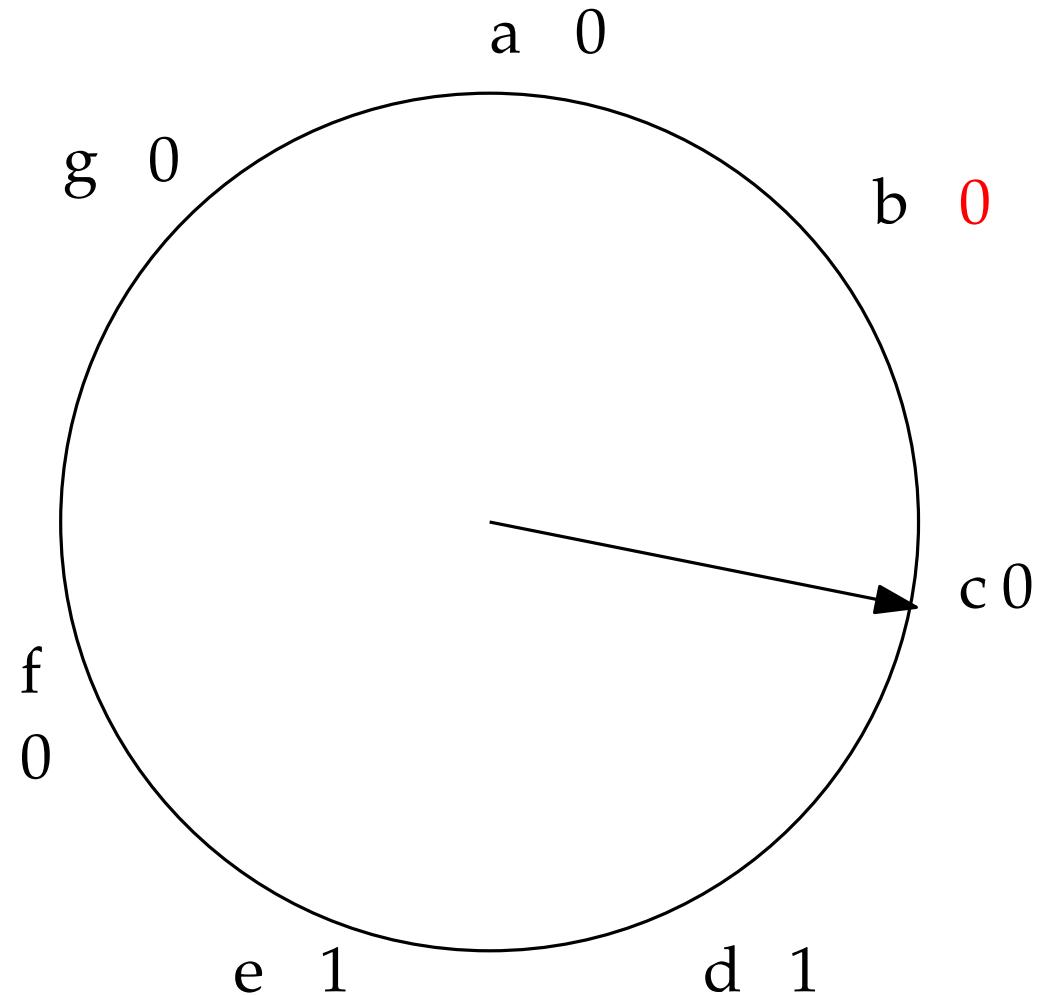
e d a e b h

clock: an approximation of LRU



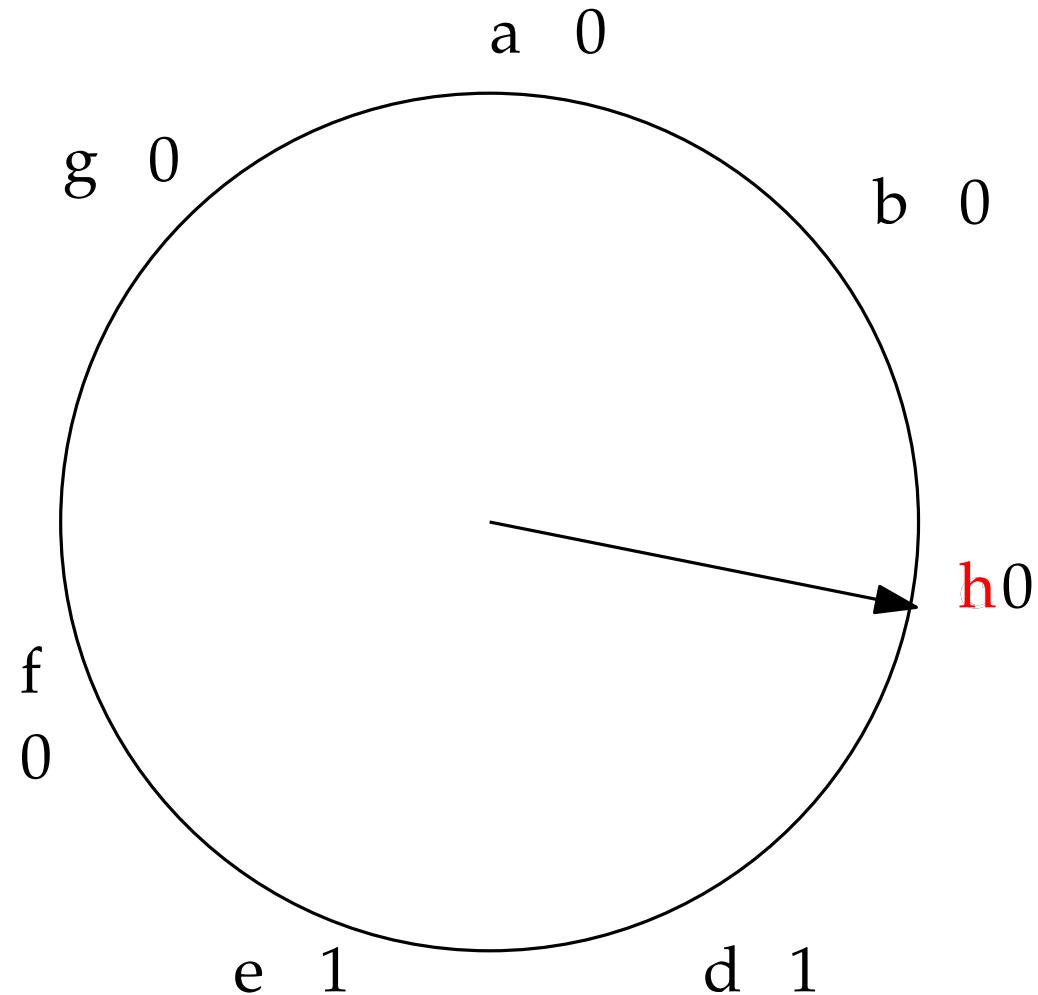
e d a e b h

clock: an approximation of LRU



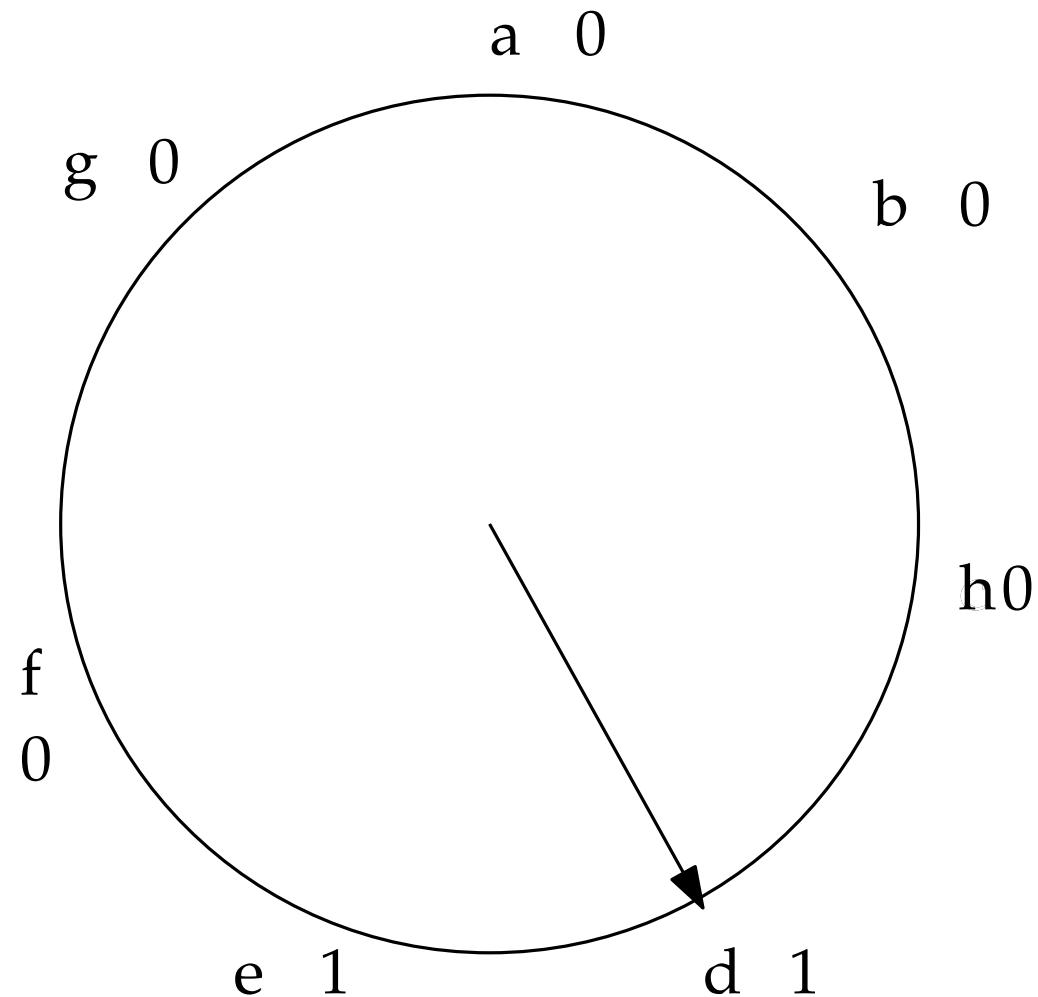
e d a e b h

clock: an approximation of LRU



e d a e b h

clock: an approximation of LRU



e d a e b h

handle-request(object o):

if o is in cache:

update priority(o)

else:

while there is not enough room in cache for o:

find and evict object in cache with smallest priority

bring o in cache

compute priority(o)

LRU

handle-request(object o):

if o is in cache:

 update priority(o) O(1)

else:

 while there is not enough room in cache for o:

 find and evict object in cache with smallest priority

 bring o in cache O(1)

 compute priority(o) O(1)

LRU

clock

handle-request(object o):

if o is in cache:

 update priority(o) O(1) move to front bit =1

else:

 while there is not enough room in cache for o:

 find and evict object in cache with smallest priority

 bring o in cache O(1) remove from tail move hand
 set bits to 0

 compute priority(o) O(1) place in front

 bit = 0

handle-request(object o):

if o is in cache:

update priority(o)

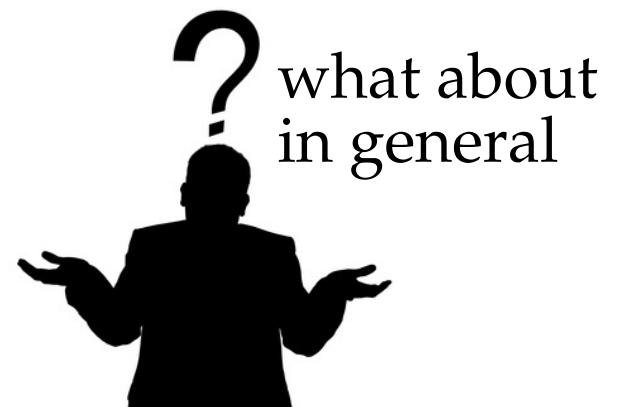
else:

while there is not enough room in cache for o:

find and evict object in cache with smallest priority

bring o in cache

compute priority(o)



what about
in general

using a priority queue (heap)

handle-request(object o):

if o is in cache:

 update priority(o) $O(\log n)$

else:

 while there is not enough room in cache for o:

 find and evict object in cache with smallest priority

 bring o in cache

$O(\log n)$

 compute priority(o)

$O(\log n)$

$n = \text{nb objects in cache}$

exact
algorithm

approximate
algorithm

A horizontal double-headed arrow connects the two algorithm names. The arrow points from the left towards the right, indicating a relationship or comparison between the two.

LRU

clock

exact
algorithm



approximate
algorithm

LRU



clock

your favorite
algorithm

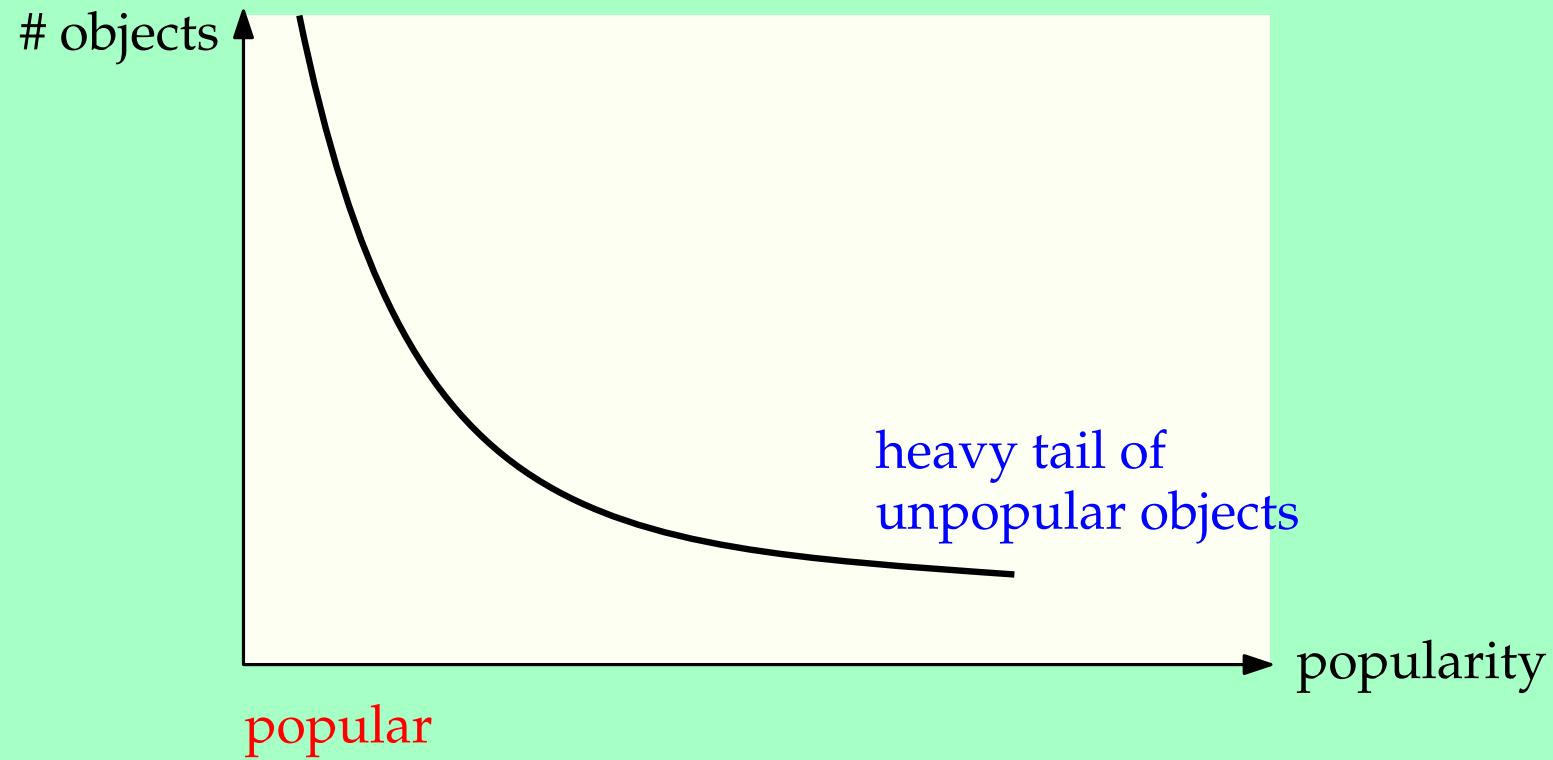


good enough
algorithm

pre-sorting as much as possible

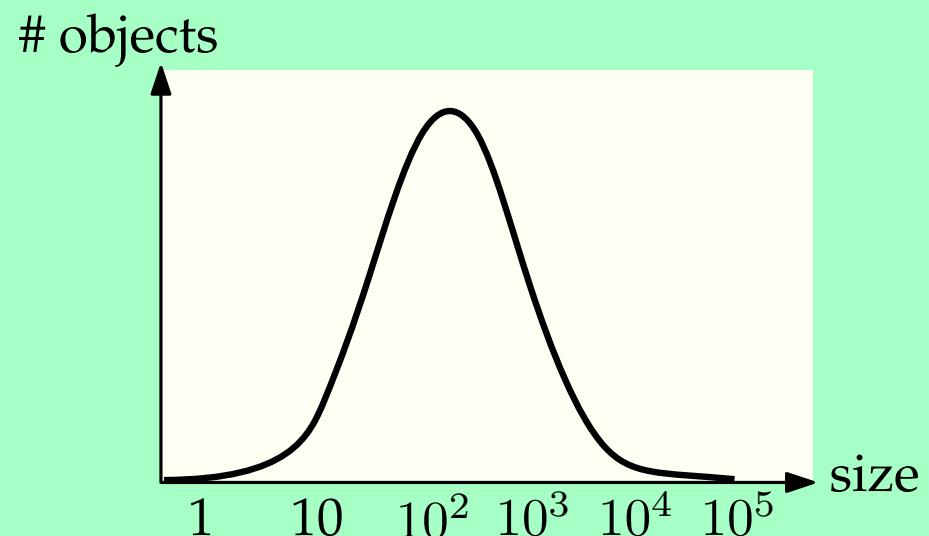
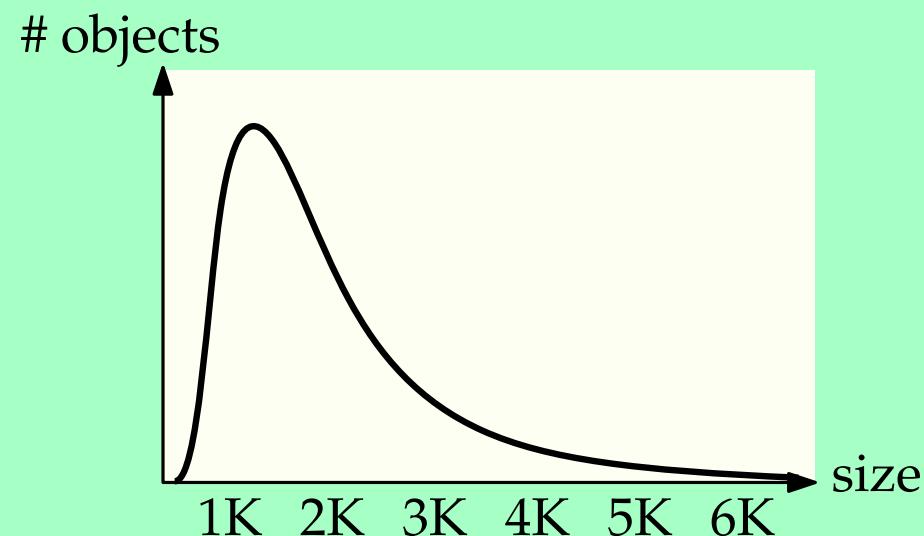
coarse graining/bucketing/rounding priorities

The popularity of objects follows Zipf's law.



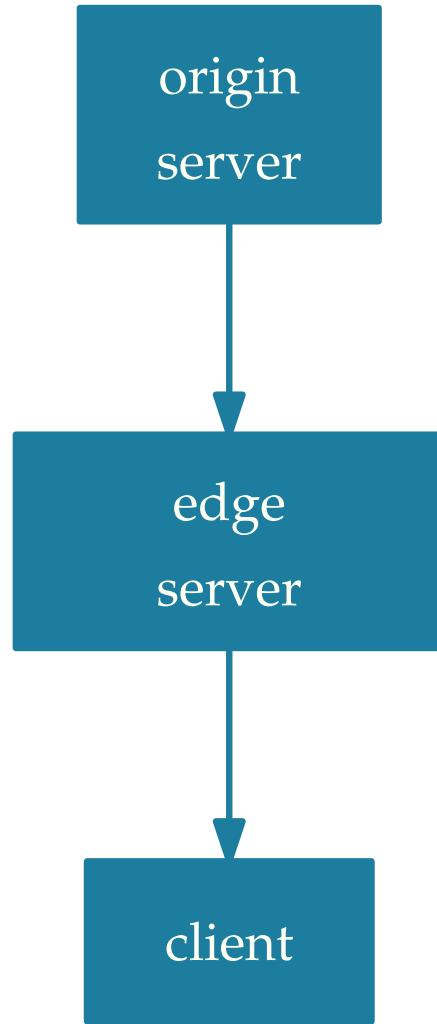
UNDERSTANDING THE WORKLOAD

The object size follows a log-normal distribution



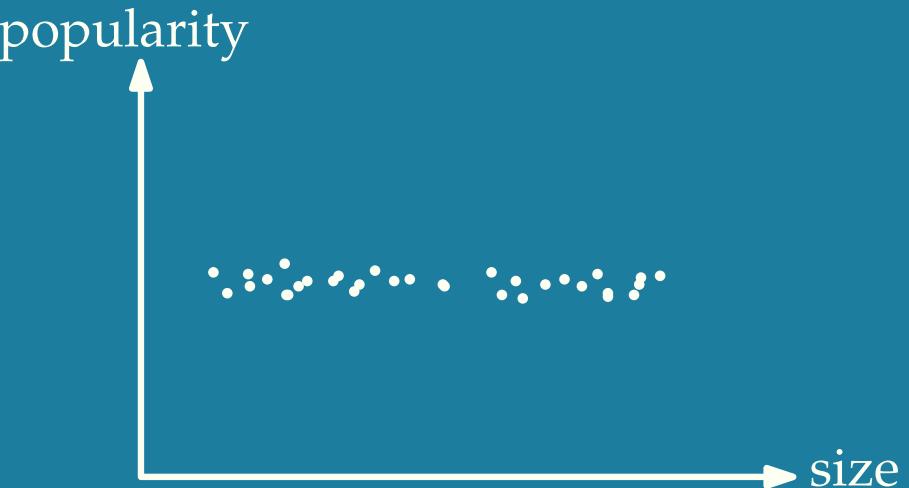
UNDERSTANDING THE WORKLOAD

FIGURES OF MERIT



- reduce load
maximize byte hit ratio
sum of sizes of objects that were request hits

- reduce perceived delay
maximize hit ratio
request hits



prefer to keep
small objects

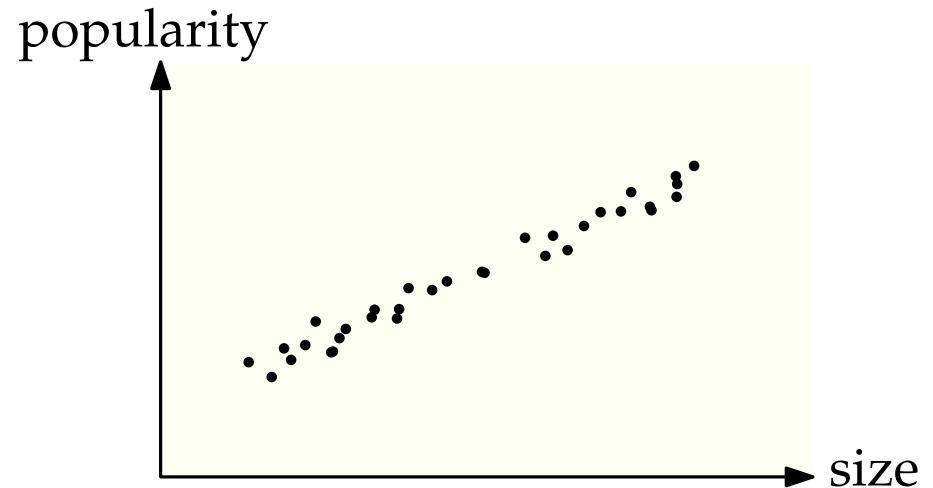
large objects

better ← $\xrightarrow{\text{hit ratio}}$ worse

worse ← $\xrightarrow{\text{byte hit ratio}}$ better

same ← $\xrightarrow{\text{byte hit ratio}}$ same

different popularity-size correlations



prefer to keep
small objects

large objects

worse ← $\xrightarrow{\text{hit ratio}}$ better

worse ← $\xrightarrow{\text{byte hit ratio}}$ better

will yield different performances

Recap

- LFD, LRU, FIFO, LFU
- provable results:
 - LFD is optimal
 - LRU and FIFO are k-competitive, LFU is not
- practice: LRU-LFU hybrids are preferred because workloads tend to exhibit locality of reference
- practice: clock approximates LRU with low overhead
- practice: best algorithm will depend on the workload

CACHE REPLACEMENT CONSIDERED LESS IMPORTANT

large caches: decreasing storage cost

reduction of cacheable traffic and rate of change
(changes to objects over time reduce
value of having a large cache to store them longer)

plenty of good enough algorithms

Extra

Locality of reference

An input sequence of references exhibits
locality of reference if
a page that is referenced is likely
to be referenced again in the near future.

Modeling locality of reference

Characteristic vector of an input sequence s

$$C = (c_0, \dots, c_{p-1})$$

p: number of distinct pages referenced

c_i: number of distance-i requests

distance-i request: a request to a page such that the number of distinct pages requested since the last time this page was requested is i

LRU beats FIFO in the presence of locality of reference

Characteristic vector of an input sequence s

$$C = (c_0, \dots, c_{p-1})$$

What does the characteristic vector of a sequence with high locality of reference look like?