


Operating Systems


Memory: replacement policies

Last time

- **caching speeds up data lookups if the data is likely to be re-requested again**
- **data structures for $O(1)$ -lookup**
 - set-associative (hardware)
 - hash tables (software)
- **cache coherence**
 - with 1 cache: buffer writes back to memory
 - with multiple caches: things get complicated
- today: **replacement policy**



Model of virtual memory




disk: slow
memory



input:
request sequence
 $p_0, p_1, p_2, p_3, \dots, p_n$


fast



cache
of size **k**

k much smaller than
of pages in disk

Model of virtual memory




goal

minimize total cost of
bringing pages from disk


slow

input:
request sequence
 $p_0, p_1, p_2, p_3, \dots, p_n$

fast



Paging problem



input

$s = p_0, p_1, p_2, p_3, \dots, p_n$

output

for each p_i , which p to
evict from cache


goal

minimize # misses

slow

input:
request sequence
 $p_0, p_1, p_2, p_3, \dots, p_n$

fast



LFD is optimal

- proposed by Belady
- LFD: longest forward distance
- synonyms: OPT, MIN, clairvoyant, Belady's optimal alg
- proof given in class

LFD proof of optimality

- let ALG be any paging algorithm, s input seq. Will construct a seq of algorithms: $ALG = ALG_0, ALG_1, ALG_2, \dots, ALG_n = LFD$ st
- $\text{cost}(ALG_0, s) \geq \text{cost}(ALG_1, s) \geq \text{cost}(ALG_2, s) \geq \dots \text{cost}(ALG_n) = LFD, n = \text{len}(s)$
- how? by induction: suppose ALG_i is determined
- construct ALG_{i+1} as follows: for the first $i-1$ requests, do the same eviction decisions as ALG_i . On request i , ALG_i evicts the page currently in its cache that will be request furthest into the future, call this page v . $ALG(i)$ may or may not evict this same page. If it does, let $ALG_{i+1} = ALG(i)$ for the future too, so then done.
- Otherwise $ALG(i)$ evicts page $u \neq v$. So after processing i , both caches are the same, but $ALG(i)$ has v and ALG_{i+1} has u , where u is requested before v .
- On future requests, ALG_{i+1} mirrors the same decisions as $ALG(i)$ except if $ALG(i)$ evicts v , then ALG_{i+1} evicts u . Since u will be requested first, $ALG(i)$ will incur a miss when u is requested, at which point both may or may not be identical (off by 1). If v is requested in the future, then the costs will be identical.

Belady's anomaly

- increasing the cache size may result in more misses
- hw: FIFO suffers from this phenomenon
- hw: LRU doesn't

Competitive analysis

Which replacement policy is better?

$$\text{competitive ratio of ALG} = \max_{\text{sequences } s} \frac{\text{misses}(\text{ALG}, s)}{\text{misses}(\text{OPT}, s)}$$

ALG is said to be **k-competitive** if
its competitive ratio is $\leq k$

LRU is k-competitive

for any input sequence s :

$$\text{misses}(\text{LRU}, s) \leq k \text{ misses}(\text{OPT}, s)$$

Proof sketch: divide s into k -phases, where each phase consists of references to exactly k distinct pages
then $\text{misses}(\text{LRU}, s) \leq k$
and $\text{misses}(\text{OPT}, s) \geq 1$

FIFO is k -competitive

for any input sequence s :

$$\text{misses}(\text{FIFO}, s) \leq k \text{ misses}(\text{OPT}, s)$$

Proof: homework

LFU is not competitive

there exists a sequence of input sequences s_1, s_2, s_3, \dots

$$\lim_{i \rightarrow \infty} \frac{\text{misses}(\text{LFU}, s_i)}{\text{misses}(\text{OPT}, s_i)} = \infty$$

Proof: let $s_i = p_0(x_{i+1}), p_1(x_{i+1}), \dots, p_{k-1}(x_{i+1}),$

$[p_k, p_{k+1}](x_i)$

then $\text{misses}(\text{LFU}, s_i) = k-1 + 2^i$

and $\text{misses}(\text{OPT}, s_i) = k$

LIFO is not competitive

there exists a sequence of input sequences s_1, s_2, s_3, \dots

$$\lim_{i \rightarrow \infty} \frac{\text{misses}(\text{LIFO}, s_i)}{\text{misses}(\text{OPT}, s_i)} = \infty$$

Proof: homework

A timeline of paging analysis

- Belady's paper shows that **MIN is optimal**
- **competitive analysis**, Sleator and Tarjan
- 1995: 1st analysis of paging with **locality of reference**, Borodin et al. using access graphs
- 1996: strongly competitiveness, Irani et al.
- 1999: LRU is better than FIFO
- 2000: markov paging, Karlin et al.
- 2007: relative worst order, Boyar et al
- 2009: relative dominance, Dorrigiv et al
- 2010: parameterized analysis, Dorrigiv et al
- 2012: access graph model under worst order analysis, Boyar et al
- 2013: access graph model under relative interval analysis, Boyar et al.
- 2015: competitiveness in terms of locality parameters, Albers et al.

over the years:

- alternative performance measures proposed
- various frameworks for locality of reference proposed

Recap

- LFD, LRU, FIFO, LFU
- provable results:
 - LFD is optimal
 - LRU and FIFO are k -competitive, LFU is not
- practice: LRU-LFU hybrids are preferred because workloads tend to exhibit locality of reference
- practice: clock approximates LRU with low overhead
- practice: best algorithm will depend on the workload

Extra

Locality of reference

An input sequence of references exhibits
locality of reference if
a page that is referenced is likely
to be referenced again in the near future.

Modeling locality of reference

Characteristic vector of an input sequence s

$$C = (c_0, \dots, c_{p-1})$$

p : number of distinct pages referenced

c_i : number of distance- i requests

distance- i request: a request to a page such that the number of distinct pages requested since the last time this page was requested is i

LRU beats FIFO in the presence of locality of reference

Characteristic vector of an input sequence s

$$C = (c_0, \dots, c_{p-1})$$

What does the characteristic vector of a sequence with high locality of reference look like?