

# CACHE OPTIMIZATION FOR THE MODERN WEB


Jenny Lam

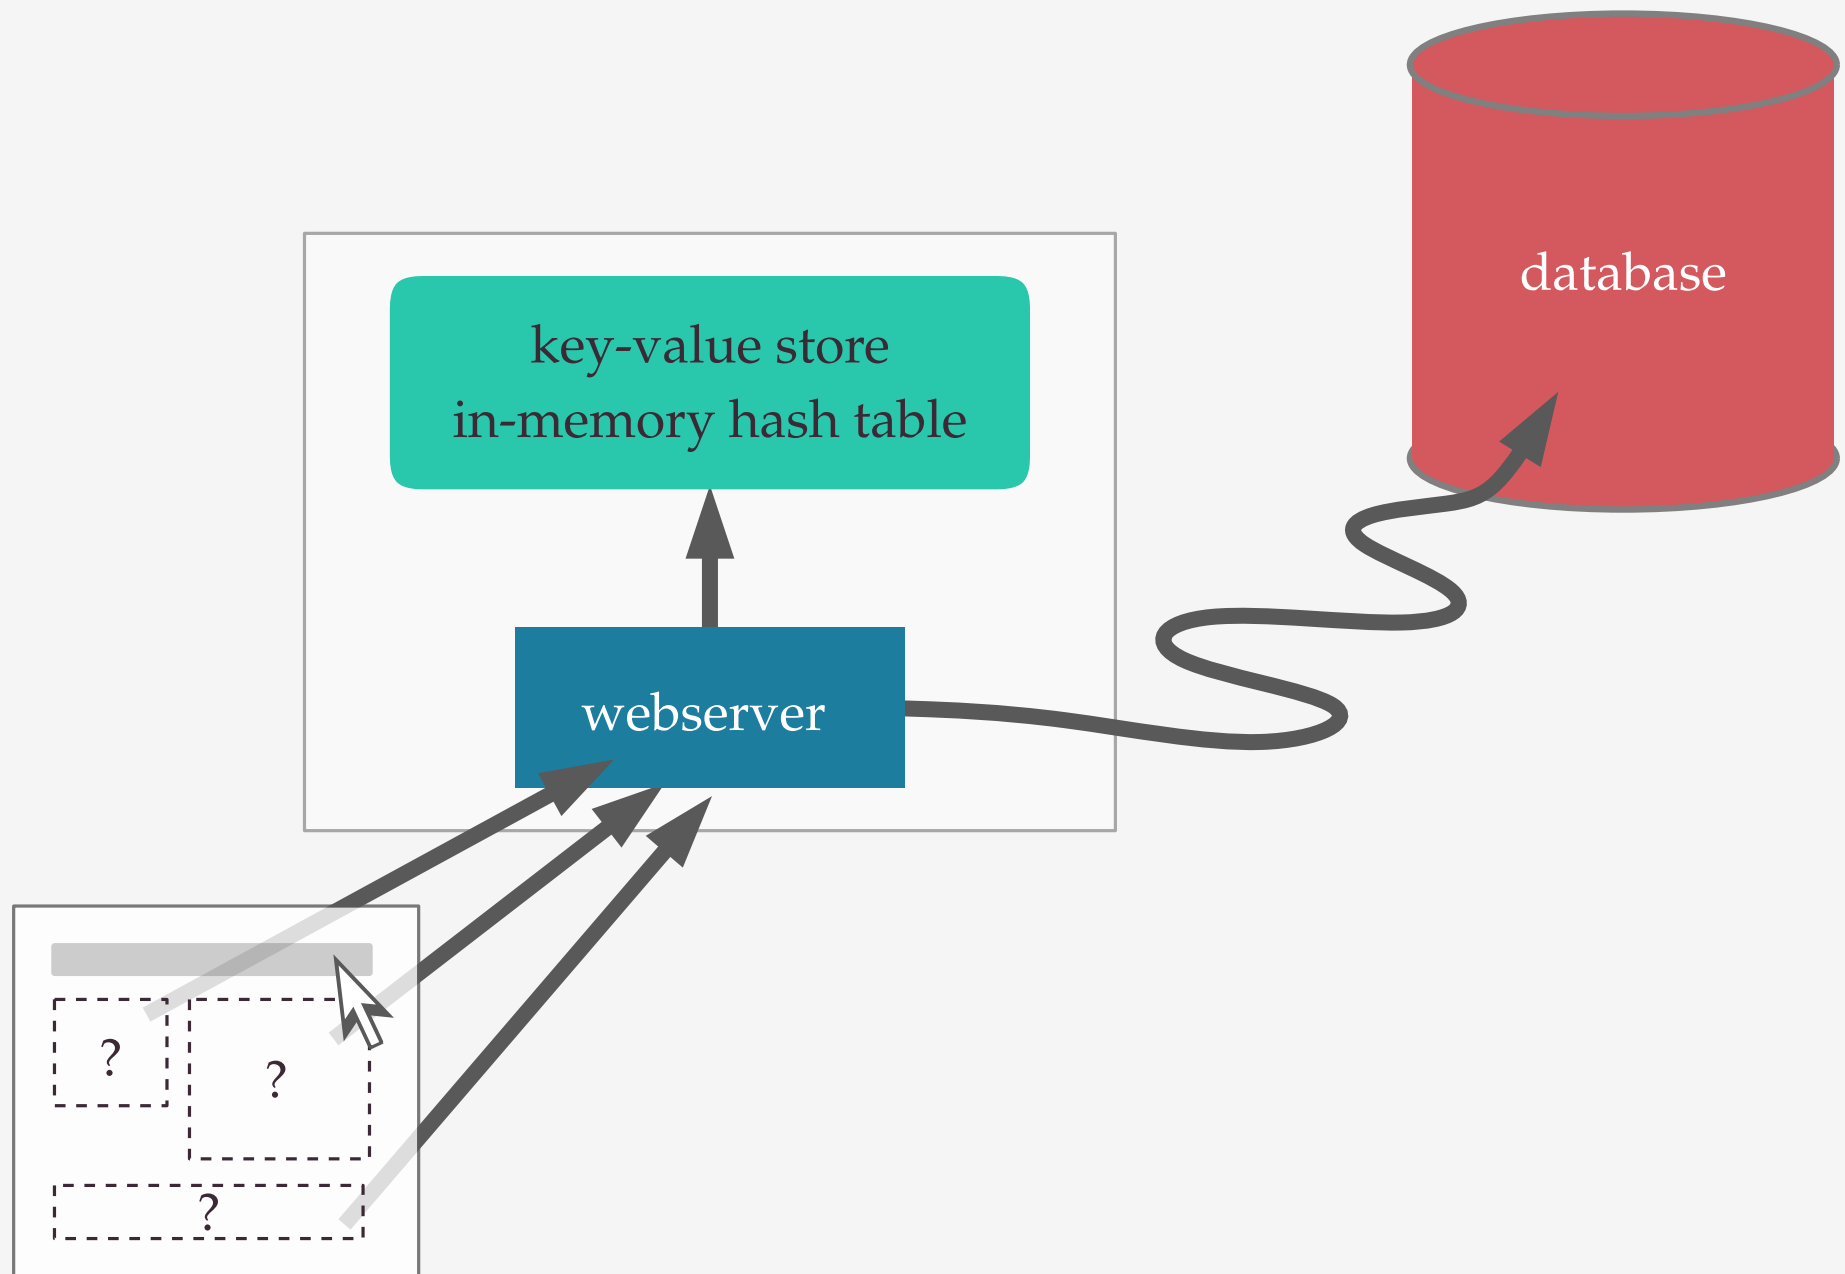
JOINT WORK WITH



Shahram Ghandeharizadeh



Sandy Irani

Jason Yap












PAGING

minimize  
number of cache misses



GENERALIZED  
CACHING


minimize  
**total cost** of cache misses



PAGING

minimize  
number of cache misses

Least Recently Used (LRU)



GENERALIZED  
CACHING

minimize  
**total cost** of cache misses

GreedyDual-Size (GDS)

## EVICTIION POLICY


GDS → CAMP

## PLACEMENT POLICY

generalized caching → managed memory caching


## MEMORY HIERARCHY

2-level cache → multi-level cache



## EVICTIION POLICY

GDS → CAMP



## PLACEMENT POLICY

generalized caching → managed memory caching

## MEMORY HIERARCHY


2-level cache → multi-level cache


Least Recently Used






Least Recently Used








pooled Least Recently Used



*Scaling Memcache at Facebook, Nishtala et al., NSDI 2013.*

pooled Least Recently Used



*Scaling Memcache at Facebook, Nishtala et al., NSDI 2013.*

need to take **recomputation cost** into consideration




GDS

---

p

GDS priority



priority  
queue




evict min

$$\frac{\text{cost}(p)}{\text{size}(p)} + \text{lowest priority}$$



# CAMP

---





# CAMP

---

$$\text{round} \left( \frac{\text{cost}(p)}{\text{size}(p)} \right)$$








# CAMP

---

$$\text{round} \left( \frac{\text{cost}(p)}{\text{size}(p)} \right)$$





# CAMP

---

$$\text{round} \left( \frac{\text{cost}(p)}{\text{size}(p)} \right)$$





# PERFORMANCE

---

log (#items) per update



$$\text{cost}(\text{GDS}) \leq k \text{ cost}(\text{OPT})$$

log (#queues) per update



$$\text{cost}(\text{CAMP}) \leq (1 + \varepsilon)k \text{ cost}(\text{OPT})$$

approximation  
parameter



20,000 items

costs

sizes

## EXPERIMENTS

key-value  
store

Twemcache

LRU

pooled LRU

CAMP

client




request  
generator

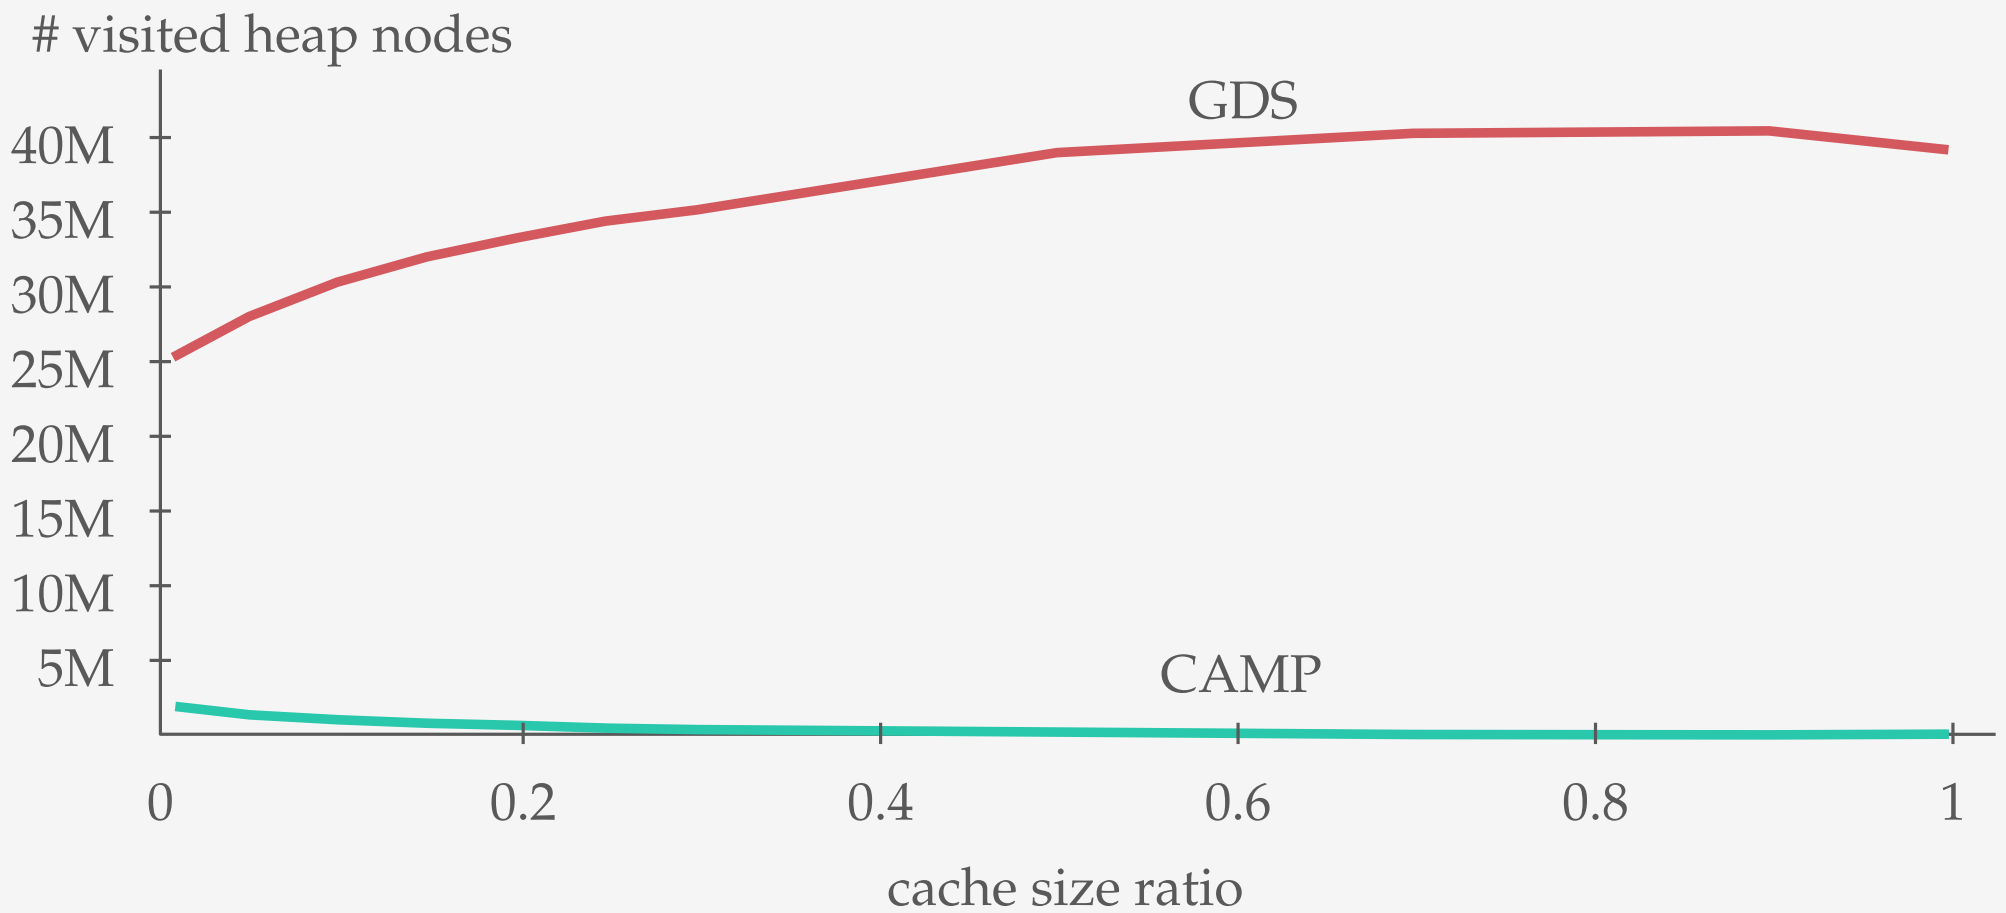
Whalin client for memcached

BG social networking benchmark

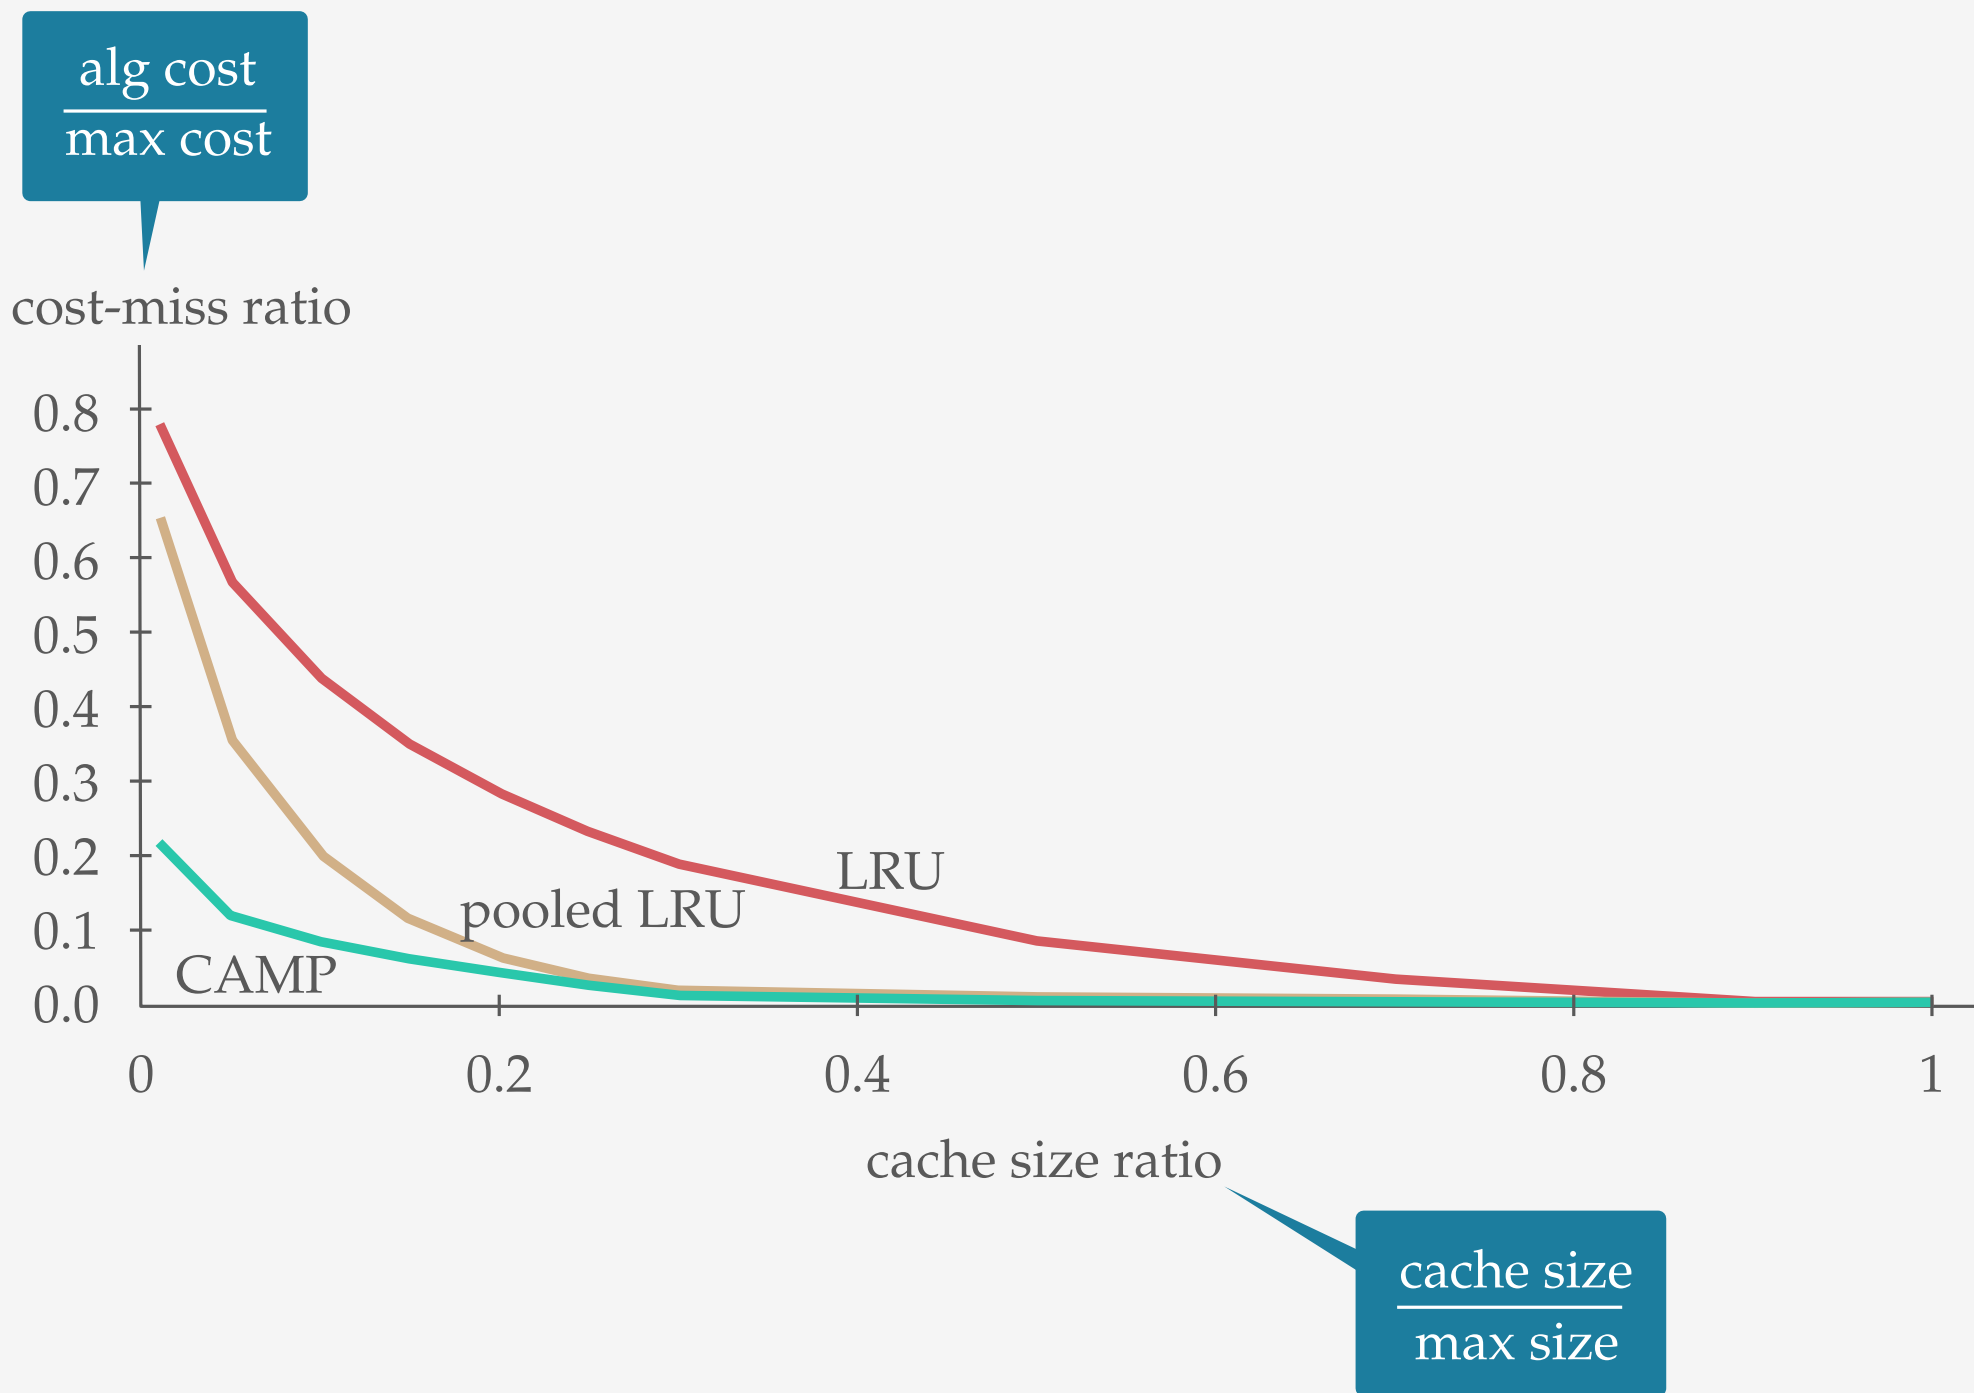
4 million requests

i.i.d. with 70% of requests to 20% of items






$$\frac{\text{cache size}}{\text{max size}}$$



miss rate




$\frac{\text{cache size}}{\text{max size}}$



## EVICTIION POLICY

GDS → CAMP




## PLACEMENT POLICY

generalized caching → managed memory caching

## MEMORY HIERARCHY

2-level cache → multi-level cache

pooled Least Recently Used



*Scaling Memcache at Facebook, Nishtala et al., NSDI 2013.*



cold expensive









cold expensive

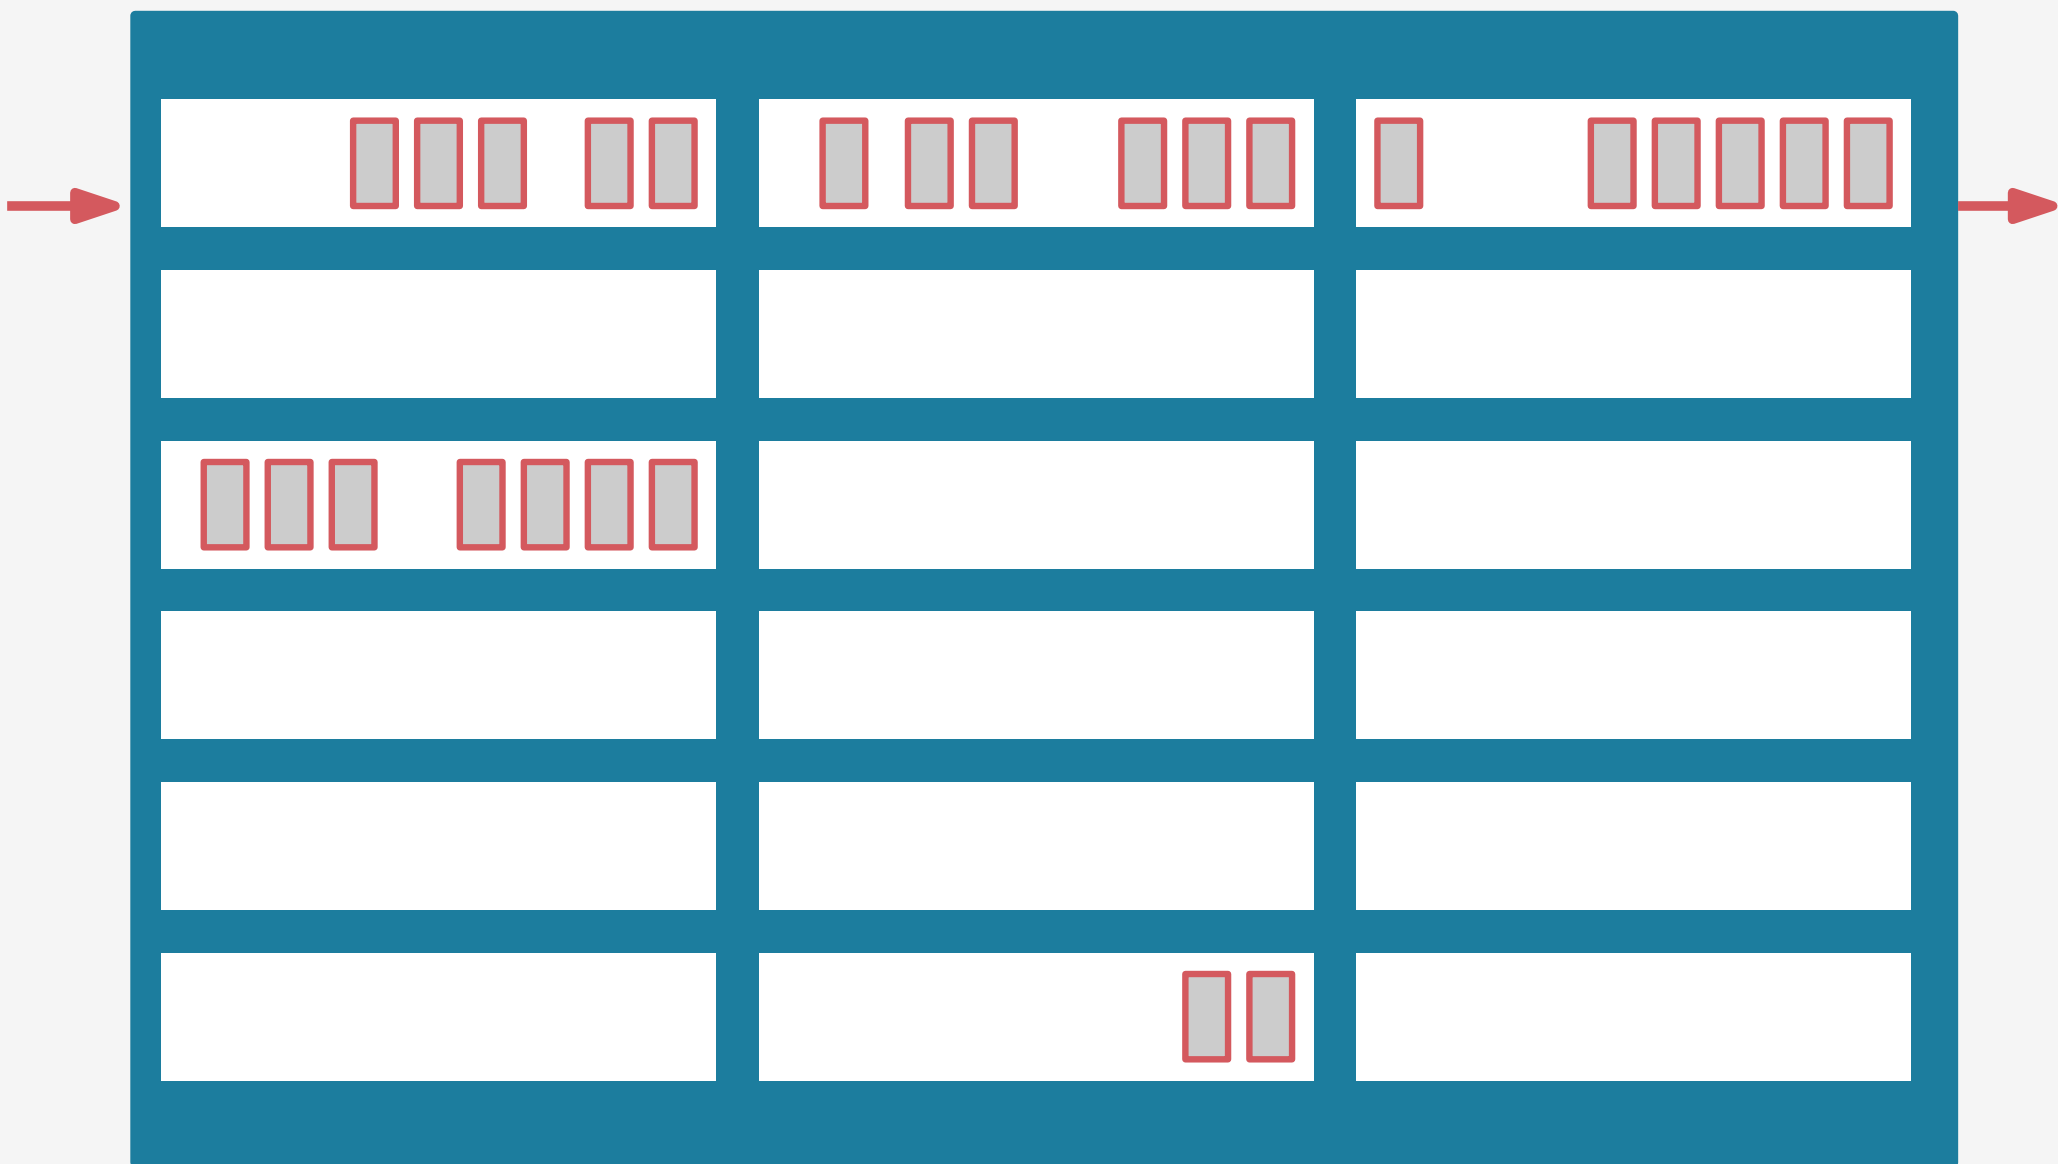



cold expensive

























EVICTIION POLICY



PLACEMENT POLICY

# THE GENERALIZED CACHING PROBLEM

variable size and cost



## GOAL

minimize **total cost** of cache misses

## SUBJECT TO

total size of items in cache  
cannot exceed the cache size



# THE MANAGED MEMORY CACHING PROBLEM

variable size and cost



every item must fit in a contiguous  
segment of memory

CACHE REPLACEMENT  
MEMORY ALLOCATION



# CAMP-MALLOC

---



LRU queues



# CAMP-MALLOC

---




FIFO queues



# CAMP-MALLOC

---




FIFO queue



# CAMP-MALLOC

---




FIFO queue



# CAMP-MALLOC

---




FIFO queue



# CAMP-MALLOC

---




FIFO queue



# CAMP-MALLOC

---




FIFO queue





# CAMP-MALLOC

---



FIFO queue



# CAMP-MALLOC

---



FIFO queue



# CAMP-MALLOC

---




FIFO queue



# CAMP-MALLOC

---



FIFO queue



# CAMP-MALLOC

---




FIFO queue



# CAMP-MALLOC

---



FIFO queue



# CAMP-MALLOC

---



FIFO queue



# CAMP-MALLOC

---



first in

FIFO queue





# CAMP-MALLOC

---



first in

FIFO queue




# CAMP-MALLOC

---



first in




FIFO queue



# CAMP-MALLOC

---



FIFO queue



# CAMP-MALLOC

---



first in




FIFO queue



# CAMP-MALLOC

---






# CAMP-MALLOC

---



first in



FIFO queue




# CAMP-MALLOC

---



first in




FIFO queue



# CAMP-MALLOC

---








# CAMP-MALLOC

---

fragmentation  $\leq 2$  (max item size)




FIFO queue



# CAMP-MALLOC


---





# CAMP-MALLOC


---





# CAMP-MALLOC


---





# CAMP-MALLOC

---





# CAMP-MALLOC


---





# CAMP-MALLOC


---





# CAMP-MALLOC

---

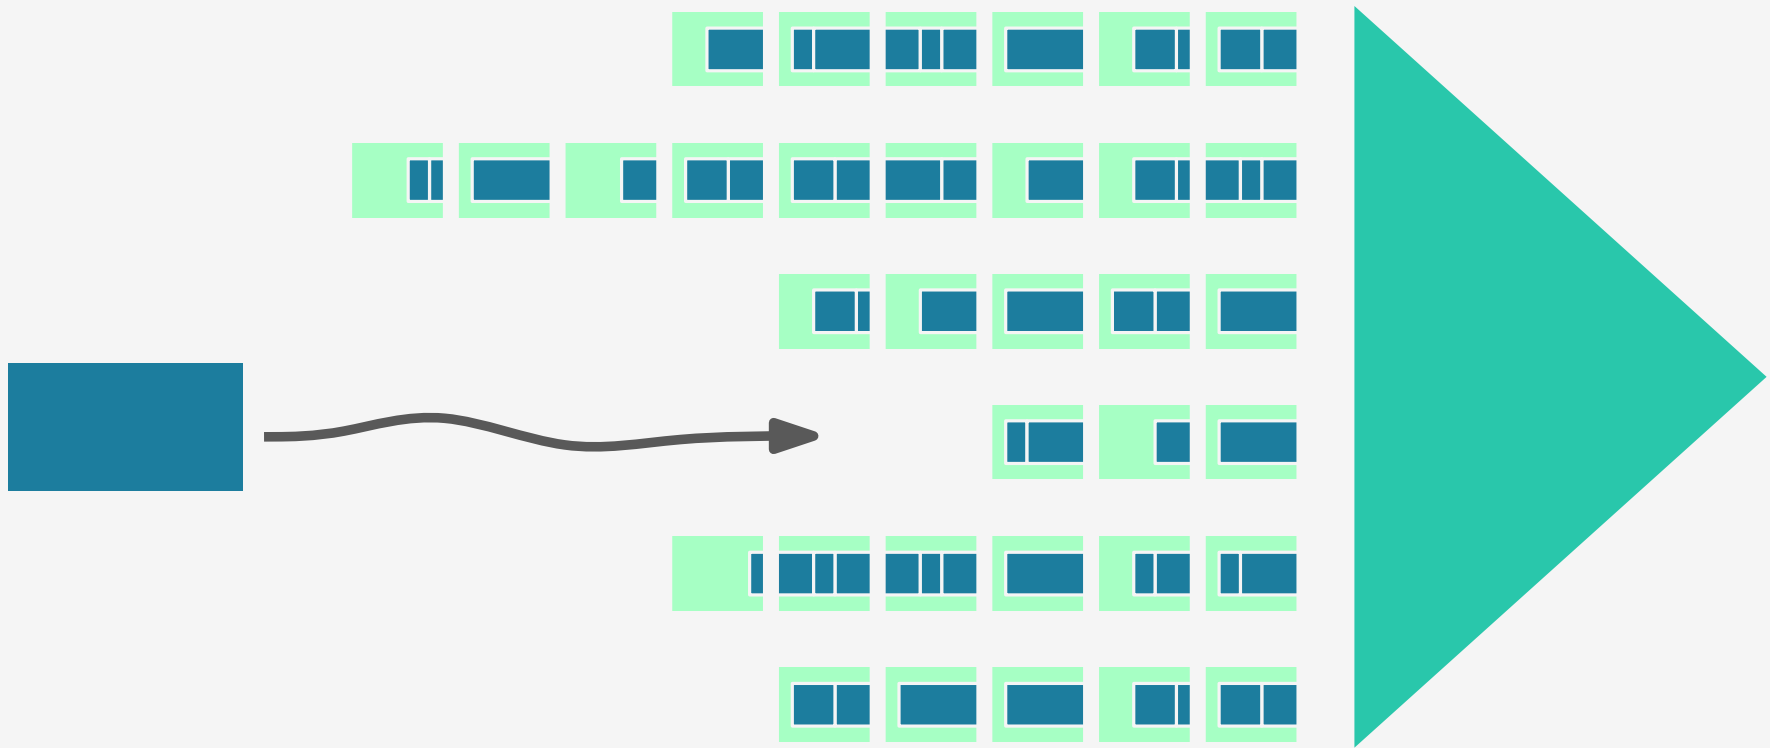






# CAMP-MALLOC

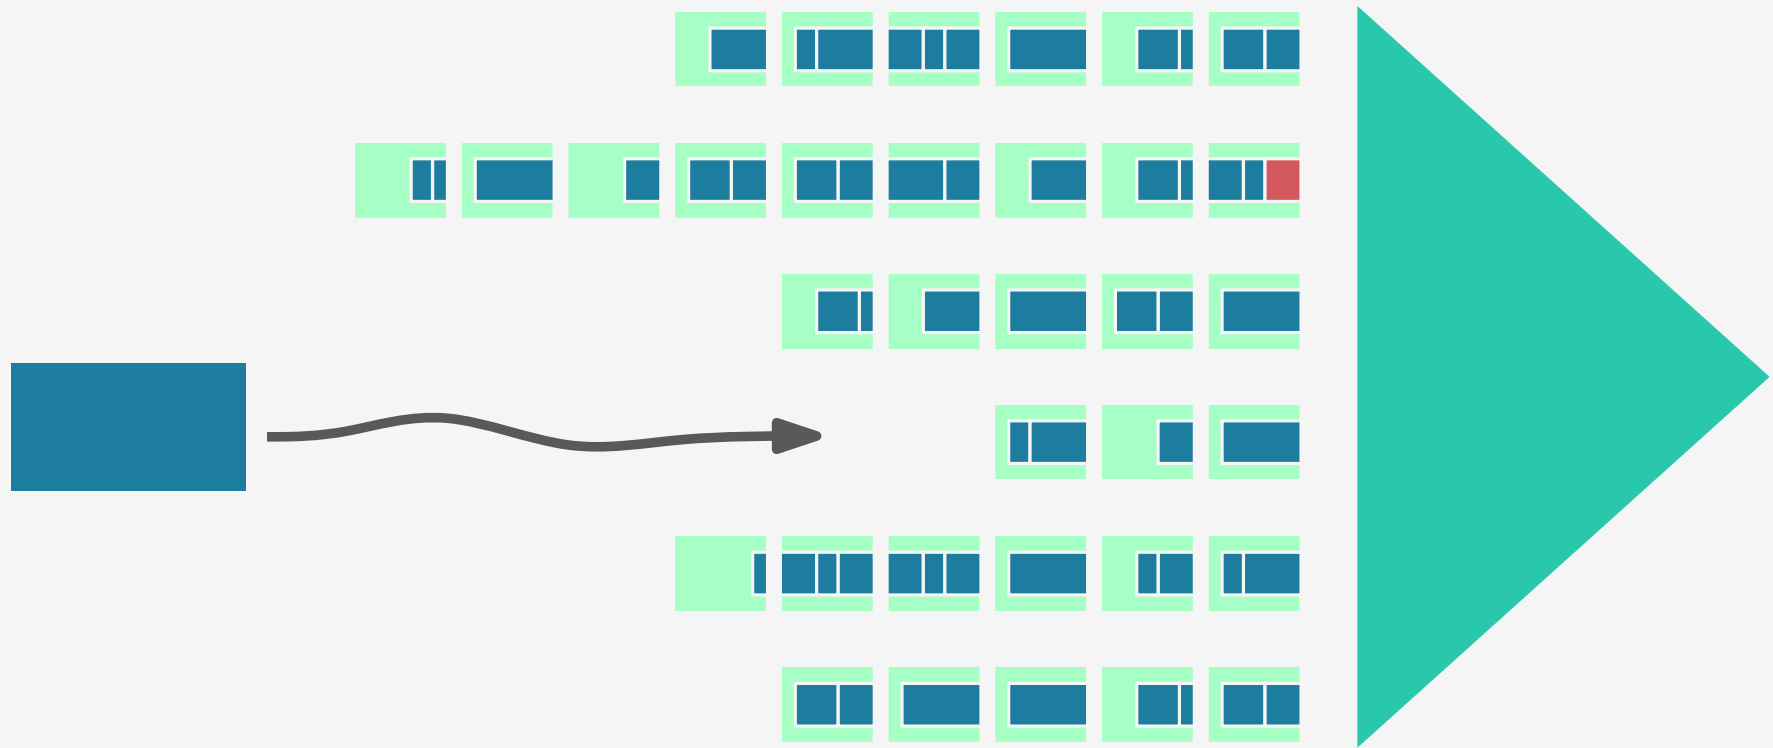
---





# CAMP-MALLOC

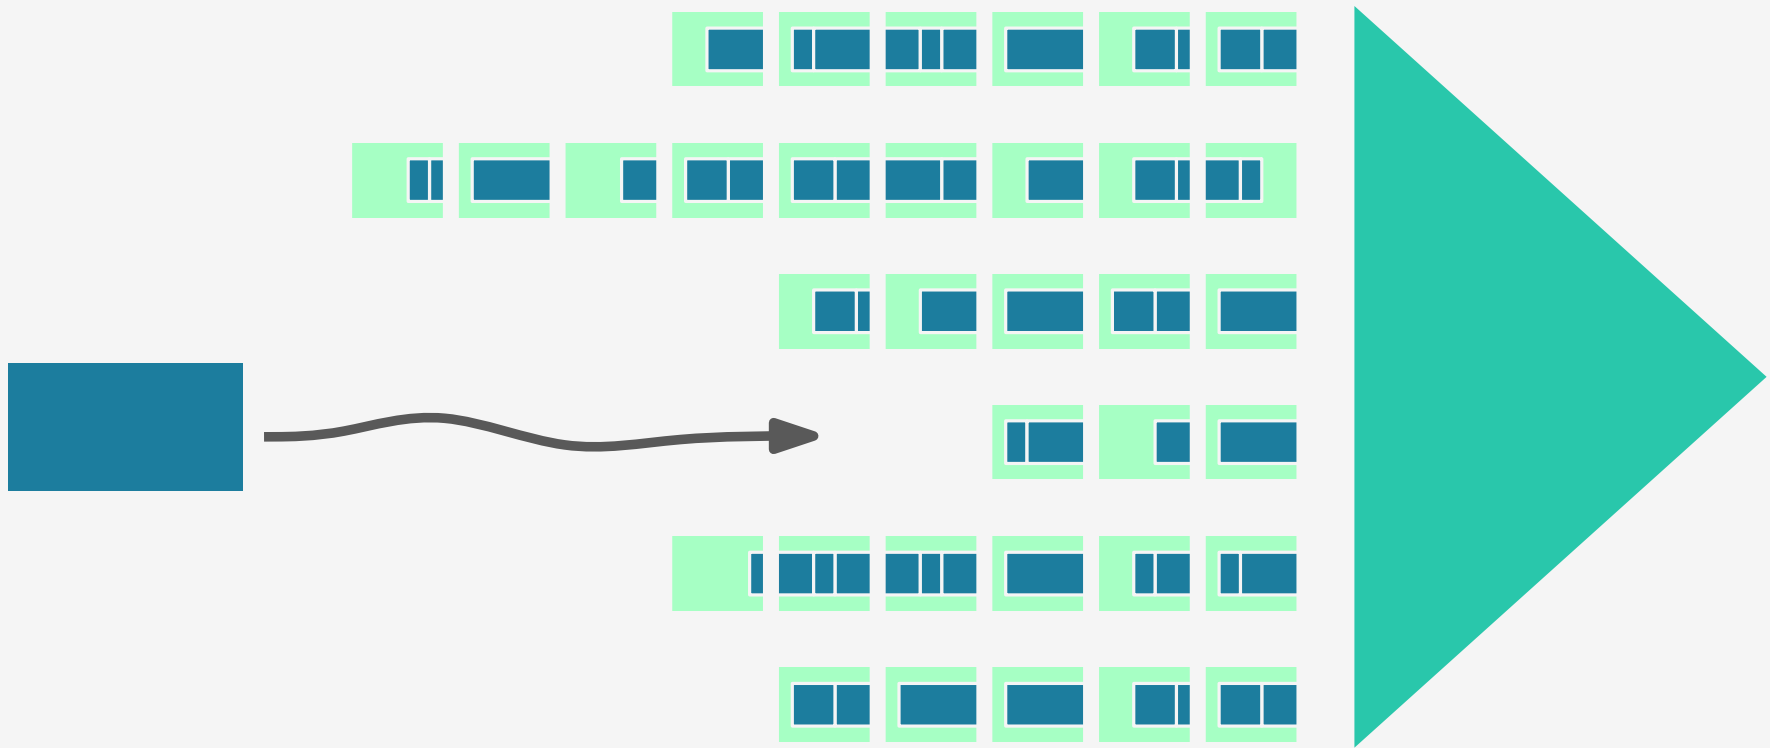
---





# CAMP-MALLOC

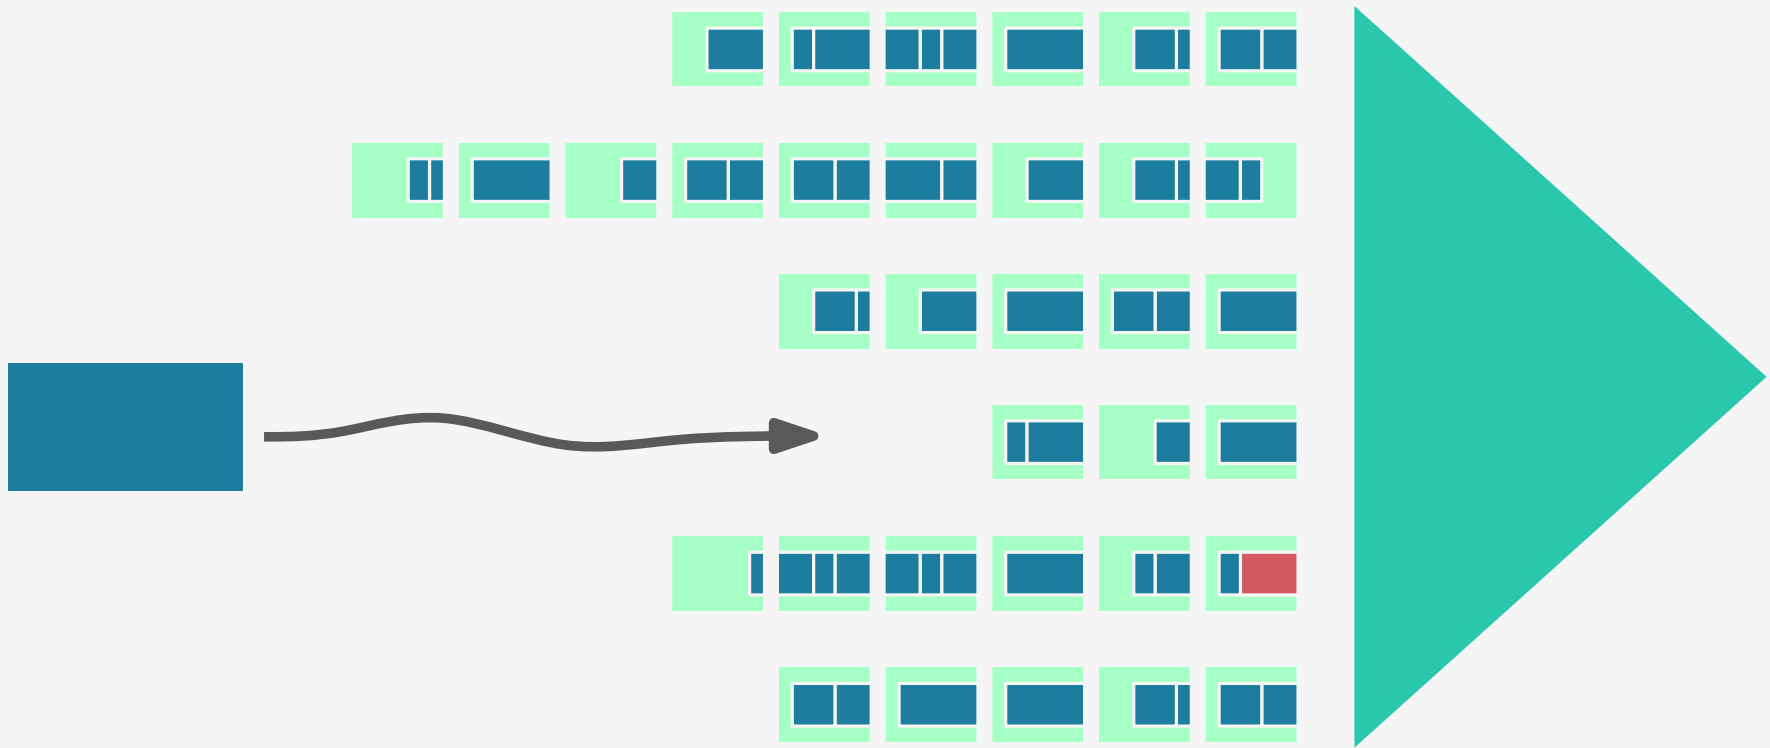
---





# CAMP-MALLOC


---





# CAMP-MALLOC


---





# CAMP-MALLOC

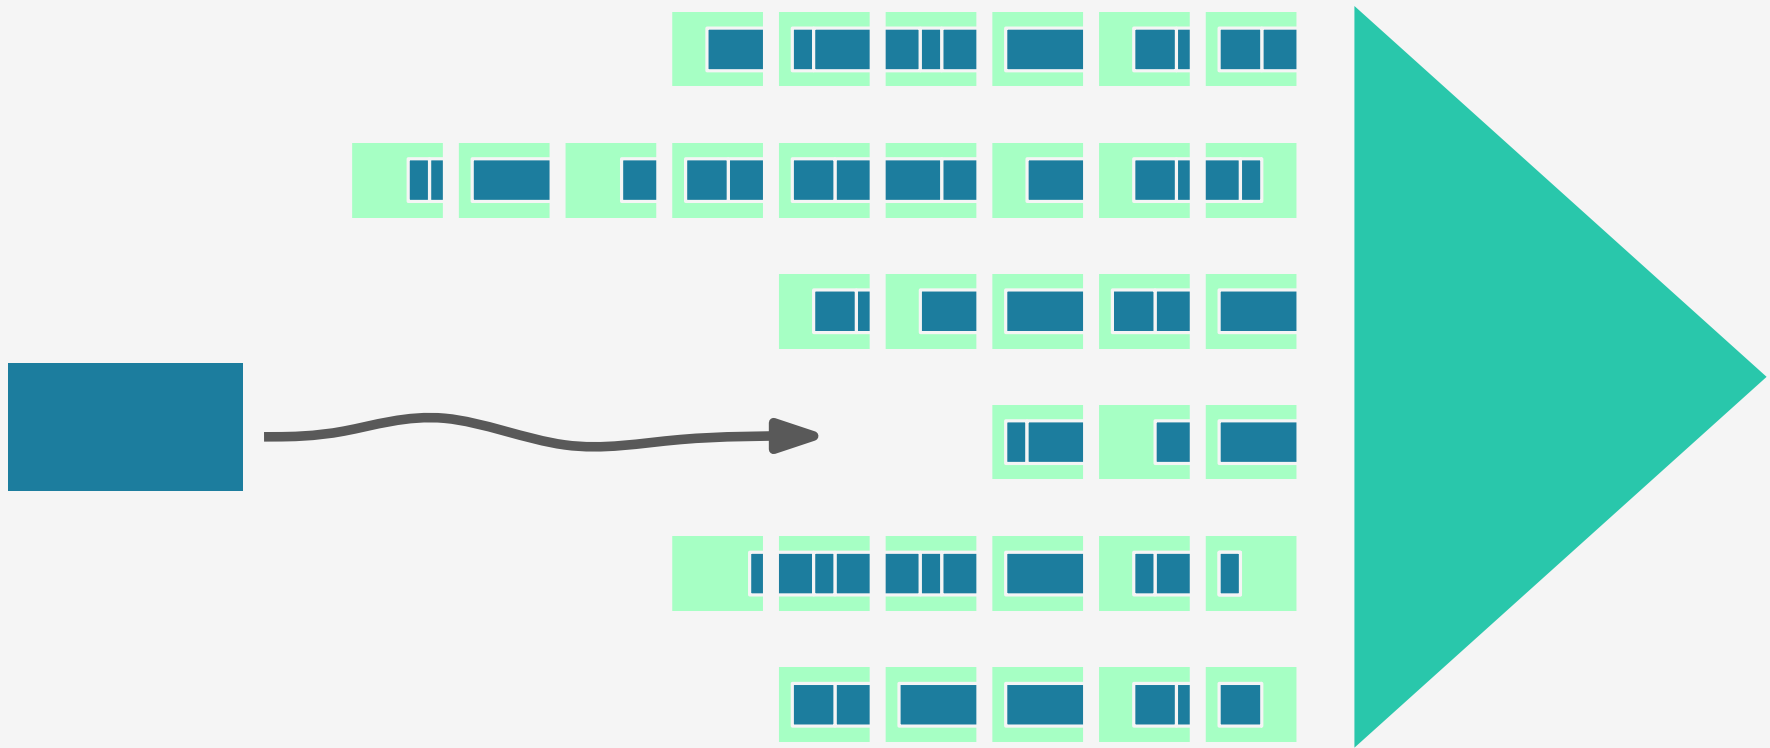
---





# CAMP-MALLOC


---





# CAMP-MALLOC

---

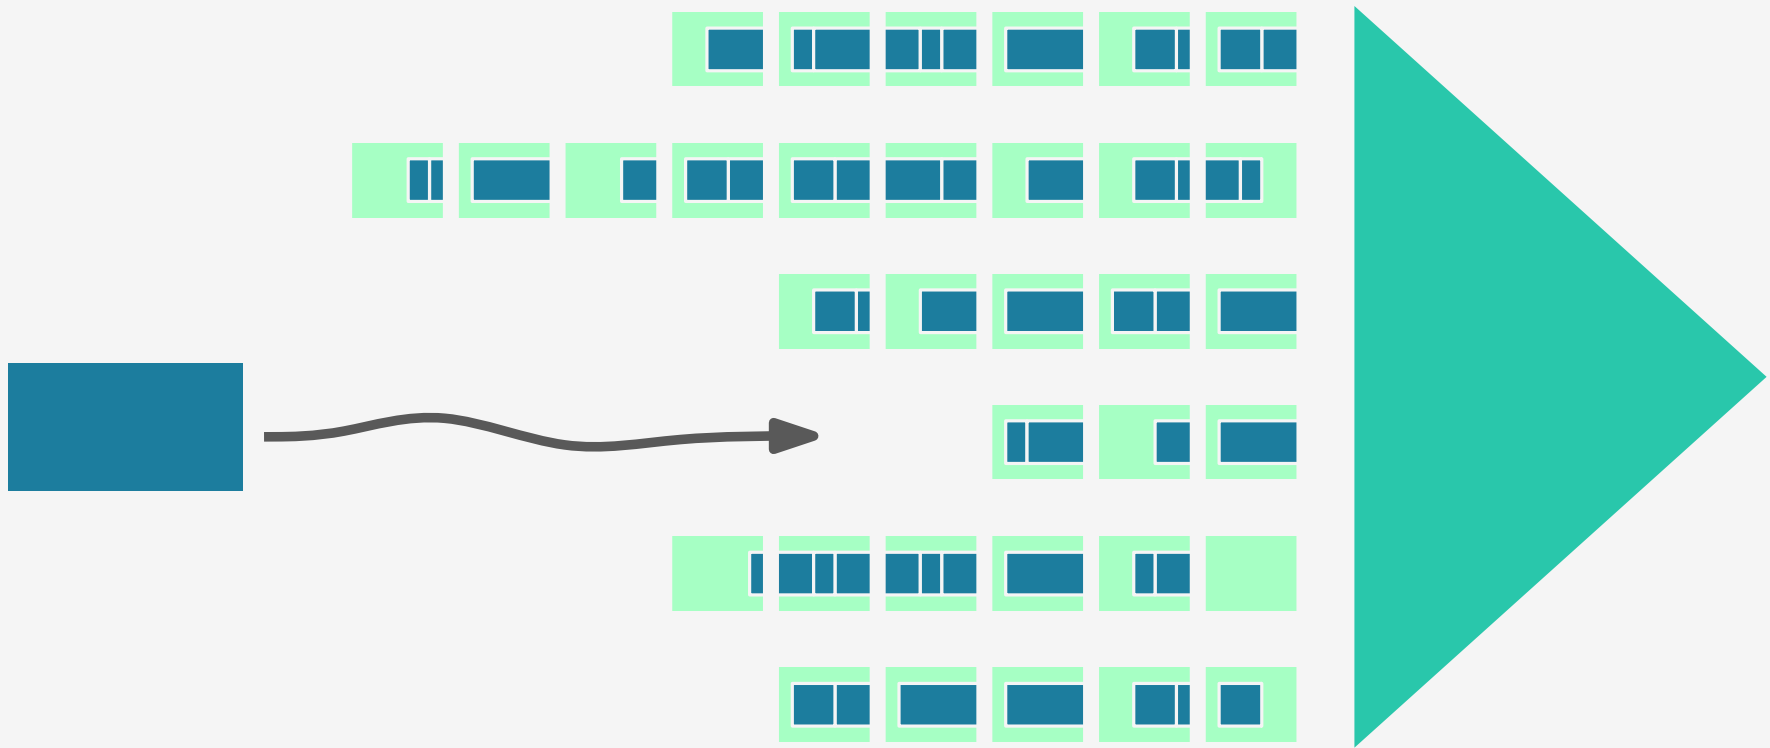






# CAMP-MALLOC


---





# CAMP-MALLOC

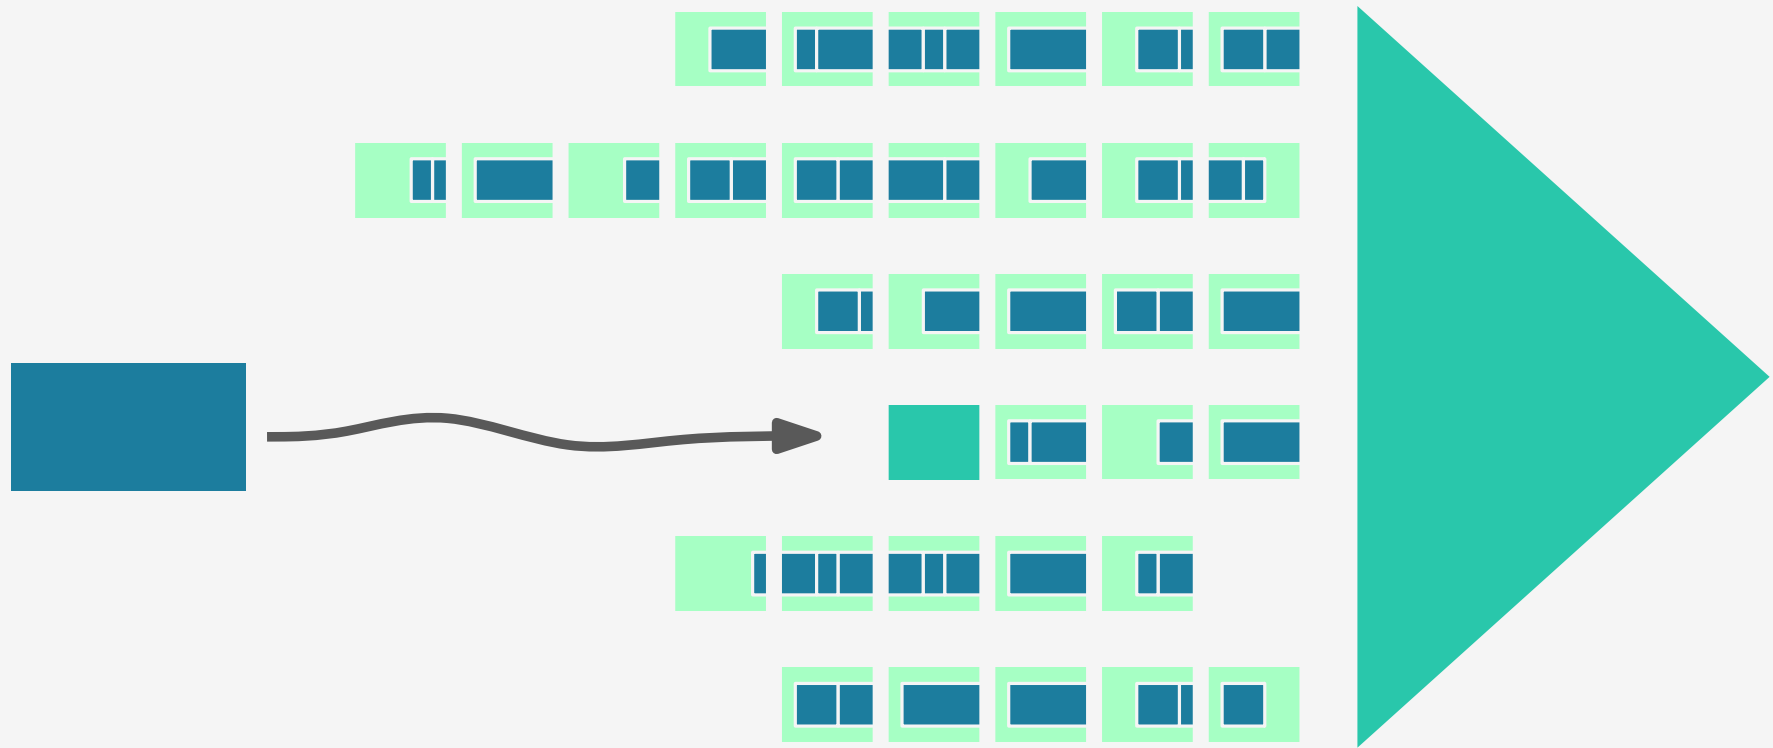
---





# CAMP-MALLOC


---





# CAMP-MALLOC


---





# CAMP-MALLOC


---






# CAMP-MALLOC

---








# CAMP-MALLOC

---



fragmentation  $\leq 2$  (num queues) (block size) + (num blocks) (max item size)





# CAMP-MALLOC

---


is competitive if memory augmented

if  $\text{OPT's cache size} \leq \text{C-M's cache size} - \text{fragmentation bound}$


then  $\text{cost(C-M)} \leq \frac{\text{C-M's cache size}}{\text{min item size}} \text{cost(OPT)}$

fragmentation  $\leq 2$  (num queues) (block size) + (num blocks) (max item size)

cost-miss ratio




percent fragmentation




block size (units of max item size)

percent fragmentation



block size (units of max item size)



## EVICTIION POLICY


GDS → CAMP

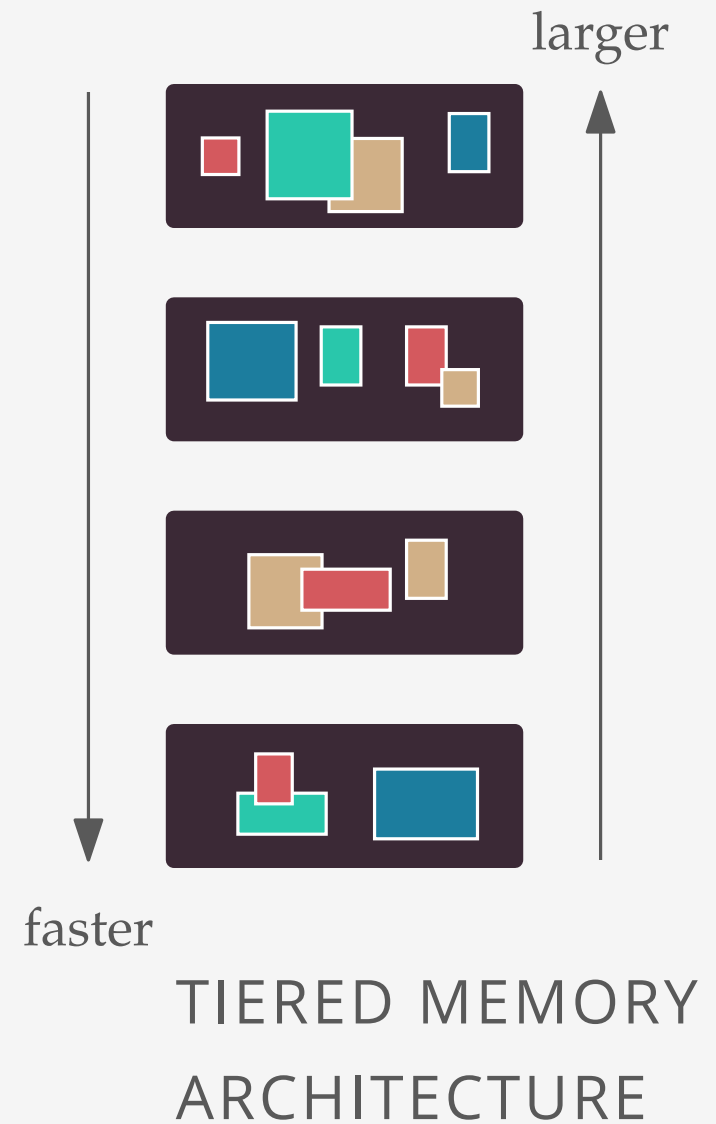

## PLACEMENT POLICY

generalized caching → managed memory caching

## MEMORY HIERARCHY

2-level cache → multi-level cache





GENERALIZED CACHING

capacity

read speed

write speed

failure rate

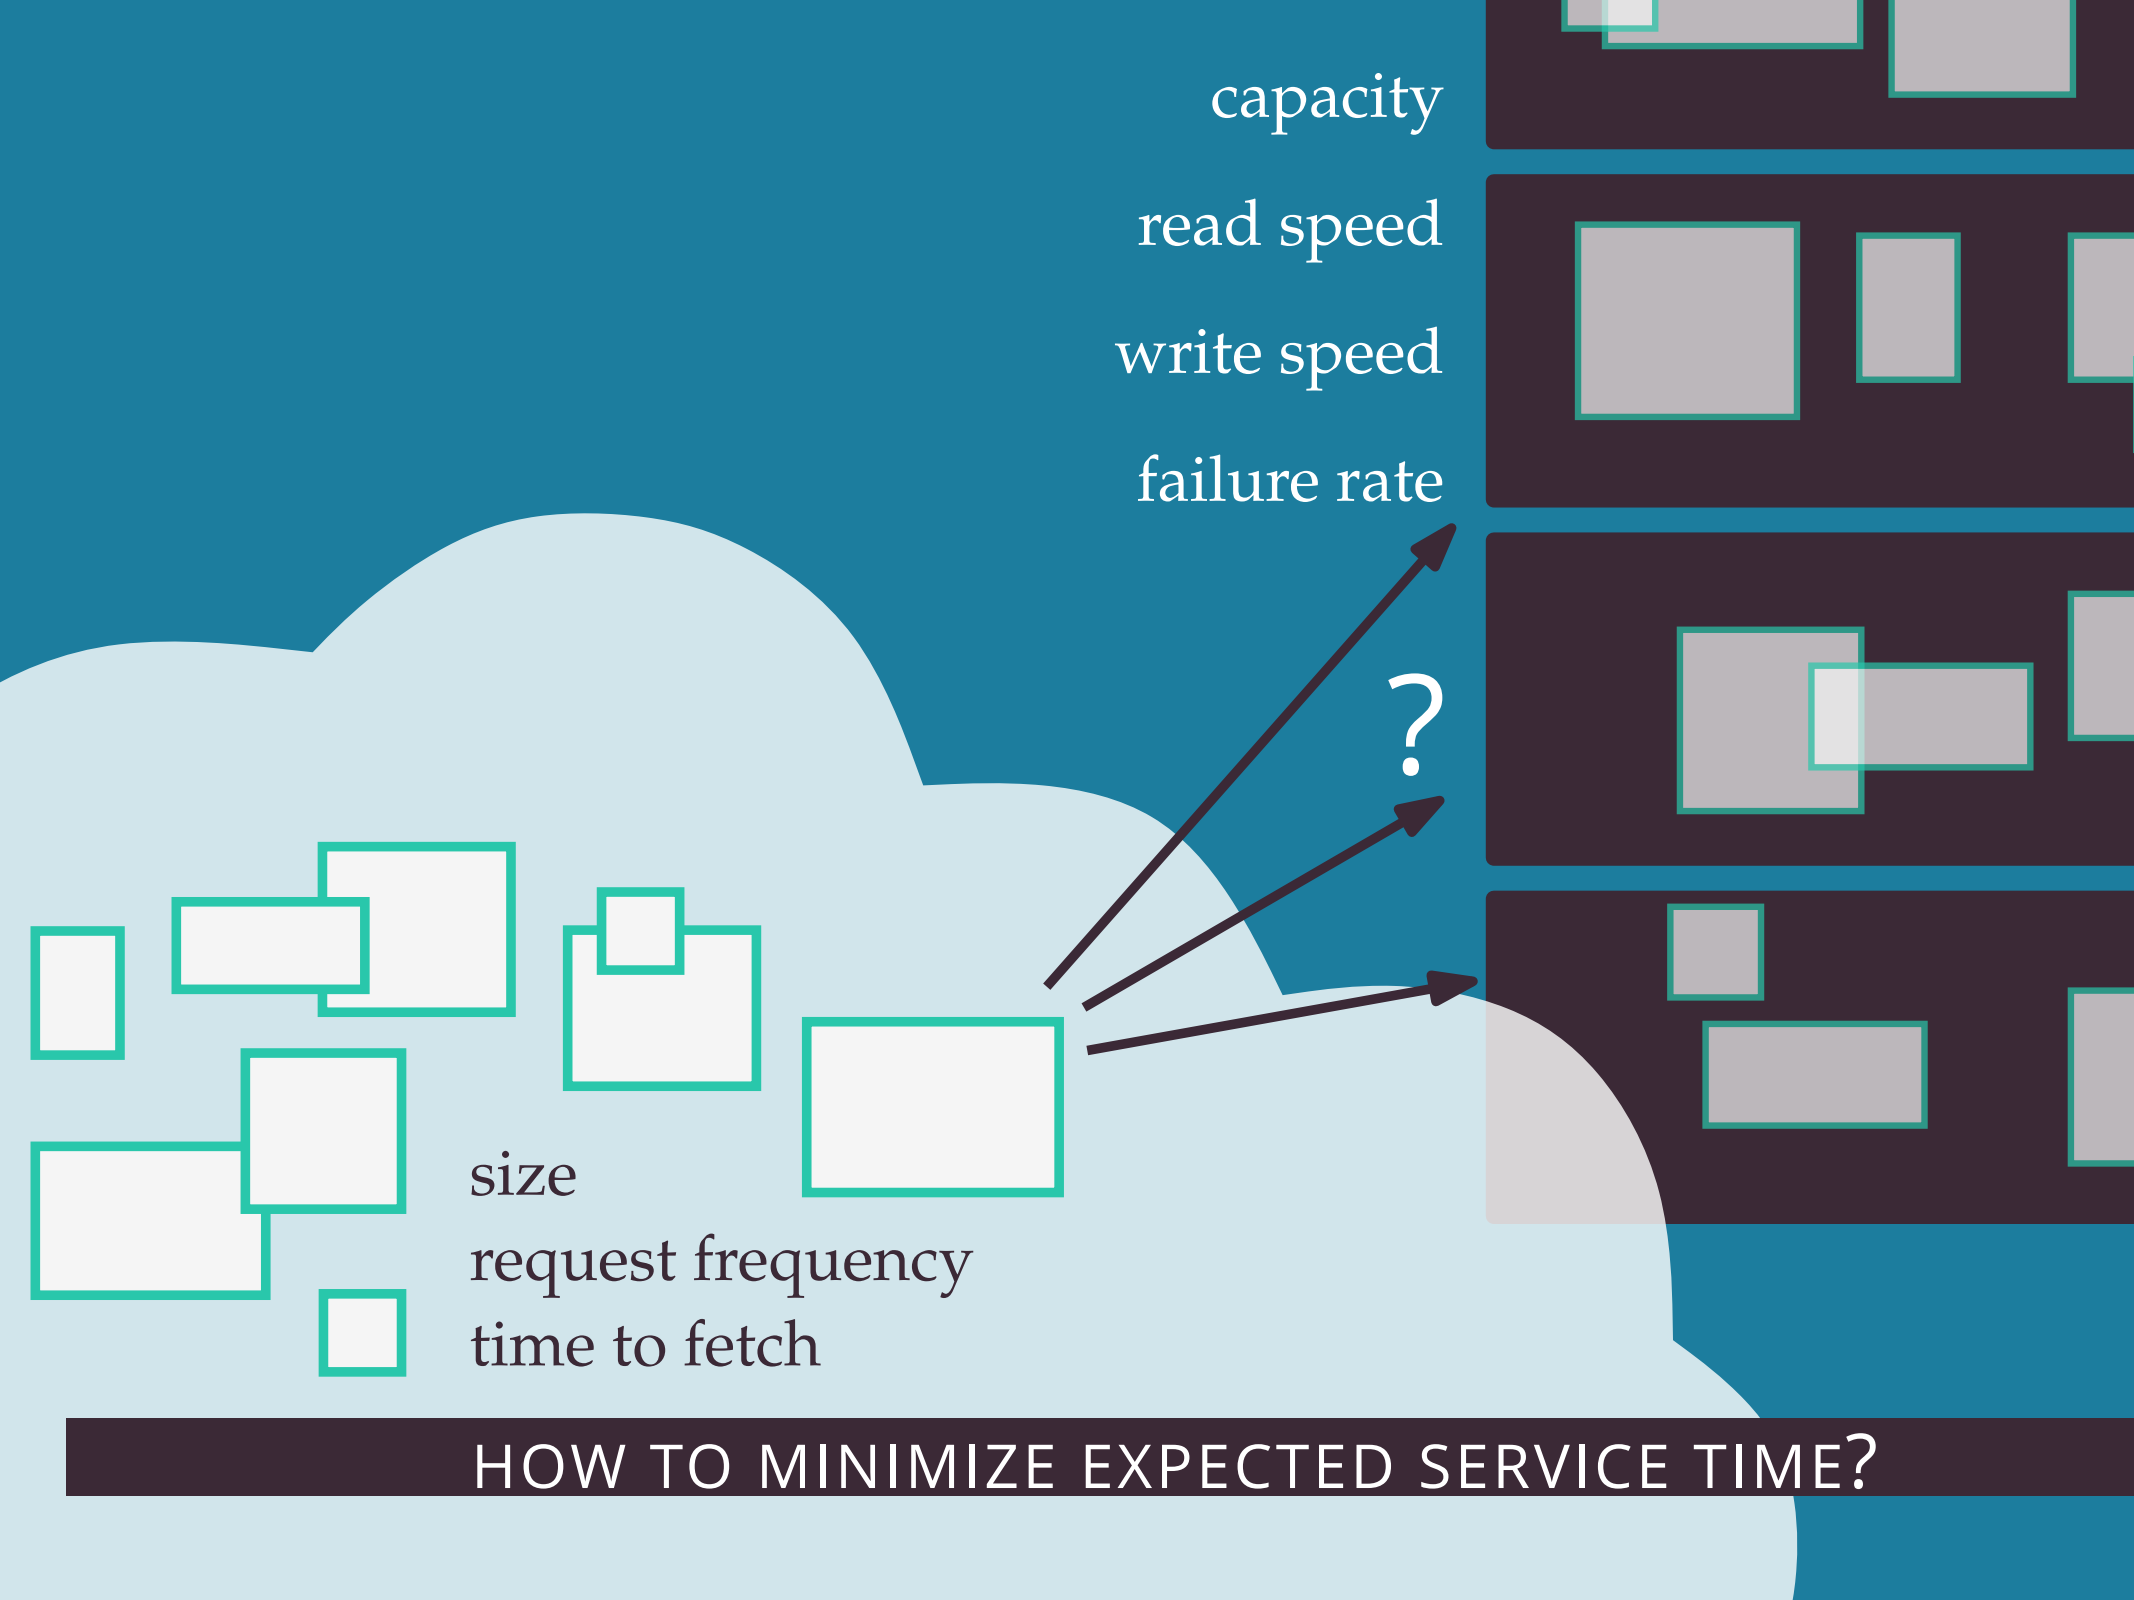
?

size

request frequency


time to fetch

HOW TO MINIMIZE EXPECTED SERVICE TIME?





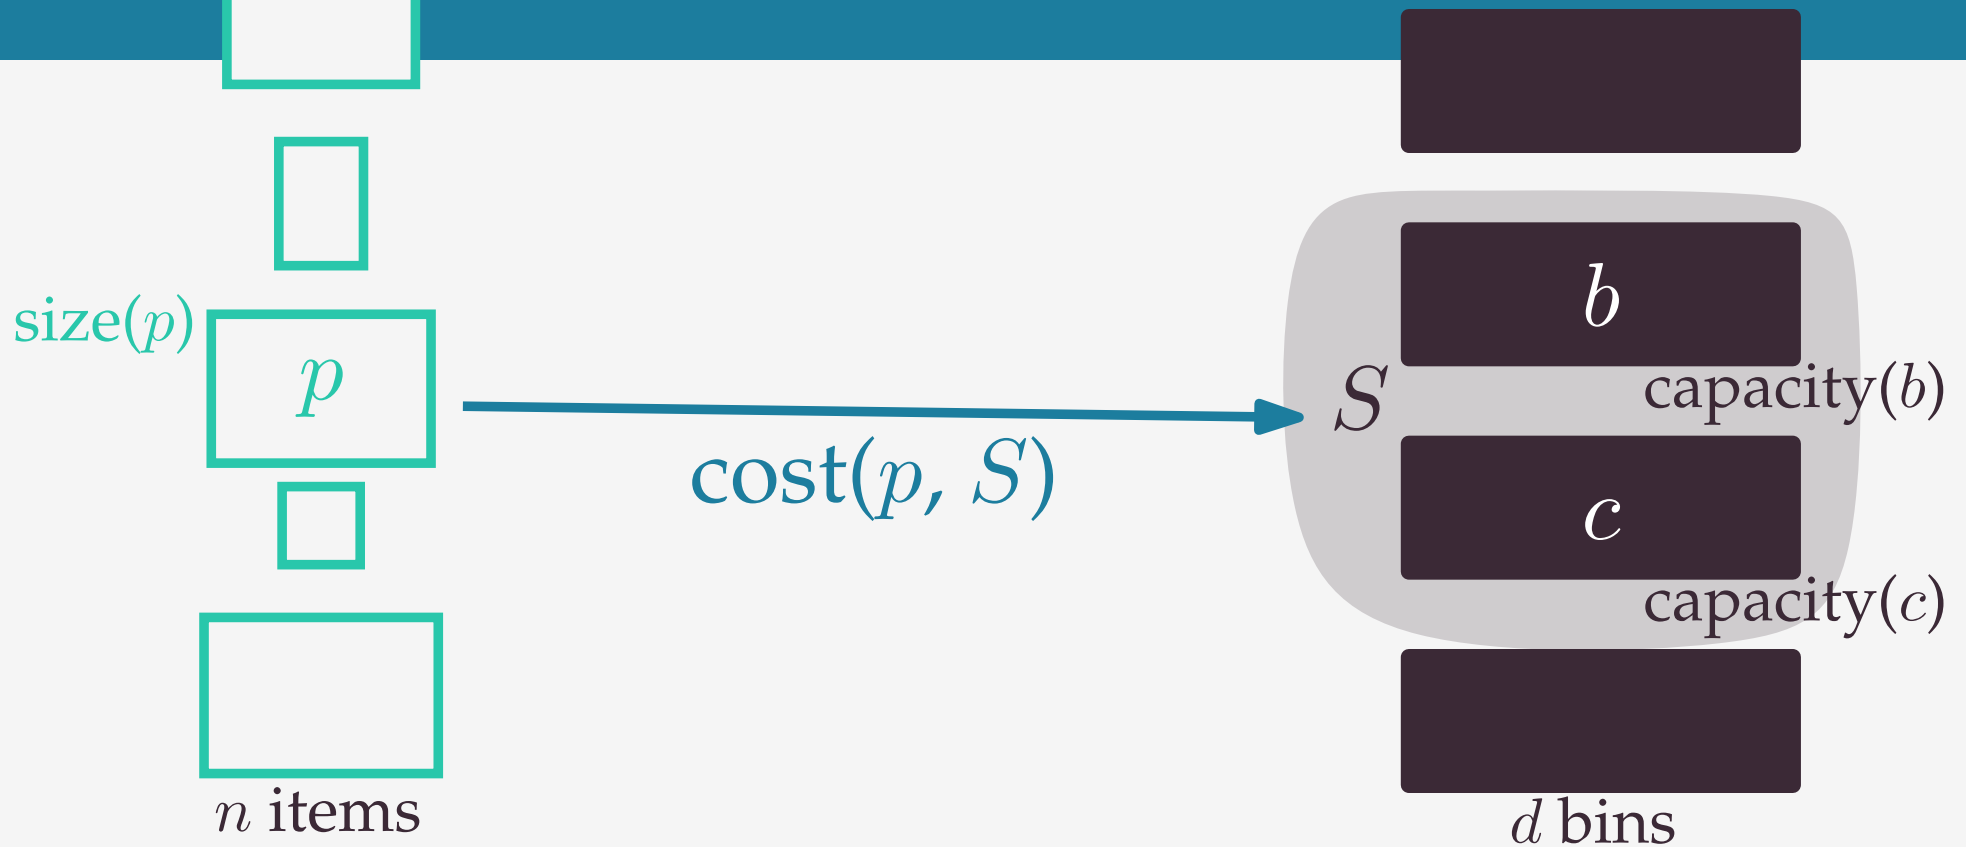
# MULTIPLE KNAPSACK PROBLEM




GOAL

minimize total cost of assignment  
subject to capacity constraints


# SUBSET ASSIGNMENT PROBLEM



GOAL  
minimize total cost of assignment  
subject to capacity constraints



$$\begin{aligned}
 \text{service-time}(p, S) = & \text{read-frequency}(p) \quad \text{read-time}(p, S) \\
 & + \text{write-frequency}(p) \quad \text{write-time}(p, S) \\
 & + \sum_{F \subseteq S} \text{fail-freq}(F) \left( \text{read-time}(p, S \setminus F) \right. \\
 & \quad \left. + \text{write-time}(p, S \cap F) \right)
 \end{aligned}$$



### cache configuration

$$\text{maximize } \sum_{p, S} \text{benefit}(p, S) x(p, S)$$

$$\sum_S x(p, S) = 1$$

$$\sum_{p, S} \text{price}(p, S) x(p, S) \leq \text{budget}$$

$$x = 0, 1$$

### subset assignment

$$\text{minimize } \sum_{p, S} \text{cost}(p, S) x(p, S)$$


$$\sum_S x(p, S) = \text{size}(p)$$

$$\sum_{p, S \ni b} x(p, S) \leq \text{capacity}(b)$$

$$x(p, S) = 0, \text{size}(p)$$



# CACHE CONFIGURATION





# SUBSET ASSIGNMENT

---

HAVE

$$d \ll n$$

sol to LP relaxation has few fractional assignments

GOAL


solve LP relaxation in  $f(d) \text{ poly}(n)$

1. cycle canceling algorithm

2. simplex algorithm




# MIN COST FLOW







# 1. cycle canceling algorithm



feasible flow  
in original network



negative cycle  
in residual network








## “cycle” in subset assignment problem

augmentation



such that

$$\sum_i \alpha_i \overrightarrow{S_i T_i} = \vec{0}$$


cost difference  
(negative)

$$\sum_i \alpha_i (\text{cost}(p_i, T_i) - \text{cost}(p_i, S_i))$$



## 2. simplex algorithm

basic feasible solution



BASIC FEASIBLE  
ASSIGNMENT

$< 2d$  fractional assignments

bound granularity of vars


$$x(p, S) = \frac{k}{\ell}$$

$< d^{d/2}$



# ALGORITHM

---






# ALGORITHM

---

preprocessing  $\sum_i \alpha_i \overrightarrow{S_i T_i} = \vec{0}$



time  $O(\exp(d(d+1)) \text{poly}(d) \ n \log(n) \log(nC) \log(S))$





EVICTIION POLICY

CAMP



PLACEMENT POLICY

CAMP-malloc



MEMORY HIERARCHY

cache configuration  
subset asssignment




EVICTIION POLICY

CAMP




PLACEMENT POLICY

CAMP-malloc



LRU-FIFO hybrid?



MEMORY HIERARCHY

cache configuration  
subset asssignment




EVICTIION POLICY

CAMP




PLACEMENT POLICY

CAMP-malloc



LRU-FIFO hybrid?



MEMORY HIERARCHY

cache configuration  
subset asssignment

continual  
updates?



## EVICTIOIN POLICY

CAMP



customer  
fairness?



## PLACEMENT POLICY

CAMP-malloc

LRU-FIFO hybrid?



## MEMORY HIERARCHY

cache configuration  
subset asssignment

continual  
updates?







CAMP: a cost adaptive multi-queue eviction policy for key-value stores. Shahram Ghandeharizadeh, Sandy Irani, Jenny Lam, and Jason Yap. In Proceedings of the 15th International Middleware Conference, 2014.

A demonstration of KOSAR: an elastic, scalable, highly available SQL middleware. Shahram Ghandeharizadeh, Connor Gorman, Sandy Irani, Shiva Jahangiri, Jenny Lam, Hieu Nguyen, Ryan Tani and Jason Yap, Middleware 2014.



Cache replacement with memory allocation. Shahram Ghandeharizadeh, Sandy Irani, and Jenny Lam. In Proceedings of the 17th Workshop on Algorithm Engineering and Experiments (ALENEX), 2015.



Memory hierarchy design for caching middleware in the age of NVM. Shahram Ghandeharizadeh, Sandy Irani, and Jenny Lam. (in submission)

The subset assignment problem for data placement in caches. Shahram Ghandeharizadeh, Sandy Irani, and Jenny Lam. (in submission)