

CACHE OPTIMIZATION FOR THE MODERN WEB


Jenny Lam

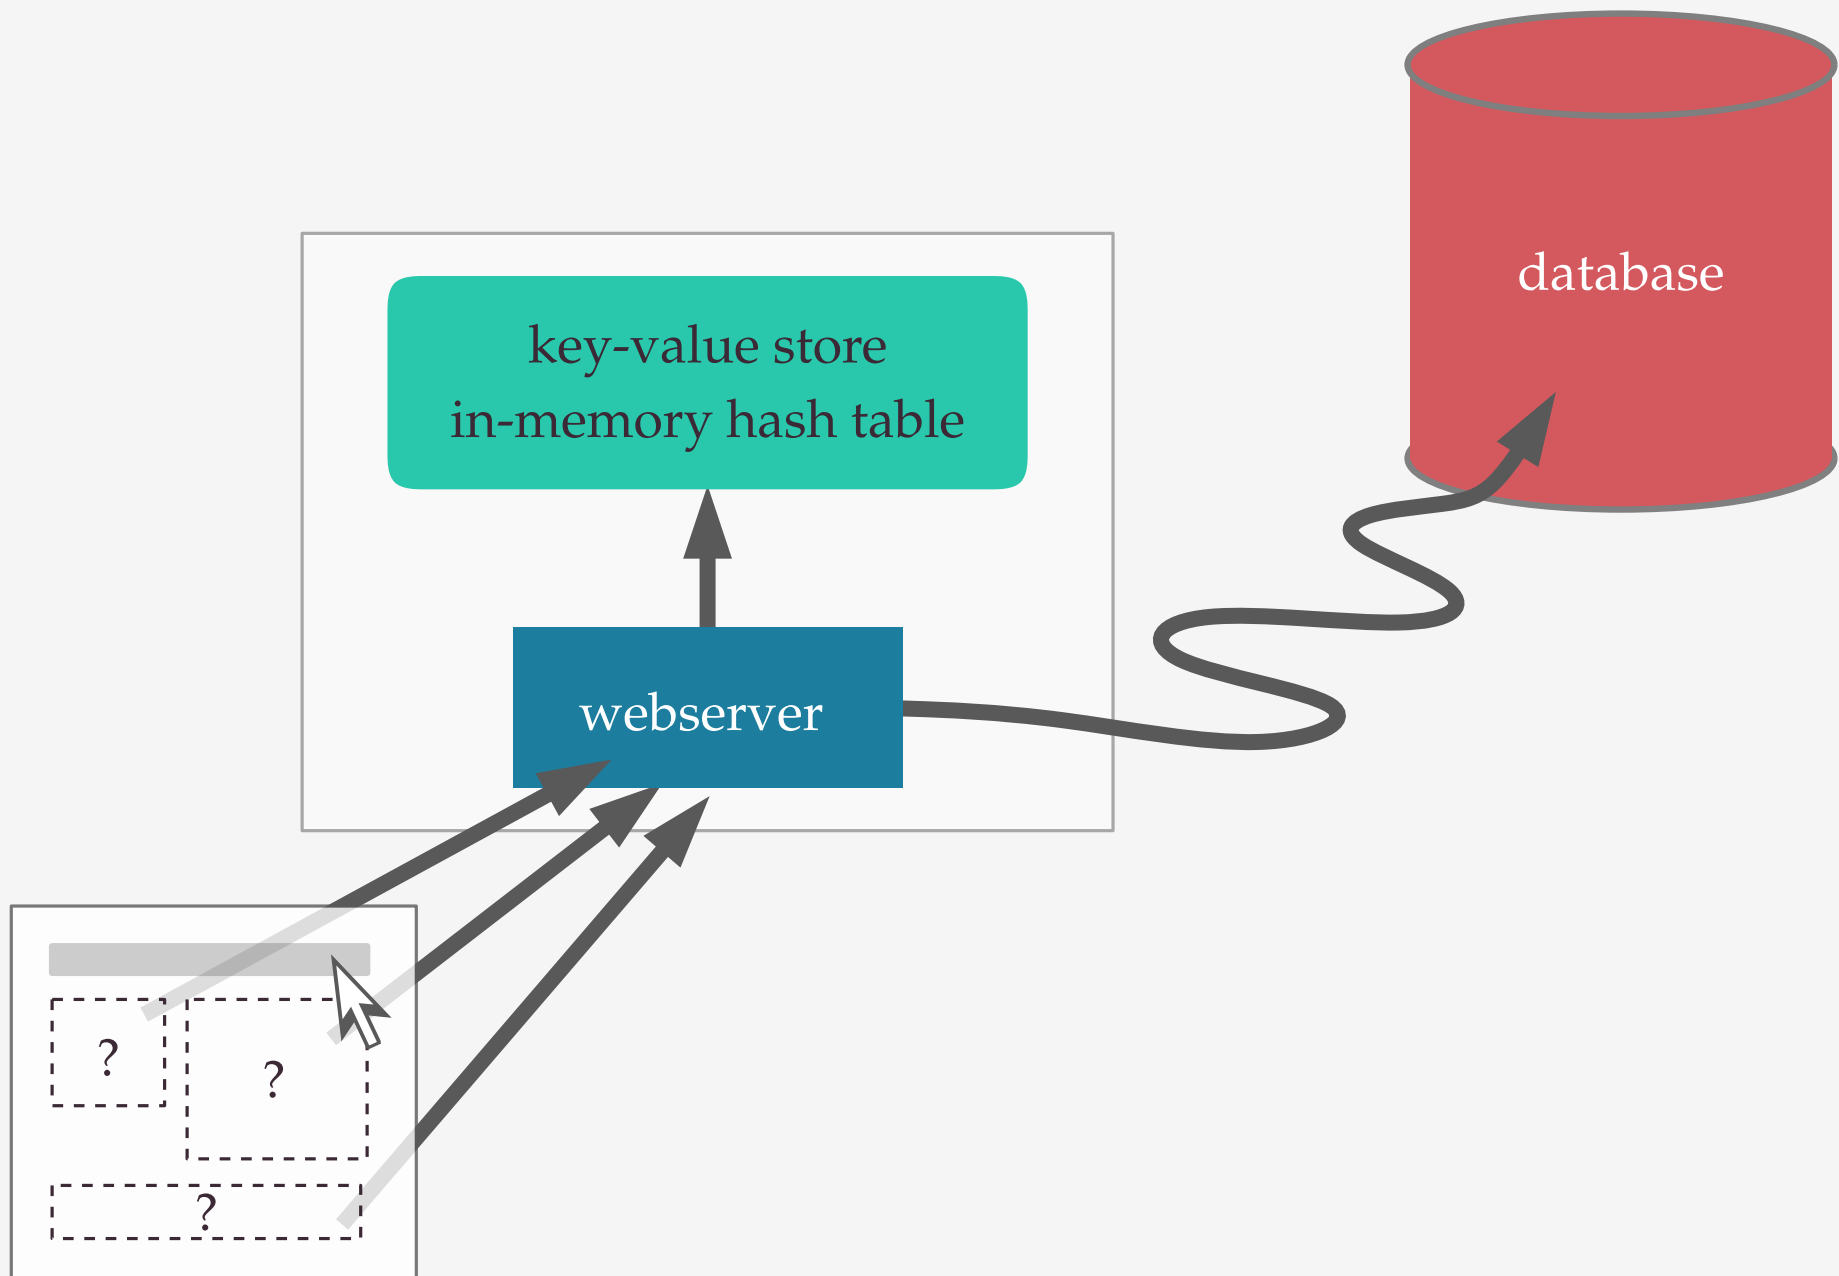
JOINT WORK WITH



Shahram Ghandeharizadeh



Sandy Irani

Jason Yap













PAGING

minimize
number of cache misses



GENERALIZED CACHING


minimize
total cost of cache misses



PAGING

minimize
number of cache misses

Least Recently Used (LRU)



GENERALIZED
CACHING

minimize
total cost of cache misses

GreedyDual-Size (GDS)

EVICTIION POLICY


GDS → CAMP

PLACEMENT POLICY

generalized caching → managed memory caching


MEMORY HIERARCHY

2-level cache → multi-level cache



EVICTIION POLICY

GDS → CAMP



PLACEMENT POLICY


generalized
caching



managed memory
caching

MEMORY HIERARCHY

2-level cache




multi-level cache

Least Recently Used




Least Recently Used








pooled Least Recently Used



Scaling Memcache at Facebook, Nishtala et al., NSDI 2013.

pooled Least Recently Used



Scaling Memcache at Facebook, Nishtala et al., NSDI 2013.


need to take **recomputation cost** into consideration



GDS

p

GDS priority



priority
queue



evict min


$$\frac{\text{cost}(p)}{\text{size}(p)} + \text{lowest priority}$$



CAMP


GDS priority

$$\frac{\text{cost}(p)}{\text{size}(p)} + \text{lowest priority}$$



LRU queues



evict min



CAMP


$$\text{round} \left(\frac{\text{cost}(p)}{\text{size}(p)} \right)$$





CAMP

$$\text{round} \left(\frac{\text{cost}(p)}{\text{size}(p)} \right)$$





CAMP

$$\text{round} \left(\frac{\text{cost}(p)}{\text{size}(p)} \right)$$





PERFORMANCE

log (#items) per update



$$\text{cost}(\text{GDS}) \leq k \text{ cost}(\text{OPT})$$

log (#queues) per update



$$\text{cost}(\text{CAMP}) \leq (1 + \varepsilon)k \text{ cost}(\text{OPT})$$

approximation
parameter



20,000 items

costs

sizes

EXPERIMENTS

key-value
store

Twemcache



LRU



pooled LRU



CAMP

client




request
generator

Whalin client for memcached


BG social networking benchmark

4 million requests

i.i.d. with 70% of requests to 20% of items



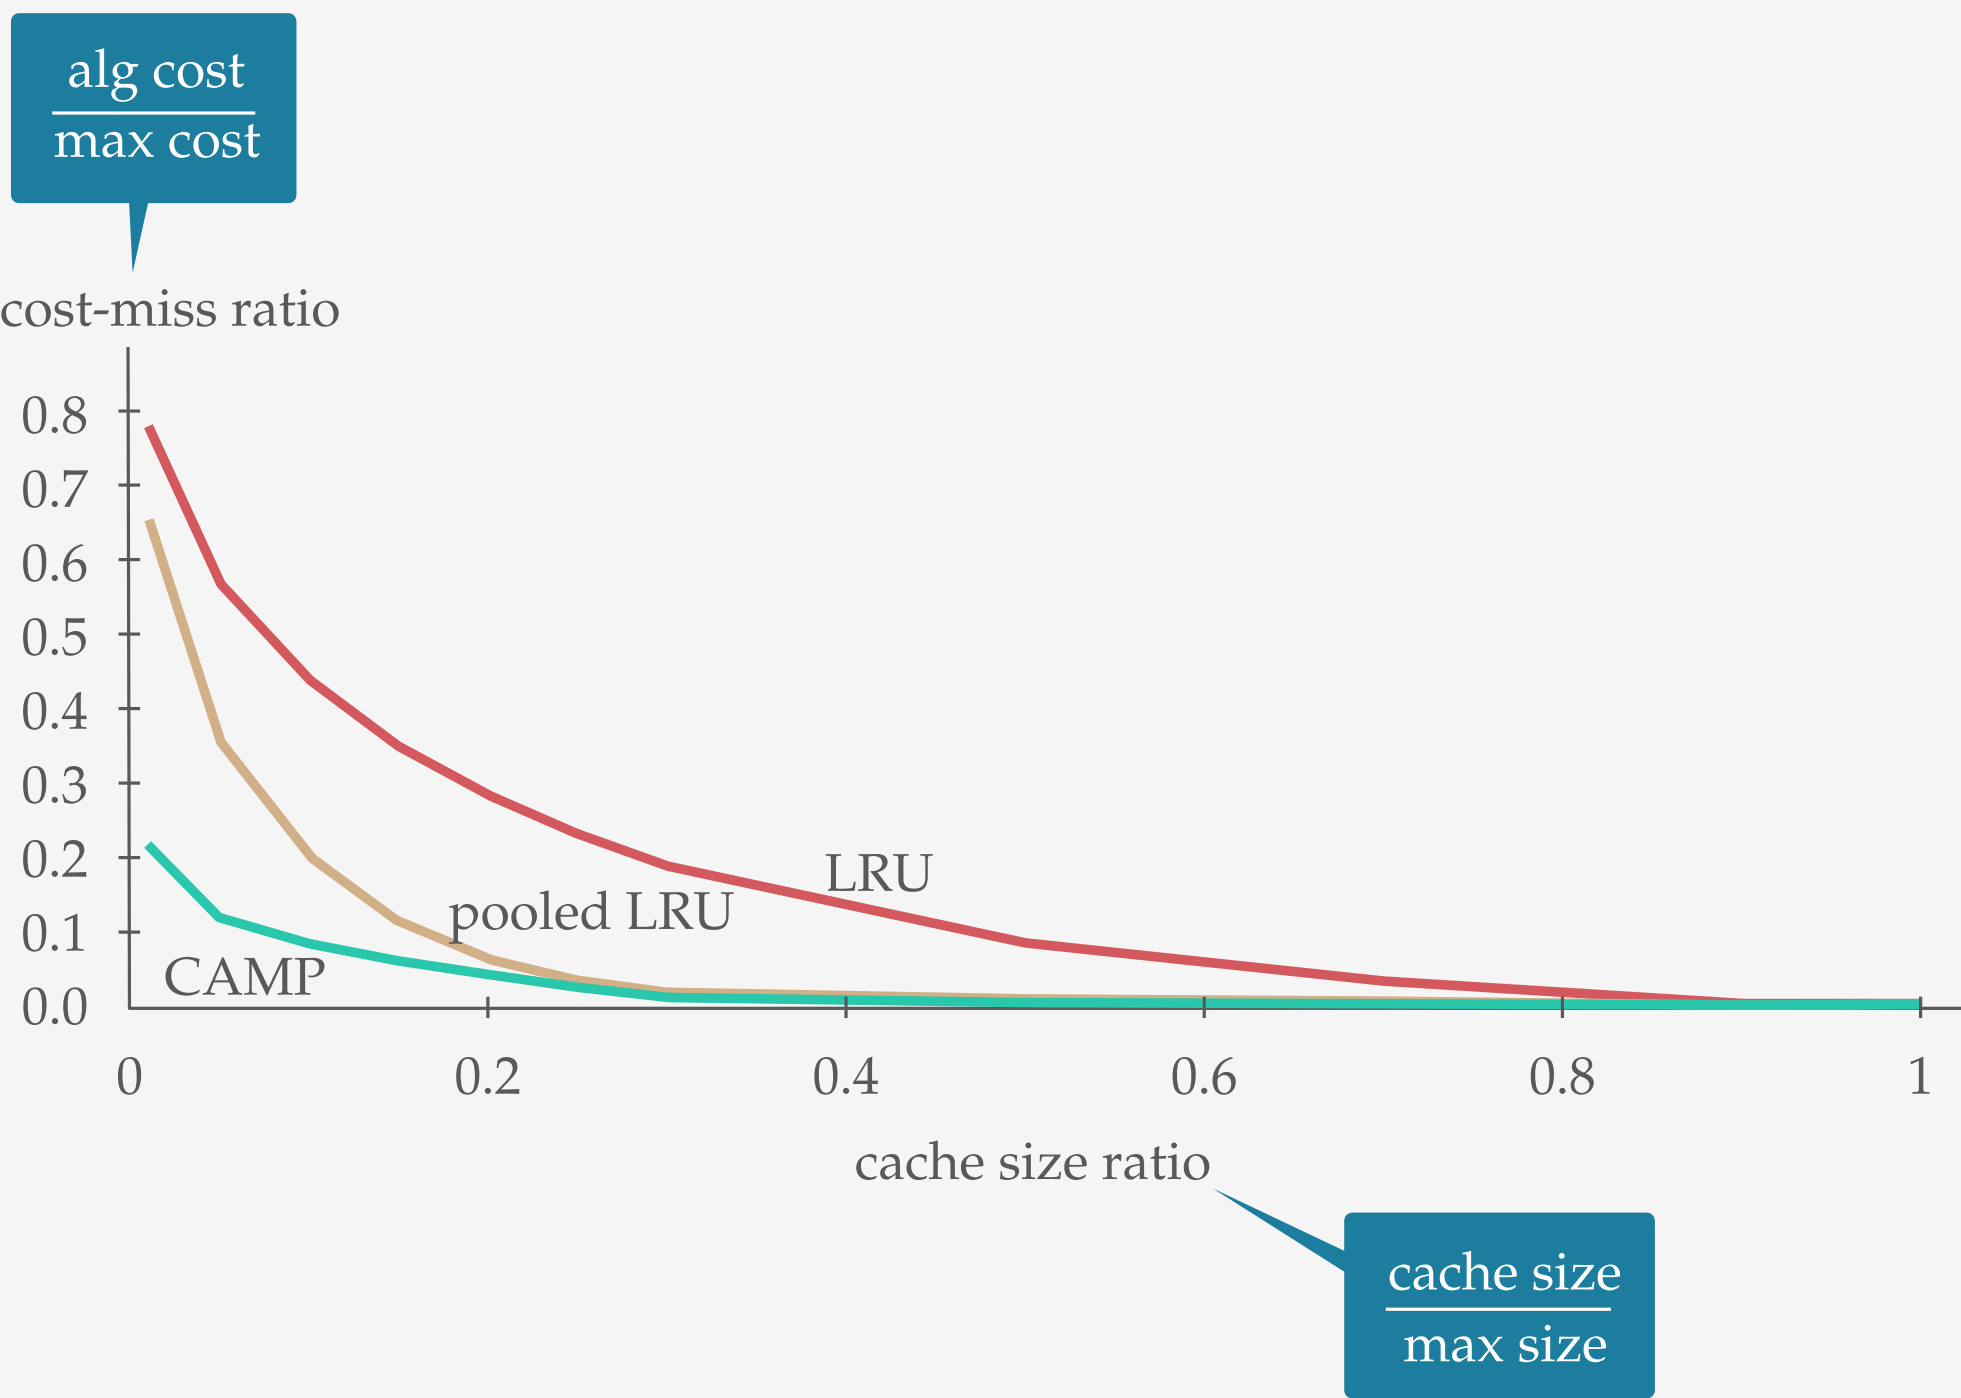
visited heap nodes




CAMP

GDS

$$\frac{\text{cache size}}{\text{max size}}$$




miss rate



$\frac{\text{cache size}}{\text{max size}}$

EVICTIION POLICY

GDS → CAMP




PLACEMENT POLICY

generalized caching → managed memory caching


MEMORY HIERARCHY

2-level cache → multi-level cache

pooled Least Recently Used



Scaling Memcache at Facebook, Nishtala et al., NSDI 2013.



cold expensive







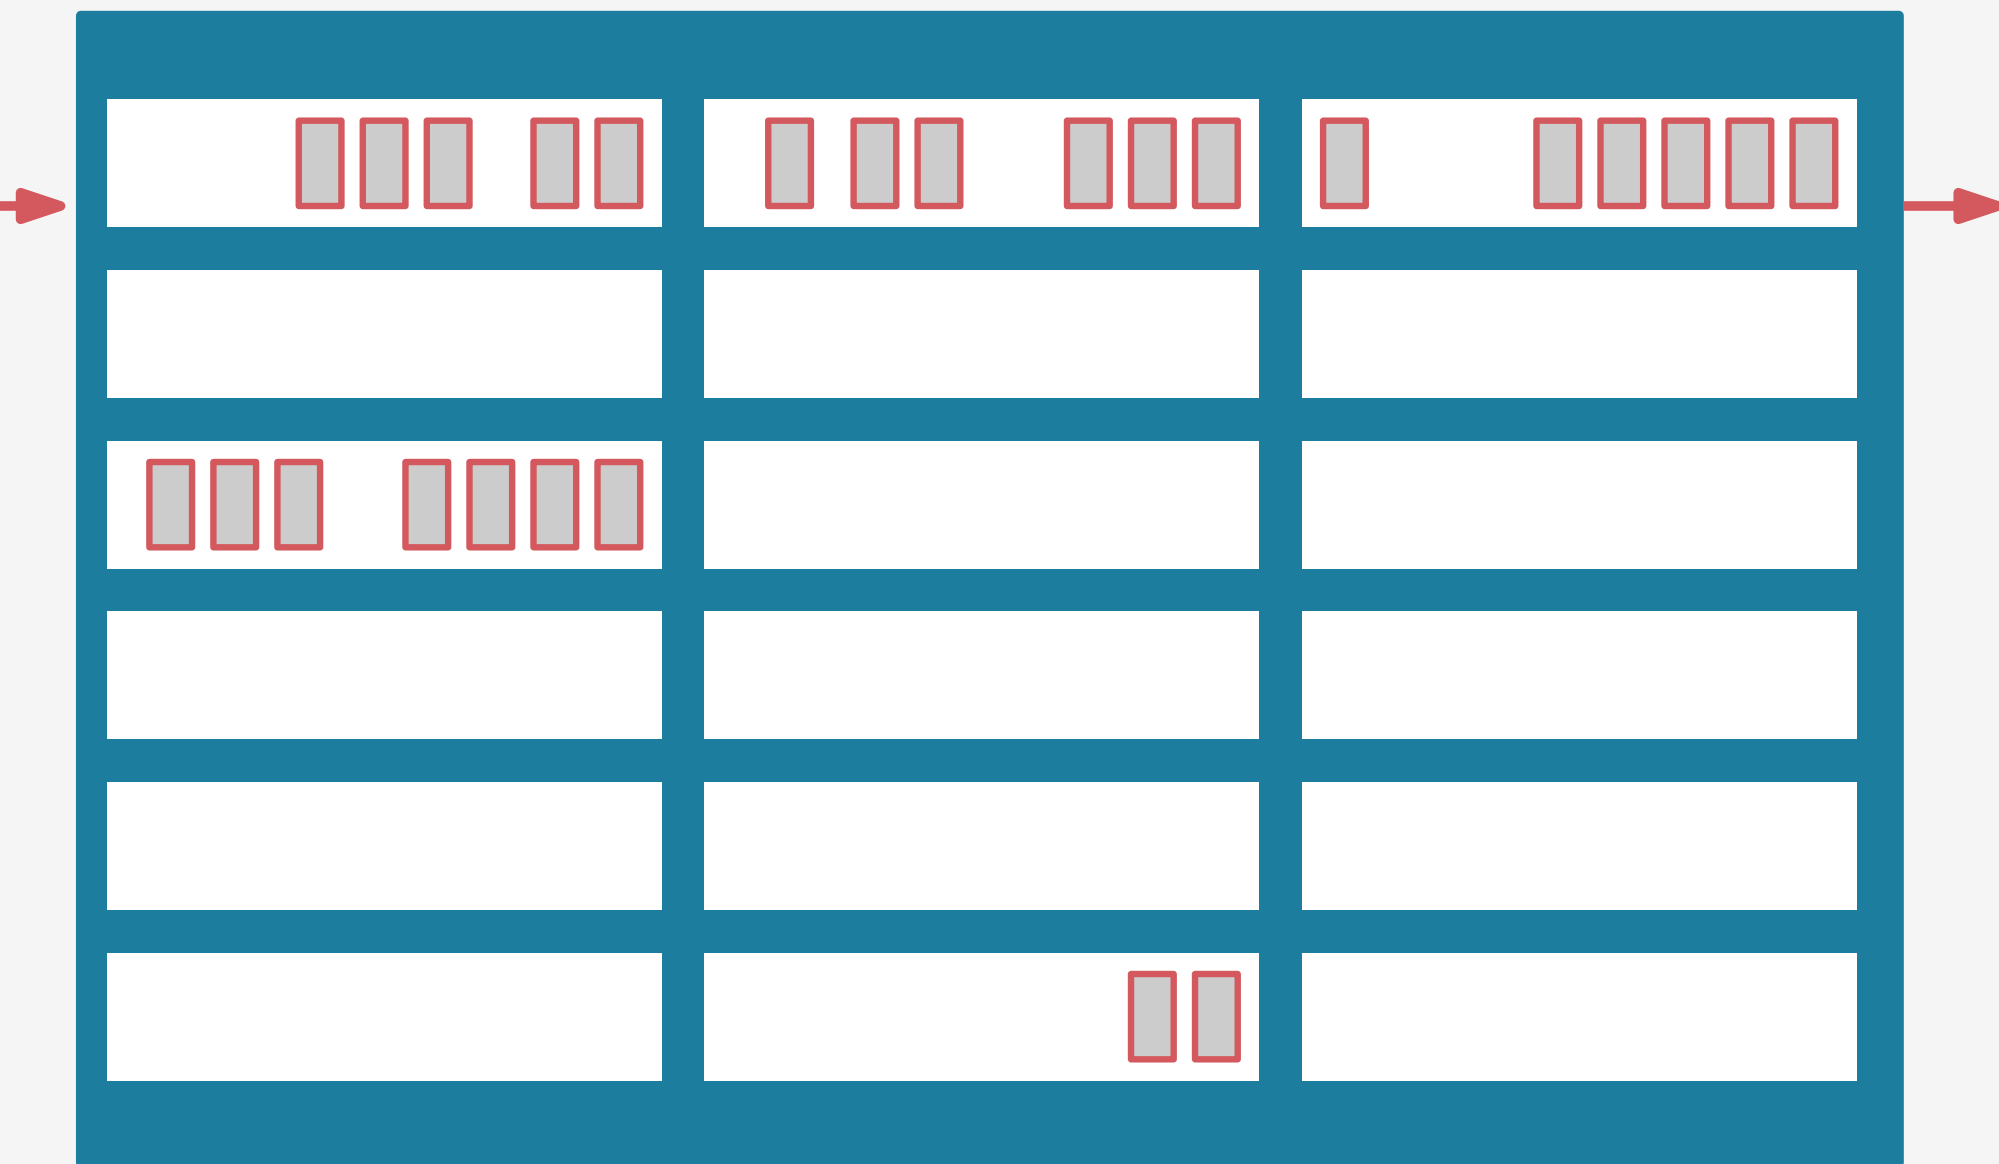



cold expensive




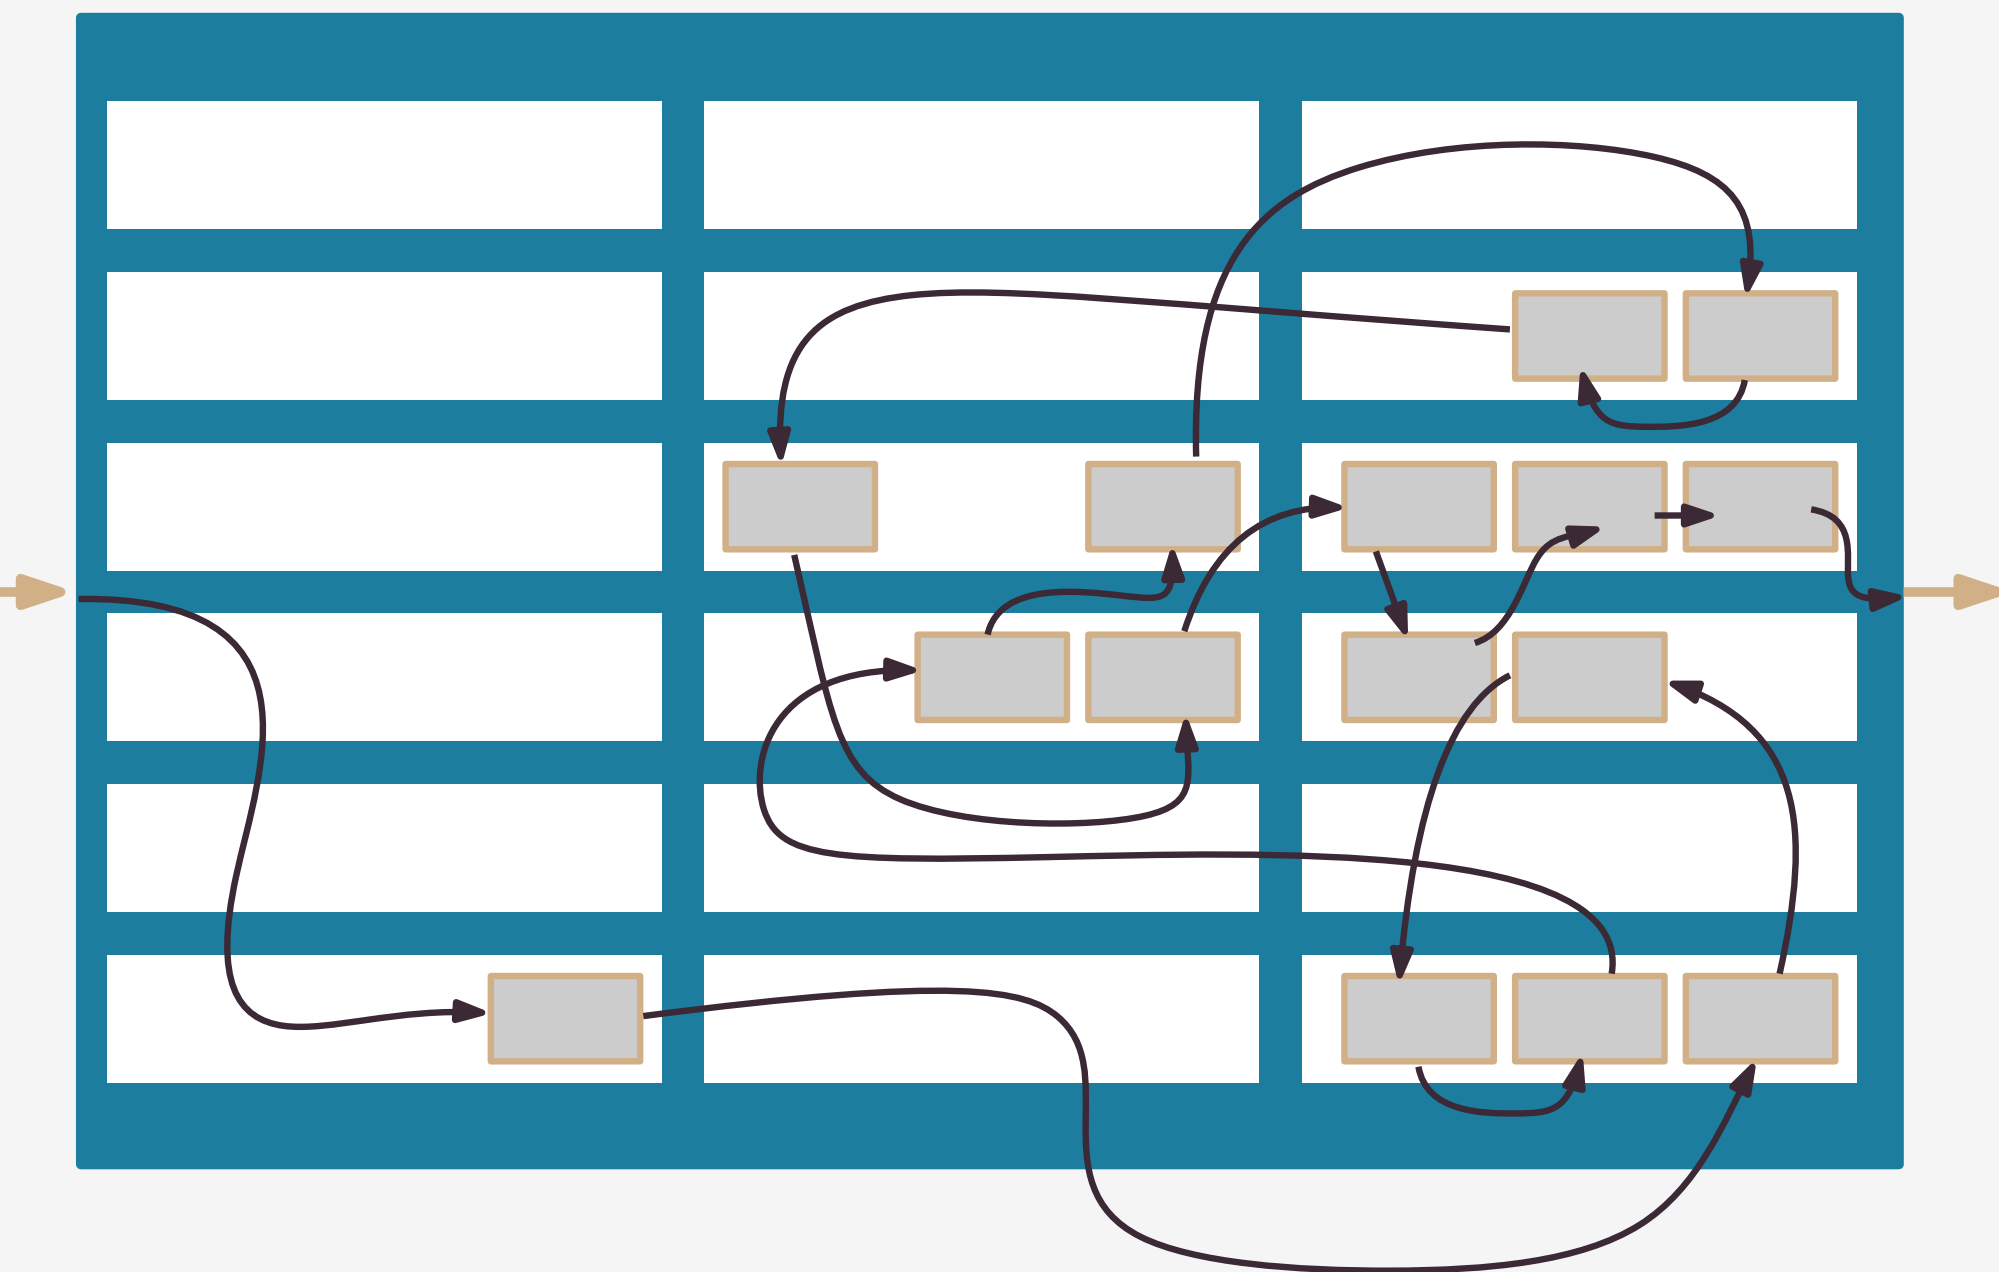
cold expensive



















EVICTIION POLICY



PLACEMENT POLICY

THE GENERALIZED CACHING PROBLEM

variable size and cost



the cache

GOAL

minimize **total cost** of cache misses

SUBJECT TO

total size of items in cache
cannot exceed the cache size

THE MANAGED MEMORY CACHING PROBLEM

variable size and cost



the cache

every item must fit in a contiguous
segment of memory

CACHE REPLACEMENT
MEMORY ALLOCATION



CAMP-MALLOC



LRU queues




CAMP-MALLOC



FIFO queues




CAMP-MALLOC



FIFO queue




CAMP-MALLOC



FIFO queue




CAMP-MALLOC



FIFO queue




CAMP-MALLOC



FIFO queue




CAMP-MALLOC



FIFO queue



CAMP-MALLOC



FIFO queue




CAMP-MALLOC



FIFO queue



CAMP-MALLOC



FIFO queue




CAMP-MALLOC



FIFO queue




CAMP-MALLOC



FIFO queue



CAMP-MALLOC



FIFO queue



CAMP-MALLOC



FIFO queue




CAMP-MALLOC



FIFO queue



CAMP-MALLOC



first in

FIFO queue



CAMP-MALLOC




first in



FIFO queue



CAMP-MALLOC





FIFO queue



CAMP-MALLOC




first in



FIFO queue



CAMP-MALLOC




FIFO queue



CAMP-MALLOC




first in



FIFO queue




CAMP-MALLOC



FIFO queue




CAMP-MALLOC





CAMP-MALLOC


fragmentation ≤ 2 (max item size)



FIFO queue




CAMP-MALLOC






CAMP-MALLOC






CAMP-MALLOC






CAMP-MALLOC






CAMP-MALLOC





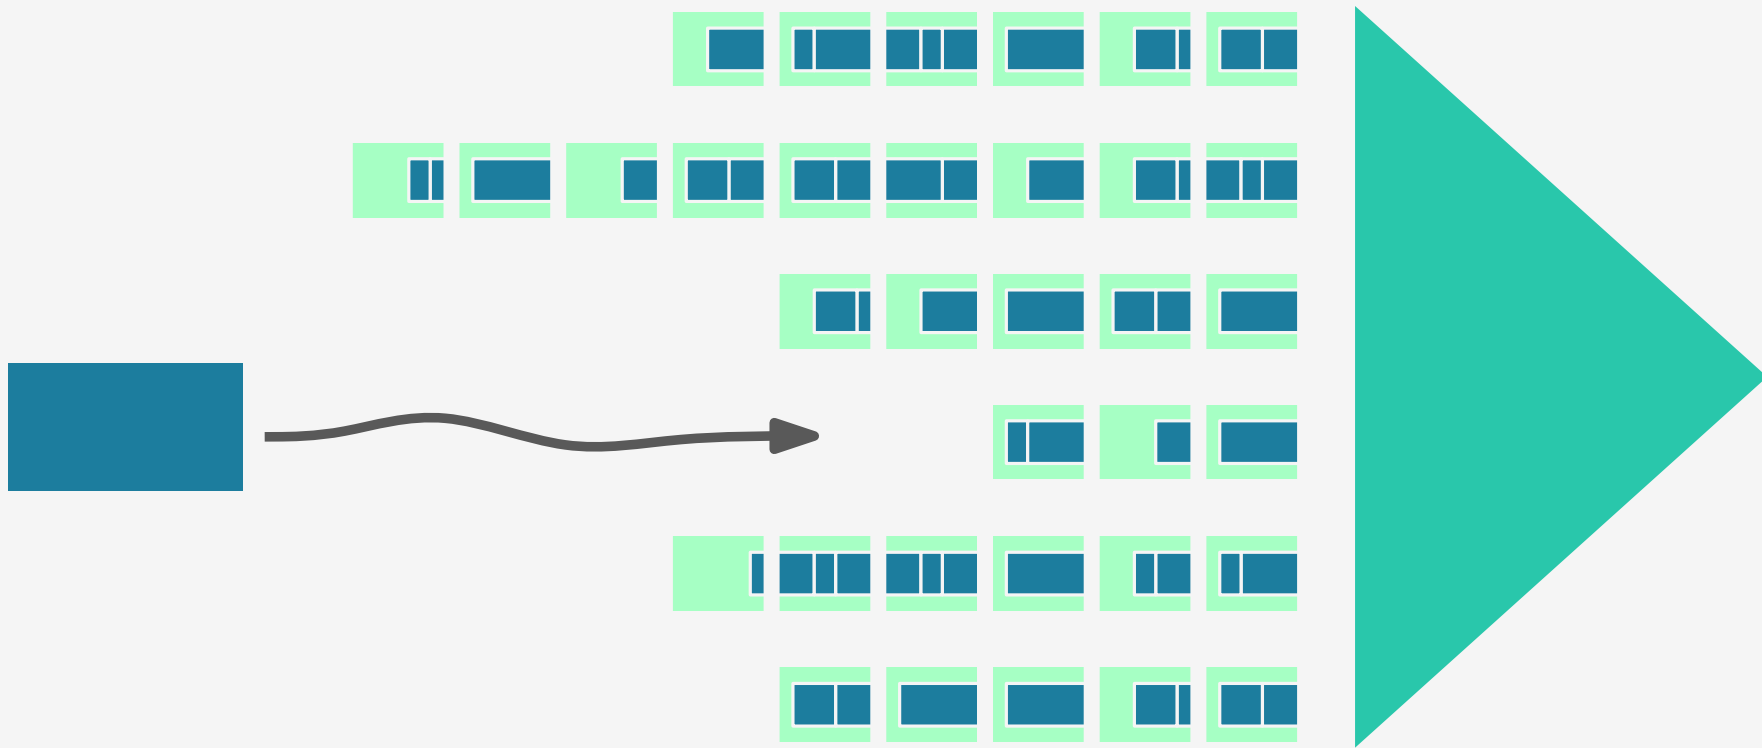


CAMP-MALLOC



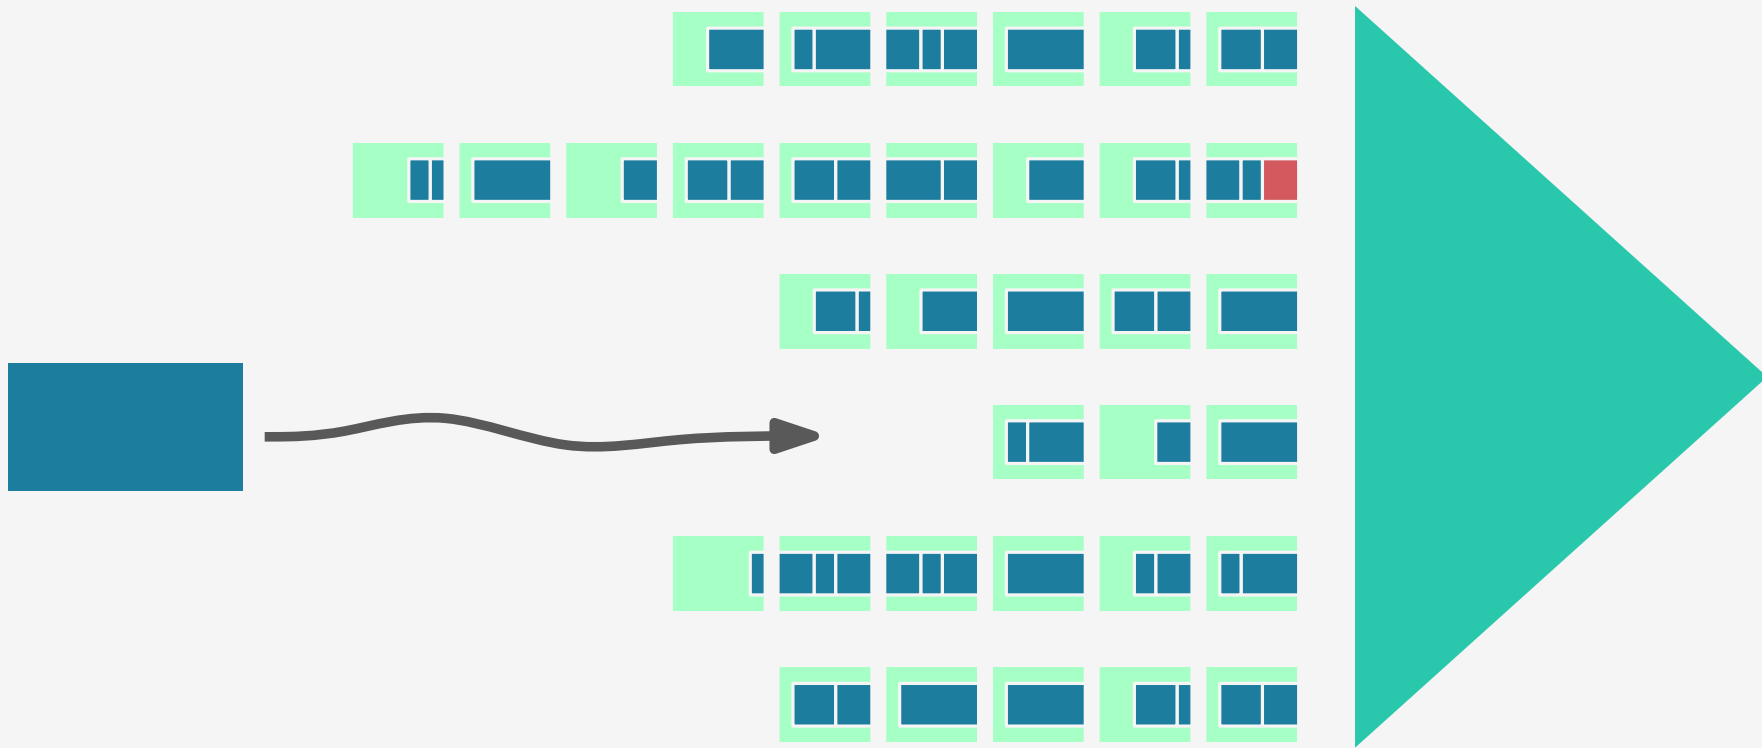


CAMP-MALLOC



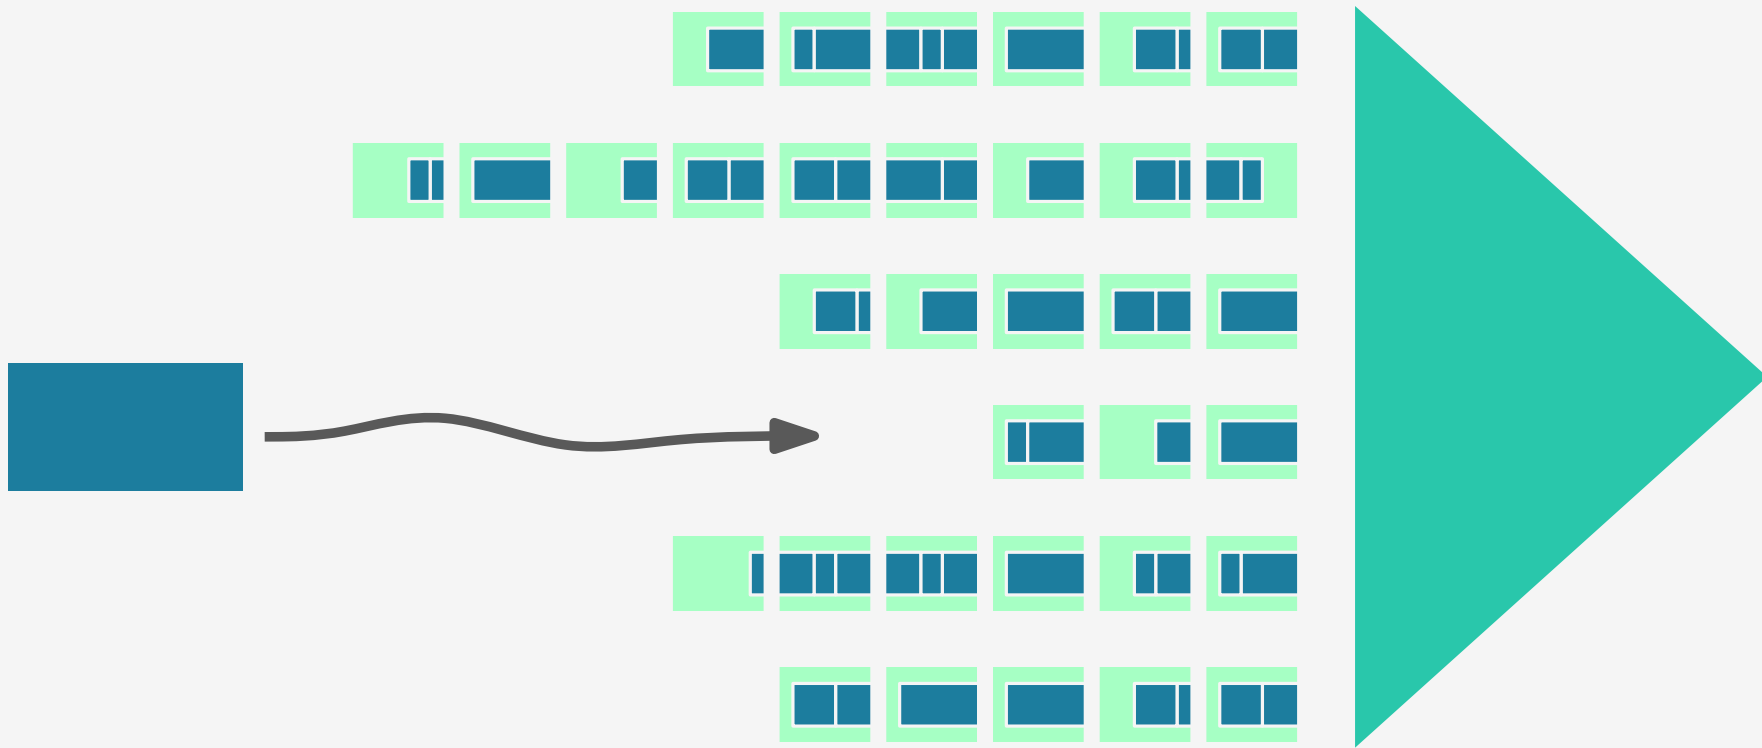


CAMP-MALLOC



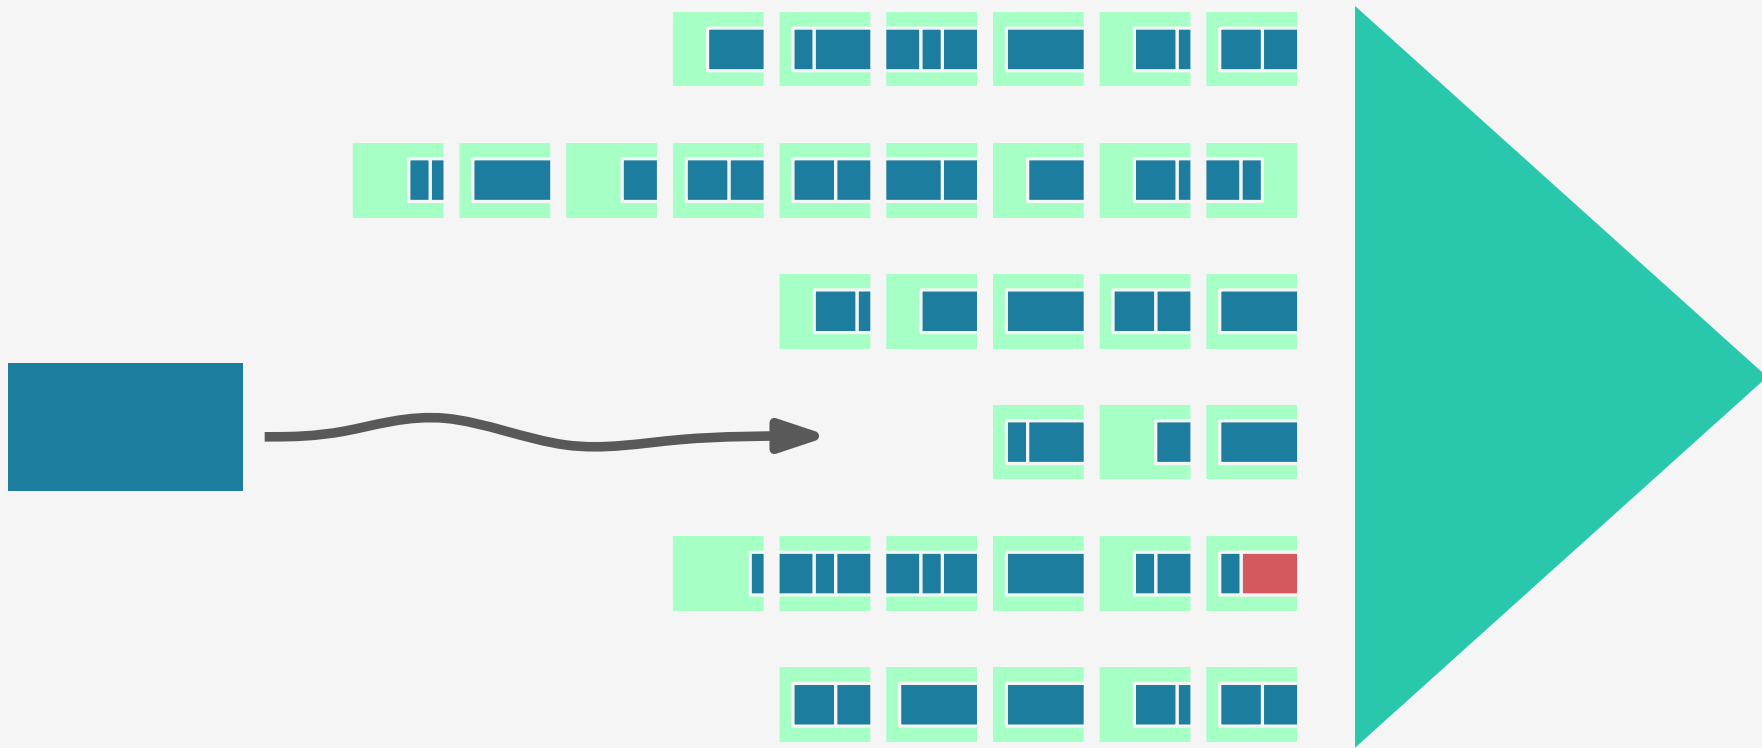


CAMP-MALLOC



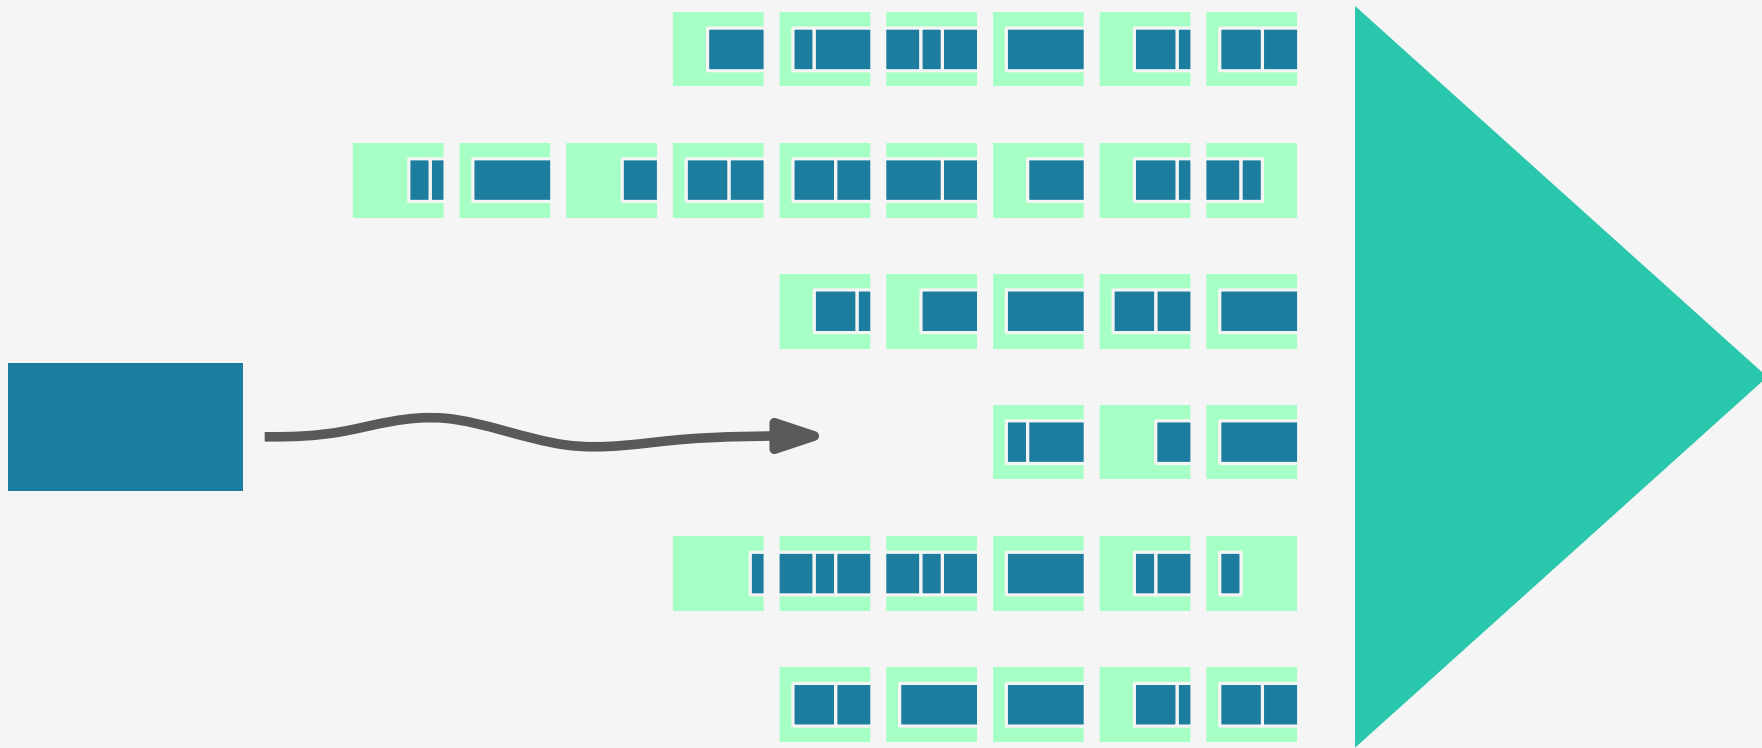


CAMP-MALLOC



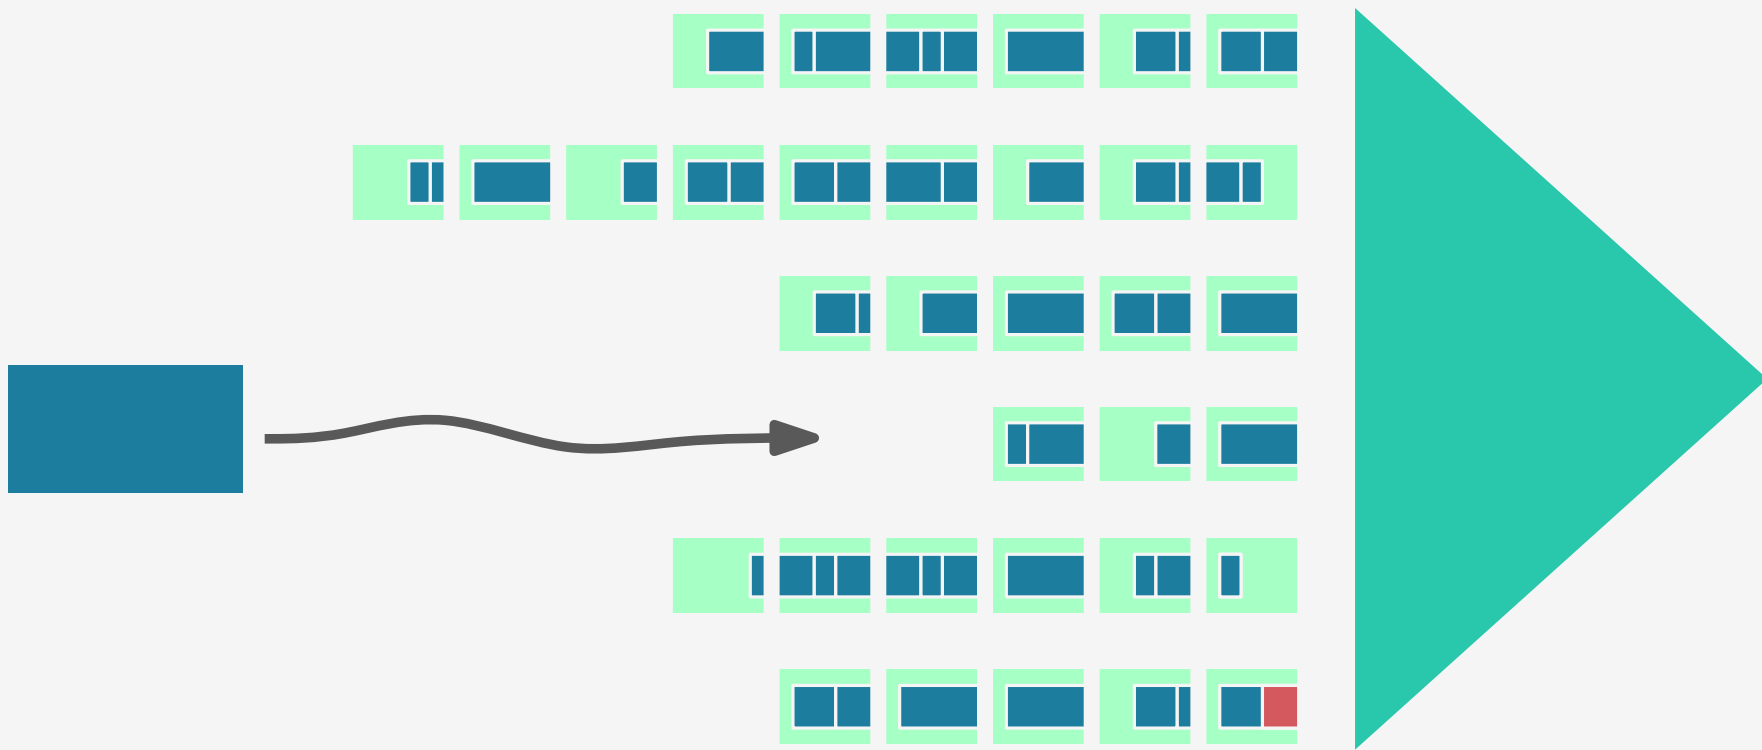


CAMP-MALLOC



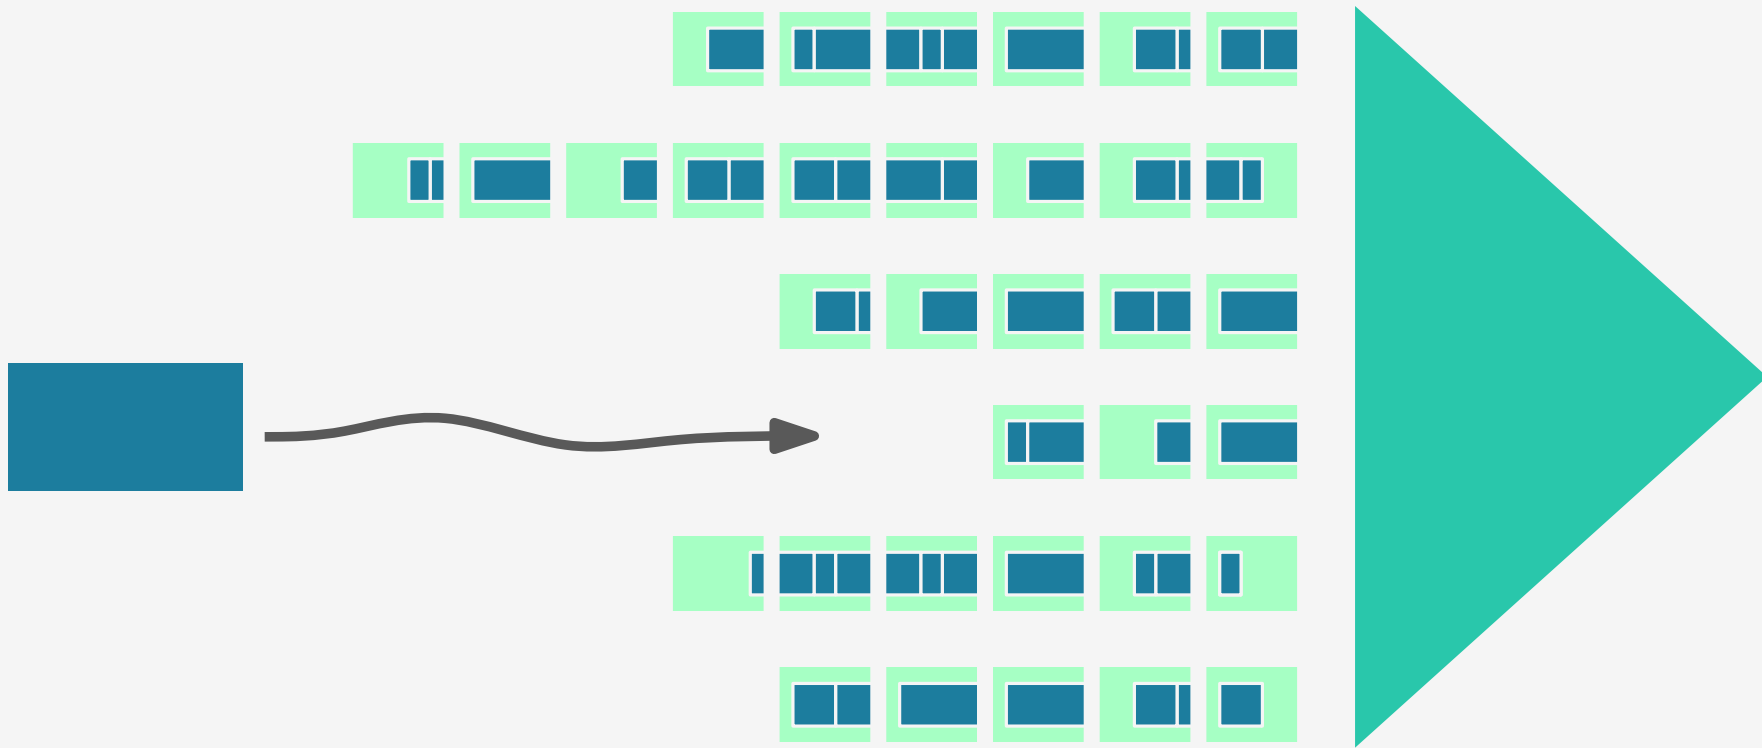


CAMP-MALLOC



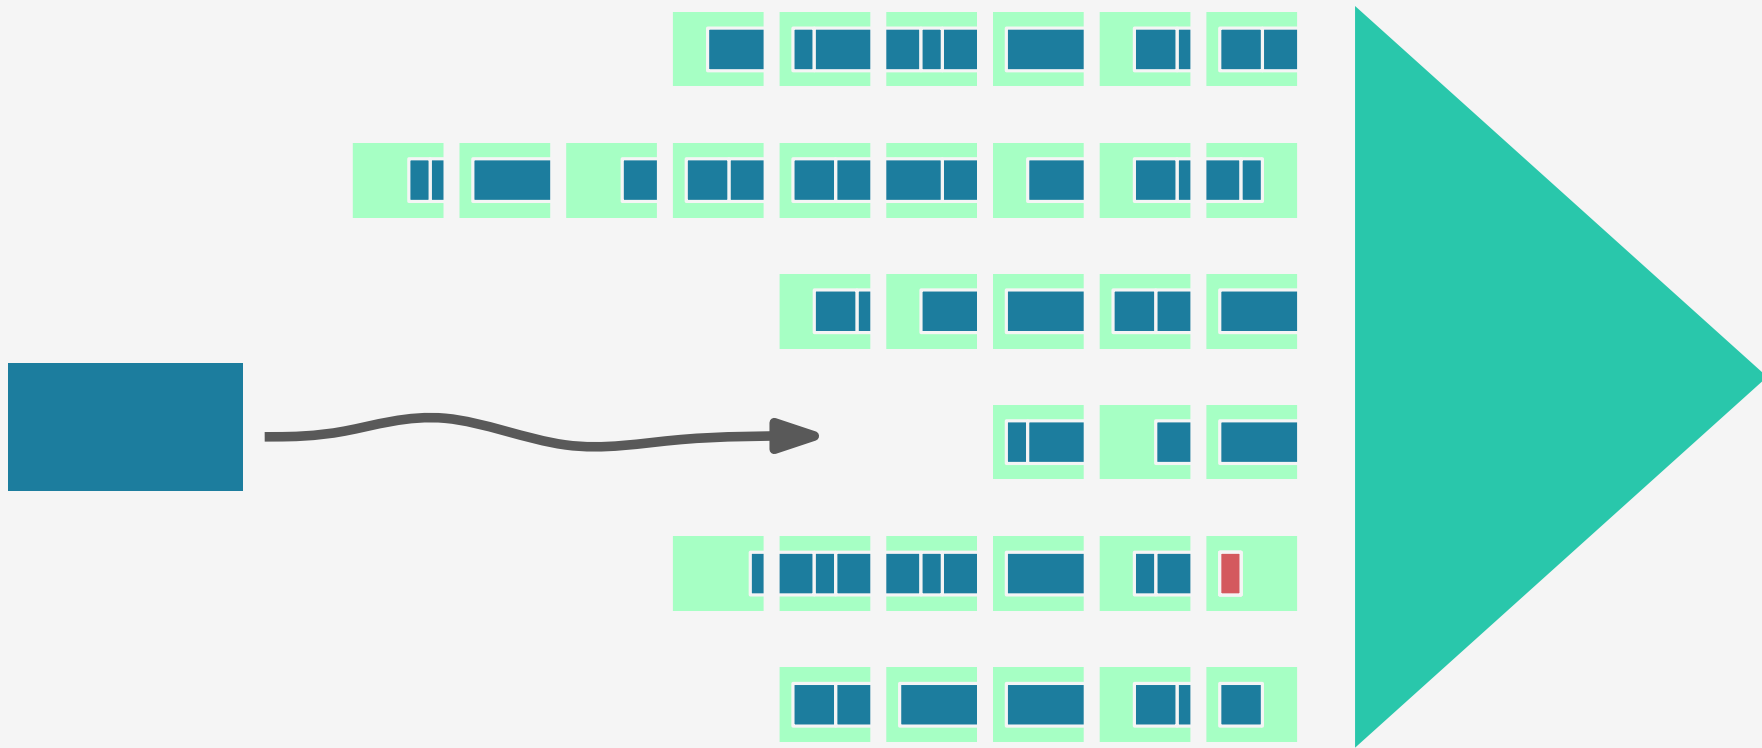


CAMP-MALLOC



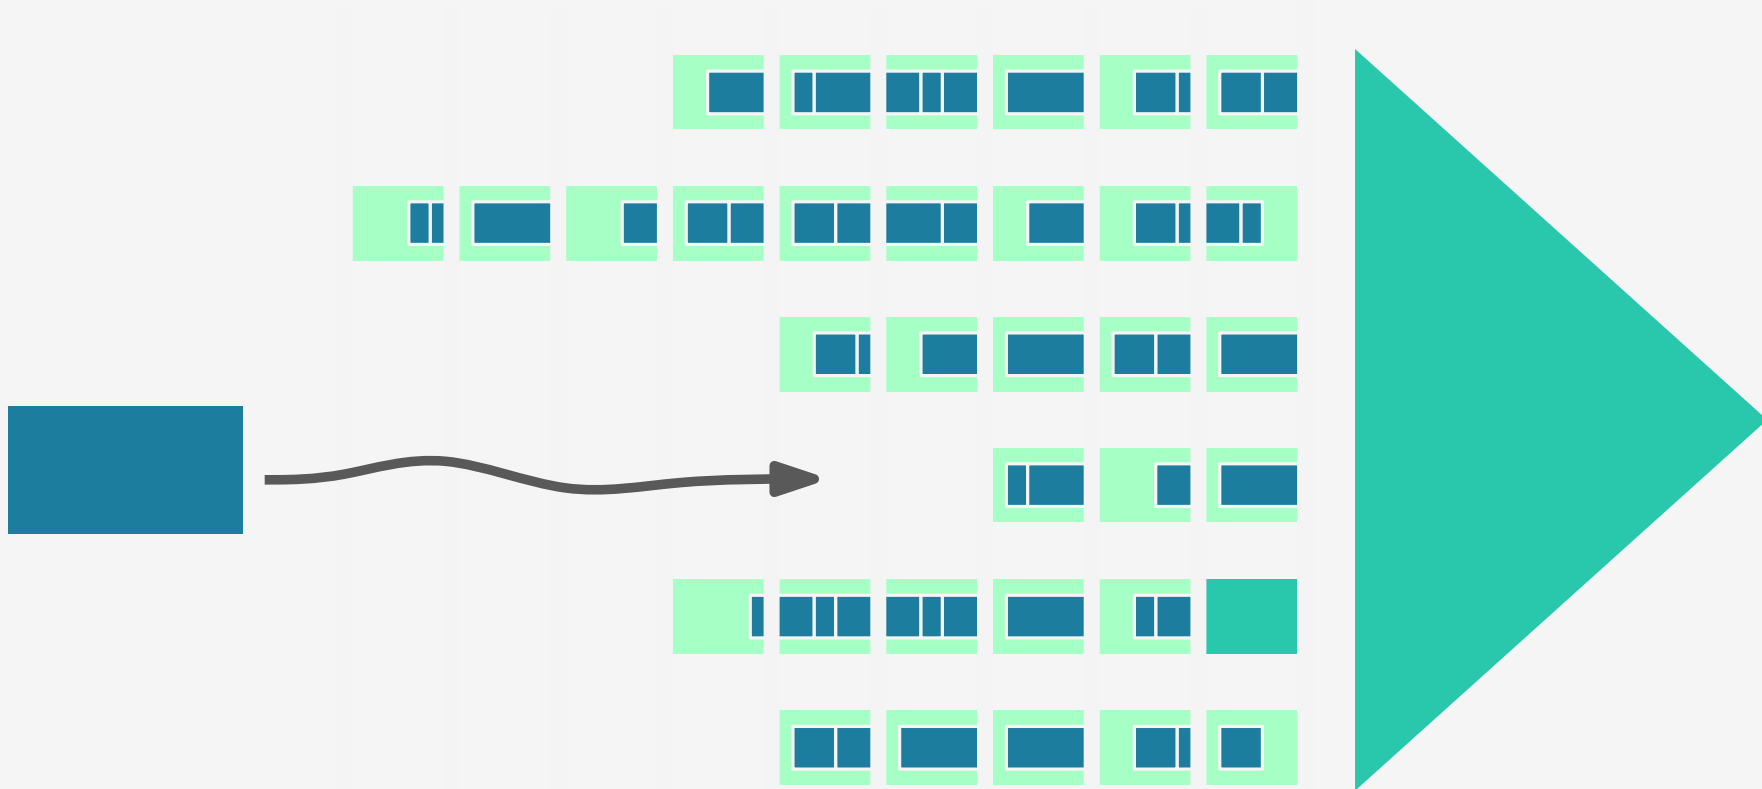


CAMP-MALLOC



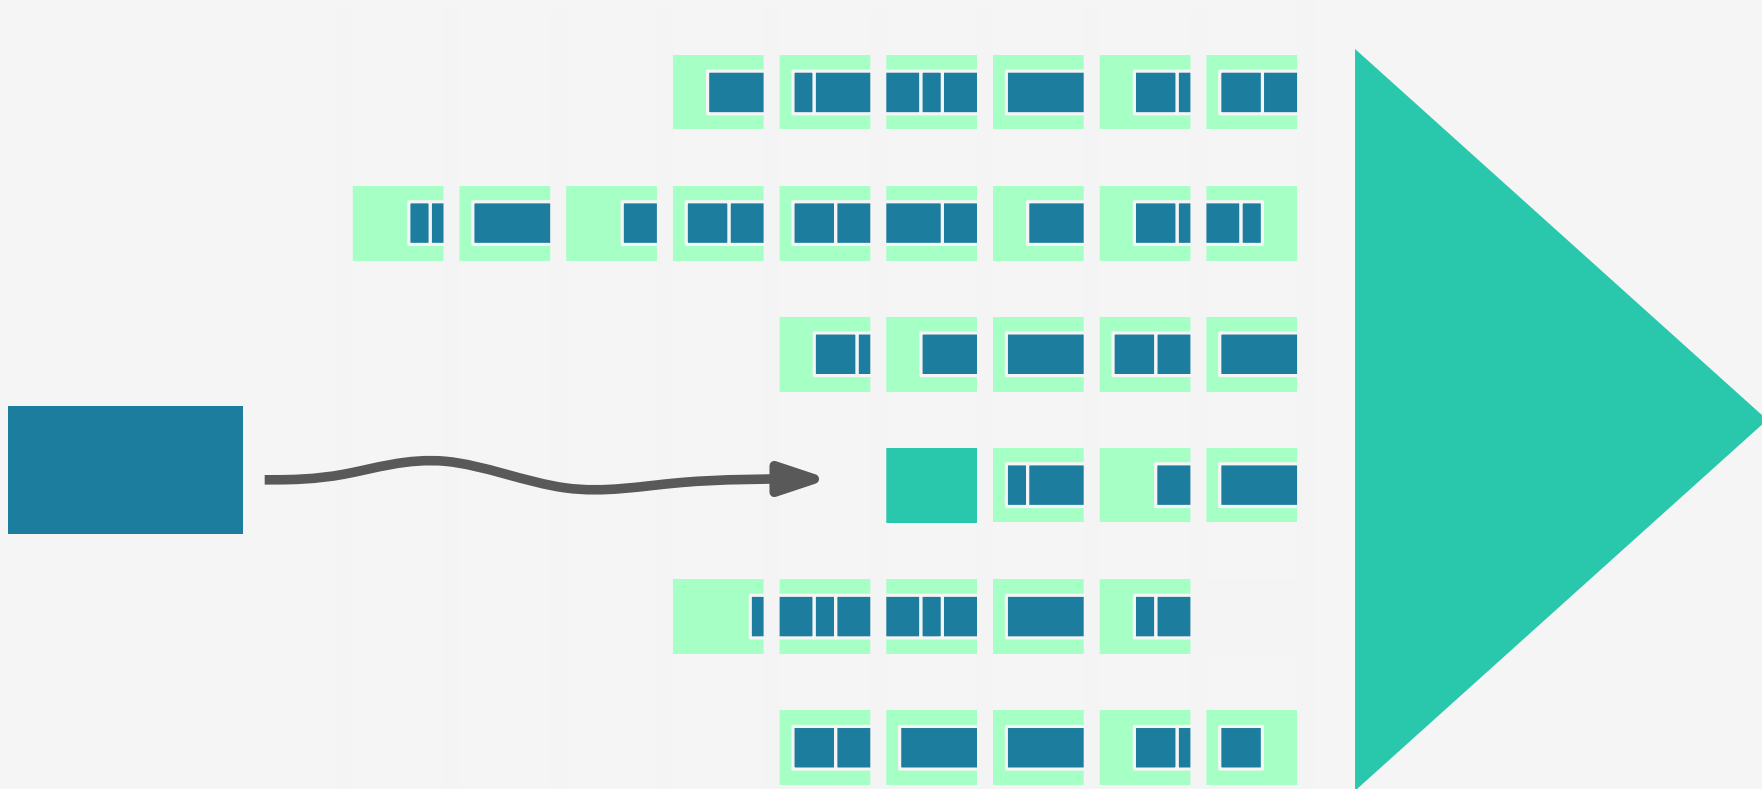


CAMP-MALLOC






CAMP-MALLOC






CAMP-MALLOC






CAMP-MALLOC






CAMP-MALLOC






CAMP-MALLOC





CAMP-MALLOC



fragmentation ≤ 2 (num queues) (block size) + (num blocks) (max item size)



CAMP-MALLOC


is competitive if memory augmented

if $\text{OPT's cache size} \leq \text{C-M's cache size} - \text{fragmentation bound}$


then $\text{cost(C-M)} \leq \frac{\text{C-M's cache size}}{\text{min item size}} \text{cost(OPT)}$

fragmentation ≤ 2 (num queues) (block size) + (num blocks) (max item size)

cost-miss ratio




percent fragmentation




block size (units of max item size)

percent fragmentation



block size (units of max item size)



EVICTIION POLICY


GDS → CAMP



PLACEMENT POLICY

generalized caching → managed memory caching

MEMORY HIERARCHY

2-level cache → multi-level cache





GENERALIZED CACHING

capacity

read speed

write speed

failure rate

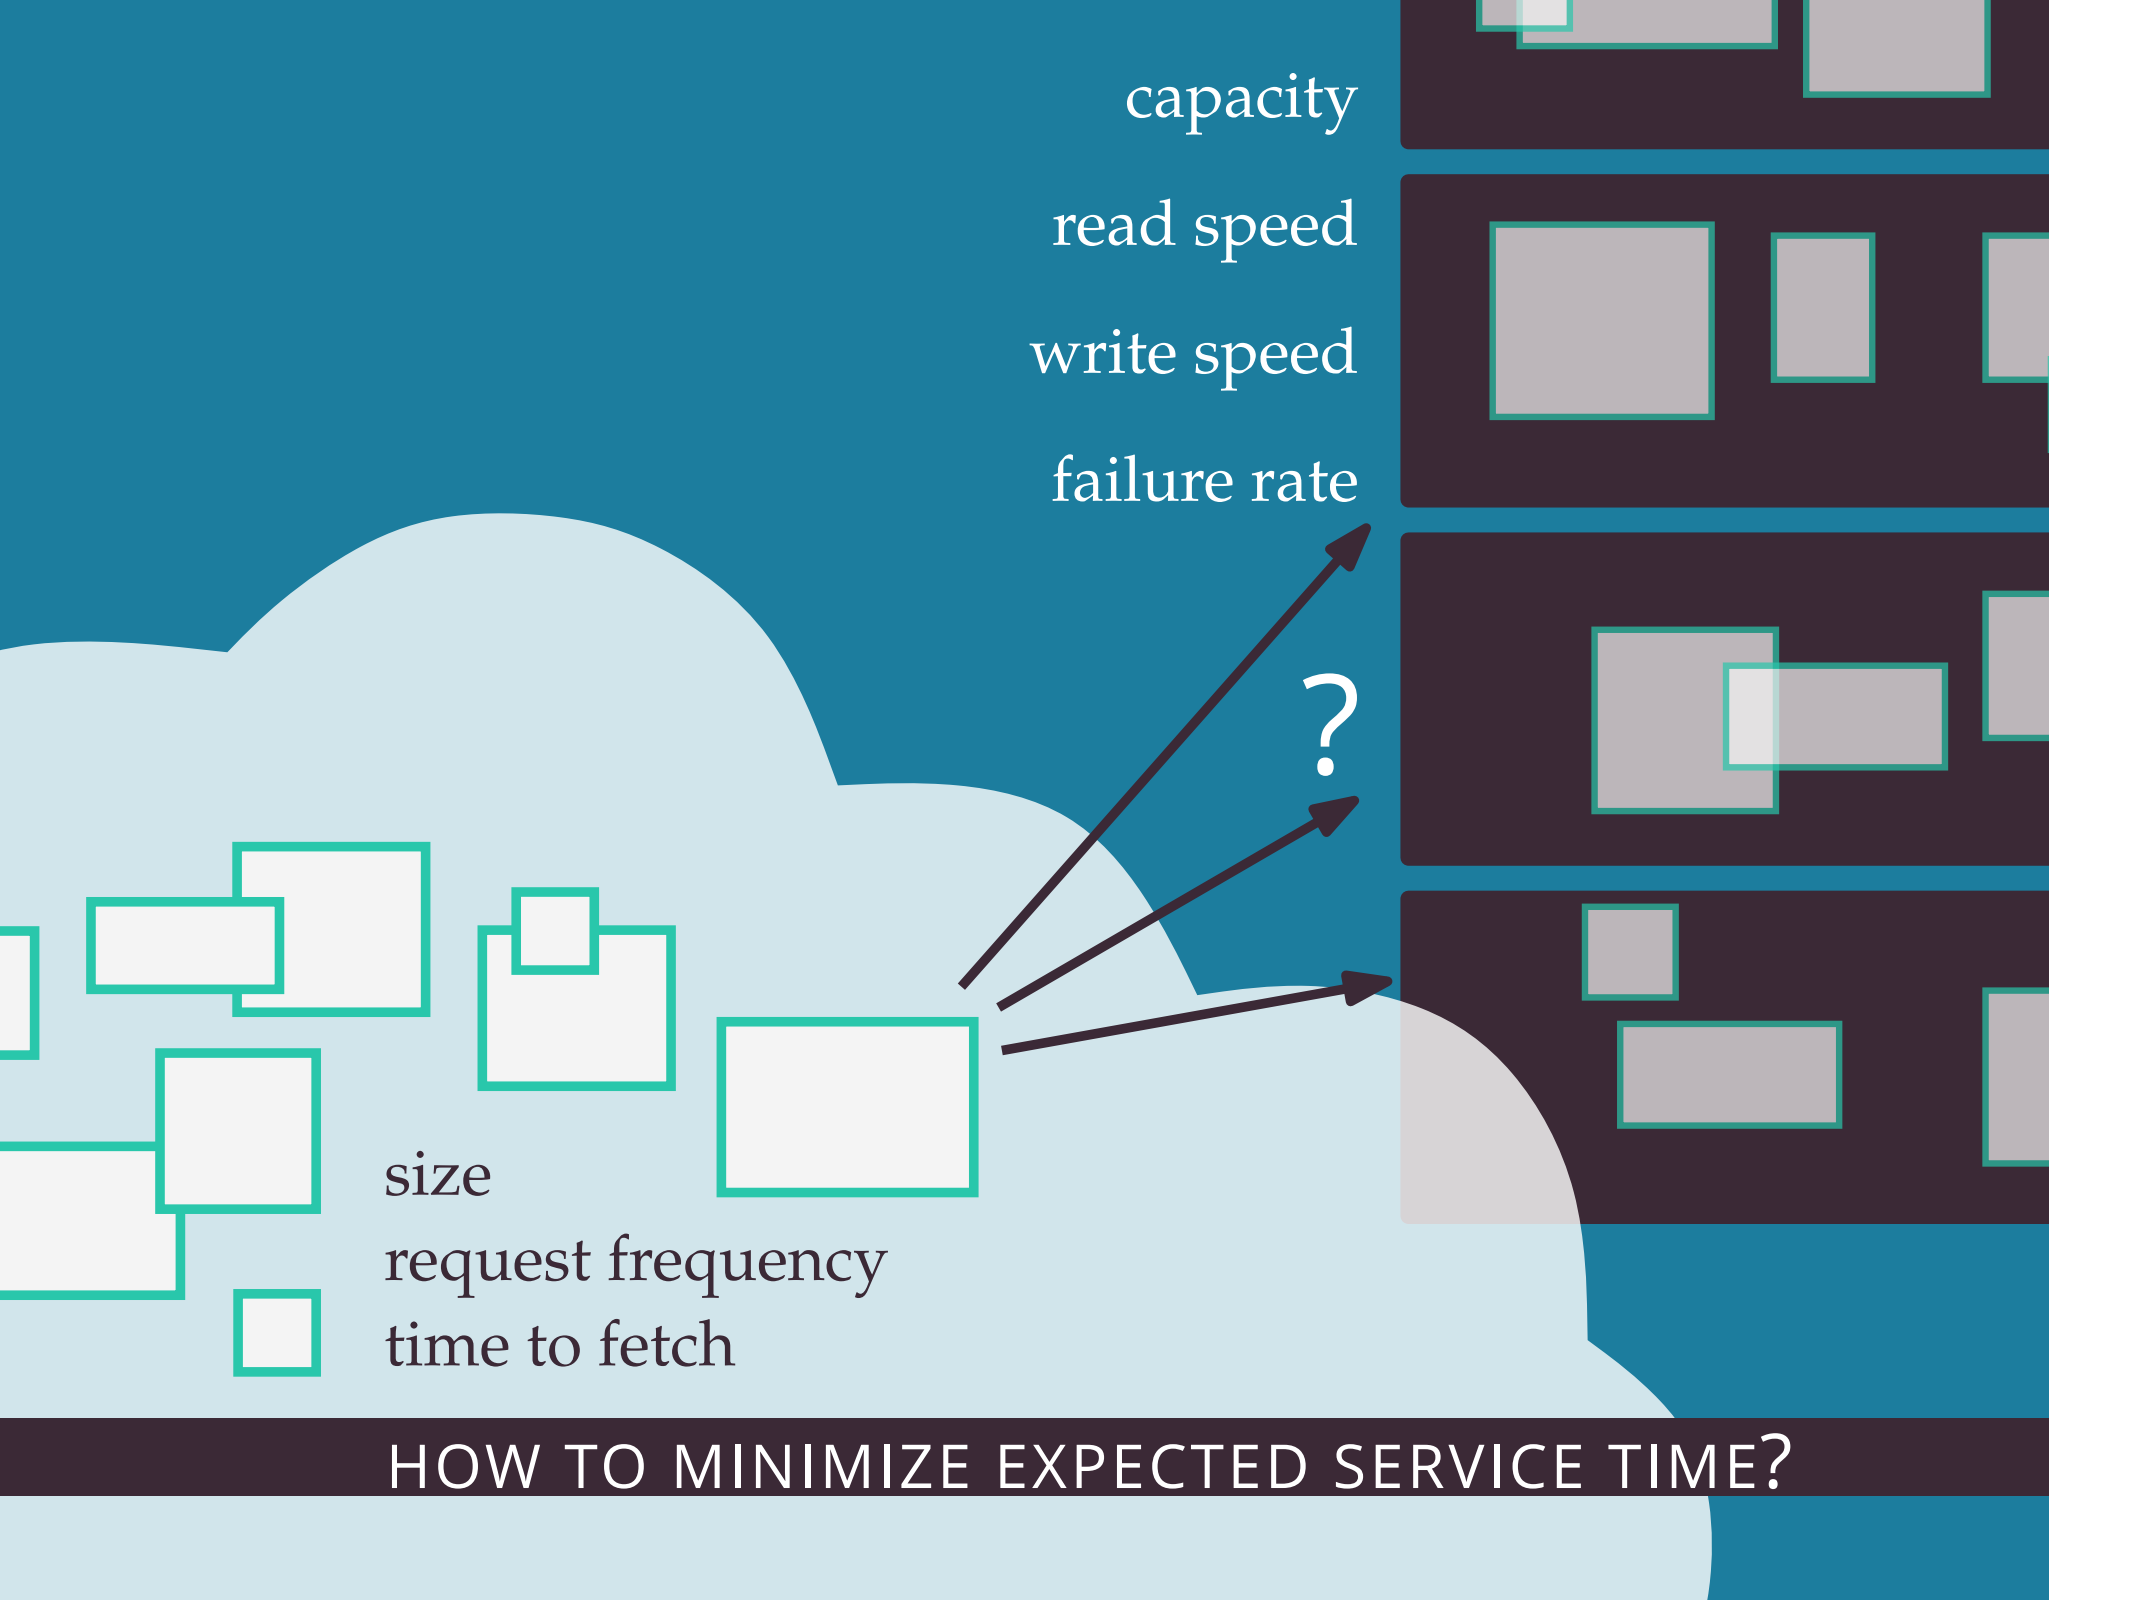
?

size

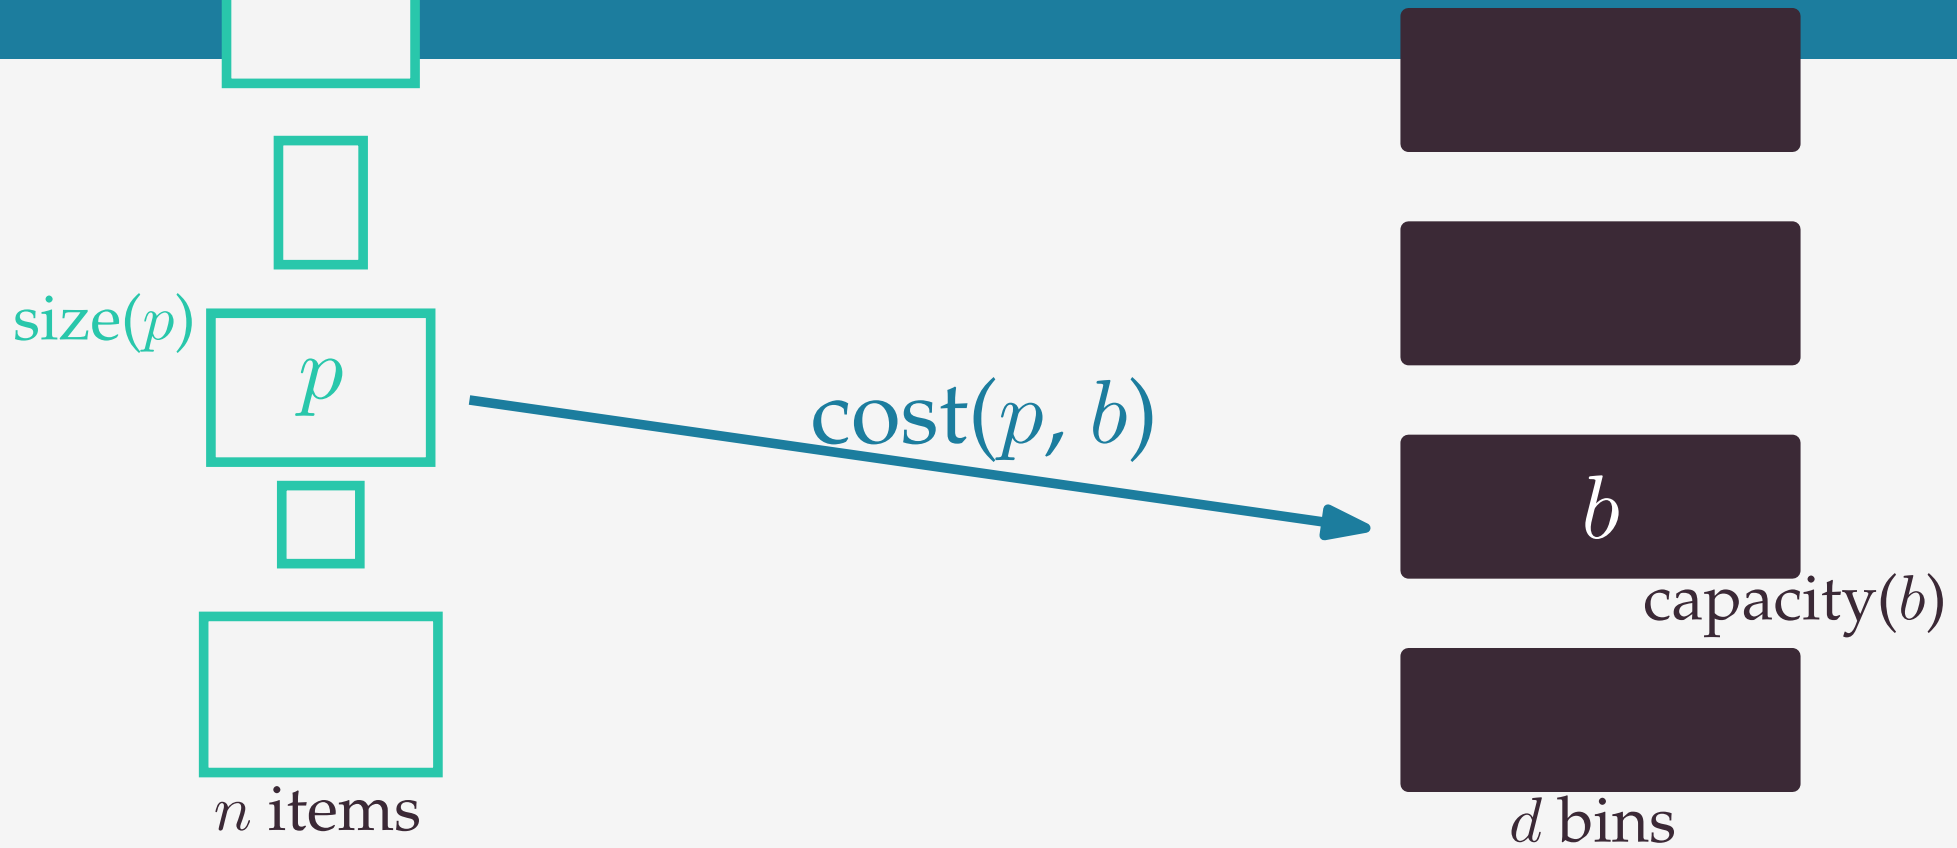
request frequency

time to fetch

HOW TO MINIMIZE EXPECTED SERVICE TIME?



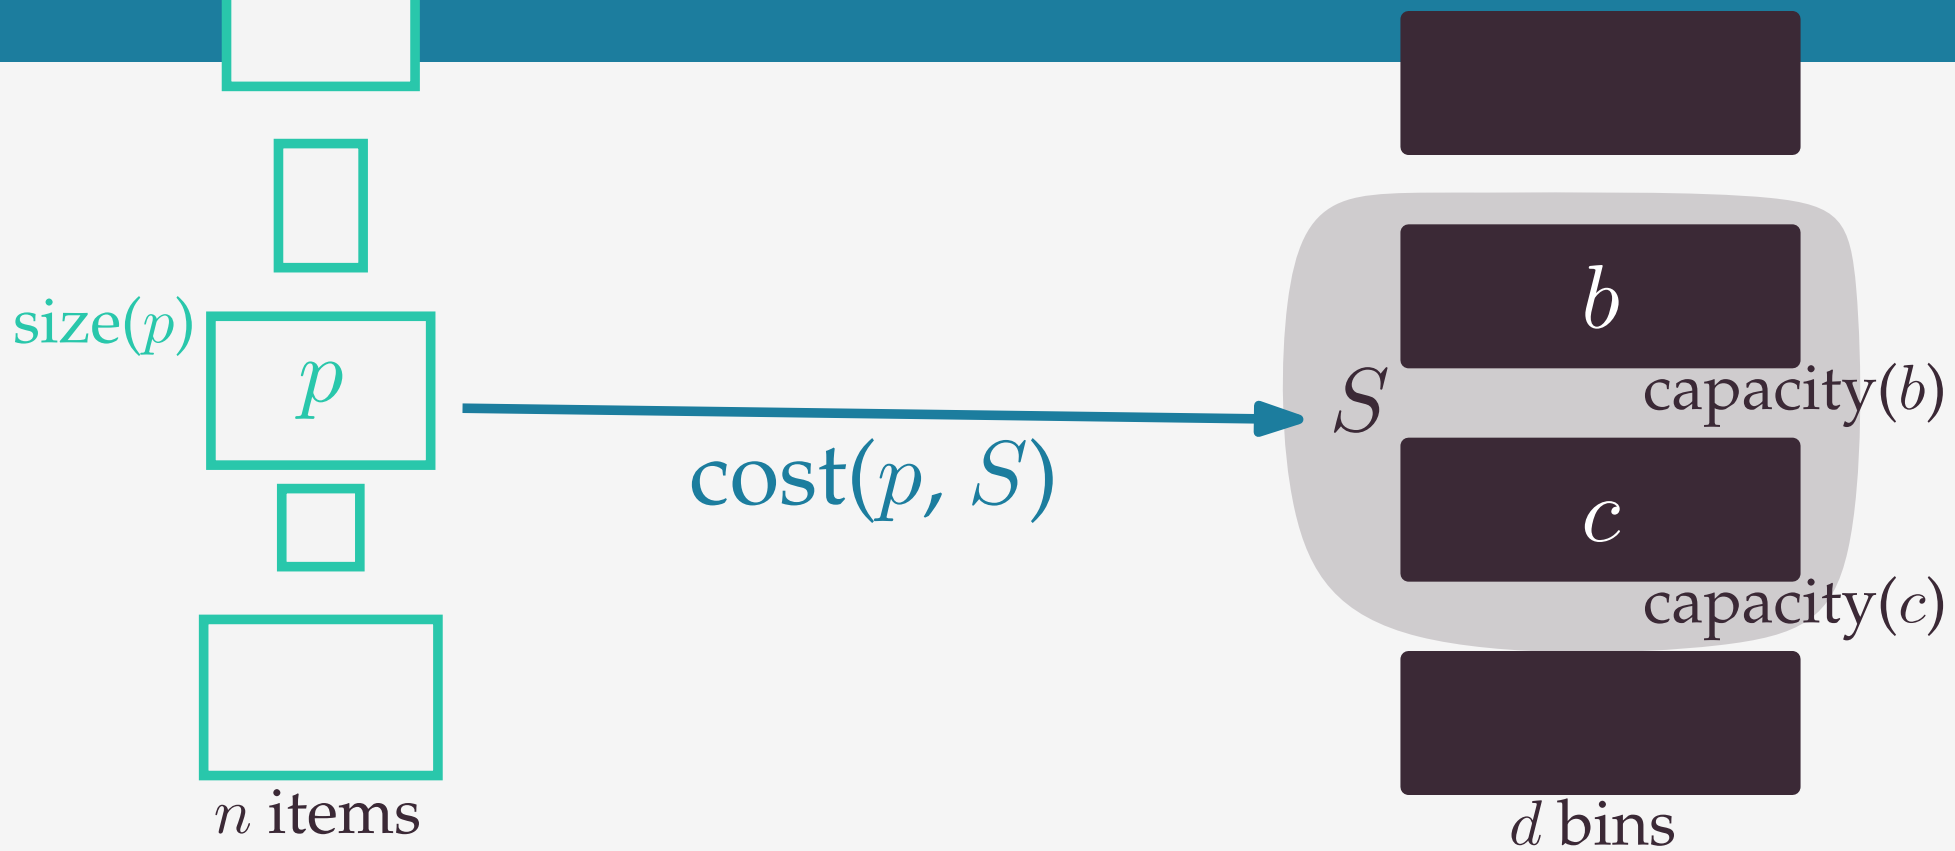
MULTIPLE KNAPSACK PROBLEM



GOAL


minimize total cost of assignment
subject to capacity constraints

SUBSET ASSIGNMENT PROBLEM




GOAL

minimize total cost of assignment
subject to capacity constraints



$$\begin{aligned}
 \text{service-time}(p, S) = & \text{read-frequency}(p) \quad \text{read-time}(p, S) \\
 & + \text{write-frequency}(p) \quad \text{write-time}(p, S) \\
 & + \sum_{F \subseteq S} \text{fail-freq}(F) \left(\text{read-time}(p, S \setminus F) \right. \\
 & \quad \left. + \text{write-time}(p, S \cap F) \right)
 \end{aligned}$$



cache configuration

$$\text{maximize } \sum_{p, S} \text{benefit}(p, S) x(p, S)$$

$$\sum_S x(p, S) = 1$$

$$\sum_{p, S} \text{price}(p, S) x(p, S) \leq \text{budget}$$

$$x = 0, 1$$

subset assignment

$$\text{minimize } \sum_{p, S} \text{cost}(p, S) x(p, S)$$


$$\sum_S x(p, S) = \text{size}(p)$$

$$\sum_{p, S \ni b} x(p, S) \leq \text{capacity}(b)$$

$$x(p, S) = 0, \text{size}(p)$$



CACHE CONFIGURATION





SUBSET ASSIGNMENT

HAVE

$$d \ll n$$

sol to LP relaxation has few fractional assignments

GOAL


solve LP relaxation in $f(d) \text{ poly}(n)$

1. cycle canceling algorithm

2. simplex algorithm




MIN COST FLOW







1. cycle canceling algorithm



feasible flow
in original network




negative cycle
in residual network





“cycle” in subset assignment problem

augmentation



such that

$$\sum_i \alpha_i \overrightarrow{S_i T_i} = \vec{0}$$


cost difference
(negative)

$$\sum_i \alpha_i (\text{cost}(p_i, T_i) - \text{cost}(p_i, S_i))$$



2. simplex algorithm

basic feasible solution



BASIC FEASIBLE
ASSIGNMENT

$< 2d$ fractional assignments


bound granularity of vars

$$x(p, S) = \frac{k}{\ell}$$

$< d^{d/2}$




ALGORITHM





ALGORITHM

preprocessing $\sum_i \alpha_i \overrightarrow{S_i T_i} = \vec{0}$



time $O(\exp(d(d+1))\text{poly}(d) \ n \log(n) \log(nC) \log(S))$





EVICTIION POLICY

CAMP



PLACEMENT POLICY

CAMP-malloc



MEMORY HIERARCHY

cache configuration
subset assigment




EVICTIION POLICY

CAMP




PLACEMENT POLICY

CAMP-malloc



LRU-FIFO hybrid?



MEMORY HIERARCHY

cache configuration
subset assnignment




EVICTIION POLICY

CAMP




PLACEMENT POLICY

CAMP-malloc



LRU-FIFO hybrid?



MEMORY HIERARCHY

cache configuration
subset assnignment

continual
updates?



EVICTIION POLICY

CAMP



customer
fairness?



PLACEMENT POLICY

CAMP-malloc

LRU-FIFO hybrid?



MEMORY HIERARCHY

cache configuration
subset assnignment

continual
updates?





CAMP: a cost adaptive multi-queue eviction policy for key-value stores. Shahram Ghandeharizadeh, Sandy Irani, Jenny Lam, and Jason Yap. In Proceedings of the 15th International Middleware Conference, 2014.

A demonstration of KOSAR: an elastic, scalable, highly available SQL middleware. Shahram Ghandeharizadeh, Connor Gorman, Sandy Irani, Shiva Jahangiri, Jenny Lam, Hieu Nguyen, Ryan Tani and Jason Yap, Middleware 2014.



Cache replacement with memory allocation.
Shahram Ghandeharizadeh, Sandy Irani, and Jenny Lam.
In Proceedings of the 17th Workshop on Algorithm
Engineering and Experiments (ALENEX), 2015.



Memory hierarchy design for caching middleware in the age of NVM. Shahram Ghandeharizadeh, Sandy Irani, and Jenny Lam. 7th Annual Non-Volatile Memories Workshop, 2016.

The subset assignment problem for data placement in caches. Shahram Ghandeharizadeh, Sandy Irani, and Jenny Lam. (in submission)